

SAS® SQL 1: Essentials

Course Notes

SAS® SQL 1: Essentials Course Notes was developed by Davetta Dunlap, Marya Ilgen-Lieth, and Richard Bell. Additional contributions were made by Randall Cates, Brian Gayle, Marty Hobbs, Mark Jordan, Kathy Kiraly, Linda Mitterling, Rich Papel, Lorilyn Russell, and Kathy Passarella. Editing and production support was provided by the Curriculum Development and Support Department.

SAS and all other SAS Institute Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc. in the USA and other countries. ® indicates USA registration. Other brand and product names are trademarks of their respective companies.

SAS® SQL 1: Essentials Course Notes

Copyright © 2016 SAS Institute Inc. Cary, NC, USA. All rights reserved. Printed in the United States of America. No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, electronic, mechanical, photocopying, or otherwise, without the prior written permission of the publisher, SAS Institute Inc.

Book code E70607, course code LWSQ1/SQ1, prepared date 09Mar2016.

LWSQ1_002

ISBN 978-1-62960-265-3

Table of Contents

Course Description	vi
Prerequisites	vii
Chapter 1 Introduction.....	1-1
1.1 Introducing the Structured Query Language.....	1-3
1.2 Course Logistics	1-8
Demonstration: Creating Course Data Files.....	1-15
Exercises.....	1-16
Chapter 2 Basic Queries.....	2-1
2.1 Overview of the SQL Procedure.....	2-3
2.2 Specifying Columns.....	2-14
Exercises.....	2-34
2.3 Specifying Rows	2-37
Exercises.....	2-47
2.4 Solutions	2-49
Solutions to Exercises	2-49
Chapter 3 Displaying Query Results.....	3-1
3.1 Presenting Data	3-3
Exercises.....	3-12
3.2 Summarizing Data	3-14
Exercises.....	3-36
3.3 Solutions	3-39
Solutions to Exercises	3-39
Chapter 4 SQL Joins	4-1
4.1 Introduction to SQL Joins	4-3
4.2 Inner Joins.....	4-9
Exercises.....	4-18

4.3	Outer Joins	4-21
	Exercises.....	4-32
4.4	Complex SQL Joins	4-34
	Exercises.....	4-37
4.5	Solutions	4-39
	Solutions to Exercises	4-39
Chapter 5 Subqueries		5-1
5.1	Noncorrelated Subqueries.....	5-3
	Exercises.....	5-18
5.2	In-Line Views.....	5-22
	Exercises.....	5-38
5.3	Solutions	5-42
	Solutions to Exercises	5-42
Chapter 6 Set Operators		6-1
6.1	Introduction to Set Operators.....	6-3
6.2	The UNION Operator	6-9
6.3	The OUTER UNION Operator.....	6-19
	Exercises.....	6-25
6.4	The EXCEPT Operator	6-27
6.5	The INTERSECT Operator.....	6-34
	Exercises.....	6-43
6.6	Solutions	6-46
	Solutions to Exercises	6-46
Chapter 7 Creating Tables and Views		7-1
7.1	Creating Tables with the SQL Procedure.....	7-3
	Exercises.....	7-16
7.2	Creating Views with the SQL Procedure	7-19
	Exercises.....	7-31

7.3	Solutions	7-34
	Solutions to Exercises	7-34
Chapter 8	Advanced PROC SQL Features.....	8-1
8.1	Dictionary Tables and Views	8-3
	Exercises.....	8-13
8.2	Using SQL Procedure Options.....	8-16
8.3	Interfacing PROC SQL with the Macro Language	8-23
	Demonstration: Using Macro Variables to Make Your Program Dynamic	8-38
	Exercises.....	8-40
8.4	Solutions	8-44
	Solutions to Exercises	8-44
Chapter 9	Learning More	9-1
9.1	SAS Resources.....	9-3
9.2	Beyond This Course.....	9-6
Appendix A	Self-Study Topics	A-1
A.1	Correlated Subqueries.....	A-3
	Exercises.....	A-25
A.2	Introduction to Indexes	A-26
A.3	Creating Indexes	A-33
	Exercises.....	A-39
A.4	Maintaining Tables with the SQL Procedure	A-42
	Exercises.....	A-58
A.5	Integrity Constraints.....	A-61
A.6	Interfacing PROC SQL with the Macro Language	A-70
	Demonstration: Creating and Using Macro Variables in PROC SQL.....	A-79
A.7	Solutions	A-80
	Solutions to Exercises	A-80

Course Description

This course teaches you how to process SAS data using Structured Query Language (SQL).

To learn more...



For information about other courses in the curriculum, contact the SAS Education Division at 1-800-333-7660, or send e-mail to training@sas.com. You can also find this information on the web at <http://support.sas.com/training/> as well as in the Training Course Catalog.



For a list of other SAS books that relate to the topics covered in this course notes, USA customers can contact the SAS Publishing Department at 1-800-727-3228 or send e-mail to sasbook@sas.com. Customers outside the USA, please contact your local SAS office.

Also, see the SAS Bookstore on the web at <http://support.sas.com/publishing/> for a complete list of books and a convenient order form.

Prerequisites

Before attending this class, you should be able to

- submit SAS programs on your operating system
- create and access SAS data sets
- use arithmetic, comparison, and logical operators
- invoke SAS procedures.

You can gain this experience from the SAS® Programming 1: Essentials course. No knowledge of SQL is necessary.

Chapter 1 Introduction

1.1	Introducing the Structured Query Language	1-3
1.2	Course Logistics.....	1-8
	Demonstration: Creating Course Data Files	1-15
	Exercises	1-16

1.1 Introducing the Structured Query Language

Objectives

- Describe the historical development of Structured Query Language (SQL).
- Explain how SQL is used.

3

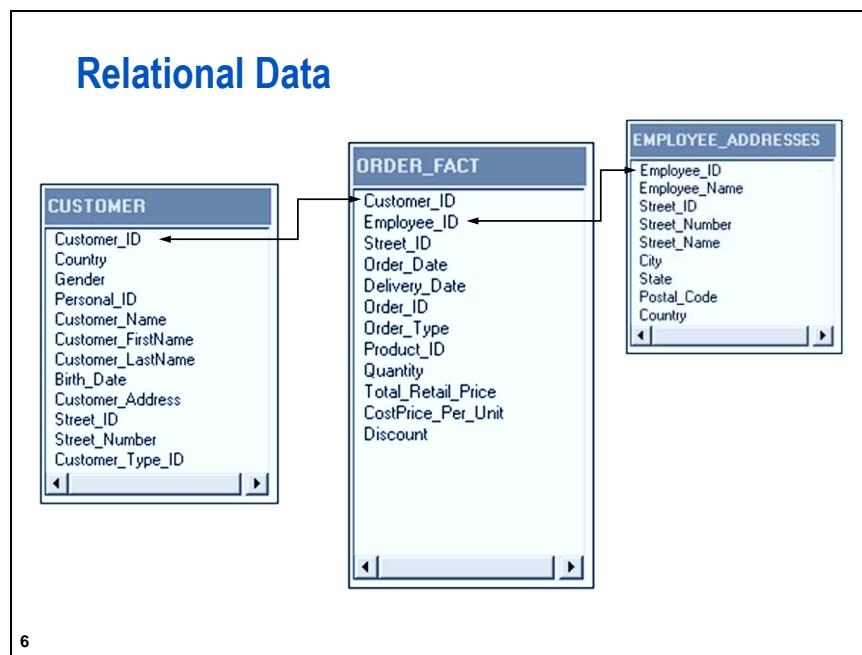
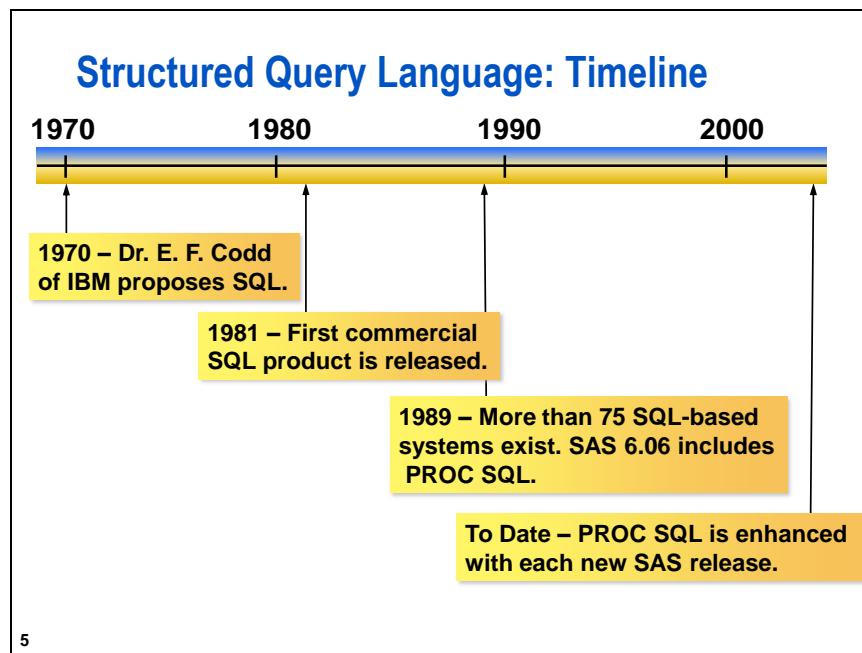
Structured Query Language

Structured Query Language (SQL) is a standardized language originally designed as a relational database query tool.

SQL is currently used in many software products to retrieve and update data.

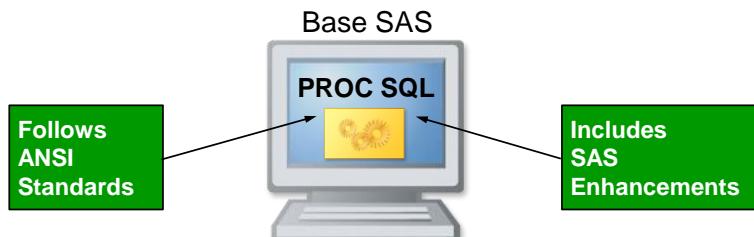


4



SQL Procedure

The SAS SQL procedure enables the use of SQL in SAS.



7

SQL Procedure



What PROC SQL Is

- A tool for querying data
- A tool for data manipulation and management
- An augmentation to the DATA step



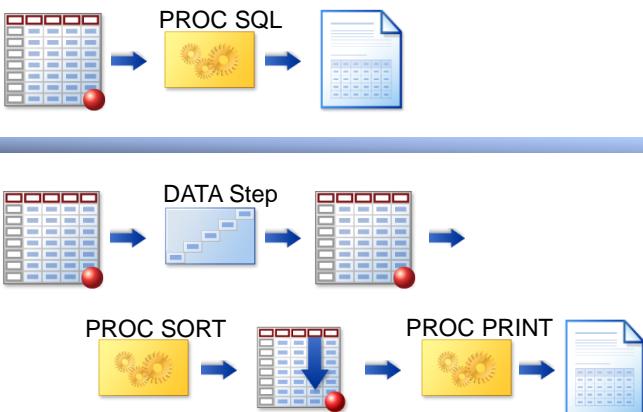
What PROC SQL Is Not

- A custom reporting tool
- A DATA step replacement

8

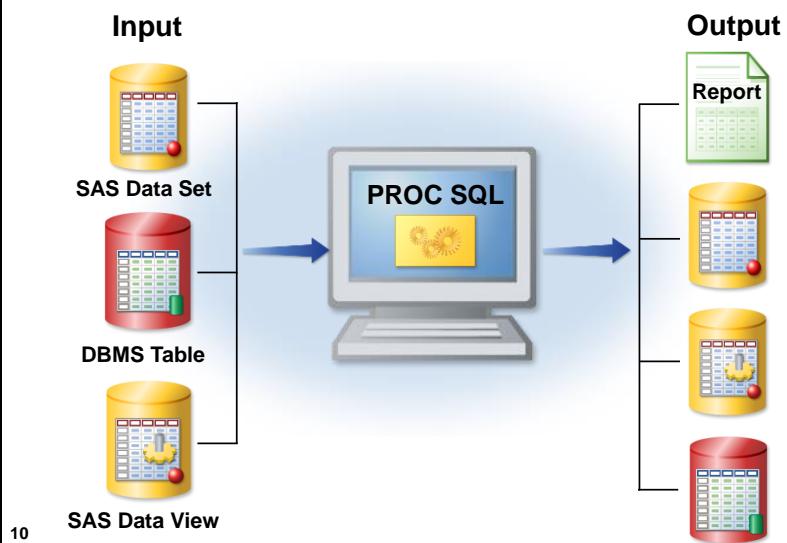
SQL Procedure versus Traditional SAS

The SQL procedure can sometimes reproduce the results of multiple DATA and procedure steps with a single query.



9

Structured Query Language



10

Terminology

Data Processing	SAS	SQL
File	Data Set	Table
Record	Observation	Row
Field	Variable	Column

11

1.01 Multiple Choice Poll

Have you worked with PROC SQL?

- a. yes, just maintaining programs
- b. yes, writing some programs
- c. no, not at all

12

1.2 Course Logistics

Objectives

- Describe the data that is used in the course.
- Identify the interfaces that are available for workshops.
- Specify the naming convention that is used for course files.
- Define the three levels of exercises.
- Explain the extended learning for this course.
- Execute a SAS program to create course data files.
- Execute a SAS program to define the data location.

15

Orion Star Sports & Outdoors

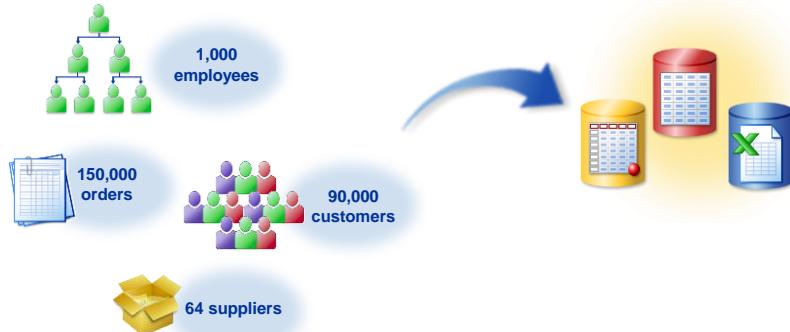
This course focuses on a fictitious global sports and outdoors retailer that has traditional stores, an online store, and a large catalog business.



16

Orion Star Data

Large amounts of data are stored in various formats.



17

SAS Programming Interfaces

In this course, three interfaces are available for you to use during workshops.

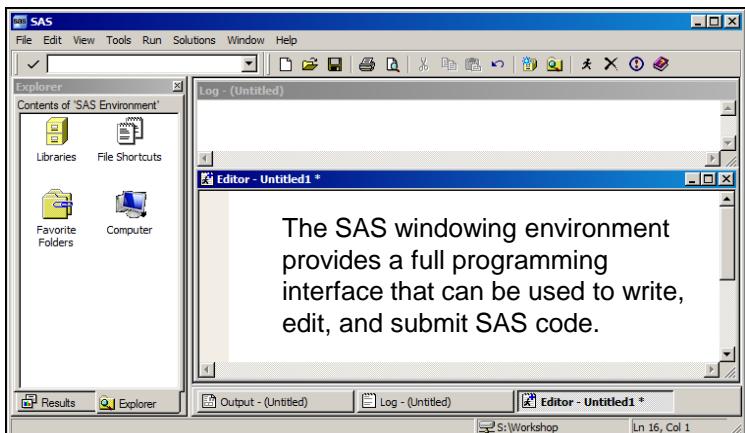


- ✍ To enable you to learn more about each of these interfaces, self-directed exercises are available in the course data location.

18

sas SAS Windowing Environment

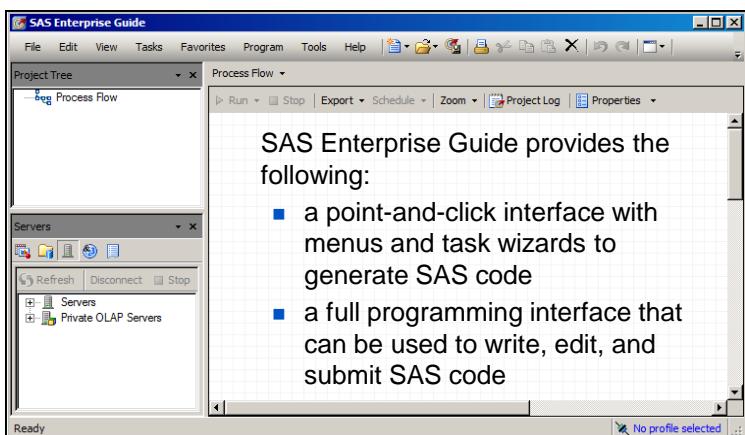
The SAS *windowing environment* is an application that is accessed from different operating environments.



19

SAS Enterprise Guide

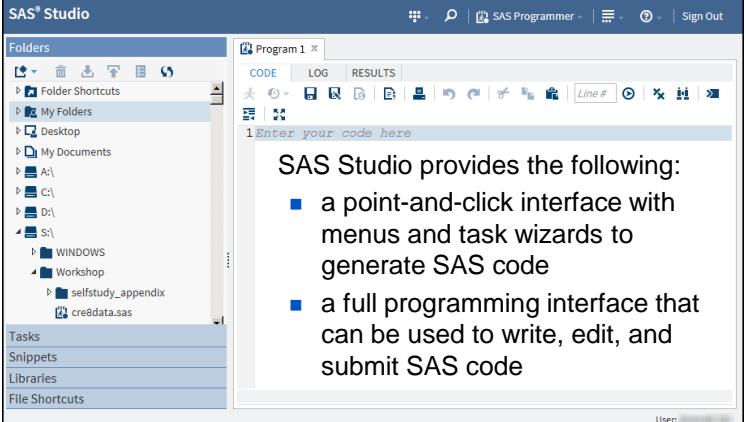
SAS *Enterprise Guide* is a client application that is accessed from the Windows operating environment.



20

 **SAS Studio**

SAS Studio is a web client that is accessed through an HTML5-compliant web browser.



SAS Studio provides the following:

- a point-and-click interface with menus and task wizards to generate SAS code
- a full programming interface that can be used to write, edit, and submit SAS code

21

1.02 Multiple Choice Poll

Which interface will you use during this course?

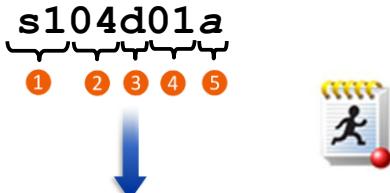
- a. SAS Studio
- b. SAS Enterprise Guide
- c. SAS windowing environment

22

Program Naming Conventions

In this course, you retrieve and save SAS programs using the structure below.

- ① course ID
- ② chapter #
- ③ type
 - a=activity
 - d=demo
 - e=exercise
 - s=solution
- ④ item #
- ⑤ placeholder



23

Three Levels of Exercises

The course is designed to have you complete only **one** set of exercises. Select the level that is most appropriate for your skill set.

Level 1	Provides step-by-step instructions.
Level 2	Provides less information and guidance.
Challenge	Provides minimal information and guidance. You might need to use SAS Help and documentation.

24

Extending Your Learning

After class, you will have access to an extended learning page that was created for this course. The page includes

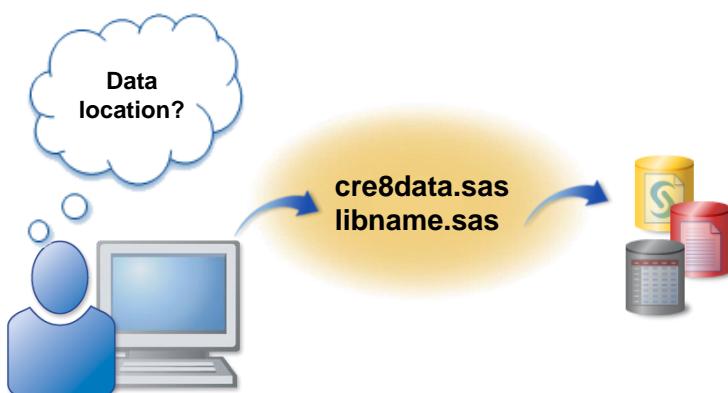
- course data and program files
- a PDF file of the course notes
- other course-specific resources.



25

Course Data Creation

You must execute programs to create the files and define the location of the course data.



26

Locating Data Files

In this course, macro variable references are used to give a more flexible approach for locating files.

Examples:

```
%let path=s:/workshop;
```



Creating Course Data Files

This demonstration illustrates how to create the course data files and define the data location.

1. In your SAS session, open the **cre8data** program.

2. Find the following %LET statement:

```
%let path=s:/workshop;
```

3. If your data files are to be created at a location other than **s:\workshop**, change the value that is assigned to the **path** macro variable to reflect the location. If your data files are created in **s:\workshop**, then no change is needed.

The **cre8data** program uses forward slashes for portability across operating environments. UNIX and Linux require forward slashes. Windows accepts forward slashes and might convert them to back slashes.

4. Submit the program.
5. View the results and verify that the output contains a list of data files.
6. Open the **libname** program, which was created by the **cre8data** program.
7. Submit the program.
8. View the log and verify that there are no errors or warnings.

End of Demonstration



Exercises



You **must** complete the exercise to create the course data files. If you do not create the data files, all programs in this course will fail.

Required Exercise

1. Creating Course Data Files

- In your SAS session, open the **cre8data** program.

- Find the following %LET statement:

```
%let path=s:/workshop;
```

- If your data files are to be created at a location other than **s:\workshop**, change the value that is assigned to the **path** macro variable to reflect the location. If your data files are created in **s:\workshop**, then no change is needed.

The **cre8data** program uses forward slashes for portability across operating environments. UNIX and Linux require forward slashes. Windows accepts forward slashes and might convert them to back slashes.

- Submit the program.
- View the results and verify that the output contains a list of data files.

The CONTENTS Procedure				
Directory				
Libref	ORION			
Engine	V9			
Physical Name	s:\workshop			
Filename	s:\workshop			
# Name Member Type File Size Last Modified				
1	CUSTOMER	DATA	131072	02/27/2015 14:18:26
2	CUSTOMER_TYPE	DATA	131072	02/27/2015 14:18:26

- Open the **libname** program, which was created by the **cre8data** program.

```
%let path=s:/workshop;
libname orion "s:/workshop";
```

- Submit the program.
- View the log and verify that there are no errors or warnings.

End of Exercises

Chapter 2 Basic Queries

2.1 Overview of the SQL Procedure	2-3
2.2 Specifying Columns	2-14
Exercises	2-34
2.3 Specifying Rows	2-37
Exercises	2-47
2.4 Solutions	2-49
Solutions to Exercises	2-49
Solutions to Student Activities (Polls/Quizzes)	2-52

2.1 Overview of the SQL Procedure

Objectives

- Identify key syntax of the SQL procedure.
- List key features of the SQL procedure.
- List key features of the SELECT statement.
- List SQL procedure statements.

3

Business Scenario

As a new Orion Star programmer, you have inherited several SQL programs. You need to learn the PROC SQL syntax to use and extend these programs and ultimately write your own.



4

The SQL Procedure

The SQL procedure is initiated with a PROC SQL statement. It is terminated with a QUIT statement.

```
proc sql;  
  select Employee_ID, Employee_Gender,  
         Salary  
    from orion.employee_information;  
quit;
```

**PROC SQL <option(s)>;
statement(s);
QUIT;**

5

s102d01

PROC SQL does not require a RUN statement. It uses the QUIT statement to explicitly terminate SQL processing. The SQL procedure, like other SAS procedures, is terminated if SAS encounters a DATA step or a PROC step.

SQL Procedure

- Multiple statements can be included in a PROC SQL step.
- Each statement defines a process and is executed immediately.

**PROC SQL <option(s)>;
statement(s);
QUIT;**

6

SELECT Statement

A *SELECT statement* is used to query one or more tables. The results of the SELECT statement are written to the default output destination.

```
proc sql;
select Employee_ID, Employee_Gender, Salary
  from orion.employee_information
 where Employee_Gender='F'
   order by Salary desc;
quit;
```

7

s102d01

SELECT Statement

A SELECT statement contains smaller building blocks called *clauses*.

```
proc sql;
select Employee_ID, Employee_Gender, Salary
  from orion.employee_information
 where Employee_Gender='F'
 clauses order by Salary desc;
quit;
```

- Although it can contain multiple clauses, each SELECT statement begins with the SELECT keyword and ends with a semicolon.

8

s102d01

Viewing the Output

Partial PROC SQL Output

The SAS System		
Employee ID	Employee Gender	Employee Annual Salary
120260	F	\$207,885
120719	F	\$87,420
120661	F	\$85,495
121144	F	\$83,505
120798	F	\$80,755

9

SELECT Statement: Required Clauses

```
SELECT object-item <, ...object-item>  
FROM from-list;
```

- The SELECT clause specifies the columns and column order.
- The FROM clause specifies the data sources.
- You can query from 1 to 256 tables.

10

SELECT Statement: Optional Clauses

```
SELECT object-item <, ...object-item>
  FROM from-list
  <WHERE sql-expression>
  <GROUP BY object-item <, ... object-item >>
  <HAVING sql-expression>
  <ORDER BY order-by-item <DESC>
    <, ...order-by-item>>;
```

- The WHERE clause specifies data that meets certain conditions.
- The GROUP BY clause groups data for processing.
- The HAVING clause specifies groups that meet certain conditions.
- The ORDER BY clause specifies an order for the data.

11

SELECT Statement Syntax

```
PROC SQL;
SELECT object-item <, ...object-item>
  FROM from-list
  <WHERE sql-expression>
  <GROUP BY object-item <, ... object-item >>
  <HAVING sql-expression>
  <ORDER BY order-by-item <DESC>
    <, ...order-by-item>>;
QUIT;
```

-  The specified order of the above clauses within the SELECT statement is required.

12

Business Scenario

You have encountered a syntax error in one of the programs. What techniques exist to check the PROC SQL code for syntax errors?



13

Setup for the Poll

- Open and submit the program **s102a01**.
- View the SAS log.
- Consider the notes from the first two steps.
 - 1) Which step generated an error?
 - 2) What error message was generated?

14

2.01 Multiple Choice Poll

Which step generated an error?

- a. Step 1
- b. Step 2

15

s102a01

2.02 Short Answer Poll

What error message was generated in Step 2?

```
/* Step 2 */
proc sql;
select Employee_ID, Employee_Gender,
       Salary
  from orion.employee_information
 order by Employee_ID
  where Employee_Gender='M';
quit;
```

17

s102a01

2.03 Quiz

- Open and submit the program **s102a02**.
- View the SAS output and log.

Did the code generate any output?

What does the note in the log indicate?

19

s102a02

Features of the VALIDATE Statement

```
proc sql;
validate
select Employee_ID, Employee_Gender,
       Salary
  from orion.employee_information
 where Employee_Gender='M'
   order by Employee_ID;
quit;
```

The VALIDATE statement

- is used only in conjunction with SELECT statements
- tests the syntax without executing the query
- checks table and column name validity
- prints error messages for invalid queries.

21

Viewing the Log

A common syntax error is to end a clause with a semicolon.

Partial SAS Log

```
39  proc sql;
40  validate
41  select Employee_ID, Employee_Gender,
42      Salary;
        -
22
76
ERROR 22-322: Syntax error, expecting one of the following:
```

22

s102d02

NOEXEC Option

To explicitly check for syntax errors without submitting the code for execution, include the NOEXEC option in the PROC SQL statement. This option applies to all statements in a PROC SQL step.

```
PROC SQL <NOEXEC>;
proc sql noexec;
select Employee_ID, Employee_Gender, Salary
  from orion.employee_information
 where Employee_Gender='F'
   order by Salary desc;
quit;
```

23

s102d03

Viewing the Log

Partial SAS Log

```
proc sql noexec;
select Employee_ID, Employee_Gender, Salary
  from orion.employee_information
 where Employee_Gender='F'
   order by Salary desc;
NOTE: Statement not executed due to NOEXEC option.
quit;
```

24

Resetting Options

After an option is specified, it remains in effect for the PROC SQL step. You can use the RESET statement to add or change PROC SQL options without restarting the procedure.

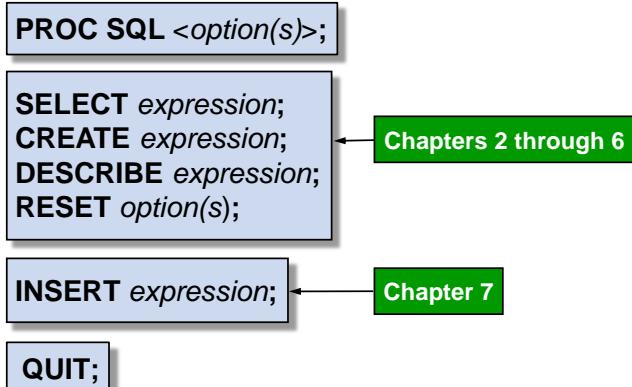
```
RESET option(s);
```

25

In a PROC SQL step, you can use options for many purposes. For details about available options, see the *SAS 9.3 SQL Procedure: User's Guide*.

Additional PROC SQL Statements

PROC SQL supports many statements in addition to the SELECT statement.



26

SELECT identifies columns to be selected.

CREATE builds new tables, views, or indexes.

DESCRIBE displays table attributes or view definitions.

INSERT adds rows of data to tables.

RESET adds to or changes PROC SQL options without re-invoking the procedure.

2.2 Specifying Columns

Objectives

- Explore unfamiliar data.
- Display columns directly from a table.
- Display columns calculated from other columns in a query.
- Calculate columns conditionally using the CASE expression.
- Create a table from query results.

29

Business Scenario

Your manager has requested analysis of a table that is unfamiliar to you. You need to learn about the PROC SQL statements and options to explore the structure of the table and the data that it contains.



30

Querying All Columns in a Table

To print all of a table's columns in the order in which they were stored, specify an asterisk in a SELECT clause.

```
proc sql;
  select *
    from orion.employee_information;
quit;
```

Partial PROC SQL Output

The SAS System								
Employee ID	Start Date	End Date	Department	Employee Annual Salary	Employee Gender	Employee Birth Date	Employee Hire Date	Employee Termination Date Manager for Employee
Employee Job Title								
120101	01JUL2007	31DEC9999	Sales Management	\$163,040	M	18AUG1980	01JUL2007	.
Director								120261

s102d04

FEEDBACK Option

When using an asterisk in the SELECT clause, add the FEEDBACK option to write the expanded SELECT statement to the SAS log.

```
PROC SQL <option(s)>;
proc sql feedback;
  select *
    from orion.employee_information;
quit;
```

Setup for the Poll

Submit the program **s102a03** and review the SAS log to answer the following question:

How are the column names represented in the expanded SELECT statement?

33

2.04 Multiple Choice Poll

How are the column names represented in the expanded SELECT statement?

- a. The column names are preceded by the table name **EMPLOYEE_INFORMATION**.
- b. The column names are preceded by the library reference **orion**.
- c. The column names are preceded by the library reference **work**.

34

Viewing the Log

The column names are preceded by the table name.

Partial SAS Log

```
proc sql feedback;
  select *
    from orion.employee_information;
NOTE: Statement transforms to:

  select EMPLOYEE_INFORMATION.Employee_ID,
EMPLOYEE_INFORMATION.Start_Date,EMPLOYEE_INFORMATION.End_Date,
EMPLOYEE_INFORMATION.Department, EMPLOYEE_INFORMATION.Job_Title,
EMPLOYEE_INFORMATION.Salary, EMPLOYEE_INFORMATION.Employee_Gender,
EMPLOYEE_INFORMATION.Birth_Date,EMPLOYEE_INFORMATION.Employee_Hire_Date,
EMPLOYEE_INFORMATION.Employee_Term_Date,EMPLOYEE_INFORMATION.Manager_ID
  from ORION.EMPLOYEE_INFORMATION;

quit;
```

36

s102a03

DESCRIBE Statement

Use the DESCRIBE statement to see the table's column names and their attributes in the SAS log.

```
proc sql;
  DESCRIBE TABLE|VIEW table-name;
  describe table orion.employee_information;
quit;
```

Partial SAS Log

```
Employee_ID num format=12. label='Employee ID',
Start_Date num format=DATE9. label='Start Date',
End_Date num format=DATE9. label='End Date',
Department char(40),
Job_Title char(25) label='Employee Job Title',
Salary num format=DOLLAR12. label='Employee Annual Salary',
Employee_Gender char(1) label='Employee Gender',
Birth_Date num format=DATE9. label='Employee Birth Date',
Employee_Hire_Date num format=DATE9. informat=DATE9.,
Employee_Term_Date num format=DATE9. informat=DATE9.,
```

37

s102d05

Business Scenario

Produce a report that contains selected information for all Orion Star employees.

`orion.employee_information`



PROC SQL



Employee_ID	Employee_Gender	Salary
120101	M	163040
120102	M	108255
120103	M	87975
120104	F	46230
120105	F	27110

39

Querying Specific Columns in a Table

List the columns that you want and the order to display them in the SELECT clause.

```
proc sql;
select Employee_ID, Employee_Gender,
      Salary
  from orion.employee_information;
quit;
```

40

s102d06



Remember to use commas to separate items in a list, such as a list of column names in the SELECT, GROUP BY, or ORDER BY clauses.

Viewing the Output

Partial PROC SQL Output

The SAS System		
Employee ID	Employee Gender	Employee Annual Salary
120101	M	\$163,040
120102	M	\$108,255
120103	M	\$87,975
120104	F	\$46,230
120105	F	\$27,110
120106	M	\$26,960
120107	F	\$30,475

41

2.05 Short Answer Poll

Submit the program **s102a04** and review the SAS log to answer the following question:

What syntax error was identified in the log?

42

Business Scenario

Modify the previous report by creating a new column, **Bonus**, which contains an amount equal to 10% of the employee's salary.

`orion.employee_information`



Employee_ID	Salary	Bonus
120101	163040	16304
120102	108255	10825.5
120103	87975	8797.5
120104	46230	4623
120105	27110	2711

44

Calculated Columns

Name the new column using the AS keyword.

```
proc sql;
  select Employee_ID, Salary,
    Salary*.10 as Bonus
  from orion.employee_information;
quit;
```

Partial PROC SQL Output

The SAS System		
Employee		
Employee ID	Annual Salary	Bonus
120101	\$163,040	16304
120102	\$108,255	10825.5
120103	\$87,975	8797.5
...		

s102d07

45

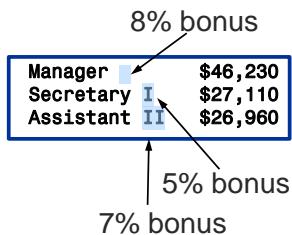


The new column name, **Bonus**, is called an *alias*. Assigning an alias to a calculated column is optional, but if an alias *is* assigned, the AS keyword is required. Omission of the alias causes the column heading in the report to be blank.

Business Scenario

You have been asked to modify the previous bonus report to conditionally calculate bonuses.

Job Title Level	Bonus Percent
I	5%
II	7%
III	10%
IV	12%
Other	8%



46

2.06 Multiple Answer Poll

Which of these SAS character functions have you used?

- a. CHAR()
- b. FIND()
- c. SCAN()
- d. SUBSTR()

47

SCAN Function

The *SCAN function* returns the *n*th word or segment from a character string after breaking it up by the delimiters.

<code>scan(Job_Title,3,' ')</code>	<code>Office Assistant III</code>
------------------------------------	-----------------------------------

<code>SCAN(string,count<,charlist><,modifier(s)>)</code>
--

<i>string</i>	a character constant, variable, or expression
<i>count</i>	an integer specifying the number of the word or segment that you want SCAN to select
<i>charlist</i>	characters used as delimiters to separate words
<i>modifier</i>	a character that modifies the action of the SCAN function

48

If the third argument (*charlist*) is omitted, the default delimiters are as shown below:

ASCII (PC, UNIX)	blank . < (+ & ! \$ *) ; - / , % ^
-------------------------	--

EBCDIC (z/0s)	blank . < (+ & ! \$ *) ; - / , % ¢ ¬
----------------------	--

The two most commonly use modifiers are I and T.

I – indicates to ignore the case of the characters.

T – indicates to trim trailing blanks from the *string* and *charlist* arguments.

Extracting the Level from Job_Title

Return the third word from **Job_Title** using a blank space as the delimiter. The SCAN function processes from left to right when the value of the count argument is positive.

<code>scan(Job_Title,3,' ')</code>

<code>Office Assistant II</code>		<code>II</code>
----------------------------------	---	-----------------



49

...

Extracting the Level from Job_Title

Some **Job_Title** values have fewer than three words. If the value of *count* is greater than the number of words in the character string, the SCAN function returns a missing value.

```
scan(Job_Title, 3, ' ')
```

Secretary I

1 2

...

50

Extracting the Level from Job_Title

The SCAN function processes from right to left when the value of the *count* argument is negative.

```
scan(Job_Title, -1, ' ')
```

Secretary I

-2 -1

Office Assistant II

-3 -2 -1

51

CASE Expression

To create new columns conditionally, use the CASE expression in the SELECT list.

```
SELECT object-item<, ...object-item>,
       CASE <case-operand>
           WHEN when-condition THEN result-expression
           <WHEN when-condition THEN result-expression>
           <ELSE result-expression>
       END <AS column>
   FROM table;
```

There are two forms of the CASE expression.

The CASE expression is evaluated for each row of a table and returns a single value.

52

Calculating Columns Conditionally: CASE

CASE Expression: Standard Form

```
select Job_Title, Salary,
       case
           when scan(Job_Title,-1,' ')='I'
               then Salary*.05
           when scan(Job_Title,-1,' ')='II'
               then Salary*.07
           when scan(Job_Title,-1,' ')='III'
               then Salary*.10
           when scan(Job_Title,-1,' ')='IV'
               then Salary*.12
           else Salary*.08
       end as Bonus
   from orion.employee_information;
```

With this form of the CASE syntax, you can use equality and non-equality tests for validity.

s102d08

53

Calculating Columns Conditionally: CASE

CASE Expression: Standard Form

```
select Job_Title, Salary,
       case
         when scan(Job_Title,-1,' ')='I'
              then Salary*.05
         when scan(Job_Title,-1,' ')='III'
              then Salary*.07
         when scan(Job_Title,-1,' ')='III'
              then Salary*.10
         when scan(Job_Title,-1,' ')='IV'
              then Salary*.12
         else Salary*.08
       end as Bonus
     from orion.emp
```

The first WHEN clause evaluated as true determines which value the CASE expression returns.
Subsequent WHEN clauses are not evaluated.

54

s102d08

Calculating Columns Conditionally: CASE

CASE Expression: Standard Form

```
select Job_Title, Salary
       case
         when scan(Job_Title,-1,' ')='I'
              then Salary*.05
         when scan(Job_Title,-1,' ')='III'
              then Salary*.07
         when scan(Job_Title,-1,' ')='III'
              then Salary*.10
         when scan(Job_Title,-1,' ')='IV'
              then Salary*.12
         else Salary*.08
       end as Bonus
     from orion.employee_information;
```

The optional ELSE expression provides an alternate action if none of the WHEN expressions are true.

If no ELSE expression is present and every WHEN condition is false, the CASE expression returns a missing value.

55

s102d08

Calculating Columns Conditionally: CASE

CASE Expression: Shortcut (Case-Operand) Form

```
select Job_Title, Salary,
       case scan(Job Title,-1,' ')
         when 'I' then Salary*.05
         when 'II' then Salary*.07
         when 'III' then Salary*.10
         when 'IV' then Salary*.12
         else Salary*.08
       end as Bonus
  from orion.employee_information;
```

Case
operand

- With this form of the CASE syntax, you can use only an equality test for validity.

56

s102d08

- The shortcut form is more efficient because the SCAN function is evaluated only once, but this method requires an EQUAL comparison operator. If you need a different operator, use the standard form.

Viewing the Output

Partial PROC SQL Output

The SAS System			
Employee Job Title	Employee Annual Salary	Bonus	
Director	\$163,040	13043.2	
Sales Manager	\$108,255	8660.4	
Sales Manager	\$87,975	7038	
Administration Manager	\$46,230	3698.4	
Secretary I	\$27,110	1355.5	
Office Assistant II	\$26,960	1887.2	
Office Assistant III	\$30,475	3047.5	
Warehouse Assistant II	\$27,660	1936.2	
Warehouse Assistant I	\$26,495	1324.75	

57

Business Scenario

Management needs a report that includes the employee identifier, gender, and age for an upcoming audit.

Here is a sketch of the desired report:

Employee_ID	Employee_Gender	Age
120101	M	32
120102	M	39

58

Business Data

Employee data is stored in the **orion.employee_information** table. An employee's age can be calculated based on his or her birthdate.

orion.employee_information



Employee_ID
Employee_Gender
Birth_Date

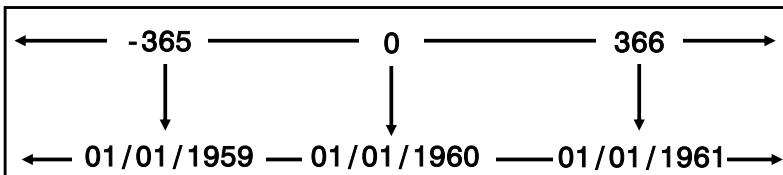
Employee_ID	Employee_Gender	Age
120101	M	32
120102	M	39

59

SAS Date Values

A SAS date is stored as the number of whole days between January 1, 1960, and the date specified.

Stored Values



Display Values (formatted MMDDYY10.)

60

SAS date constants are used throughout this course.

Syntax	Used to	Example
'ddMONyyyy'd	Convert a calendar date into a SAS date value	'15JUN2008'd as Date

Selected SAS Numeric Functions

The following SAS numeric functions are frequently used when you work with SAS dates.

Function	Used To Return	Example
TODAY()	Today's date as a SAS date value (integer)	<code>today() as Date</code>
MONTH(arg)	The month portion of a SAS date variable as an integer between 1-12	<code>month(Birth_Date) as Birth_Month</code>
INT(arg)	The integer portion of a numeric value	<code>int(fullage) as Age</code>

61



PROC SQL supports most SAS DATA step functions. However, some are not supported (for example, the LAG and DIF functions). For a complete list of functions that are not supported in SQL, see the most recent documentation.

Calculating Columns Using SAS Dates: Step 1

Calculate **Age** based on today's date being 02JAN2013 and a **Birth_Date** value of 18AUG1980.

```
19360          7535
proc sql;
select Employee_ID, Employee_Gender,
       int((today()-Birth_Date)/365.25)
     as Age
  from orion.employee_information;
quit;
```

62

s102d09

Calculating Columns Using SAS Dates: Step 2

Calculate **Age** based on today's date being 02JAN2013 and a **Birth_Date** value of 18AUG1980.

```
proc sql;
select Employee_ID, Employee_Gender,
       int((today()-Birth_Date)/365.25)
     as Age
  from orion.employee_information; 32.3751
quit;
```

63

s102d09

Calculating Columns Using SAS Dates: Step 3

Calculate **Age** based on today's date being 02JAN2013 and a **Birth_Date** value of 18AUG1980.

```
proc sql;
select Employee_ID, Employee_Gender,
32      int((today()-Birth_Date)/365.25)
      as Age
from orion.employee_information;
quit;
```

64

s102d09

Viewing the Output

Partial PROC SQL Output

The SAS System		
Employee ID	Employee Gender	Age
120101	M	32
120102	M	39
120103	M	59
120104	F	54
120105	F	34
120106	M	64
120107	F	59

The values of **Age** vary based on the date that the program is executed.

65

2.07 Short Answer Poll

What date function would you use to create the **Birth_Month** column for the following rows from the **employee_information** table?

The SAS System			
Employee_ID	Birth_Date	Birth_Month	Employee_Gender
120101	7535	8	M
120102	4971	8	M
120103	-2535	1	M
120104	-600	5	F
120105	6929	12	F
120106	-4026	12	M
120107	-2536	1	F

66

Business Scenario

To support management's requests for reports, you need to create a new table that contains selected columns from an existing table and a calculated column named **Birth_Month**.



68

Creating and Populating a Table

To define the structure of the **work.birth_months** table, use the SELECT list.

```

CREATE TABLE table-name AS
query-expression;

proc sql;
create table work.birth_months as
select Employee_ID, Birth_Date,
month(Birth_Date) as Birth_Month, ← column alias
Employee_Gender
from orion.employee_information;
describe table work.birth_months;
select * from work.birth_months;
quit;

```

69

s102d11

 If a new column is computed, the column name in the new table is the name of the alias.

If a table is created and an alias is not assigned, the omission of the AS keyword causes the new column name to be **_TEMA001**. If other new columns are created without the AS keyword, the new names would be **_TEMA002**, **_TEMA003**, and so on.

Viewing the Log

Partial SAS Log

```

121 proc sql;
122 create table work.birth_months as
123 select Employee_ID, Birth_Date,
124 month(Birth_Date) as Birth_Month,
125 Employee_Gender
126 from orion.employee_information;
NOTE: Table WORK.BIRTH_MONTHS created, with 424 rows and 4 columns.

127 describe table work.birth_months;
NOTE: SQL table WORK.BIRTH_MONTHS was created like:

create table WORK.BIRTH_MONTHS( bufsize=4096 )
(
  Employee_ID num format=12.,
  Birth_Date num,
  Birth_Month num,
  Employee_Gender char(1)
);

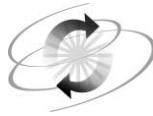
```

70

Viewing the Output

Partial PROC SQL Output

The SAS System			
Employee_ID	Birth_Date	Birth_Month	Employee_Gender
120101	7535	8	M
120102	4971	8	M
120103	-2535	1	M
120104	-600	5	F
120105	6929	12	F
120106	-4026	12	M
120107	-2536	1	F



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Querying a Table

- a. Write a query that displays all rows and all columns from the **orion.employee_payroll** table.

Partial PROC SQL Output

Employee_ID	Employee_Gender	Salary	Birth_Date	Hire_Date	Employee_Term_Date	Marital_Status	Dependents
120101	M	163040	7535	17384	.	S	0
120102	M	108255	4971	12205	.	O	2
120103	M	87975	-2535	6575	.	M	1
120104	F	46230	-600	9132	.	M	1
120105	F	27110	6929	15826	.	S	0

- b. Modify the previous query so that only the columns for **Employee_ID**, **Employee_Gender**, **Marital_Status**, and **Salary** are displayed.

Partial PROC SQL Output

Employee_ID	Employee_Gender	Marital_Status	Salary
120101	M	S	163040
120102	M	O	108255
120103	M	M	87975
120104	F	M	46230
120105	F	S	27110

Level 2

2. Calculating a Column

Write a query that generates the report below. The report should do the following:

- display **Employee_ID**, **Employee_Gender**, **Marital_Status**, **Salary**, and a new column (**Tax**) that is one-third of the employee's salary
- use the **orion.employee_payroll** table

Partial PROC SQL Output

Employee_ID	Employee_Gender	Marital_Status	Salary	Tax
120101	M	S	163040	54346.67
120102	M	O	108255	36085
120103	M	M	87975	29325
120104	F	M	46230	15410
120105	F	S	27110	9036.667
120106	M	M	26960	8986.667

3. Creating a New Table

Create a new table named **work.bonus**.

- The new table should include these columns: **Employee_ID**, **Salary**, a new column representing a 4% bonus (**Bonus**).
- Use the **orion.employee_payroll** table as the basis.
- In the same SQL step, generate a report of the new table.

SAS Log Output

```
NOTE: Table WORK.BONUS created, with 424 rows and 3 columns.
```

Partial PROC SQL Output

Employee_ID	Salary	Bonus
120101	163040	6521.6
120102	108255	4330.2
120103	87975	3519
120104	46230	1849.2
120105	27110	1084.4
120106	26960	1078.4

Challenge

4. Conditional Processing

Create a report that displays **Employee_ID**, **Level**, **Salary**, and **Salary_Range** using the **orion.staff** table. **Level** and **Salary_Range** are two new columns in the report. The report should also contain salary information for only the Orion Star executives. Conditionally, assign values to the two new columns as follows:

Job_Title (Last Word)	Level	Salary Ranges		
		Low	Medium	High
Manager	Manager	< 52,000	52,000-72,000	> 72,000
Director	Director	<108,000	108,000-135,000	> 135,000
Officer, President	Executive	<240,000	240,000-300,000	> 300,000
Other	not applicable	DO NOT INCLUDE IN THE REPORT.		

Partial PROC SQL Output

		Employee Annual		
Employee	ID	Level	Salary	Salary_Range
120101	Director	\$163,040	High	
120659	Director	\$161,290	High	
121142	Director	\$156,065	High	
120800	Director	\$80,210	Low	
120270	Director	\$48,435	Low	
120259	Executive	\$433,800	High	

Why are you seeing column headings instead of column names in your output?

End of Exercises

2.3 Specifying Rows

Objectives

- Eliminate duplicate rows in a query.
- Select a subset of rows in a query.

75

Business Scenario

Display the names of the Orion Star departments using the **orion.employee_information** table.

Desired Report

The SAS System	
Department	
<hr/>	
Accounts	
Accounts Management	
Administration	
Concession Management	
...	

No duplicates should appear in the report.

76

Displaying All Rows

```
proc sql;
  select Department
    from orion.employee_information;
quit;
```

Partial PROC SQL Output

The SAS System
Department
Sales Management
Sales Management
Sales Management
Administration

s102d12

77

Remove duplicates.



Eliminating Duplicate Rows

Use the *DISTINCT* keyword to eliminate duplicate rows.

```
proc sql;
  select distinct Department
    from orion.employee_information;
quit;
```

- ✍ The DISTINCT keyword applies to all columns in the SELECT list. One row is displayed for each unique combination of values.

78

s102d13



- ✍ UNIQUE is an alternative keyword for DISTINCT. It is a SAS enhancement.

Viewing the Output

Partial PROC SQL Output

The SAS System	
Department	
Accounts	one row per department
Accounts Management	
Administration	
Concession Management	
Engineering	
Executives	
Group Financials	
Group HR Management	
IS	
Logistics Management	
Marketing	
Purchasing	
Sales	

79

2.08 Multiple Choice Poll

Which partial SELECT statement selects only the unique combinations of **Employee_Gender** and **Job_Title**?

- a. `select distinct Employee_Gender,
distinct Job_Title...`
- b. `select distinct Employee_Gender,
Job_Title...`
- c. `select distinct Employee_Gender
Job_Title...`

80

Business Scenario

Management requested a list of employees whose salaries exceed \$112,000.

`orion.employee_information`



Employee ID	Employee Job Title	Employee Annual Salary
120101	Director	\$163,040
120259	Chief Executive Officer	\$433,800
120260	Chief Marketing Officer	\$207,885
120261	Chief Sales Officer	\$243,190
120262	Chief Financial Officer	\$268,455

82

Subsetting with the WHERE Clause

Use a WHERE clause to specify a condition that the data must satisfy before being selected.

```

proc sql;
select Employee_ID, Job_Title, Salary
  from orion.employee_information
 where Salary > 112000;
quit;
  
```

WHERE sql-expression

83

s102d14

Viewing the Output

PROC SQL Output

The SAS System		Employee Annual Salary
Employee ID	Employee Job Title	
120101	Director	\$163,040
120259	Chief Executive Officer	\$433,800
120260	Chief Marketing Officer	\$207,885
120261	Chief Sales Officer	\$243,190
120262	Chief Financial Officer	\$268,455
120659	Director	\$161,290
121141	Vice President	\$194,885
121142	Director	\$156,065

84

Subsetting: Comparison Operators

You can use all common comparison operators in a WHERE clause.

Mnemonic	Symbol	Definition
LT †	<	Less than
GT †	>	Greater than
EQ †	=	Equal to
LE †	<=	Less than or equal to
GE †	>=	Greater than or equal to
NE †	< >	Not equal to
	¬= †	Not equal to (EBCDIC)
	¬= †	Not equal to (ASCII)

85

† SAS enhancement

Subsetting: Logical Operators

Use only one WHERE clause in a SELECT statement.
To specify multiple subsetting criteria, combine
expressions with logical operators.

Mnemonic	Symbol	Definition
OR	†	or, either
AND	& †	and, both
NOT	¬ †	not, negation (EBCDIC)
NOT	^ †	not, negation (ASCII)

86

† SAS enhancement

Subsetting: Special Operators

Common WHERE clause operators with examples:

Operator	Example
IN	<code>where JobCategory in ('PT', 'NA', 'FA')</code>
CONTAINS or ? †	<code>where word ? 'LAM'</code>
IS NULL or IS MISSING †	<code>where Product_ID is missing</code>
BETWEEN – AND	<code>where Salary between 70000 and 80000</code>
SOUNDS LIKE (=*) †	<code>where LastName =* 'SMITH'</code>
LIKE using % or _	<code>where Employee_Name like 'H%'</code> <code>where JobCategory like '__1'</code>

87

† SAS enhancement

Alternative statements for using the IS NULL or IS MISSING operator are as follows:

- `where Product_ID = ''`
- `where Product_ID = .`

With the = operator, you must know whether **Product_ID** is character or numeric. However, if you use IS MISSING or IS NULL, you do not need advance knowledge of the column type.

2.09 Short Answer Poll

Modify program **s102a05** to provide a WHERE expression that selects only those rows in which the employees' first names begin with N.

Desired PROC SQL Output

The SAS System	
Employee_Name	Employee_ID
Apr, Nishan	120759
James, Narelle	120155
Kokoszka, Nikeisha	120765
Plybon, Nicholas	120276
Post, Nahliah	120748
Smith, Nasim	121032

88

Business Scenario

Management requested a report that includes only those employees who receive bonuses less than \$3000.

orion.employee_information



Employee_ID	Employee_Gender	Salary	Bonus
120105	F	27110	2711
120106	M	26960	2696
120108	F	27660	2766
120109	F	26495	2649.5
120110	M	28615	2861.5

90

Subsetting with Calculated Values

First attempt:

```
proc sql;
  select Employee_ID, Employee_Gender,
         Salary, Salary*.10 as Bonus
    from orion.employee_information
   where Bonus<3000;
quit;
```

A **WHERE** clause is evaluated before the **SELECT** clause. Therefore, columns used in the WHERE clause must exist in the table.

Partial SAS Log

```
ERROR: The following columns were not found in the contributing
tables: Bonus.
```

91

s102d15

Subsetting with Calculated Values

One solution is to repeat the calculation in the WHERE clause.

```
proc sql;
  select Employee_ID, Employee_Gender,
         Salary, Salary*.10 as Bonus
    from orion.employee_information
   where Salary*.10<3000;
quit;
```

ANSI standard

92

s102d15

Subsetting with Calculated Values

An alternate method is to use the CALCULATED keyword in the WHERE clause.

```
proc sql;
select Employee_ID, Employee_Gender,
       Salary, Salary*.10 as Bonus
  from orion.employee_information
 where calculated Bonus<3000;
quit;
```

SAS enhancement

93

s102d15

Viewing the Output

Partial PROC SQL Output

The SAS System			
Employee ID	Employee Gender	Employee	
		Annual Salary	Bonus
120105	F	\$27,110	2711
120106	M	\$26,960	2696
120108	F	\$27,660	2766
120109	F	\$26,495	2649.5
120110	M	\$28,615	2861.5
120111	M	\$26,895	2689.5
120112	F	\$26,550	2655

94

Using the Calculated Keyword

You can also use the CALCULATED keyword in other parts of a query.

```
proc sql;
select Employee_ID, Employee_Gender,
       Salary, Salary*.10 as Bonus,
       calculated Bonus/2 as Half
  from orion.employee_information
 where calculated Bonus<3000;
quit;
```

95

s102d15



The CALCULATED keyword or repeated calculated column expression is required when referring to any calculated column, character or numeric, in the SELECT or WHERE clause, but it is not necessary with the ORDER BY or HAVING clause.

Example:

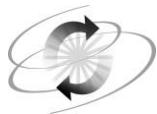
```
select Employee_ID, Salary,
       (scan(Job_Title,-1,' ')) as Job_Level
  from orion.Staff
 where calculated Job_Level='IV';
```

Viewing the Output

Partial PROC SQL Output

The SAS System				
Employee ID	Employee Gender	Employee	Annual	
		Salary	Bonus	Half
120105	F	\$27,110	2711	1355.5
120106	M	\$26,960	2696	1348
120108	F	\$27,660	2766	1383
120109	F	\$26,495	2649.5	1324.75
120110	M	\$28,615	2861.5	1430.75
120111	M	\$26,895	2689.5	1344.75
120112	F	\$26,550	2655	1327.5

96



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

5. Eliminating Duplicates

Write a query that generates a report that displays the cities where the Orion Star employees reside. The report should do the following:

- display one unique row per **City**
- use the **orion.employee_addresses** table

PROC SQL Output

city
Melbourne
Miami-Dade
Philadelphia
San Diego
Sydney

Level 2

6. Subsetting Data

Write a query that generates a report that displays Orion Star employees whose charitable contributions exceed \$90.00. The report should have the following characteristics:

- display **Employee_ID**, **Recipients**, and the new column **Total** that represents the total charitable contribution for each employee over the four quarters
- use the **orion.employee_donations** table
- include only employees whose charitable contribution **Total** for all four quarters exceeds \$90.00

Hint: The total charitable contribution is calculated by adding the amount of **Qtr1**, **Qtr2**, **Qtr3**, and **Qtr4**. Use the SUM function to ensure that missing values are ignored.

Partial PROC SQL Output

Employee ID	Recipients	Total
120660	Disaster Assist, Inc.	100
120677	EarthSalvors 60%, Vox Victimas 40%	100
120753	Conserve Nature, Inc. 50%, AquaMissions International 50%	100
120766	Mitleid International 80%, Save the Baby Animals 20%	100
120791	Child Survivors	120

Challenge

7. Subsetting Data Using the ESCAPE Clause

Create a report that displays the **Employee_ID** and **Recipients** for all employees who contributed 90% of their charitable contributions to a single company that was incorporated (Inc.) and is the last recipient in the list to receive the 90% contribution. Use the **orion.employee_donations** table.

Hint: Use the ESCAPE clause in the WHERE clause to solve this problem.



Alternative methods can be used to solve this problem, but for this exercise, use the LIKE operator with an ESCAPE clause.

PROC SQL Output

Employee ID	Recipients
120783	Disaster Assist, Inc. 10%, Cancer Cures, Inc. 90%
121012	Child Survivors 10%, Disaster Assist, Inc. 90%
121136	Disaster Assist, Inc. 10%, Cancer Cures, Inc. 90%

End of Exercises

2.4 Solutions

Solutions to Exercises

1. Querying a Table

a.

```
*** s102s01 ***
proc sql;
select *
  from orion.employee_payroll;
quit;
```

b.

```
proc sql;
select Employee_ID, Employee_Gender, Marital_Status, Salary
  from orion.employee_payroll;
quit;
```

2. Calculating a Column

```
*** s102s02 ***
proc sql;
select Employee_ID, Employee_Gender, Marital_Status,
       Salary, Salary/3 as Tax
  from orion.employee_payroll;
quit;
```

3. Creating a New Table

```
*** s102s03 ***
proc sql;
create table work.bonus as
select Employee_ID, Salary,
       Salary * .04 as Bonus
  from orion.employee_payroll;
select * from work.bonus;
quit;
```

4. Conditional Processing

```
*** s102s04 ***
proc sql;
select Employee_ID,
       case (scan(Job_Title,-1," "))
         when "Manager" then "Manager"
         when "Director" then "Director"
         when "Officer" then "Executive"
         when "President" then "Executive"
         else "N/A"
       end as Level,
       Salary,
```

```

case (calculated Level)
when "Manager" then
  case
    when (Salary>72000) then "High"
    when (Salary>52000) then "Medium"
    else "Low"
  end
when "Director" then
  case
    when (Salary>135000) then "High"
    when (Salary>108000) then "Medium"
    else "Low"
  end
when "Executive" then
  case
    when (Salary>300000) then "High"
    when (Salary>240000) then "Medium"
    else "Low"
  end
else "N/A"
end as Salary_Range
from orion.Staff
where calculated level ne "N/A"
order by Level, Salary desc;
quit;

```

Why are you seeing column headings instead of column names in your output?

The label attribute is associated with these columns in the descriptor portion of the input SAS data set.

These labels are automatically used by PROC SQL.

5. Eliminating Duplicates

```

*** s102s05 ***;
proc sql;
select distinct City
  from orion.employee_addresses;
quit;

```

6. Subsetting Data

```

*** s102s06 ***;
proc sql;
select Employee_ID, Recipients,
      sum(Qtr1,Qtr2,Qtr3,Qtr4) as Total
    from orion.employee_donations
   where calculated Total>90;
quit;

```

7. Subsetting Data Using the ESCAPE Clause

```
*** s102s07 ***;
proc sql;
select Employee_ID, Recipients
  from orion.employee_donations
 where Recipients like "% Inc. 90~%" ESCAPE "~";
quit;
```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

2.01 Multiple Choice Poll – Correct Answer

Which step generated an error?

- a. Step 1
- b.** Step 2

16

s102a01

2.02 Short Answer Poll – Correct Answer

What error message was generated in Step 2?

```
/* Step 2 */
proc sql;
select Employee_ID, Employee_Gender,
       Salary
  from orion.employee_information
 order by Employee_ID
where Employee_Gender='M' ;
quit;
```

ERROR: Syntax error, expecting one of the
following....

18

s102a01

2.03 Quiz – Correct Answer

Did the code generate any output?

no

What does the note in the log indicate?

Valid syntax – the VALIDATE statement checks the SELECT statement syntax.

Partial SAS Log

```
proc sql;
validate
select Employee_ID, Employee_Gender,
       Salary
  from orion.employee_information
 where Employee_Gender='M'
 order by Employee_ID;
NOTE: PROC SQL statement has valid syntax.
```

20

s102a02

2.04 Multiple Choice Poll – Correct Answer

How are the column names represented in the expanded SELECT statement?

- a. The column names are preceded by the table name **EMPLOYEE_INFORMATION**.
- b. The column names are preceded by the library reference **orion**.
- c. The column names are preceded by the library reference **work**.

35

2.05 Short Answer Poll – Correct Answer

What syntax error was identified in the log?

Syntax error, expecting one of the following:

The commas used to separate the items in a list were omitted.

- Remember to use commas to separate items in a list, such as a list of column names in a SELECT, GROUP BY, or ORDER BY clause.

43

2.07 Short Answer Poll – Correct Answer

What date function would you use to create the **Birth_Month** column for the following rows from the **employee_information** table? **MONTH** function

```
proc sql;
select Employee_ID, Birth_Date,
       month(Birth_Date) as Birth_Month,
       Employee_Gender
  from orion.employee_information;
quit;
```

67

s102d10

2.08 Multiple Choice Poll – Correct Answer

Which partial SELECT statement selects only the unique combinations of Employee_Gender and Job_Title?

- a. `select distinct Employee_Gender,
distinct Job_Title...`
- b.** `select distinct Employee_Gender,
Job_Title...`
- c. `select distinct Employee_Gender
Job_Title...`

81

2.09 Short Answer Poll – Correct Answer

Modify program **s102a05** to provide a WHERE expression that selects only those rows in which the employees' first names begin with N.

One possible solution:

```
select Employee_Name, Employee_ID  
from orion.employee_addresses  
where Employee_Name contains ', N';
```

89

s102a05s

Chapter 3 Displaying Query Results

3.1 Presenting Data	3-3
Exercises	3-12
3.2 Summarizing Data.....	3-14
Exercises	3-36
3.3 Solutions	3-39
Solutions to Exercises	3-39
Solutions to Student Activities (Polls/Quizzes)	3-42

3.1 Presenting Data

Objectives

- Display a query's results in a specified order.
- Use SAS formats, labels, and titles to enhance the appearance and usability of a query's output.
- Use PROC SQL options to control the query's output.

3

Business Scenario

Management wants to analyze employee donation information. They requested two reports.

Report 1: first quarter donations in descending order

Employee ID	Qtr1
121143	35
121142	35
121145	35
120791	30
...	

Report 2: maximum quarterly donations in descending order with employee IDs in ascending order

Employee ID
120265
120736
120270
120679
...

4

Ordering Rows

Report 1: Use the ORDER BY clause to order the query results.

```
proc sql;
  select Employee_ID, Qtr1
    from orion.employee_donations
      order by Qtr1 desc;
quit; ORDER BY order-by-item <DESC>
          <,...order-by-item <DESC>>
```

The default sort order when using an ORDER BY clause is ascending (no keyword or **ASC**). Use the **DESC** keyword following the column name to reverse the order.

5

s103d01

- The order of your output is guaranteed only if you include an ORDER BY clause in your query and the order is guaranteed only for the columns specified in the ORDER BY clause.

Viewing the Output

Partial PROC SQL Output

Employee ID	Qtr1
121143	35
121142	35
121145	35
120791	30
120677	25
120736	25



descending order

6

Ordering Rows

In an ORDER BY clause, *order-by-item* is one of the following:

```
ORDER BY order-by-item <DESC>
    <,...order-by-item <DESC>>
```

- a column name from any table in the FROM clause, even if the column is not in the SELECT clause
- a column alias
- an integer representing the position of an item in the SELECT clause
- an sql-expression

If more than one *order-by-item* is specified, then the first one determines the major sort order.

7

3.01 Multiple Choice Poll

Which ORDER BY clause orders a report by descending **State** and descending **City**?

- a. `order by desc State, City`
- b. `order by State, City desc`
- c. `order by State desc, City desc`
- d. `order by desc State, desc City`

8

Ordering Rows

Report 2: Order the results in descending order of maximum quarterly donation and then by employee ID in ascending order.

```
proc sql;
  select Employee_ID,
         max(Qtr1,Qtr2,Qtr3,Qtr4)
    from orion.employee_donations
   where Paid_By="Cash or Check"
  order by 2 desc, Employee_ID;
quit;
```

Mix and match!

10

s103d02

Viewing the Output

Partial PROC SQL Output

Employee ID	
120265	25
120736	25
120270	20
120679	20
120759	20
120760	20
120681	15
120734	15

11

Business Scenario

Enhance the previous report by adding descriptive headings, displaying the donation column with dollar signs, adding titles and constant text, and by numbering the rows.

The diagram illustrates a transformation process. On the left, a table titled "Employee ID" lists eight employee IDs with their corresponding maximum quarterly donations: 120265 (\$25), 120736 (\$25), 120270 (\$20), 120679 (\$20), 120759 (\$20), 120760 (\$20), 120681 (\$15), and 120734 (\$15). A blue arrow points from this table to a larger table on the right. The right table is titled "Maximum Quarterly Donation" and includes a header row with "Row", "Employee ID", and "Maximum". Below this, eight rows of data are shown, each mapping an employee ID to its maximum donation amount, such as Row 1 for Employee ID 120265 with a maximum donation of \$25.

Maximum Quarterly Donation			
Row	Employee ID		Maximum
1	120265	Maximum Donation is:	\$25
2	120736	Maximum Donation is:	\$25
3	120270	Maximum Donation is:	\$20
4	120679	Maximum Donation is:	\$20
5	120759	Maximum Donation is:	\$20
6	120760	Maximum Donation is:	\$20
7	120681	Maximum Donation is:	\$15
8	120734	Maximum Donation is:	\$15

12

Adding Formats and Column Labels

Column labels and formats must follow the column name and precede the comma.

```
proc sql;
select Employee_ID 'Employee ID',
       max(Qtr1,Qtr2,Qtr3,Qtr4)
         label='Maximum' format=dollar5.
      from orion.employee_donations
     where Paid_By="Cash or Check"
       order by 2 desc, Employee_ID;
quit;
```

The code is annotated with two green boxes. One box, labeled "ANSI standard", points to the line "label='Maximum' format=dollar5.". Another box, labeled "SAS enhancements", points to the line "from orion.employee_donations".

13

s103d03

Viewing the Output

Partial PROC SQL Output

Employee ID	Maximum
120265	\$25
120736	\$25
120270	\$20
120679	\$20
120759	\$20
120760	\$20
120681	\$15
120734	\$15

labels

format

14

Adding Titles and Constant Text

```
proc sql;
  title 'Maximum Quarterly Donation';
  select Employee_ID 'Employee ID',
         'Maximum Donation is:',
         max(Qtr1,Qtr2,Qtr3,Qtr4)
           label='Maximum' format=dollar5.
    from orion.employee_donations
   where Paid_By="Cash or Check"
   order by 3 desc, Employee_ID;
quit;
```

TITLE 'text';

'constant text' <AS alias> <column label>

15

s103d04

- ✍ TITLE and FOOTNOTE statements must precede the SELECT statement.
- ✍ When text is the only column specification, PROC SQL interprets it as a value, not a label.

Viewing the Output

Partial PROC SQL Output

Maximum Quarterly Donation	
Employee ID	Maximum
120265	Maximum Donation is: \$25
120736	Maximum Donation is: \$25
120270	Maximum Donation is: \$20
120679	Maximum Donation is: \$20
120759	Maximum Donation is: \$20
120760	Maximum Donation is: \$20
120681	Maximum Donation is: \$15
120734	Maximum Donation is: \$15

16

Adding a Row Number

```

PROC SQL <option(s)>;
proc sql number;
title 'Maximum Quarterly Donation';
select Employee_ID 'Employee ID',
      'Maximum Donation is:',
      max(Qtr1,Qtr2,Qtr3,Qtr4)
      label='Maximum' format=dollar5.
      from orion.employee_donations
      where Paid_By="Cash or Check"
      order by 3 desc, Employee_ID;
quit;

```

s103d05

17



There are some PROC SQL statement options that affect the SAS LISTING (Output) destination only: NOFLOW, FLOW=n, FLOW<=n, <m>> and DOUBLE. These options are ignored in other output destinations such as HTML.

Viewing the Output

Partial PROC SQL Output

Maximum Quarterly Donation		
Row	Employee ID	Maximum
1	120265	Maximum Donation is: \$25
2	120736	Maximum Donation is: \$25
3	120270	Maximum Donation is: \$20
4	120679	Maximum Donation is: \$20
5	120759	Maximum Donation is: \$20
6	120760	Maximum Donation is: \$20
7	120681	Maximum Donation is: \$15
8	120734	Maximum Donation is: \$15

18

PROC SQL Options

Options in the PROC SQL statement, such as line numbers, remain in effect until they are changed, deleted, or until a new PROC SQL statement is executed. Options can be added, removed, or changed between PROC SQL statements with the RESET statement.

```
PROC SQL option(s);
SELECT ...;

RESET option(s);

SELECT ...;
QUIT;
```

19

Setup for the Poll

Open the program file **s103a01**.

Submit the program and view the results.

Modify the program as follows:

- Add an option in the PROC SQL statement so that the rows of the first report are numbered.
- Add a new statement with the appropriate option after the first SELECT statement so that the rows of the second report are not numbered.

Resubmit the program and view the results.

20

3.02 Short Answer Poll

Which option, NUMBER or NONNUMBER, applies to the third report?

21



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Enhancing Output by Ordering the Rows and Adding Titles and Formats

Open program **s103e01** and modify the query.

- Select only the **Employee_ID**, **Salary**, and **Tax** columns.
- Display the **Tax** and **Salary** columns using the COMMA10.2 format.
- Add labels to match the output below.
- Order the report by **Salary** in descending order.
- Add this title to the report: **Single Male Employee Salaries**.

Partial PROC SQL Output

Single Male Employee Salaries			
Employee ID	Annual Salary	Federal Tax	
		Withdrawn	
121141	194,885.00	64,961.67	
120101	163,040.00	54,346.67	
120268	76,105.00	25,368.33	
120724	63,705.00	21,235.00	
120660	61,125.00	20,375.00	
...			
121011	25,735.00	8,578.33	
121044	25,660.00	8,553.33	
121010	25,195.00	8,398.33	

Level 2

2. Using Formats to Limit the Width of Columns in the Output

Write a query that retrieves the **Supplier_Name**, **Product_Group**, and **Product_Name** columns from the table **orion.product_dim**.

- Add this title to the report: **Australian Clothing Products**.
- Include only rows where the following are true: **Product_Category="Clothes"** and **Supplier_Country="AU"** (Australia).
- To enable the report to be printed in portrait orientation, use formats to limit the width of the **Supplier_Name** column to 18 characters, **Group** to 12, and **Product_Name** to 30.

- d. Label the columns **Supplier**, **Group**, and **Product**, respectively.
- e. Order the report by **Product_Name**.

PROC SQL Output

Supplier	Group	Product
Typhoon Clothing	Street Wear	Tyfoon Flex Shorts
Typhoon Clothing	Street Wear	Tyfoon Ketch T-Shirt
Typhoon Clothing	Street Wear	Tyfoon Oliver Sweatshirt

Challenge

3. Enhancing Output with Multiple Techniques

Create a report that displays **Customer_ID** and the customer's name displayed as **LastName**, **FirstName**, and **Gender**, as well as the customer's age as of 02FEB2013. Use the data contained in the **orion.customer** table. Include only U.S. customers who were more than 50 years old on 02FEB2013. Present the data ordered by descending age and customer name. Give the report an appropriate title. Limit the space used to display the customer's name to a maximum of 20 characters so that the report can be printed in portrait orientation. The **Customer_ID** values must be displayed with leading zeros as shown in this sample report.

PROC SQL Output

US Customers >50 Years Old as of 02FEB2013			
Customer ID	Last Name, First Name	Gender	Age
0000056	Siferd, Roy	M	75
0000089	Lewis, Wynella	F	74
0000092	Celii, Lendon	M	64
0000023	Devereaux, Tulio	M	59
0000018	Asmussen, Tonie	M	55
0000017	Evans, Jimmie	M	54

End of Exercises

3.2 Summarizing Data

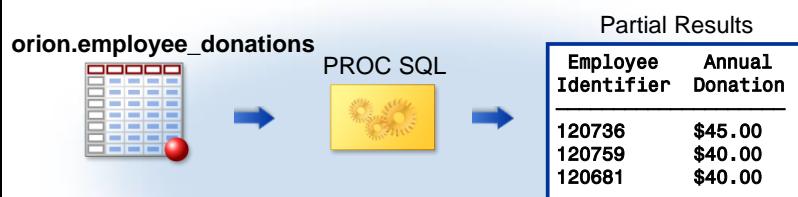
Objectives

- Use functions to create summary queries.
- Group data and produce summary statistics for each group.
- Subset a query on summarized values.

26

Business Scenario

Management requested a report containing the total annual donations for each employee.



27

Summary Functions: Across a Row

For a summary function with multiple arguments, nonmissing values are totaled across a row.

```
sum(Qtr1,Qtr2,Qtr3,Qtr4)
```

Partial orion.employee_donations

SAS Data Set

Employee_ID	Qtr1	Qtr2	Qtr3	Qtr4
120736	25	.	.	20
120759	15	20	5	.
120681	10	10	5	15
120679	.	20	5	15
120777	5	15	5	15

28

Summary Functions: Across a Row

Total each employee's annual cash donations.

```
proc sql;
  select Employee_ID      SUM(col1, ..., coln)
        ,label='Employee Identifier',
        ,Qtr1,Qtr2,Qtr3,Qtr4,
        ,sum(Qtr1,Qtr2,Qtr3,Qtr4)
        ,label='Annual Donation'
        ,format=dollar5.
  from orion.employee_donations
  where Paid_By="Cash or Check"
  order by 6 desc;
quit;
```

29

s103d06

Viewing the Output

Partial PROC SQL Output

Employee Identifier	Qtr1	Qtr2	Qtr3	Qtr4	Annual Donation
120736	25	.	.	20	\$45
120759	15	20	5	.	\$40
120681	10	10	5	15	\$40
120679	.	20	5	15	\$40
120777	5	15	5	15	\$40

30

3.03 Quiz

The following SELECT statements return the same results.

```
select sum(Qtr1,Qtr2,Qtr3,Qtr4)
      from orion.employee_donations;
```

```
select Qtr1+Qtr2+Qtr3+Qtr4
      from orion.employee_donations;
```

- a. True
- b. False

SAS Data Set

Employee_ID	Qtr1	Qtr2	Qtr3	Qtr4
120736	25	.	.	20
120759	15	20	5	.

31

Business Scenario

Management requested a report containing the total contributions for all employees in the first quarter.



33

Summary Functions: Down a Column

For a summary function with a single argument, nonmissing values are totaled down a column.

```
sum(Qtr1)
```

Partial orion.employee_donations

SAS Data Set

Employee_ID	Qtr1	Qtr2	Qtr3	Qtr4
120736	25	.	.	20
120759	15	20	5	.
120681	10	10	5	15
120679	.	20	5	15
120777	5	15	5	15

34

Summary Functions: Down a Column

Calculate the total of all charitable donations in quarter 1.

```
proc sql;
select sum(Qtr1)           SUM(argument)
      'Total Quarter 1 Donations'
   from orion.employee_donations;
quit;
```

PROC SQL Output

Total
Quarter 1
Donations

1515

35

s103d07

Alternative Method: MEANS Procedure

Calculate the total all charitable donations in quarter 1.

```
proc means data=orion.employee_donations
            sum maxdec=0;
var Qtr1;
run;
```

PROC MEANS Output

Analysis Variable : Qtr1
Sum

1515

36

s103d07

Summary Functions: Recap

How a summary function works in SQL depends on the number of columns specified in the argument list.

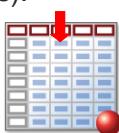
- If the summary function specifies more than one column, the statistic is calculated for the row (using values from the listed columns).

`sum(Qtr1, Qtr2, Qtr3, Qtr4)`



- If the summary function specifies only one column, the statistic is calculated for the column (using values from one or more rows).

`sum(Qtr1)`



37

Commonly Used Summary Functions

Both ANSI SQL and SAS functions can be used in PROC SQL.

SQL	SAS	Description
AVG	MEAN	Returns the mean (average) value.
COUNT	FREQ, N	Returns the number of nonmissing values.
MAX	MAX	Returns the largest value.
MIN	MIN	Returns the smallest nonmissing value.
SUM	SUM	Returns the sum of nonmissing values.
	NMISS	Counts the number of missing values.
	STD	Returns the standard deviation.
	VAR	Returns the variance.

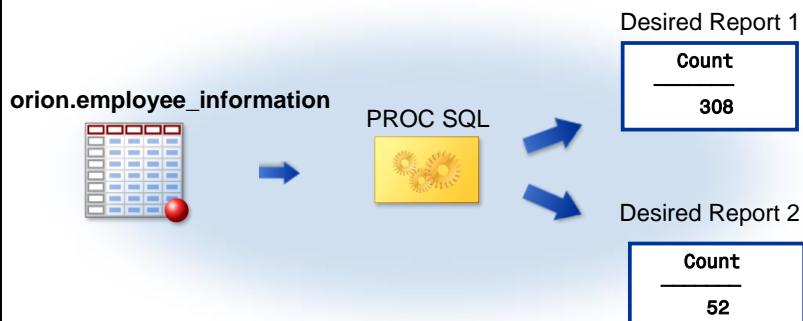
38



The AVG and COUNT functions do not accept multiple arguments and can summarize data only for a single column.

Business Scenario

Create a report that shows the total number of current Orion Star employees and a report that shows the total number of current Orion Star managers.



39

Summary Functions: COUNT Function

The COUNT function counts the number of rows returned by a query.

```
proc sql;
  select count(*) as Count
    from orion.employee_information
    where Employee_Term_Date is missing;
quit;
```

COUNT(argument)

Argument value	Counts
* (asterisk)	All rows in a table or group
A column name	The number of nonmissing values in that column

40

s103d08a

Viewing the Output

The COUNT function returns the number of distinct nonmissing **Manager_ID** values returned by a query.

```
proc sql;
  select count(distinct Manager_ID) as Count
    from orion.employee_information
      where Employee_Term_Date is missing;
quit;
```

PROC SQL Output

Count
52

41

s103d08b

3.04 Quiz

Open the program file **s103a02**. Submit the program and view the output.

1. How many rows did the first query create?
2. How many rows did the second query create?
3. In the second query's results, was the value in the **Average** column different for every gender listed?

43

s103a02

Remerging Summary Statistics

When a SELECT clause contains a column created by a summary function and a column that is *not* summarized, the summarized data is automatically appended to each row of the original data table (or *remerged*).

Employee_Gender	Average
M	40476.92
M	40476.92
F	40476.92

NOTE: The query requires remerging summary statistics back with the original data.

47

- Most Database Management Systems do not enable remerging summary statistics. They generate an error instead.

Business Scenario

Calculate each male employee's salary as a percentage of all male employees' salaries.

`orion.employee_information`



PROC SQL



Partial Results

Male Employee Salaries

Employee_ID	Salary	
120259	433,800	5.9%
120262	268,455	3.7%
120261	243,190	3.3%

48

Using Remerged Summary Statistics

Calculate each male employee's salary as a percentage of all male employees' salaries. Display **Employee_ID**, **Salary**, and percentage in decreasing order of percentage.

```
proc sql;
title "Male Employee Salaries";
select Employee_ID, Salary format=comma12.,
       Salary / sum(Salary)
          'PCT of Total' format=percent6.2
     from orion.employee_information
    where Employee_Gender="M"
          and Employee_Term_Date is missing
      order by 3 desc;
quit;
title;
```

Select only the group of rows that you want to analyze.

49

...

Using Remerged Summary Statistics

Calculate each male employee's salary as a percentage of all male employees' salaries. Display **Employee_ID**, **Salary**, and percentage in decreasing order of percentage.

```
proc sql;
title "Male Employee Salaries";
select Employee_ID, Salary format=comma12.,
       Salary / sum(Salary)
          'PCT of Total' format=percent6.2
     from orion.employee_information
    where Employee_Gender="M"
          and Employee_Term_Date is missing
      order by 3 desc;
quit;
title;
```

individual salary value for each row

divided by a remerged summary value (sum of all salaries)

50

s103d09

Viewing the Output

Partial PROC SQL Output

Male Employee Salaries			
Employee_ID	Employee Annual Salary	PCT of Total	
120259	433,800	5.9%	
120262	268,455	3.7%	
120261	243,190	3.3%	
121141	194,885	2.7%	
120101	163,040	2.2%	

Partial SAS Log

NOTE: The query requires remerging summary statistics back with the original data.

51

This is a good example of SQL simplifying code. The following traditional SAS code produces similar results:

```
proc means data=orion.employee_information (where=(Employee_Gender="M"
                                                and Employee_Term_Date is missing))
    sum noplay;
output out=work.summary sum=TotalSalary;
var salary;
run;
data report;
merge orion.employee_information (where=(Employee_Gender="M"
                                            and Employee_Term_Date is missing))
    work.summary(keep=TotalSalary);
retain Total 0;
if _n_=1 then Total=TotalSalary;
Percent=Salary / Total;
keep Employee_ID Salary Percent;
format salary comma12. Percent percent6.2;
run;
proc sort data=report;
by descending Percent;
run;
title "Male Employee Salaries - Traditional SAS Programming";
proc print data=report noobs split='*';
label Percent='*';
run;
title;
```

Remerging Summary Statistics

To change the default behavior, use either of the following:

- NOSQLREMERGE SAS system option
- PROC SQL NOREMERGE option

Resubmitting the query with the NOREMERGE option in effect produces no output and results in this SAS log error message:

ERROR: The query requires remerging summary statistics back with the original data. This is disallowed due to the NOREMERGE proc option or NOSQLREMERGE system option.

52

s103d09a

3.05 Quiz

Open the program file **s103a03**. Submit the query and review the output and the log. Answer the following questions:

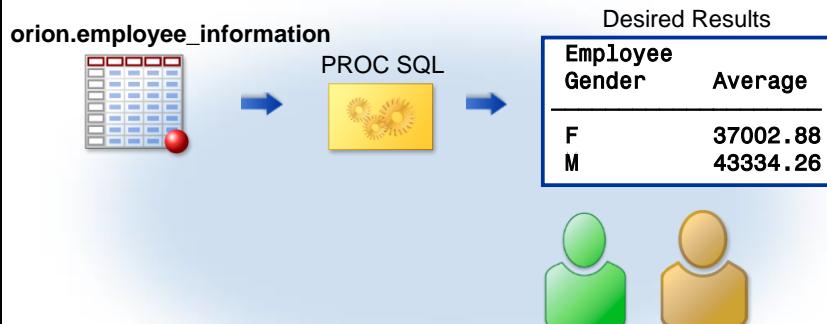
1. How many rows of output are created?
2. What major difference is there in the log between this query's results and the second query in the previous activity?

53

s103a03

Business Scenario

Produce a report that determines the average salary by gender.



57

Grouping Data

You can use the GROUP BY clause to do the following:

- classify the data into groups based on the values of one or more columns
- calculate statistics for each unique value of the grouping columns

```
proc sql;
  title "Average Salary by Gender";
  select Employee_Gender as Gender,
         avg(Salary) as Average
    from orion.employee_information
   where Employee_Term_Date is missing
  group by Employee_Gender;
quit; GROUP BY group-by-item<..., group-by-item>
```

s103d10

58

Viewing the Output

PROC SQL Output

Average Salary by Gender

Employee Gender	Average
F	37002.88
M	43334.26

59

3.06 Poll

Can you group by more than one column?

- Yes
- No

60

Business Scenario

Produce a report showing the count of employees in those departments that have at least 25 people. Display the results in descending order by count.

`orion.employee_information`



Desired Results

Department	Count
Sales	201
Administration	34
Stock & Shipping	26
IS	25

62

Step 1

Produce a report showing the number of employees in each department.

`orion.employee_information`



PROC SQL



Partial Results

Department	Count
Accounts	17
Accounts Management	9
Administration	34

63

Analyzing Groups of Data

```
proc sql;
select Department, count(*) as Count
  from orion.employee_information
 group by Department;
quit;
```

64

s103d11

Step 2

Control the result to include only the departments that have at least 25 people, with the departments in decreasing order.

Department	Count
Accounts	17
Accounts Management	9
Administration	34
Concession Management	11
Engineering	9
Executives	4
Group Financials	3
Group HR Management	18
IS	25
Logistics Management	
Marketing	
Purchasing	
Sales	
...	

Department	Count
Sales	201
Administration	34
Stock & Shipping	26
IS	25

65

Selecting Groups with the HAVING Clause

The *HAVING* clause subsets groups based on the expression value.

```
proc sql;
  select Department, count(*) as Count
    from orion.employee_information
   group by Department
   having Count ge 25
  order by Count desc;
quit;
```

GROUP BY group-by-item <,...,group-by-item>
HAVING sql-expression

66

s103d12

Viewing the Output

PROC SQL Output

Department	Count
Sales	201
Administration	34
Stock & Shipping	26
IS	25

67

WHERE Clause versus HAVING Clause

- The WHERE clause is evaluated **before** a row is available for processing and determines which individual rows are available for grouping.

WHERE sql-expression

- The HAVING clause is processed **after** the GROUP BY clause and determines which groups are displayed.

HAVING sql-expression

68

3.07 Quiz

Which syntax selects **Job_Title** values with a total **Bonus** value greater than \$10,000?

- ```
select Job_Title, sum(Salary*0.1) as Bonus
 from orion.employee_information
 group by Job_Title
 having Bonus > 10000;
```
- ```
select Job_Title, sum(Salary*0.1) as Bonus
  from orion.employee_information
 where Bonus > 10000
 group by Job_Title;
```
- both of the above
- neither of the above

69

Business Scenario

Create a report that lists the following for each department:

- total number of managers
- total number of non-manager employees
- manager-to-employee (M/E) ratio

Below is a rough sketch of the desired report.

Department	Managers	Employees	M/E Ratio
Accounts	1	5	20%
Administration	2	20	10%

71

Business Data

Determine whether an employee is a manager or a non-manager.

The **Job_Title** column contains the information about each employee.

Department	Job_Title
Administration	Administration Manager
Administration	Secretary I
Administration	Office Assistant II

72

Counting Rows That Meet a Specified Criterion

How do you determine the rows that *do* have *Manager* in **Job_Title**, as well as rows that *do not*? You cannot use a WHERE clause to exclude either group.

Department	Job_Title
Administration	Administration Manager
Administration	Secretary I
Administration	Office Assistant II

Use the FIND function in a Boolean expression to identify rows that contain *Manager* in the **Job_Title** column.

73

FIND Function

The *FIND* function returns the starting position of the first occurrence of a substring within a string (character value).

Find the starting position of the substring *Manager* in the character variable **Job_Title**.

```
find(Job_Title,"manager","i")
```

Job_Title	1	2
1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5		
A d m i n i s t r a t i o n M a n a g e r		

The value returned by the FIND function is 16.

```
FIND(string, substring<,modifier(s)><,startpos>)
```

74

string constant, variable, or expression to be searched

substring constant, variable, or expression sought within the string

modifiers I=ignore case, T=trim trailing blanks

startpos an integer specifying the start position and direction of the search

Using Boolean Expressions

Part 1: Use a Boolean expression to determine whether an employee is a manager.

```
proc sql;
  select Department, Job_Title,
    (find(Job_Title,"manager","i")>0)
    "Manager"
  from orion.employee_information;
quit;
```

- ✍ Boolean expressions evaluate to true (1) or false (0).
 - If **Job_Title** contains *Manager*, the value is 1.
 - If **Job_Title** does not contain *Manager*, the value is 0.

75

s103d13

Viewing the Output

Partial PROC SQL Output

Department	Job_Title	Manager
Administration	Administration Manager	1
Administration	Secretary I	0
Administration	Office Assistant II	0
Administration	Office Assistant III	0
Administration	Warehouse Assistant II	0
Administration	Warehouse Assistant I	0
Administration	Warehouse Assistant III	0
Administration	Security Guard II	0
Administration	Security Guard I	0
Administration	Security Guard II	0
Administration	Security Manager	1

76

Using Boolean Expressions

Part 2: Calculate the statistics requested.

```
proc sql;
title "Manager-to-Employee Ratios";
select Department,
       sum((find(Job_Title,"manager","i")>0))
           as Managers,
       sum((find(Job_Title,"manager","i")=0))
           as Employees,
       calculated Managers/calculated Employees
           "M/E Ratio" format=percent8.1
from orion.employee_information
group by Department;
quit;
```

77

s103d14

Viewing the Output

PROC SQL Output

Manager-to-Employee Ratios			
Department	Managers	Employees	M/E Ratio
Accounts	3	14	21.4%
Accounts Management	1	8	12.5%
Administration	5	29	17.2%
Concession Management	1	10	10.0%
Engineering	1	8	12.5%
Executives	0	4	0.0%
Group Financials	0	3	0.0%
Group HR Management	3	15	20.0%
IS	2	23	8.7%
Logistics Management	6	8	75.0%
Marketing	6	14	42.9%
Purchasing	3	15	20.0%
Sales	0	201	0.0%
Sales Management	5	6	83.3%
Secretary of the Board	0	2	0.0%
Stock & Shipping	5	21	23.8%
Strategy	0	2	0.0%

78



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

4. Summarizing Data

Create a report that displays the number of employees residing in each city.

- a. Use the **orion.employee_addresses** table.
- b. Use the **City** column and the COUNT(*) function.
- c. Group the data and order the output by **City**.
- d. Add this title to the report: **Cities Where Employees Live**.

PROC SQL Output

Cities Where Employees Live	
City	Count
Melbourne	41
Miami-Dade	109
Philadelphia	95
San Diego	112
Sydney	67

5. Using SAS Functions

Create a report that includes each employee's age at time of employment.

- a. The report should contain the columns **Employee_ID**, **Birth_Date**, **Employee_Hire_Date**, and **Age**.
- b. Obtain the data for the report from the **orion.employee_payroll** table.
- c. Calculate **Age** as `int((Employee_Hire_Date-Birth_Date)/365.25)`.
- d. Add this title to the report: **Age at Employment**.
- e. Display **Birth_Date** and **Employee_Hire_Date** values using the MMDDYY10. format.
- f. Label each column as shown in the following sample report:

Partial PROC SQL Output

Age at Employment				
Employee ID	Birth Date	Hire Date	Age	
120101	08/18/1980	07/01/2007	26	
120102	08/11/1973	06/01/1993	19	
120103	01/22/1953	01/01/1978	24	
120104	05/11/1958	01/01/1985	26	
120105	12/21/1978	05/01/2003	24	
120106	12/23/1948	01/01/1978	29	
120107	01/21/1953	02/01/1978	25	

Level 2**6. Summarizing Data**

- a. Using data contained in the **orion.customer** table, create a report that shows the following statistics for each country:
 - 1) total number of customers
 - 2) total number of male customers
 - 3) total number of female customers
 - 4) percent of all customers that are male (**Percent Male**)
- b. Add this title to the report: **Customer Demographics: Gender by Country**.
- c. Arrange the report by value of **Percent Male** so that the country with the lowest value is listed first, with the remaining countries following in ascending order.

Customer Demographics: Gender by Country				
Customer Country	Customers	Men	Women	Percent Male
ZA	4	1	3	25%
CA	15	7	8	47%
US	28	15	13	54%
AU	8	5	3	63%
DE	10	7	3	70%
IL	5	5	0	100%
TR	7	7	0	100%

Hint: The Boolean expression (**Gender="M"**) evaluates as 1 when the value of **Gender** is *M* and 0 when the value of **Gender** is *F*.

7. Summarizing Data in Groups

Use the **orion.customer** table to determine the number of Orion Star customers of each gender in each country. Display columns named **Country**, **Male Customers**, and **Female Customers**. Display only those countries that have more female customers than male customers. Order the report by descending female customers. Add this title to the report: **Countries with More Female than Male Customers**.

PROC SQL Output

Countries with More Female than Male Customers		
Country	Male Customers	Female Customers
CA	7	8
ZA	1	3

Country	Male Customers	Female Customers
CA	7	8
ZA	1	3

Challenge**8. Advanced Summarizing Data in Groups**

Use the **orion.employee_addresses** table to create a report that displays the countries and cities where Orion Star employees reside, and the number of employees in each city. Include only one row per country/city combination. Display the values in country/city order, and give the report an appropriate title.

PROC SQL Output

Countries and Cities Where Employees Live		
Country	City	Employees
AU	Melbourne	41
AU	Sydney	67
US	Miami-Dade	109
US	Philadelphia	95
US	San Diego	112

Country	City	Employees
AU	Melbourne	41
AU	Sydney	67
US	Miami-Dade	109
US	Philadelphia	95
US	San Diego	112

Hint: Some data might not have consistent capitalization.

End of Exercises

3.3 Solutions

Solutions to Exercises

1. Enhancing Output by Ordering the Rows and Adding Titles and Formats

```
*** s103s01 ***
proc sql;
title "Single Male Employee Salaries";
select Employee_ID 'Employee ID',
       Salary format=comma10.2 'Annual Salary',
       Salary/3 format=comma10.2 as Tax
          'Federal Tax Withdrawn'
from orion.employee_payroll
where Marital_Status="S"
      and Employee_Gender ="M"
      and Employee_Term_Date is missing
order by Salary desc;
quit;
title;
```

2. Using Formats to Limit the Width of Columns in the Output

```
*** s103s02 ***
proc sql;
title "Australian Clothing Products";
select Supplier_Name 'Supplier'format=$18.,
       Product_Group 'Group' format=$12.,
       Product_Name 'Product' format=$30.
from orion.product_dim
where Supplier_Country="AU"
      and Product_Category="Clothes"
order by Product_Name;
quit;
title;
```

3. Enhancing Output with Multiple Techniques

```
*** s103s03 ***
proc sql;
title "US Customers >50 Years Old as of 02FEB2013";
select Customer_ID format=z7. 'Customer ID',
       catx(' ', Customer_LastName,Customer_FirstName)
           format=$20. 'Last Name, First Name' as Name,
       Gender 'Gender',
       int((02feb2013'd-Birth_Date)/365.25) as Age
from orion.customer
where Country="US"
      and calculated Age>50
order by Age desc, Name;
quit;
title;
```

4. Summarizing Data

```
*** s103s04 ***;

proc sql;
title "Cities Where Employees Live";
select City, Count(*) as Count
  from orion.employee_addresses
 group by City
 order by City;
quit;
title;
```

5. Using SAS Functions

```
*** s103s05 ***;

proc sql;
title "Age at Employment";
select Employee_ID 'Employee ID',
      Birth_Date format=mmddyy10. 'Birth Date',
      Employee_Hire_Date format=mmddyy10. 'Hire Date',
      int((Employee_Hire_Date-Birth_Date)/365.25)as Age 'Age'
    from orion.employee_payroll;
quit;
title;
```

6. Summarizing Data

```
*** s103s06 ***;

proc sql;
title "Customer Demographics: Gender by Country";
select Country, Count(*) as Customers,
      sum(Gender="M") as Men,
      sum(Gender ="F") as Women,
      calculated Men/calculated Customers 'Percent Male'
      format=percent6.1
    from orion.customer
 group by Country
 order by 5;
quit;
title;
```

7. Summarizing Data in Groups

```
*** s103s07 ***;

proc sql;
title "Countries with more Female than Male Customers";
select Country 'Country',
      sum(Gender="M") as M "Male Customers",
      sum(Gender="F") as F "Female Customers"
    from orion.customer
 group by Country
```

```

    having F > M
    order by F desc;
quit;
title;

```

8. Advanced Summarizing Data in Groups

```

*** s103s08 ***;

proc sql;
title "Countries and Cities Where Employees Live";
select upcase(Country) 'Country',
       propcase(City) 'City',
       count(*) 'Employees'
  from orion.employee_addresses
 group by 1,2
 order by 1,2;
quit;
title;

```

Items of interest:

- **Country** and **City** columns – The case of the data values is inconsistent. Sorting by column position uses the data in the intermediate results set, which was already corrected for case using SAS functions.
- Sorting by column name, even if you use column aliases (as **Country**, as **City**), produces duplicates in the output, although the data displayed appears to be identical. This happens because, by default, the ORDER BY and GROUP BY clauses operate on the **Country** and **City** data values found in the original table, before case correction. To ensure that you get the desired results, specify column position numbers in the GROUP BY and ORDER BY clauses instead of column names. This causes SQL to group by and sort by the values in the intermediate result set (corrected data) instead of the values in the original table.

Code demonstrating *incorrect* results:

```

*** s103s08 ***;
proc sql;
title "Countries and Cities Where Employees Live";
title2 "Incorrect Solution: Improper grouping due to Case issues";
title3 "Even with DISTINCT specified, duplicates remain";
select upcase(Country),
       propcase(City),
       count(*) 'Employees'
  from orion.employee_addresses
 group by country, city
 order by country, city;
quit;
title;

```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

3.01 Multiple Choice Poll – Correct Answer

Which ORDER BY clause orders a report by descending **State** and descending **City**?

- a. `order by desc State, City`
- b. `order by State, City desc`
- c. `order by State desc, City desc`
- d. `order by desc State, desc City`

9

3.02 Short Answer Poll – Correct Answer

Which option, NUMBER or NONNUMBER, applies to the third report? **The NONNUMBER option applies to the third report.**

22

3.03 Quiz – Correct Answer

The following SELECT statements return the same results.

```
select sum(Qtr1,Qtr2,Qtr3,Qtr4)
      from orion.employee_donations;
```

```
select Qtr1+Qtr2+Qtr3+Qtr4
      from orion.employee_donations;
```

- a. True
- b. False**

The addition operator returns a missing value if any operand is missing.

32

3.04 Quiz – Correct Answer

1. How many rows did the first query create?

```
proc sql;
  select 'The Average Salary is:',
         avg(Salary)
    from orion.employee_information
   where Employee_Term_Date is missing;
quit;
```

One row

```
_____
```

The Average Salary is: 40476.92

44

s103a02

3.04 Quiz – Correct Answer

2. How many rows did the second query create?

```
proc sql;
select Employee_Gender,
       avg(Salary) as Average
  from orion.employee_information
 where Employee_Term_Date is missing;
quit;
```

The output contains 308 rows. This is the number of rows returned by the COUNT(*) function in program s103d08.

45

s103a02

3.04 Quiz – Correct Answer

3. In the second query's results, was the value in the average column different for every gender listed? **no**

Every row contained the same Average value, which is the overall average salary for the entire table.

Employee_Gender	Average
M	40476.92
M	40476.92
M	40476.92
F	40476.92
F	40476.92

46

s103a02

3.05 Quiz – Correct Answer

1. How many rows of output are created?

```
proc sql number;
select Employee_Gender,
       avg(Salary) as Average
  from orion.employee_information
 where Employee_Term_Date is missing
 group by Employee_Gender;
quit;
```

two rows

Employee_	Row	Gender	Average
	1	F	37002.88
	2	M	43334.26

54

3.05 Quiz – Correct Answer

2. What major difference is there in the log between this query's results and the second query in the previous activity?

SAS log notes from the previous activity:

```
NOTE: The query requires remerging summary statistics back
with the original data.
NOTE: PROCEDURE SQL used (Total process time):
      real time            0.01 seconds
      cpu time             0.01 seconds
```

SAS log notes from this activity:

```
NOTE: PROCEDURE SQL used (Total process time):
      real time            0.01 seconds
      cpu time             0.01 seconds
```

There was no note about remerging statistics.

55

3.06 Poll – Correct Answer

Can you group by more than one column?

- Yes
 No

61

3.07 Quiz – Correct Answer

Which syntax selects **Job_Title** values with a total **Bonus** value greater than \$10,000?

1.

```
select Job_Title, sum(Salary*0.1) as Bonus
      from orion.employee_information
            group by Job_Title
                  having Bonus > 10000;
```
2.

```
select Job_Title, sum(Salary*0.1) as Bonus
      from orion.employee_information
            where Bonus > 10000
                  group by Job_Title;
```
3. both of the above
4. neither of the above

70

Chapter 4 SQL Joins

4.1	Introduction to SQL Joins	4-3
4.2	Inner Joins	4-9
	Exercises	4-18
4.3	Outer Joins.....	4-21
	Exercises	4-32
4.4	Complex SQL Joins.....	4-34
	Exercises	4-37
4.5	Solutions	4-39
	Solutions to Exercises	4-39
	Solutions to Student Activities (Polls/Quizzes)	4-42

4.1 Introduction to SQL Joins

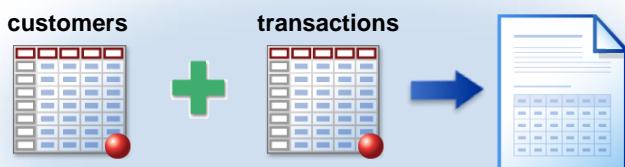
Objectives

- Identify different ways to combine data horizontally from multiple tables.
- Distinguish between inner and outer SQL joins.
- Understand the Cartesian product.

3

Business Scenario

Management has requested multiple reports. You have to understand how to combine the data from the tables to complete these requests.



4

Exploring the Data

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

The **customers** table is representative of a customer dimension table. There would be additional columns with data about our customers including address, age, and so on.

The **transactions** table is representative of a fact table. There would be columns holding all the key column data, **Product_ID**, **Employee_ID**, and so on.

5

Combining Data from Multiple Tables

SQL uses *joins* to combine tables horizontally. Requesting a join involves matching data from one row in one table with a corresponding row in a second table. Matching is typically performed on one or more columns in the two tables.



6

4.01 Multiple Choice Poll

Which of these DATA step statements is used to combine tables horizontally?

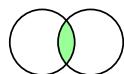
- a. SET
- b. APPEND
- c. MERGE
- d. INPUT
- e. INFILE

7

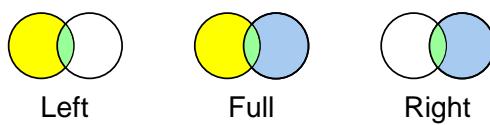
Types of Joins

PROC SQL supports two types of joins:

Inner joins return only matching rows.



Outer joins return all matching rows, plus nonmatching rows from one or both tables.



9



Cross joins, natural joins, and union joins are special cases of the standard join types.

A *cross join* is a Cartesian product.

A *natural join* automatically selects columns with the same name and type to use in determining matching rows.

A *union join* is similar to combining two tables with the OUTER UNION set operator.

Cartesian Product

A query that lists multiple tables in the FROM clause without a WHERE clause produces all possible combinations of rows from all tables. This result is called a *Cartesian product*.

```
proc sql;
  select *
    from customers, transactions;
quit;
```

SELECT ...
 FROM table-name, table-name
 <, ...,table-name >;

To understand how SQL processes a join, it is helpful to understand the concept of the Cartesian product.

10

s104d01

Building the Cartesian Product

customers		transactions		
ID	Name	ID	Action	Amount
101	Smith	102	Purchase	\$100
104	Jones	103	Return	\$52
102	Blank	105	Return	\$212

Result Set

ID	Name	ID	Action	Amount
101	Smith	102	Purchase	\$100
101	Smith	103	Return	\$52
101	The Cartesian product	105	Return	\$212
104	is rarely the desired			
104	result of a query.			
102				
102	Blank	103	Return	\$52
102	Blank	105	Return	\$212

20

Non-Matching Data in the Cartesian Product

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

Result Set

Non-matching
IDs

ID	Name	ID	Action	Amount
101	Smith	102	Purchase	\$100
101	Smith	103	Return	\$52
101	Smith	105	Return	\$212
104	Jones	102	Purchase	\$100
104	Jones	103	Return	\$52
104	Jones	105	Return	\$212
102	Blank	102	Purchase	\$100
102	Blank	103	Return	\$52
102	Blank	105	Return	\$212

21

Size of the Cartesian Product

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

Result Set

ID	Name	ID	Action	Amount
101	Smith	102	Purchase	\$100
101	Smith	103	Return	\$52
101	Smith	105	Return	\$212
104	Jones	102	Purchase	\$100
104	Jones	103	Return	\$52
104	Jones	105	Return	\$212
102	Blank	102	Purchase	\$100
102	Blank	103	Return	\$52
102	Blank	105	Return	\$212

9 rows

22

Size of the Cartesian Product

The number of rows in a Cartesian product is the product of the number of rows in the contributing tables.

$$3 \times 3 = 9$$

$$1,000 \times 1,000 = 1,000,000$$

$$100,000 \times 100,000 = 10,000,000,000$$

Partial SAS Log

NOTE: The execution of this query involves performing one or more Cartesian product joins that cannot be optimized.

23

4.02 Short Answer Poll

How many rows and columns are returned from this query?

```
select *
  from customer2, transaction2;
```

customer2		transaction2		
ID	Name	ID	Action	Amount
101	Jones	102	Purchase	\$376
101	Jones	102	Return	\$119
102	Kent	103	Purchase	\$57
102	Kent	105	Purchase	\$98
104	Avery			

24

4.2 Inner Joins

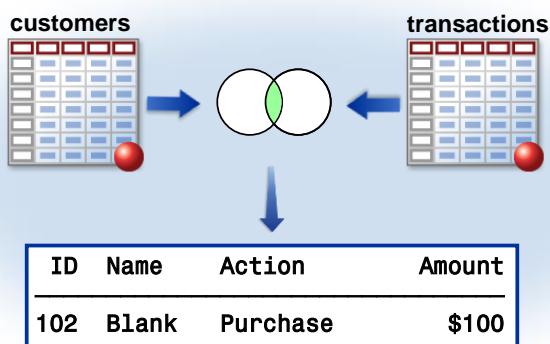
Objectives

- Join two or more tables on matching columns.
- Qualify column names to identify specific columns.
- Use a table alias to simplify the SQL code.
- Join two tables with alternative join syntax.

28

Report 1: Inner Join

Management has requested a report showing all valid order information.



29

Inner Join

Specify the matching criteria in the WHERE clause.

```
proc sql;
  select *
    from customers, transactions
   where customers.ID=
         transactions.ID;
quit;
```

```
SELECT object-item<, ...object-item>
  FROM table-name, ... table-name
 WHERE join condition
       <AND sql-expression>
       <other clauses>;
```

PROC SQL Output

ID	Name	ID	Action	Amount
102	Blank	102	Purchase	\$100

s104d02

30

-  The SQL Procedure Optimizer has algorithms from which it selects to optimize a join. For example, the optimizer might use a hashing algorithm when joining a small table with a large table.

A detailed paper on PROC SQL join optimization can be found here:
<http://support.sas.com/techsup/technote/ts553.html>

4.03 Short Answer Poll

Why are there two ID columns in the result?

ID	Name	ID	Action	Amount
102	Blank	102	Purchase	\$100

31

Qualifying the Column Names

When you specify columns with the same name from more than one table, it is necessary to qualify the column name.

```
proc sql;
select *
  from customers, transactions
 where customers.ID=
       transactions.ID;
quit;
```

table-name.column-name

33

s104d02

Setup for the Poll

Run program **s104a02** and review the results.

What message is displayed in the log when the program is executed?

```
proc sql;
select *
  from customers, transactions
 where ID=ID;
quit;
```

34

4.04 Short Answer Poll

What message is displayed in the log when the program is executed?

35

Completed Code for Report 1

To display the ID column only once in the results, qualify the ID column in the SELECT clause.

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

```
select customers.ID, Name, Action, Amount
      from customers, transactions
     where customers.ID=transactions.ID;
```

ID	Name	Action	Amount
102	Blank	Purchase	\$100

37

s104d03

Equijoins are inner joins where the join criteria specify equality for the identifying variables. Natural joins are equijoins in which SQL compares the values of all same-name columns in both tables and produces output containing **only one column** for results from same-name columns in the source tables. The rows in the tables are joined, based on matching values in same-name columns.

For example, the following natural join code would produce results identical to the query above:

```
proc sql;
  select *
    from customers natural join transactions;
quit;
```

Abbreviating the Code

A *table alias* is a temporary, alternative name for a table. You can make the query easier to read by using table aliases.

```
SELECT alias-1.object-item<, ...alias-2.object-item>
FROM table-name <AS> alias-1,
      table-name <AS> alias-2
WHERE join-condition(s)
      <other clauses>;
```

The AS keyword is optional in the table alias syntax.

38

Abbreviating the Code with a Table Alias

```
proc sql;
select c.ID, Name, Action, Amount
  from customers as c, transactions as t
  where c.ID=t.ID;
quit;
```

PROC SQL Output

ID	Name	Action	Amount
102	Blank	Purchase	\$100

39

s104d04

SQL Inner Join versus DATA Step Merge

A PROC SQL inner join and the DATA step match merge can return the same results.

```
proc sql;
  select c.ID, Name, Action,
         Amount
    from customers as c,
         transactions as t
   where c.ID=t.ID;
quit;

data orders;
  merge customers(in=c)
        transactions(in=t);
  by ID;
  if c=1 and t=1;
run;
proc print data=orders;
run;
```

40

s104d04

SQL Inner Join versus DATA Step Merge

A PROC SQL inner join and the DATA step match merge will not always return the same results.

customer2

ID	Name
101	Jones
101	Jones
102	Kent
102	Kent
104	Avery

transaction2

ID	Action	Amount
102	Purchase	\$376
102	Return	\$119
103	Purchase	\$57
105	Purchase	\$98

```
select *
  from customer2 as c2, transaction2 as t2
 where c2.ID=t2.ID;
```

ID	Name	ID	Action	Amount
102	Kent	102	Purchase	\$376
102	Kent	102	Purchase	\$376
102	Kent	102	Return	\$119
102	Kent	102	Return	\$119

s104d05

41

Setup for the Poll

Run program **s104a03** and review the results to determine how many rows (observations) the DATA step MERGE statement produces in the output table.

customer2		transaction2		
ID	Name	ID	Action	Amount
101	Jones	102	Purchase	\$376
101	Jones	102	Return	\$119
102	Kent	103	Purchase	\$57
102	Kent	105	Purchase	\$98
104	Avery			

```
data work.new;
  merge customer2 in=(InCust)
    transaction2 in=(InTrans);
  by ID;
  if InCust=1 and InTrans=1;
run;

proc print data=work.new;
run;
```

s104a03

42

4.05 Multiple Choice Poll

How many rows (observations) result from the DATA step MERGE statement in program **s104a03**?

- a. 2
- b. 4
- c. 6
- d. 20
- e. none of the above

43

Compare SQL Join and DATA Step Merge

Key Points	SQL Join	DATA Step Merge
Explicit sorting of data before join/merge	Not required	Required
Same-name columns in join/merge expressions	Not required	Required
Equality in join or merge expressions	Not required	Required

45

Alternative Join Syntax

This alternative syntax names the join type and includes an ON clause.

```
select c.ID, Name, Action, Amount
      from customers as c
            inner join
                  transactions as t
            on c.ID=t.ID;
```

```
SELECT object-item <, ...object-item>
      FROM table-name <> AS alias>
            INNER JOIN
                  table-name <> AS alias>
            ON join-condition(s)
            WHERE sql-expression
            <other clauses>;
```

46

s104d06



Up to 256 tables can be joined using the alternative join syntax. The join processes two tables at a time, until all tables have been joined.

```
select *
      from one inner join two
            on one.id=two.id
            inner join three
            on one.id=three.id
            inner join four
            on one.id=four.id;
```

4.06 Short Answer Poll

How many tables can be combined using the following PROC SQL syntax?

```
proc sql;
  select *
    from table-name, ..., table-name
   where sql-expression;
quit;
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Inner Joins

Produce a report containing **Employee_Name** and calculated years of service (YOS) as of February 1, 2013, by joining **orion.employee_addresses** and **orion.employee_payroll** on **Employee_ID**. Label the columns and provide two title lines, as shown in the sample output. Limit the report to employees where YOS > 30. Order the output alphabetically by **Employee_Name**.

- The **orion.employee_addresses** table contains the **Employee_Name** column.
- The **orion.employee_payroll** table contains the **Employee_Hire_Date** column.
- Both **orion.employee_addresses** and **orion.employee_payroll** contain columns named **Employee_ID**.
- Use TITLE1 and TITLE2 statements to produce title lines as indicated in the sample report:

Partial PROC SQL Output

Employees With More Than 30 Years of Service
As of February 1, 2013

Name	Years of Service
Abbott, Ray	33
Banchi, Steven	35
Blackley, James	35
Bleu, Henri Le	35
Branly, Wanda	35
Buddery, Jeannette	35
Campbell, Carston	35
Capps, Ramond	35



Age and years of service calculations can be difficult to render precisely. In this course, you use the following:

`int('1FEB2013'd - Employee_Hire_Date)/365.25) as YOS`

For a more in-depth discussion, see “Accurately Calculating Age with Only One Line of Code” at http://staff.washington.edu/billyk/TechTips_SC4Q98.pdf.

Level 2

2. Creating a Summary Report from Two Tables

The head of the Sales Department wants to know how many of each product was sold since the beginning of 2010. The report should include the product ID, the product name, and the total sold for that product and ordered to match the output shown below.

The data that you need can be found in the listed columns of the following tables:

- **orion.product_dim** contains
 - **Product_ID**
 - **Product_Name.**
- **orion.order_fact** contains
 - **Product_ID**
 - **Quantity.**

Partial PROC SQL Output

Total Quantities Sold by Product ID and Name		
Product ID	Product Name	Total Sold
230100600016	Expedition Zero,Medium,Right,Charcoal	6
230100700011	Hurricane 4	6
230100600031	Outback Sleeping Bag, Large,Right, Blue/Black	6
240800200035	Shine Black PRO	6
240700100001	Armour L	5
220101400088	Casual Genuine Polo-Shirt	5
240200100118	Hi-fly Intrepid Stand 8 Black	5
240100100312	Hot Mini Backboard Bulls	5
230100500082	Lucky Tech Intergal Wp/B Rain Pants	5

Challenge

3. Joining Multiple Tables

Create a report showing Orion Star catalog and Internet customers residing in the U.S. or Australia who purchased foreign-manufactured products (that is, a product that was not manufactured in their country of residence). The report should be titled **US and Australian Catalog and Internet Customers Purchasing Foreign Manufactured Products** and should display the customers' names and the number of foreign purchases made. Present the information so that those with the largest number of purchases appear at the top of the report, and customers who have the same number of purchases are displayed in alphabetical order.

Employee_ID 99999999 is a dummy ID that can be used to identify catalog and Internet orders. The data that you need can be found in the listed columns of the following tables:

- **orion.product_dim** contains
 - **Product_ID**
 - **Supplier_Country.**
- **orion.order_fact** contains
 - **Product_ID**
 - **Employee_ID**
 - **Customer_ID.**
- **orion.customer** contains
 - **Customer_ID**
 - **Country.**

Partial PROC SQL Output

US and Australian Catalog and Internet Customers
Purchasing Foreign Manufactured Products

Name	Purchases
Candy Kinsey	10
Phenix Hill	7
Cynthia McCluney	5
Korolina Dokter	5
Najma Hicks	4
Robert Bowerman	4

End of Exercises

4.3 Outer Joins

Objectives

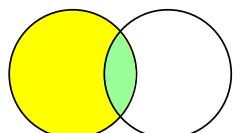
- Join two tables on matching columns and include non-matching rows from one table.
- Join two tables on matching columns and include non-matching rows from both tables.
- Compare outer joins and DATA step merges.

52

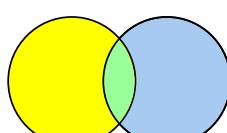
Outer Joins

You can retrieve both non-matching and matching rows using an outer join.

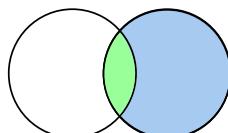
Outer joins include left, full, and right outer joins. Many tables can be referenced in outer joins. The tables are processed two tables at a time.



Left



Full

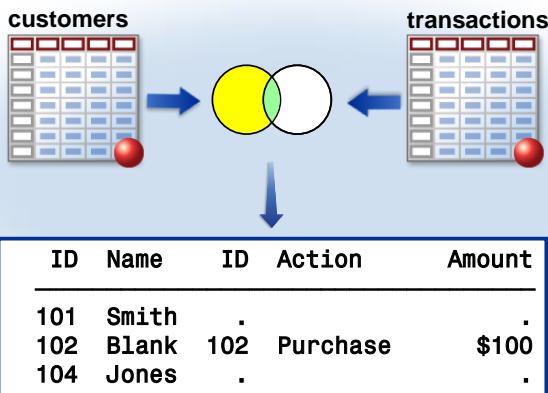


Right

53

Report 2: Outer Joins

You have been asked for a report that displays **all** customers and any recent transactions that they have completed.



54

Outer Joins

Outer join syntax is similar to the alternate inner join syntax.

```
proc sql;
  title 'All Customers';
  select *
    from customers as c
    left join
      transactions as t
    on c.ID=t.ID;
  quit;
```

The ON clause specifies the join criteria in outer joins.

```
SELECT object-item <, ...object-item>
  FROM table-name <<AS> alias>
  LEFT|RIGHT|FULL JOIN
  table-name <<AS> alias>
  ON join-condition(s)
  <other clauses>;
```

55

s104d07

The WHERE and ON clauses operate independently.

- The ON clause specifies the rows to be joined.
- The WHERE clause produces a subset of the results.

Viewing the Output

PROC SQL Output

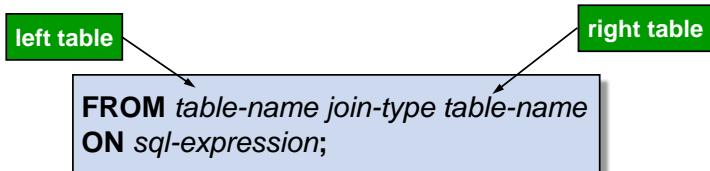
All Customers				
ID	Name	ID	Action	Amount
101	Smith	.	.	.
102	Blank	102	Purchase	\$100
104	Jones	.	.	.

56

Determining Left and Right

Consider the position of the tables in the FROM clause.

- Left joins return all matching and non-matching rows from the *left* table and the matching rows from the *right* table.
- Right joins return all matching and non-matching rows from the *right* table and the matching rows from the *left* table.
- Full joins return all matching and non-matching rows from all of the tables.



57

Left Join

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

```
select *
  from customers c left join transactions t
  on c.ID = t.ID;
```

ID	Name	ID	Action	Amount
101	Smith	.	.	.
102	Blank	102	Purchase	\$100
104	Jones	.	.	.

Includes all rows from the left table, even if there are no matching rows in the right table.

58

s104d08

Right Join

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

```
select *
  from customers c right join transactions t
  on c.ID = t.ID;
```

ID	Name	ID	Action	Amount
102	Blank	102	Purchase	\$100
.	.	103	Return	\$52
.	.	105	Return	\$212

Includes all rows from the right table, even if there are no matching rows in the left table.

59

s104d09

Full Join

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

```
select *
  from customers c full join transactions t
  on c.ID = t.ID;
```

ID	Name	ID	Action	Amount
101	Smith	.	.	.
102	Blank	102	Purchase	\$100
.	.	103	Return	\$52
104	Jones	.	.	.
.	.	105	Return	\$212

Includes all rows from both tables, even if there are no matching rows in either table

s104d10

60

4.07 Quiz

Given these two SELECT statements:

```
select *
  from customers left join transactions
  on customers.ID=transactions.ID;
```

```
select *
  from transactions right join customers
  on customers.ID=transactions.ID;
```

1. Will the result sets contain the same data?
2. What is the difference in the result sets?

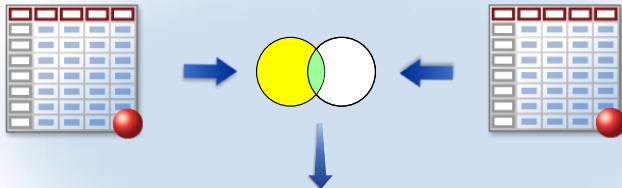
s104d11

61

Report 3

Management is considering a matching spouse donation program. A report is needed showing employees who are married and make contributions to company-sponsored charities.

orion.employee_donations **orion.employee_payroll**



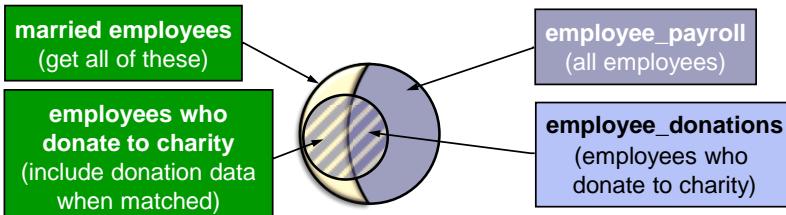
Employee_ID Recipients

121128	Cancer Cures, Inc.
121131	Vox Victimas 40%, Conserve Nature, Inc. 60%
121132	EarthSalvors 50%, Vox Victimas 50%

63

Business Data: Part 4

- The table **orion.employee_payroll** contains gender and marital status information.
- The table **orion.employee_donations** contains records only for those employees who donate to a charity via the company program.
- About half of all employees are married.



67

4.08 Multiple Choice Poll

For the report, you need the data for all married employees from **orion.employee_payroll**. You also want to include the charity names from the **orion.employee_donations** table if **Employee_ID** matches. What type of join should you use to combine the information from these two tables?

- a. inner join
- b. left join
- c. full join
- d. none of the above

68

Completed Code for Report 3

This is the PROC SQL code that generates the manager's requested report.

```
proc sql;
  select d.Employee_ID, Recipients
    from orion.employee_donations as d
      left join
        orion.employee_payroll as p
      on d.Employee_ID=
        p.Employee_ID
      where Marital_Status="M";
quit;
```

-  Output order is not guaranteed unless an ORDER BY clause is included.

s104d12

70

Viewing the Output

Partial PROC SQL Output (Rows 54–64)

Employee_ID	Recipients
121102	Cancer Cures, Inc.
121126	Save the Baby Animals 60%, EarthSalvors 40%
121128	Cancer Cures, Inc.
121131	Vox Victimas 40%, Conserve Nature, Inc. 60%
121132	EarthSalvors 50%, Vox Victimas 50%
121133	Disaster Assist, Inc.
121138	Cuidadores Ltd.
121142	AquaMissions International 10%, Child Survivors 90%
121143	Mitleid International 60%, Save the Baby Animals 40%
121145	Save the Baby Animals
121147	Cuidadores Ltd. 50%, Mitleid International 50%

71

DATA Step Merge

A DATA step with MERGE and BY statements combines tables similar to an SQL join. Columns are treated differently in the results.

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

Customers must be sorted or indexed on column ID before a merge can be performed.

72

DATA Step Merge

A DATA step merge automatically overlays same-name columns.

```
data merged;
  merge customers transactions;
  by ID;
run;
proc print data=merged;
run;
```

ID	Name	Action	Amount
101	Smith		.
102	Blank	Purchase	\$100
103		Return	\$52
104	Jones		.
105		Return	\$212

73

s104d13

SQL Join versus DATA Step Merge

SQL joins do not overlay same-name columns.

customers

ID	Name
101	Smith
104	Jones
102	Blank

transactions

ID	Action	Amount
102	Purchase	\$100
103	Return	\$52
105	Return	\$212

```
proc sql;
select *
  from customers c full join transactions t
  on c.ID=t.ID;
quit;
```

ID	Name	ID	Action	Amount
101	Smith	.	.	.
102	Blank	102	Purchase	\$100
.		103	Return	\$52
104	Jones	.	.	.
.		105	Return	\$212

74

s104d13

COALESCE Function

You can use the *COALESCE function* to overlay columns. The COALESCE function returns the value of the first nonmissing argument.

```
proc sql;
select coalesce(c.ID,t.ID) as ID,
       Name, Action, Amount
  from customers c full join transactions t
    on c.ID=t.ID;
quit;
```

COALESCE(argument-1,argument-2<, ...argument-n>)

75

s104d13



The arguments in the COALESCE function can be a variable name, a constant, or an expression. When all arguments are missing, COALESCE returns a missing value. All arguments must be of the same type (character or numeric).

Viewing the Output

The COALESCE function returns a single ID column in the output.

ID	Name	Action	Amount
101	Smith		.
102	Blank	Purchase	\$100
103		Return	\$52
104	Jones		.
105		Return	\$212

76

Comparing Inner Joins and Outer Joins

Key Points	Inner Join	Outer Join
Table Limit	256	256
Join Behavior	Returns matching rows only	Returns matching and nonmatching rows
Join Options	Matching rows only	LEFT, FULL, RIGHT
Syntax changes	<ul style="list-style-type: none"> ■ Multiple tables, separated by commas, in the FROM clause ■ WHERE clause that specifies join criteria 	ON clause that specifies join criteria

77

Tables can be joined on inequalities—for example:

```
proc sql;
  title "List of things I could buy";
  select Item_name, Price
    from budget, wish_list
   where budget.Cash_Available > wish_list.Price;
quit;
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

4. Outer Join

Join **orion.sales** and **orion.employee_addresses** on **Employee_ID** to create a report showing the names and cities of all Orion Star employees. If an employee is in the Sales Department, also show the job title. Present the report in alphabetical order by city, job title, and name.

- The **orion.sales** table contains a record for every employee in the Sales Department and includes columns **Employee_ID** and **Job_Title**.
- The **orion.employee_addresses** table contains a record for every employee and includes **Employee_ID**, **Employee_Name**, and **City**.

Partial PROC SQL Output

Name	City	Job_Title
Blanton, Brig	Melbourne	
Catenacci, Reyne	Melbourne	
Dillin, Kerrin	Melbourne	
Fiocca, Jina	Melbourne	
Fouche, Madelaine	Melbourne	
Glattback, Ellis	Melbourne	
Graham-Rowe, Jannene	Melbourne	
Gromek, Gladys	Melbourne	
Harwood, Reece	Melbourne	
Hieds, Merle	Melbourne	
Horsey, Riu	Melbourne	
McCleary, Bill	Melbourne	
Moffat, Trent	Melbourne	
Pettolino, Peter	Melbourne	
Povey, Liz	Melbourne	
Santomaggio, Pearl	Melbourne	
Sheedy, Sherie	Melbourne	
Streit, Russell	Melbourne	
Zhou, Tom	Melbourne	Sales Manager
Barcoe, Selina	Melbourne	Sales Rep. I
Chantharasy, Judy	Melbourne	Sales Rep. I
Duckett, Shani	Melbourne	Sales Rep. I
Osborn, Hernani	Melbourne	Sales Rep. I
Pa, Koavea	Melbourne	Sales Rep. I
Scordia, Randal	Melbourne	Sales Rep. I
Simms, Doungkamol	Melbourne	Sales Rep. I
Aisbitt, Sandy	Melbourne	Sales Rep. II
George, Vino	Melbourne	Sales Rep. II
Magrath, Brett	Melbourne	Sales Rep. II

Level 2

5. Outer Join

Join the **orion.order_fact** and **orion.product_dim** tables to determine whether there are any items in the product dimension table that have never been sold.

End of Exercises

4.4 Complex SQL Joins

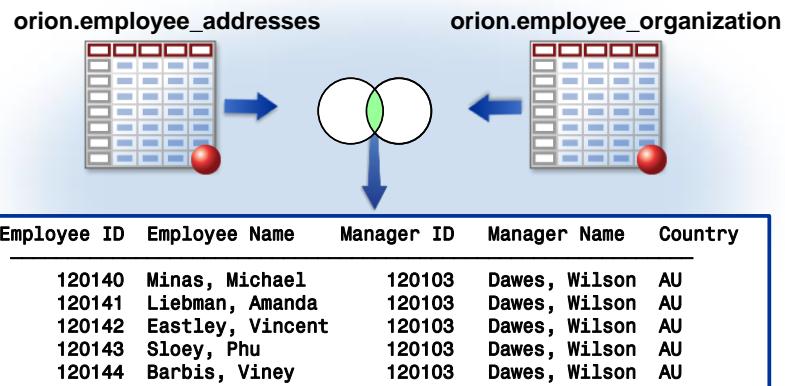
Objectives

- Perform a self-join (reflexive join).

81

Business Scenario

The chief sales officer wants to have a report with the name of all sales employees and the name of each employee's direct manager.



82

Business Data

To return the employee name and the manager name, you need to read the **addresses** table twice.

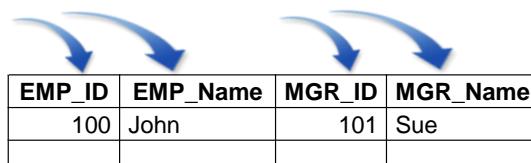
1. Return the employee's ID and name.
2. Determine the ID of the employee's manager.
3. Return the manager's name.

addresses

EMP_ID	EMP_NAME
100	John
101	Sue

organization

EMP_ID	MGR_ID
100	101
101	57



85

Required Table Aliases

In order to read from the same table twice, it must be listed in the FROM clause twice. Here, a different table alias is required to distinguish the different uses.

```
from orion.employee_addresses as e,
      orion.employee_addresses as m,
```

```
FROM table-name-1 <AS> alias-1,
      table-name-1 <AS> alias-2
```

86

Required Table Aliases

```

select e.Employee_ID "Employee ID",
       e.Employee_Name "Employee Name",
       m.Employee_ID "Manager ID",
       m.Employee_Name "Manager Name",
       e.Country
  from orion.employee_addresses as e,
       orion.employee_addresses as m,
       orion.employee_organization as o
 where e.Employee_ID=o.Employee_ID and
       o.Manager_ID=m.Employee_ID and
       Department contains 'Sales'
 order by Country,4,1;

```

87

s104d14

Viewing the Output

Partial PROC SQL Output

Employee ID	Employee Name	Manager ID	Manager Name	Country
120140	Minas, Michael	120103	Dawes, Wilson	AU
120141	Liebman, Amanda	120103	Dawes, Wilson	AU
120142	Eastley, Vincent	120103	Dawes, Wilson	AU
120143	Sloey, Phu	120103	Dawes, Wilson	AU
120144	Barbis, Viney	120103	Dawes, Wilson	AU

88



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

6. Creating a Report with a Self-Join

Create a report of all trainees and temporary workers to include their ID, name, and job title and their manager's ID and name. Order the report by the employee's ID. Label the columns and title the report as shown in the sample output.

The data that you need can be found in the listed columns of the following tables:

- **orion.employee_addresses** contains
 - **Employee_ID**
 - **Employee_Name**.
- **orion.staff** contains
 - **Employee_ID**
 - **Job_Title**
 - **Manager_ID** (**Employee_ID** of the person's manager).

Partial PROC SQL Output

Trainee and Temporary Employees					
Employee		Employee Job Title		Manager	
ID	Employee_Name	ID	Job Title	ID	Manager Name
120181	Cantatore, Lorian		Temp. Sales Rep.	120103	Dawes, Wilson
120182	Barreto, Geok-Seng		Temp. Sales Rep.	120103	Dawes, Wilson
120183	Blanton, Brig		Temp. Sales Rep.	120103	Dawes, Wilson
120184	Moore, Ari		Temp. Sales Rep.	120103	Dawes, Wilson
120185	Bahlman, Sharon		Temp. Sales Rep.	120103	Dawes, Wilson
120186	Quinby, Merryn		Temp. Sales Rep.	120103	Dawes, Wilson
120187	Catenacci, Reyne		Temp. Sales Rep.	120103	Dawes, Wilson
120188	Baran, Shanmuganathan		Temp. Sales Rep.	120103	Dawes, Wilson
120189	Lachlan, Mihailo		Temp. Sales Rep.	120103	Dawes, Wilson
120190	Czernezkyi, Ivor		Trainee	120103	Dawes, Wilson
120191	Graham-Rowe, Jannene		Trainee	120103	Dawes, Wilson

Level 2

7. Joining Multiple Tables

Create a report of Orion Star employees with more than 30 years of service as of February 1, 2013. Display the employee's name, years of service, and the name of the employee's manager. Order the report alphabetically by manager name, by descending years of service, and then alphabetically by employee name. Label the columns and title the report as shown in the sample output.

The data that you need can be found in the listed columns of the following tables:

- **orion.employee_addresses** contains
 - **Employee_ID**
 - **Employee_Name**.
- **orion.employee_payroll** contains
 - **Employee_ID**
 - **Employee_Hire_Date**.
- **orion.employee_organization** contains
 - **Employee_ID**
 - **Manager_ID** (**Employee_ID** of the person's manager).

Partial PROC SQL Output

Employees with More Than 30 Years of Service As of February 1, 2013		
Employee Name	Years of Service	Manager Name
Marion, Chiorene	35	Ardskin, Elizabeth
Zied, Ahmed	35	Ardskin, Elizabeth
Droste, Victor	31	Ardskin, Elizabeth
Parker, U'Vonda	31	Ardskin, Elizabeth
Buddery, Jeannette	35	Billington, Karen

End of Exercises

4.5 Solutions

Solutions to Exercises

1. Inner Joins

```
*** s104s01 ***;

proc sql;
title "Employees With More Than 30 Years of Service";
title2 "As of February 1, 2013";
select Employee_Name 'Name',
       int((1FEB2013'd-Employee_Hire_Date)/365.25)
            as YOS 'Years of Service'
  from orion.employee_addresses as a,
       orion.employee_payroll as p
 where a.Employee_ID=p.Employee_ID
       and calculated YOS gt 30
  order by Employee_Name;
quit;
title;
```

2. Creating a Summary Report from Two Tables

```
*** s104s02 ***;

proc sql;
title 'Total Quantities Sold by Product ID and Name';
select p.Product_ID,
       Product_Name,
       sum(Quantity) 'Total Sold'
  from orion.product_dim as p,
       orion.order_fact as o
 where p.Product_ID = o.Product_ID and Order_Date>='01Jan2010'd
  group by p.Product_ID, Product_Name
  order by 3 desc, Product_Name;
quit;
title;
```

3. Joining Multiple Tables

```
*** s104s03 ***;

proc sql;
title "US and Australian Catalog and Internet Customers";
title2 "Purchasing Foreign Manufactured Products";
select Customer_Name 'Name', Count(*) 'Purchases'
       as Count
  from orion.product_dim as p,
       orion.order_fact as o,
       orion.customer as c
 where p.Product_ID=o.Product_ID
```

```

        and o.Customer_ID=c.Customer_ID
        and Employee_ID=99999999
        and p.Supplier_Country ne Country
        and Country in ('US','AU')
    group by Customer_Name
    order by Count desc, Customer_Name;
quit;
title;

```

4. Outer Join

```

*** s104s04 ***;

proc sql;
select Employee_Name 'Name' format=$35.,
       City, Job_Title
  from orion.employee_addresses as a
       left join
       orion.sales as s
      on a.Employee_ID=s.Employee_ID
  order by City, Job_Title, Employee_Name;
quit;

```

5. Outer Join

```

*** s104s05 **;

proc sql;
title 'Products That Have Not Been Sold';
select Product_Name,
       p.Product_ID,
       Quantity
  from orion.product_dim as p
       left join
       orion.order_fact as o
      on p.Product_ID=o.Product_ID
     where Order_ID is missing;
quit;
title;

```

6. Creating a Report with a Self-Join

```

*** s104s06 ***;

proc sql;
title 'Trainee and Temporary Employees';
select a1.Employee_ID 'Employee ID',
       a1.Employee_Name 'Employee Name',
       Job_Title,
       a2.Employee_ID 'Manager ID',
       a2.Employee_Name 'Manager Name'
  from orion.employee_addresses as a1,
       orion.employee_addresses as a2,
       orion.staff as s

```

```

where a1.Employee_ID=s.Employee_ID and
      s.Manager_ID=a2.Employee_ID and
      (Job_Title contains 'Trainee' or
       Job_Title contains 'Temp')
order by a1.Employee_ID;
quit;
title;

```

7. Joining Multiple Tables

```

*** s104s07 ***;

proc sql;
title "Employees With More Than 30 Years of Service";
title2 "As of February 1, 2013";
select emp.Employee_Name 'Employee Name' format=$35.,
      int((1FEB2013'd-Employee_Hire_Date)/365.25)
            as YOS 'Years of Service',
      mgr.Employee_Name 'Manager Name' as Manager_Name
/* Employee Addresses:
   First copy is required to
   look up Employee information Employee's
   Employee_ID */
/* Employee_Organization:
   Links Employee_ID to Manager_ID */
/* Employee Addresses:
   Second copy is required to
   look up Manager information using Manager's
   Employee_ID */
from orion.employee_addresses as emp,
     orion.employee_organization as org,
     orion.employee_payroll as pay,
     orion.employee_addresses as mgr
where emp.Employee_ID=pay.Employee_ID
  and emp.Employee_ID=org.Employee_ID
  and org.Manager_ID=mgr.Employee_ID
  and calculated YOS gt 30
order by Manager_Name, YOS desc, Employee_Name;
quit;
title;

```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

4.01 Multiple Choice Poll – Correct Answer

Which of these DATA step statements is used to combine tables horizontally?

- a. SET
- b. APPEND
- c. MERGE**
- d. INPUT
- e. INFILE

8

4.02 Short Answer Poll – Correct Answer

How many rows and columns are returned from this query?

20 rows

5 columns

ID	Name	ID	Action	Amount
101	Jones	102	Purchase	\$376
101	Jones	102	Return	\$119
101	Jones	103	Purchase	\$57
101	Jones	105	Purchase	\$98
101	Jones	102	Purchase	\$376
101	Jones	102	Return	\$119
101	Jones	103	Purchase	\$57
101	Jones	105	Purchase	\$98
102	Kent	102	Purchase	\$376
102	Kent	102	Return	\$119
102	Kent	103	Purchase	\$57
102	Kent	105	Purchase	\$98
102	Kent	102	Purchase	\$376
102	Kent	102	Return	\$119
102	Kent	103	Purchase	\$57
102	Kent	105	Purchase	\$98
104	Avery	102	Purchase	\$376
104	Avery	102	Return	\$119
104	Avery	103	Purchase	\$57
104	Avery	105	Purchase	\$98

25

4.03 Short Answer Poll – Correct Answer

Why are there two ID columns in the result?

ID	Name	ID	Action	Amount
102	Blank	102	Purchase	\$100

The * was used to select all columns from both tables.
Therefore, you have two ID columns in the results.

32

4.04 Short Answer Poll – Correct Answer

What message is displayed in the log when the program is executed?

Error messages are shown.

```
ERROR: Ambiguous reference, column ID is in more than one table.  
ERROR: Ambiguous reference, column ID is in more than one table.  
      quit;  
NOTE: The SAS System stopped processing this step because of errors.
```

36

4.05 Multiple Choice Poll – Correct Answer

How many rows (observations) result from the DATA step MERGE statement in program **s104a03**?

- a. 2
- b. 4
- c. 6
- d. 20
- e. none of the above

Obs	ID	Name	Action	Amount
1	102	Kent	Purchase	\$376
2	102	Kent	Return	\$119

44

customer2 **transaction2**

ID	Name	ID	Action	Amount
101	Jones	102	Purchase	\$376
101	Jones	102	Return	\$119
102	Kent	103	Purchase	\$57
102	Kent	105	Purchase	\$98
104	Avery			

4.06 Short Answer Poll – Correct Answer

How many tables can be combined using the following PROC SQL syntax?

```
proc sql;
  select *
    from table-name,..., table-name
    where sql-expression;
quit;
```

Up to 256 tables can be joined with this syntax.

48

4.07 Quiz – Correct Answer

Given these two SELECT statements:

```
select *
  from customers left join transactions
  on customers.ID=transactions.ID;
```

```
select *
  from transactions right join customers
  on customers.ID=transactions.ID;
```

1. Will the result sets contain the same data? Yes
2. What is the difference in the result sets? The columns are in a different order.

Left Join - C and then T

ID	Name	ID	Action	Amount
001	Smith	.	.	.
002	Blank	002	Purchase	\$100
004	Jones	.	.	.

62

Right Join - T and then C

ID	Action	Amount	ID	Name
001	Smith	.	001	Smith
002	Purchase	\$100	002	Blank
.	.	.	004	Jones

4.08 Multiple Choice Poll – Correct Answer

For the report, you need the data for all married employees from **orion.employee_payroll**.

You also want to include the charity names from the **orion.employee_donations** table if **Employee_ID** matches. What type of join should you use to combine the information from these two tables?

- a. inner join
- b. left join**
- c. full join
- d. none of the above

69

Chapter 5 Subqueries

5.1 Noncorrelated Subqueries	5-3
Exercises	5-18
5.2 In-Line Views.....	5-22
Exercises	5-38
5.3 Solutions	5-42
Solutions to Exercises	5-42
Solutions to Student Activities (Polls/Quizzes)	5-48

5.1 Noncorrelated Subqueries

Objectives

- Define PROC SQL subqueries.
- Differentiate between correlated and noncorrelated subqueries.
- Subset data based on values returned from a subquery.

3

Business Scenario

HR and Payroll managers requested a report that displays **Job_Title** for job groups with an average salary greater than the average salary of the company as a whole.



4

Step 1

Calculate the company's average salary.

```
proc sql;
select avg(Salary) as CompanyMeanSalary
  from orion.staff;
quit;
```

Company MeanSalary
38041.51

5

s105d01

Step 2

Determine the job titles whose average salary exceeds the company's average salary.

```
proc sql;
select Job_Title,
       avg(Salary) as MeanSalary
  from orion.staff
 group by Job_Title
 having MeanSalary>38041.51;
quit;
```

Partial PROC SQL Output

Employee Job Title	MeanSalary
Account Manager	46090
Administration Manager	47415
Applications Developer I	42760

6

s105d01

Step 3

Write the program as a single step using a subquery.

A *subquery* is a query that resides within an outer query.

```
proc sql;
  select Job_Title, avg(Salary) as MeanSalary
    from orion.staff
   group by Job_Title
   having avg(Salary) >
      (select avg(Salary)
       from orion.staff);
quit;
```

outer query

subquery

- ✎ The subquery must be resolved before the outer query can be resolved

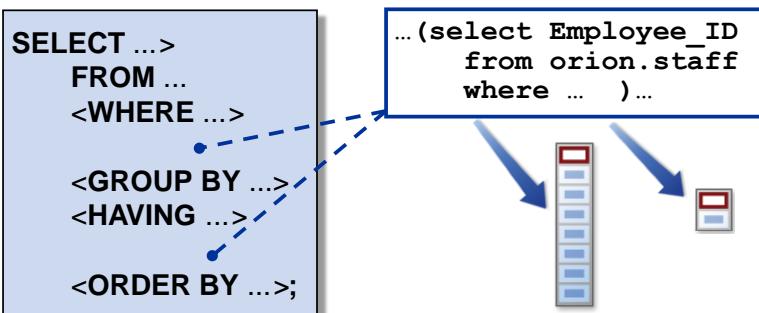
7

s105d02

Subqueries

A subquery

- returns values to be used in the outer query's WHERE or HAVING clause
- must return only a single column
- can return multiple values or a single value.



11

Subqueries are also known as *nested queries*, *inner queries*, and *sub-selects*.

Setup for the Poll

Is the following SELECT statement valid as a subquery?

```
(select Employee_ID,  
      Gender  
   from orion.staff)
```

12

5.01 Poll

Is the SELECT statement valid as a subquery?

- Yes
- No

13

Subqueries: Noncorrelated

There are two types of subqueries:

- A *noncorrelated subquery* is a self-contained query.
It executes independently of the outer query.

```
proc sql;
select Job_Title, avg(Salary) as MeanSalary
  from orion.staff
 group by Job_Title
 having avg(Salary) >
       (select avg(Salary)
        from orion.staff);
quit;
```

This query is a stand-alone query.

15

Subqueries: Correlated

- A *correlated subquery* requires a value or values to be passed to it by the outer (main) query before it can be successfully resolved.

```
proc sql;
select Employee_ID, avg(Salary) as MeanSalary
  from orion.employee_addresses
 where 'AU'=
       (select Country
        from work.supervisors
        where employee_addresses.Employee_ID=
              supervisors.Employee_ID);
quit;
```

This query is not stand-alone.
It needs additional information
from the main query.

16

Noncorrelated Subquery

```
proc sql;
  select Job_Title,
         avg(Salary) as MeanSalary
    from orion.staff
   group by Job_Title
  having avg(Salary) >
        (select avg(Salary)
         from orion.staff);
quit;
```

Evaluate the subquery first.

17

s105d02

Noncorrelated Subquery

```
proc sql;
  select Job_Title,
         avg(Salary) as MeanSalary
    from orion.staff
   group by Job_Title
  having avg(Salary) >
        (38041.51);
quit;
```

Then pass the results to the outer query.

Partial PROC SQL Output

Employee	Job Title	MeanSalary
	Account Manager	46090
	Administration Manager	47415
	Applications Developer I	42760

18

s105d02

5.02 Poll

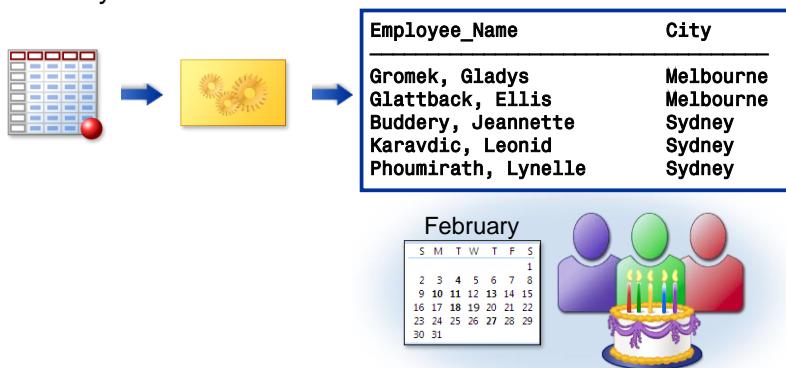
Can a subquery contain a subquery?

- Yes
- No

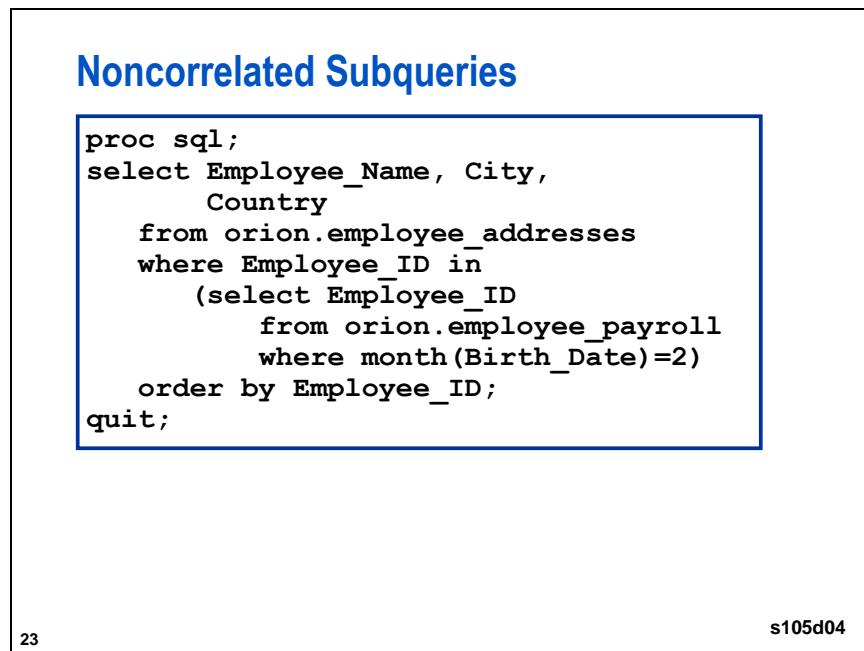
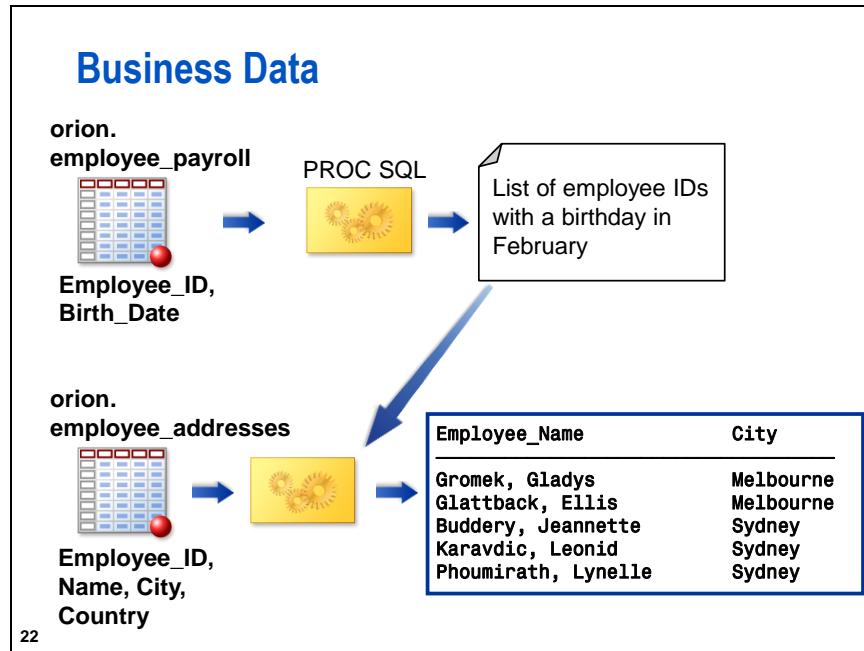
19

Business Scenario

The CEO sends a birthday card to each employee having a birthday in the current month. Create a report listing the names, cities, and countries of employees with February birthdays.



21



✍ To make this code dynamic, replace the month value, 2, with **month(today())**. Now the program will return a list of employees with a birthday in the current month.

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
select Employee_Name, City,
       Country
  from orion.employee_addresses
 where Employee_ID in
       (select Employee_ID
          from orion.employee_payroll
         where month(Birth_Date)=2)
  order by Employee_ID;
quit;
```

Step 1: Evaluate the inner query and create a list of values.

Partial
orion.employee_payroll

Employee_ID	Birth_Date
...	...
120106	23DEC1948
120107	21JAN1953
120108	23FEB1988
120109	15DEC1990
120110	20NOV1953
120111	23JUL1953
120112	17FEB1973
120113	10MAY1948
...	...

24

...

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
select Employee_Name, City,
       Country
  from orion.employee_addresses
 where Employee_ID in
       (120108,120112,120114,120157,
        120159, 120170,...)
  order by Employee_ID;
quit;
```

**list returned by
the inner query**

Partial
orion.employee_payroll

Employee_ID	Birth_Date
...	...
120106	23DEC1948
120107	21JAN1953
120108	23FEB1988
120109	15DEC1990
120110	20NOV1953
120111	23JUL1953
120112	17FEB1973
120113	10MAY1948
...	...

25

...

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
select Employee_Name, City,
       Country
  from orion.employee_addresses
 where Employee_ID in
       (120108,120112,120114,120157,
        120159, 120170,...)
  order by Employee_ID;
quit;
```

Step 2: Pass the values to the outer query for use in its WHERE clause.

Partial
orion.employee_addresses

Employee_ID	Employee_Name
...	...
120106	Hornsey, John
120107	Sheedy, Sherie
120108	Gromek, Gladys
120109	Baker, Gabriele
120110	Entwistle, Dennis
120111	Spillane, Ubaldo
120112	Glattback, Ellis
120113	Horsey, Riu
...	...

s105d04

26

Viewing the Output

Partial PROC SQL Output

Employee_Name	City	Country
Gromek, Gladys	Melbourne	AU
Glattback, Ellis	Melbourne	AU
Buddery, Jeannette	Sydney	AU
Karavdic, Leonid	Sydney	AU
Phoumirath, Lynelle	Sydney	AU

27

Setup for the Poll

- Open the program **s105a01**.
- Change the IN operator to an equal sign in the code.
- Run the changed program and review the SAS log for messages.

```
proc sql;
  select Employee_Name, City, Country
    from orion.employee_addresses
   where Employee_ID =
         (select Employee_ID
          from orion.employee_payroll
         where month(Birth_Date)=2)
      order by 1;
quit;
```

28

5.03 Multiple Choice Poll

What happens when you change the comparison operator to an equal sign?

- a. Nothing special, the program runs fine.
- b. Multiple rows are returned in your output.
- c. You get an error message.
- d. a and b

29

Subqueries That Return Multiple Values

When a subquery returns multiple values and the EQUAL operator is used, an ERROR message is generated.

ERROR: Subquery evaluated to more than one row.

- ☞ If the subquery returns multiple values, you must use the IN operator or a comparison operator with the ANY or ALL keywords.

31

ANY Keyword

The ANY expression is true if it is true for any of the values returned by the subquery.

Keyword ANY	Signifies...
= ANY(20,30,40)	=20 or =30 or =40
> ANY(20,30,40)	> 20
< ANY(20,30,40)	< 40

- ☞ The values 20,30,40 represent values returned from a subquery.

32

ALL Keyword

The ALL expression is true if it is true for all the values returned by the subquery.

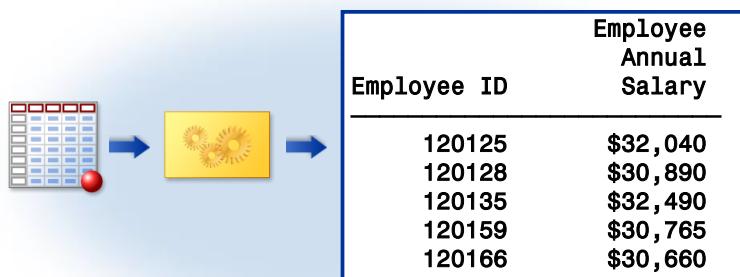
Keyword ALL	Signifies
> ALL(20,30,40)	> 40
< ALL(20,30,40)	< 20

- The values 20,30,40 represent values returned from a subquery.

33

Business Scenario

The senior sales executive asked the following: “Do any Level IV sales representatives have a salary that is lower than any of the lower-level sales representatives?”



34

Solution 1: ANY Keyword

```
proc sql;
select Employee_ID, Salary
  from orion.staff
 where Job_Title='Sales Rep. IV'
       and Salary < any
 (select Salary
   from orion.staff
  where Job_Title in
      ('Sales Rep. I','Sales Rep. II',
       'Sales Rep. III'));
quit;
```

35

s105d05

Solution 2: MAX Statistic

```
proc sql;
select Employee_ID, Salary
  from orion.staff
 where Job_Title='Sales Rep. IV'
       and Salary <
 (select max(Salary)
   from orion.staff
  where Job_Title in
      ('Sales Rep. I','Sales Rep. II',
       'Sales Rep. III'));
quit;
```

36

s105d05

Viewing the Output

Partial PROC SQL Output

Level IV Sales Reps Who Earn Less Than
Any Lower Level Sales Reps.

Employee ID	Employee Annual Salary
120125	\$32,040
120128	\$30,890
120135	\$32,490
120159	\$30,765
120166	\$30,660



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Using a Noncorrelated Subquery

The **orion.order_fact** table contains information about orders that were placed by customers. Create a report that lists the retail customers whose average retail price exceeds the company average retail sales.

- Write a query that displays the average of **Total_Retail_Price** for all retail prices in the table.
 - Use **AVG(Total_Retail_Price)** to calculate the average.
 - Subset the rows so that only the retail sales are included (**Order_Type=1**).

PROC SQL Output

MeanSales
137.8227

- Write a query that displays **Customer_ID** and **AVG(Total_Retail_Price)** for those customers whose average retail price exceeds the company average retail price. The query should do the following:
 - Display the values for **Customer_ID** and **AVG(Total_Retail_Price)**. Name the second column **MeanPrice**.
 - Subset the rows do only the retail sales are included (**Order_Type=1**).
 - Group the data by **Customer_ID**.
 - Include only groups where the employee's average retail price exceeds the company average. Use the query from 1.a. as a subquery in the HAVING clause.
 - Add a title to the report as shown.

Partial PROC SQL Output

Customers Whose Average Retail Price Exceeds the Average Retail Price for All Customers
Customer ID MeanSales
<hr/> 10 141.91 17 139.03 31 158.11 45 178.03 49 161.66

Level 2

2. Using a Noncorrelated Subquery

Each month a memo is that lists the employees who have employment anniversaries for that month is posted. Create the report for the current month and list **Employee_ID** and the first and last names for all employees hired during the current month of any year.

You can find **Employee_Name** in the **orion.employee_addresses** table and **Employee_Hire_Date** in the **orion.employee_payroll** table. Both tables contain the column **Employee_ID**. Order the report by an employee's last name.

- a. Create a query that returns a list of employee IDs for employees with a current anniversary. The query should do the following:
 - Display **Employee_ID** numbers.
 - Use the **orion.employee_payroll** table.
 - Return only employees whose hire date (**Employee_Hire_Date**) is in the current month.
 - Add a title to the report as shown.

Partial PROC SQL Output (Generated when current month = February)

Employee IDs for Current Month Anniversaries
Employee_ID
120107
120116
120136
120162
120164

- b. Using the query in 2.a. as a noncorrelated subquery, write a query that displays the employee IDs and names of employees who have current month anniversaries. The final query should do the following:
 - Display **Employee_ID** and split **Employee_Name** into two new columns: **FirstName** and **LastName**. Both new columns should have a length of \$15 and appropriate labels. (See the report below.) The original **Employee_Name** is stored as **Lastname, Firstname**.
 - Use the **orion.employee_addresses** table.
 - Select **Employee_ID** only for employees who had current month anniversaries.
 - Order the final results by **LastName**.
 - Create an appropriate title.

Partial PROC SQL Output (Generated when current month = February)

Employees with Current Month Anniversaries		
Employee_ID	First Name	Last Name
121030	Jeryl	Areu
121007	John	Banaszak
120667	Edwin	Droste
120778	Angela	Gardner

3. Creating Subqueries Using the ALL Keyword

In most companies, you can assume that the higher-level job titles have employees that are older than employees with a lower-level job title. Using the **orion.staff** table, determine whether there are any lower-level purchasing agents (Purchasing Agent I and Purchasing Agent II) that are older than all the higher-level purchasing agents (Purchasing Agent III). The final report should display **Employee_ID**, **Job_Title**, **Birth_Date**, and a calculated **Age** column for the employee as of 02FEB2013.

Hint: Use the SAS date constant ('02FEB2013'd) in the calculation for **Age**. If you use the TODAY function to calculate the age, the values will differ from the results below:

Level I or II Purchasing Agents Who are older than ALL Purchasing Agent III's			
Employee ID	Employee Job Title	Birth Date	Age
120742	Purchasing Agent I	04FEB1948	64

Challenge

4. Using Nested Subqueries

The Marketing Department is interested in developing a campaign pointed toward the Orion Club members low activity group. They specifically want to identify the customer IDs of this group whose last order was placed before 1Jan2012. The requested report should contain the customer ID and the date of the customer's last order.

The following tables contain the data that you need:

orion.customer_type contains the columns **Customer_Type** and **Customer_Type_ID**.

orion.customer contains **Customer_Type_ID** and **Customer_ID**.

orion.order_fact contains **Customer_ID** and **Order_Date**.

PROC SQL Output**Latest Order Date for
Orion Club Low Activity Members****Customer ID Order Date**

4	18-DEC-2008
18	16-FEB-2011
34	18-NOV-2011
65	03-AUG-2011
69	18-OCT-2011
75	10-NOV-2008
92	01-MAY-2011
183	21-APR-2011
195	12-JUN-2010
928	04-SEP-2011
1033	06-MAY-2011
1100	10-DEC-2010
1684	01-AUG-2010
70046	05-FEB-2008
70108	22-JUL-2009
70165	28-JUL-2011

End of Exercises

5.2 In-Line Views

Objectives

- Create and use in-line views.
- Use in-line views and subqueries to simplify coding a complex query.

41

Business Scenario

List all active Sales Department employees who have annual salaries significantly lower (less than 95%) than the average salary for everyone with the same job title.



Employees with Salaries less than 95% of the Average for their Job

Employee_Name	Employee_Job_Title	Annual_Salary	Job_Avg
Ould, Tulsidas	Sales Rep. I	22,710	26,576
Polky, Asishana	Sales Rep. I	25,110	26,576
Voron, Tachaun	Sales Rep. I	25,125	26,576

42

Step 1

Calculate the average salaries for active employees in the Sales Department, grouped by job title.

```
title  'Sales Department Average Salary';
title2 'By Job Title';
select Job_Title,
       avg(Salary) as Job_Avg
  from orion.employee_payroll as p,
       orion.employee_organization as o
 where p.Employee_ID=o.Employee_ID
       and Employee_Term_Date is missing
       and Department="Sales"
 group by Job_Title;
```

43

s105d06

Viewing the Output

PROC SQL Output

Sales Department Average Salary by Job Title	
Job_Title	Job_Avg
Sales Rep. I	26575.76
Sales Rep. II	27347.97
Sales Rep. III	29213.62
Sales Rep. IV	31588.5

44

Setup for the Poll

Can this SELECT statement be used as a subquery?

```
select Job_Title,
       avg(Salary) as Job_Avg
       format=comma7.
  from orion.employee_payroll as p,
       orion.employee_organization as o
 where p.Employee_ID=o.Employee_ID
       and Employee_Term_Date is missing
       and Department="Sales"
 group by Job_Title;
```

45

5.04 Poll

Can this SELECT statement be used as a subquery?

- Yes
- No

46

In-Line Views

An *in-line view* is a query expression (SELECT statement) that resides in a FROM clause. It acts as a virtual table, used in place of a physical table in a query.

```
PROC SQL;
SELECT *
FROM
  (in-line view query expression)
  ...
QUIT;
```

In-line views are often useful when you build complex SQL queries.

48

- ✍ An in-line view consists of any valid SQL query, except that it must not contain an ORDER BY clause. In-line views can be independently tested while building a complex query, which can simplify troubleshooting.

Step 2

Match each employee to a job title group and compare the employee's salary to the group's average to determine whether it is less than 95% of the group average.

```
select Employee_Name, emp.Job_Title,
       Salary format=comma7., Job_Avg format=comma7.
  from (select Job_Title,
              avg(Salary) as Job_Avg format=comma7.
            from orion.employee_payroll as p,
                 orion.employee_organization as o
           where p.Employee_ID=o.Employee_ID
             and Employee_Term_Date is missing
             and Department="Sales"
        group by Job_Title) as job,
       orion.salesstaff as emp
  where emp.Job_Title=job.Job_Title
    and Salary<Job_Avg*.95
    and Emp_Term_Date is missing
  order by Job_Title, Employee_Name;
```

s105d07

49

- ✍ You can assign an alias to an in-line view as if it were a table.

Viewing the Output

PROC SQL Output

Employees with Salaries less than 95% of the Average for their Job				
Employee_Name	Employee	Job Title	Annual Salary	Job_Avg
Ould, Tulsidas	Sales Rep. I		22,710	26,576
Polky, Asishana	Sales Rep. I		25,110	26,576
Voron, Tachaun	Sales Rep. I		25,125	26,576

50

Setup for the Poll

```
select Name, emp.Job_Title, Job_Avg
  from (in-line view) as job,
       orion.staff as emp
 where emp.Job_Title=job.Job_Title;
select Job_Title, Job_Avg
  from job;
```

The first SELECT statement uses an in-line view to create data that is referenced with the alias **job**. Can **job** also be referenced in the second SELECT statement?

51

5.05 Poll

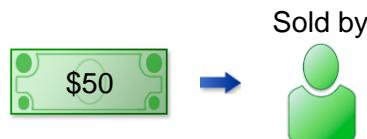
The first SELECT statement uses an in-line view to create data that is referenced with the alias **job**. Can **job** also be referenced in the second SELECT statement?

- Yes
- No

52

Business Scenario

Management wants to incent more sales of Expedition Zero sleeping bags by rewarding each sales associate who sold these items in 2011 with a \$50 certificate.



The Task:

Prepare a list of the managers' names and cities for those employees who qualify for the certificate.

55

Business Data

Because this query involves four tables, it might not be easy to code all at once. To simplify the task, split the query into small parts. Test each part individually, and test the overall query each time that a new segment is added.

These tables are required for this query:

`orion.order_fact`



`orion.product_dim`



`orion.employee_organization`



`orion.employee_addresses`



56

Planning the Complex Query

This query involves multiple tables. To simplify the task, split the query into parts and test as you progress.

Part 1 Identify the employees who sold Expedition Zero merchandise in 2011.

Part 2 Find the employee identifier for the managers of these employees.

Part 3 Obtain the managers' names and city information.

57

Business Data: Part 1

Part 1 Get employee IDs for employees who sold Expedition Zero merchandise in 2011.

Select the employee's identifier (**Employee_ID**) from the results of joining the **order_fact** and **product_dim** tables on **Product_ID**, where **Product_Name** contains *Expedition Zero*. Exclude Internet and catalog orders (**Employee_ID NE 99999999**).

58

Coding the Complex Query

Part 1 Write a query to obtain the employee ID of all employees who sold Expedition Zero merchandise in 2011.

```
select distinct Employee_ID
  from orion.order_fact as of,
       orion.product_dim as p
 where of.Product_ID=p.Product_ID
   and year(Order_Date)=2011
   and Product_Name contains
     'Expedition Zero'
   and Employee_ID ne 99999999;
```

PROC SQL Output

Employee ID
121065

s105d08

59

Business Data: Part 2

Part 2 Find the employee identifier for the managers of these employees.

Select the manager's identifier (**Manager_ID**) from the results of joining the **employee_organization** table with the first query's results on **Employee_ID**.

60

Setup for the Poll

To join the **employee_organization** table with the Step 1 query results, you use the query from Step 1 as which of the following?

```
select distinct Employee_ID
  from orion.order_fact as of,
       orion.product_dim as p
 where of.Product_ID=p.Product_ID
   and year(Order_Date)=2011
   and Product_Name contains
     'Expedition Zero'
   and Employee_ID ne 99999999;
```

61

5.06 Multiple Choice Poll

To join the **employee_organization** table with the Step 1 query results, you use the query from Step 1 as which of the following?

- a. an in-line view
- b. a noncorrelated subquery
- c. either a or b
- d. neither a or b

62

Coding the Complex Query

Part 2 Write a query to obtain the manager ID of the employee's manager.

```
select Manager_ID  
  from orion.employee_organization as o,  
       (<Step 1 query results>) as ID  
 where o.Employee_ID=ID.Employee_ID;
```

Employee_ID
121065

64

Coding the Complex Query

Part 2 Write a query to obtain the manager ID of the employee's manager.

```
select Manager_ID
  from orion.employee_organization as o,
       (select distinct Employee_ID
        from orion.order_fact as of,
             orion.product_dim as p
       where of.Product_ID=p.Product_ID
         and year(Order_Date)=2011
         and Product_Name
           contains 'Expedition Zero'
         and Employee_ID ne 99999999)as ID
 where o.Employee_ID=ID.Employee_ID;
```

PROC SQL Output

Manager_ID
121145

s105d09

65

Business Data: Part 3

Part 3 Find the managers' names and cities.

Select the employee's name (**Employee_Name**) and city from the **employee_addresses** table, where **Employee_ID** matches **Manager_ID** in the results of the previous query.

66

Setup for the Poll

Is it possible to use the entire query in Step 2 as a subquery?

```
select Manager_ID
  from orion.employee_organization as o,
       (select distinct Employee_ID
        from orion.order_fact as of,
             orion.product_dim as p
       where of.Product_ID=p.Product_ID
         and year(Order_Date)=2007
         and Product_Name
           contains 'Expedition Zero'
         and Employee_ID ne 99999999)as ID
  where o.Employee_ID=ID.Employee_ID;
```

67

5.07 Poll

Is it possible to use the entire query in Step 2 as a subquery?

- Yes
- No

68

Coding the Complex Query

Part 3 Write a query to obtain the managers' names and city information.

```
proc sql;
select Employee_Name format=$25. as Name, City
  from orion.employee_addresses
 where Employee_ID in
      (<Step 2 query results>);
```

Manager_ID
121145

70

Coding the Complex Query

```
proc sql;
select Employee_Name format=$25. as Name, City
  from orion.employee_addresses
 where Employee_ID in
      (select Manager_ID
        from orion.employee_organization as o,
             (select distinct Employee_ID
               from orion.order_fact as of,
                    orion.product_dim as p
                  where of.Product_ID=p.Product_ID
                  and year(Order_Date)=2011
                  and Product_Name contains
                      'Expedition Zero'
                  and Employee_ID ne 99999999) as ID
       where o.Employee_ID=ID.Employee_ID);
```

Part 3

71

s105d10

Viewing the Output

Part 3

PROC SQL Output

Name	City
Lansberry, Dennis	Miami-Dade

72

Coding the Complex Query

You can also solve this problem using a multiway join.

```
proc sql;
select distinct Employee_Name format=$25. as Name, City
  from orion.order_fact as of,
       orion.product_dim as pd,
       orion.employee_organization as eo,
       orion.employee_addresses as ea
 where of.Product_ID=pd.Product_ID
   and of.Employee_ID=eo.Employee_ID
   and ea.Employee_ID=eo.Manager_ID
   and Product_Name contains 'Expedition Zero'
   and year(Order_Date)=2011
   and eo.Employee_ID ne 99999999;
quit;
```

73

s105d11

Comparison with Traditional SAS Programs

```
*** s105d12 ***;

/* Traditional SAS Program Code */
/* for Complex Join */

***** Step 1: Identify the employees who sold Expedition Zero
merchandise in 2003.
*****
```

```

proc sort data=orion.order_fact
    (keep=Product_ID Employee_ID Order_Date
    where=(YEAR(Order_Date)=2011 and Employee_ID ne 99999999))
    out=orders_2011 (Drop=Order_Date);
    by Product_ID;
run;
proc sort data=orion.product_dim (keep=Product_ID Product_Name)
    out=products;
    by Product_ID;
run;
data employees (Keep=Employee_ID);
    merge orders_2011 (In=KeepMe)
        products (where=(Product_Name contains 'Expedition Zero'));
    by Product_ID;
    if KeepMe and Product_Name ne '';
run;

proc sort data=employees nodup;
    by Employee_ID;
run;

*****  

Step 2: Find the employee identifier for the managers of  

these employees  

*****  

data manager_id (rename=(Manager_ID=Employee_ID));
    merge employees (in=KeepMe)
        orion.employee_organization (keep=Employee_ID Manager_ID);
    by Employee_ID;
    if KeepMe;
    drop Employee_ID;
run;
proc sort data=manager_id nodup;
    by Employee_ID;
run;

*****  

Step 3: Obtain the managers' names and city information  

*****  

proc sort data=orion.employee_addresses (Keep=Employee_ID
Employee_Name City)
    out=employees;
    by Employee_ID;
run;

data managers;
    length Manager $28.;
    merge manager_id (in=KeepMe)
        employees;
    by Employee_ID;
    if KeepMe;

```

```
Manager=catx(' ',scan(Employee_Name,2,' ',''),
            scan(Employee_Name,1,' '));
drop Employee_ID Employee_Name;
run;
proc print data=managers noobs;
run;
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

5. Using In-Line Views

Produce a report of Orion Star sales force employees' aggregate sales in 2011.

- a. Select the **Country**, **First_Name**, **Last_Name**, **Value_Sold**, **Orders**, and **Avg_Order** columns by joining the **orion.order_fact** and **orion.sales** tables on **Employee_ID**. Group the report by **Country**, **First_Name**, and **Last_Name**. Include only employees having an aggregate **Value_Sold** of \$200.00 or more. Order the results by **Country**, **Value_Sold** (descending), and **Orders** (descending).
- Calculate **Value_Sold** by summing **Total_Retail_Price**.
 - Calculate **Orders** by using the **count(distinct Order_ID)** function to count the number of unique rows for each employee.



An employee can have multiple **Order_ID** row values for different products. Rows with the same order ID do not count as a separate order.

- Calculate **Avg_Order** by dividing **Value_Sold** by **Orders**.
- Subset the data to return only the rows for sales in 2011.
- Title the report as indicated in the sample output.

Partial PROC SQL Output

2011 Sales Force Sales Statistics For Employees With \$200.00 or More In Sales						
Country	First_Name	Last_Name	Value_Sold	Orders	Avg_Order	
AU	Lucian	Daymond	880.10	4	220.03	
AU	Ranj	Stamalis	697.60	1	697.60	
AU	Sharryn	Clarkson	400.40	3	133.47	
AU	Marinus	Surawski	398.80	2	199.40	
AU	Sian	Shannan	306.20	1	306.20	
AU	Monica	Kletschkus	239.30	1	239.30	
US	Tywanna	Mcdade	1,387.90	1	1387.90	

- b. Rewrite the query created in part a above and use it as an in-line view. Select **Country** and the maximum **Value_Sold**, **Orders**, and **Avg_Order** values, as well as the minimum **Avg_Order** value for each country. Name the report **2011 Sales Summary by Country**.

Hint: An in-line view cannot use the ORDER BY clause.

PROC SQL Output

2011 Sales Summary By Country					
Country	Max Value Sold	Max Orders	Max Average	Min Average	
AU	880.10	4.00	697.60	133.47	
US	1,387.90	4.00	1387.90	133.00	

Level 2**6. Building Complex Queries with In-Line Views**

Your ultimate goal in this exercise is to create a report showing each employee's salary expressed as a percentage of the total salary for that employee's department. The report should be sorted by department and, within each department, in descending order of salary percentage.

- The **orion.employee_payroll** table contains **Salary**.
- The **orion.employee_addresses** table contains **Employee_Name**.
- The **orion.employee_organization** table contains **Department**.

Sketch of desired report:

Employee Salaries as a percent of Department Total			
Department	Employee_Name	Salary	Percent
Accounts	Mea, Azavious	58,200.00	8.6%
Accounts	Miller, Pamela	53,475.00	7.9%
Accounts	Asta, Wendy	52,295.00	7.7%

- Create a report aggregating the sum of all salaries for each department. The report should include **Department** and the sum of all associated salary values as **Dept_Salary_Total**. Join **orion.employee_payroll** and **orion.employee_organization** by **Employee_ID** to obtain the information that you need.
 - The **orion.employee_payroll** table contains salary values.
 - The **orion.employee_organization** table contains department information.

Partial PROC SQL Output

Department	Dept_Salary_Total
Accounts	680440
Accounts Management	397175
Administration	1009850
Concession Management	372225
Engineering	276285

- b. Create a report that includes the employee ID, name, and department. Join **orion.employee_addresses** and **orion.employee_organization** by **Employee_ID** to obtain the information that you need.
- The **orion.employee_addresses** table contains **Employee_Name** and **Employee_ID**.
 - The **orion.employee_organization** table contains **Employee_ID** and **Department**.

Partial PROC SQL Output

Employee_ID	Employee_Name	Department
121044	Abbott, Ray	Sales
120145	Aisbitt, Sandy	Sales
120761	Akinfolarin, Tameaka	Accounts
120656	Amos, Salley	Logistics Management
121107	Anger, Rose	Sales
121038	Anstey, David	Sales

- c. Use the two queries that you created in **a** and **b** as in-line views. Join the views with **orion.employee_payroll** by either **Employee_ID** or **Department** to create the final report.
- The query from step **a** contains **Department** and **Dept_Salary_Total**.
 - The query from step **b** contains **Employee_ID**, **Employee_Name**, and **Department**.
 - The **orion.employee_payroll** table contains **Employee_ID** and individual **Salary** values.

Partial PROC SQL Output

Employee Salaries as a percent of Department Total				
Department	Employee_Name	Salary	Percent	
Accounts	Mea, Azavious	58,200.00	8.6%	
Accounts	Miller, Pamela	53,475.00	7.9%	
Accounts	Asta, Wendy	52,295.00	7.7%	
Accounts	Post, Nahliah	48,380.00	7.1%	
Accounts	Ferrari, Ralph	47,000.00	6.9%	
Accounts	Kimmerle, Kevie	46,090.00	6.8%	
Accounts	Farthing, Zashia	43,590.00	6.4%	
Accounts	Knopfmacher, Paul	38,545.00	5.7%	
Accounts	Thoits, Elizabeth	36,440.00	5.4%	
Accounts	Apr, Nishan	36,230.00	5.3%	
Accounts	Atkins, John	34,760.00	5.1%	
Accounts	Voltz, Sal	34,040.00	5.0%	
Accounts	Woods, Connie	32,675.00	4.8%	
Accounts	Akinfolarin, Tameaka	30,960.00	4.5%	
Accounts	Leone, Marvin	30,625.00	4.5%	
Accounts	Van Damme, Jean-Claude	30,590.00	4.5%	
Accounts	Niemann, Kevin	26,545.00	3.9%	
Accounts Management	Kempster, Janelle	53,400.00	13%	
Accounts Management	Kokoszka, Nikeisha	51,950.00	13%	
Accounts Management	Lightbourne, Abelino	47,990.00	12%	

Challenge

7. Building a Complex Query Using a Multi-Way Join

Create a report using a multi-way inner join, which produces the total of the 2011 sales figures for each Orion Star employee. The report should be titled **2011 Total Sales Figures** and must include both the managers' and employees' names (displayed as first name followed by last name), and the total retail value of all sales made by each employee in 2011. Present the information as follows:

- Use one row per employee.
- Organize the report so that the following standards are observed:
 - Employees under one manager are adjacent to each other (grouped together) on the report.
 - Within each manager's group, employees are listed in decreasing order of total sales.
 - The Australian groups are listed first, followed by the U.S. groups.
 - Manager names are in alphabetical order by last name and then first name.

Remember that you can group and order by columns that are not included in the SELECT statement list.

The data that you need can be found in the following tables (variables of interest in parentheses):

- **orion.order_fact (Employee_ID, Total_Retail_Price)**
- **orion.employee_organization (Employee_ID, Manager_ID)**
- **orion.employee_addresses (Employee_ID, Employee_Name)**

Partial PROC SQL Output

2011 Total Sales Figures		
Manager	Employee	Total_Sales
Wilson Dawes	Jina Fiocca	223.80
Wilson Dawes	Phu Sloey	17.60
Wilson Dawes	Amanda Liebman	6.40
Tom Zhou	Lucian Daymond	880.10
Tom Zhou	Ranj Stamalis	697.60
Tom Zhou	Sharryn Clarkson	400.40
Tom Zhou	Marinus Surawski	398.80
Tom Zhou	Sian Shannan	306.20
Tom Zhou	Monica Kletschkus	239.30
Tom Zhou	Fancine Kaiser	147.10
Tom Zhou	Shani Duckett	101.50
Tom Zhou	Atul Leyden	92.50
Tom Zhou	Kevin Lyon	73.99
Tom Zhou	Andrew Conolly	60.80
Tom Zhou	Sean Dives	19.10
Tom Zhou	Lynelle Phoumirath	19.10
Renee Capachietti	Brienne Darrohn	533.40
Renee Capachietti	Michael Westlund	366.00

End of Exercises

5.3 Solutions

Solutions to Exercises

1. Using a Noncorrelated Subquery

a.

```
*** s105s01 ***
/* a */

title;
proc sql;
select avg(Total_Retail_Price) as MeanPrice
  from orion.order_fact
 where Order_Type=1;
quit;
```

b.

```
/* b */

title "Customers Whose Average Retail Price";
title2 "Exceeds the Average Retail Price";
title3 "for All Retail Customers";
proc sql;
select Customer_ID,
      avg(Total_Retail_Price) as MeanPrice format=8.2
    from orion.order_fact
   where Order_Type=1
 group by Customer_ID
 having MeanPrice >
      (select avg(Total_Retail_Price)
        from orion.order_fact
       where Order_Type=1);
quit;
title;
```

2. Using a Noncorrelated Subquery

a.

```
*** s105s02***
/* a */

proc sql;
title "Employee IDs for Current Month Anniversaries";
select Employee_ID
  from orion.employee_payroll
 where month(Employee_Hire_Date)=month(today());
quit;
title;
```

b.

```
/* b */

proc sql;
title "Employees with Current Month Anniversaries";
select Employee_ID,
       scan(Employee_name,2,',') format=$15.
               as FirstName 'First Name',
       scan(Employee_name,1,',') format=$15.
               as LastName 'Last Name'
from orion.employee_addresses
where Employee_ID in
  (select Employee_ID
   from orion.employee_payroll
   where month(Employee_Hire_Date)=month(today()))
order by LastName;
quit;
title;
```

3. Creating Subqueries Using the ALL Keyword

```
*** s105s03***;

proc sql;
title "Level I or II Purchasing Agents";
title2 "Who are older than ALL Purchasing Agent III's";
select Employee_ID, Job_Title, Birth_Date,
       int((02Feb2013'd-Birth_Date)/365.25) as Age
from orion.staff
where Job_Title in ('Purchasing Agent I',
                     'Purchasing Agent II')
      and Birth_Date < all
      (select Birth_Date
       from orion.staff
       where Job_Title='Purchasing Agent III');
quit;
title;
```

Alternate Solution

```
proc sql;
title "Level I or II Purchasing Agents";
title2 "Who are older than ALL Purchasing Agent III's";
select Employee_ID, Job_Title, Birth_Date,
       int((02Feb2013'd-Birth_Date)/365.25) as Age
from orion.staff
where Job_Title in ('Purchasing Agent I',
                     'Purchasing Agent II')
      and Birth_Date <
      (select min(Birth_Date)
       from orion.staff
       where Job_Title='Purchasing Agent III');
quit;
```

```
title;
```

4. Using Nested Subqueries

```
*** s105s04 ***;

title 'Latest Order Date for';
title2 'Orion Club Low Activity Members';
proc sql;
select distinct Customer_ID,
      max(Order_Date) 'Order Date' format=date11.
from orion.order_fact
where Order_Date < '01Jan2012'd and
      Customer_ID in (
            select Customer_ID
            from orion.customer
            where Customer_Type_ID =
                  (select Customer_Type_ID
                  from orion.customer_type
                  where Customer_Type =
                        'Orion Club members low activity'))
group by Customer_ID;
quit;
title;
```

Alternate Solution

```
title 'Inner Join with In-Line View';
proc sql;
select c.Customer_ID,
      Date
from orion.customer_type as t,
      orion.customer as c,
      (select Customer_ID, max(Order_Date) format=date11. as Date
       from orion.order_fact
       where Order_Date < '01JAN2012'd
       group by 1) as of
where t.Customer_Type_ID=c.Customer_Type_ID
      and c.Customer_ID=of.Customer_ID
      and Customer_Type='Orion Club members low activity';
quit;
title;
```

5. Using In-Line Views

a.

```
*** s105s05 ***;

/* Summarizing by employee */
proc sql;
title "2011 Sales Force Sales Statistics";
title2 "For Employees With $200.00 or More In Sales";
select Country, First_Name, Last_Name,
      sum(Total_Retail_Price) as Value_Sold format=comma9.2,
```

```

count(distinct Order_ID) as Orders,
calculated Value_Sold / calculated Orders as Avg_Order
format=comma7.2
from orion.order_fact as of,
orion.sales as s
where of.Employee_ID=s.Employee_ID
and year(Order_Date)=2011
group by Country, First_Name, Last_Name
having Value_Sold ge 200
order by Country, Value_Sold desc, Orders desc;
quit;

```

b.

```

/* Further summarize by country using
the query from part a. as an in-line view */
proc sql;
title "2011 Sales Summary By Country";
select Country,
max(Value_Sold) 'Max Value Sold' format=comma9.2,
max(Orders) 'Max Orders' format=comma7.2,
max(Avg_Order) 'Max Average' format=comma7.2,
min(Avg_Order) 'Min Average' format=comma7.2
/* Begin in-line view */
from (select Country, First_Name, Last_Name,
sum(Total_Retail_Price) as Value_Sold,
count(distinct Order_ID) as Orders,
calculated Value_Sold / calculated Orders
as Avg_Order
from orion.order_fact as of,
orion.sales as s
where of.Employee_ID=s.Employee_ID
and year(Order_Date)=2011
group by Country, First_Name, Last_Name
having Value_Sold ge 200) /* End in-line view */
group by Country
order by Country;
quit;
title;

```

6. Building Complex Queries with In-Line Views

a.

```

*** s105s06 ***;

proc sql;
select Department, sum(Salary) as Dept_Salary_Total
from orion.employee_payroll as pay,
orion.employee_organization as org
where org.Employee_ID=pay.Employee_ID
group by Department;
quit;

```

b.

```
proc sql;
  select adr.Employee_ID, Employee_Name, org.Department
    from orion.employee_addresses as adr,
         orion.employee_organization as org
   where adr.Employee_ID=org.Employee_ID;
quit;
```

c.

```
proc sql;
  title "Employee Salaries as a Percent of Department Total";
  select emp.Department format=$22.,
         emp.Employee_Name format=$28.,
         Salary format=comma9.2,
         Salary/Dept_Salary_Total as Percent
         format=percent6.2
    from orion.employee_payroll as pay,
     /* In-line View: Employee ID, name and department */
     (select adr.Employee_ID, Employee_Name,
            org.Department
       from orion.employee_addresses as adr,
            orion.employee_organization as org
      where adr.Employee_ID=org.Employee_ID)
     as emp,
     /* In-line View: Aggregate sum of salary by department */
     (select Department, sum(Salary) as Dept_Salary_Total
       from orion.employee_payroll as pay,
            orion.employee_organization as org
      where org.Employee_ID=pay.Employee_ID
      group by Department)
     as sum
   where sum.Department=emp.Department and
         pay.Employee_ID=emp.Employee_ID
   order by Department, Percent desc;
quit;
```

7. Building a Complex Query Using a Multi-Way Join

```
*** s105s07 ***;

proc sql;
  title "2011 Total Sales Figures";
  select catx(' ',scan(mgr.Employee_Name,2,',',''),
             scan(mgr.Employee_Name,1,',')) format=$27.
             as Manager,
             catx(' ',scan(emp.Employee_Name,2,',',''),
                   scan(emp.Employee_Name,1,',')) format=$27.
                   as Employee,
                   Sum(Total_Retail_Price) format=comma9.2
                   as Total_Sales
  from orion.order_fact as order,
       orion.employee_organization as org,
```

```
orion.employee_addresses as emp,
orion.employee_addresses as mgr
where order.Employee_ID=org.Employee_ID
and order.Employee_ID=emp.Employee_ID
and mgr.Employee_ID=org.Manager_ID
and year(Order_Date)=2011
and order.Employee_ID ne 99999999
group by mgr.Country, mgr.Employee_Name, emp.Employee_Name
order by mgr.Country, mgr.Employee_Name, Total_Sales desc;
quit;
title;
```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

5.01 Poll – Correct Answer

Is the SELECT statement valid as a subquery?

- Yes
- No

A subquery must return only values in a single column.

14

5.02 Poll – Correct Answer

Can a subquery contain a subquery?

- Yes
- No

```
title 'Latest Order Date for';
title2 'Orion Club Low Activity Members';
proc sql;
select Customer_ID,
       max(Order_Date) 'Order Date' format=date11.
  from orion.order_fact
 where Order_Date < '1Jan2010'd and
       Customer_ID in
        (select Customer_ID
          from orion.customer
         where Customer_Type ID =
              (select Customer_Type_ID
                from orion.customer_type
               where Customer_Type =
                     'Orion Club members low activity'))
   group by Customer_ID;
quit;
title;
```

20

s105d03

5.03 Multiple Choice Poll – Correct Answer

What happens when you change the comparison operator to an equal sign?

- a. Nothing special, the program runs fine.
- b. Multiple rows are returned in your output.
- c. You get an error message.
- d. a and b

ERROR: Subquery evaluated to more than one row.

30

5.04 Poll – Correct Answer

Can this SELECT statement be used as a subquery?

- Yes
- No

A subquery must return values from a single column.

47

5.05 Poll – Correct Answer

The first SELECT statement uses an in-line view to create data that is referenced with the alias **job**. Can **job** also be referenced in the second SELECT statement?

- Yes
- No

An in-line view can be referenced only in the SELECT statement where it is defined.

53

5.06 Multiple Choice Poll – Correct Answer

To join the **employee_organization** table with the Step 1 query results, you use the query from Step 1 as which of the following?

- a. an in-line view
- b. a noncorrelated subquery
- c. either a or b
- d. neither a or b

63

5.07 Poll – Correct Answer

Is it possible to use the entire query in Step 2 as a subquery?

- Yes
- No

A subquery can return values for multiple rows, but must return values for only one column. When submitted on its own, the query in Step 2 returns one row and one column, so it can be used as a noncorrelated subquery.

Chapter 6 Set Operators

6.1	Introduction to Set Operators	6-3
6.2	The UNION Operator.....	6-9
6.3	The OUTER UNION Operator	6-19
	Exercises	6-25
6.4	The EXCEPT Operator	6-27
6.5	The INTERSECT Operator	6-34
	Exercises	6-43
6.6	Solutions	6-46
	Solutions to Exercises	6-46
	Solutions to Student Activities (Polls/Quizzes)	6-52

6.1 Introduction to Set Operators

Objectives

- Describe the business scenarios.
- Explore the example data.
- Describe the SQL set operators.

3

Business Scenario

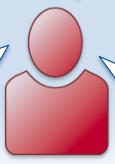
Your manager has requested reports that answer questions, including the following:

Which employees have completed training A or B?

Which employees have completed training A, but not training B?

Which employees have completed training A and/or B, and on what dates?

Which employees have completed both classes?



4

Business Data

The data required to answer the questions is stored in two tables.

Partial train_a

ID	Name	Date
11	Bob	15JUN2012
16	Sam	5JUN2012
14	Pete	21JUN2012

Training class A is completed in a single session. Date represents the date of training.

Partial train_b

Name	ID	SDate	EDate
Bob	11	9JUL2012	13JUL2012
Pam	15	25JUL2012	27JUL2012
Kyle	19	12JUL2012	20JUL2012
Chris	21	29JUL2012	.

Training class B is a multi-session class. SDate is recorded on the first training day. EDate is recorded when the course is complete.

5

Discussion

Can you answer any of the four questions by querying only one table?

- 1 Which employees have completed training A or B?
- 2 Which employees have completed training A and/or B, and on what dates?
- 3 Which employees have completed training A, but not training B?
- 4 Which employees have completed both classes?

Partial train_a

ID	Name	Date
11	Bob	15JUN2012
16	Sam	5JUN2012
14	Pete	21JUN2012

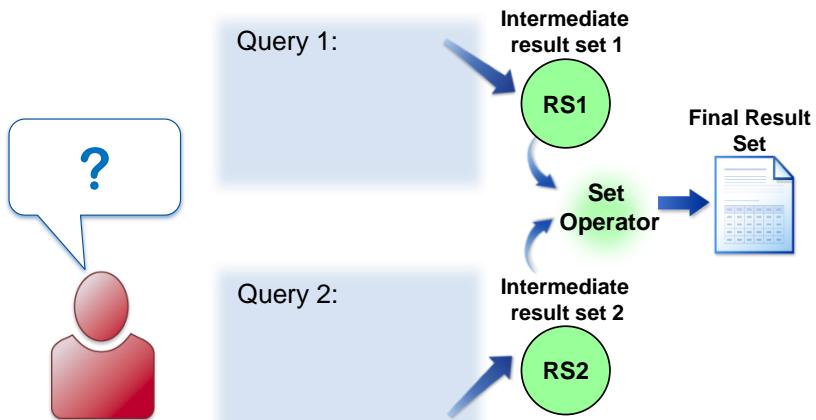
Partial train_b

Name	ID	SDate	EDate
Bob	11	9JUL2012	13JUL2012
Pam	15	25JUL2012	27JUL2012
Kyle	19	12JUL2012	20JUL2012
Chris	21	29JUL2012	.

6

Using Set Operators

Set operators use the intermediate result sets from two queries to create a final result set.



8



The set operator operates on the result sets produced by the table queries (select queries), not on the actual tables themselves.

Question 1: UNION Set Operator

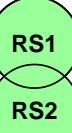
Which employees have completed training A or B?



Query 1:
List employees that have completed train_a.

Query 2:
List employees that have completed train_b.

UNION



Final Result Set

9

Question 2: OUTER UNION Set Operator

Which employees have completed training A and/or B and on what dates?



Query 1:
List employees that have completed **train_a** and the completion date.

OUTER UNION

RS1

Query 2:
List employees that have completed **train_b** and the completion date.

RS2

Final Result Set

10

Question 3: EXCEPT Set Operator

Which employees have completed training A, but not training B?



Query 1:
List employees that have completed **train_a**.

EXCEPT

RS1
RS2

Final Result Set

Query 2:
List employees that have completed **train_b**.

11

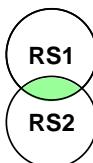
Question 4: INTERSECT Set Operator

Which employees have completed both training A and B?



Query 1:
List employees that have completed train_a.

INTERSECT



Final Result Set

Query 2:
List employees that have completed train_b.

12

Default Behavior of Set Operators

Set Operator	Rows	Columns
UNION	Unique rows from both result sets	Aligned by column position in both result sets
OUTER UNION	All rows from both result sets	All columns from both result sets
EXCEPT	Unique rows from the first result set, that are not in the second result set	Aligned by column position in both result sets
INTERSECT	Unique rows from the first result set that are in the second result set	Aligned by column position in both result sets

13

Using Set Operators

```
SELECT ...
UNION | OUTER UNION | EXCEPT | INTERSECT
<ALL><CORR>
SELECT ...;
```

The modifiers ALL and CORR change the default behavior of the set operators.

ALL modifies the default behavior for rows.

CORR modifies the default behavior for columns.

14

6.01 Poll

By default, the UNION, INTERSECT, and EXCEPT set operators remove duplicate rows from the query output.

- True
- False

15

6.2 The UNION Operator

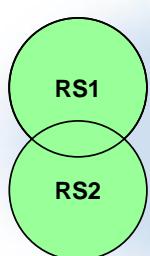
Objectives

- Describe the SQL process when you use the UNION set operator.
- Use the UNION set operator.
- Use the modifiers ALL and CORR.

19

Business Scenario: Question 1

The Education manager needs to know which employees have completed training class A or B.



Which Employees Have Completed Training A or B?	
ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
...	

20

UNION Operator

The manager requested a report that shows employees from both result sets; the UNION operator is appropriate.

```
title 'Which Employees Have Completed';
title2 'Training A or B?';
proc sql;
select ID, Name from work.train_a
union
select ID, Name from work.train_b
  where EDate is not missing;
quit;
```

```
SELECT ...
UNION <ALL><CORR>
SELECT ...;
```

21

s106d01

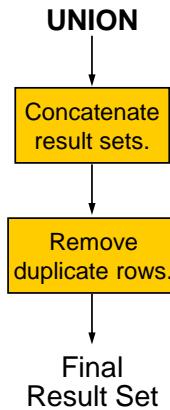
Default Behavior of the UNION Operator

Rows

- Rows from both intermediate result sets are concatenated.
- Duplicate rows are removed from the final result set.

Columns

- Columns are matched by position in the select clauses and must be the same data type.
- Column names in the final result set are determined by the first result set.



22

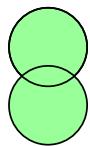


- If the two Intermediate result sets have a different number of columns, then SAS will extend one table with null columns so that the two intermediate result sets will have the same number of columns. If Result Set 1 is extended with null columns, then the name of the column in Result Set 2 will be used in the final results.

Viewing the Intermediate Result Sets

UNION

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

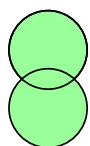
Concatenated Result Set	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand
11	Bob
15	Pam
19	Kyle
87	Ted

23

Removing Duplicates

UNION

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

Concatenated Result Set	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand
11	Bob
15	Pam
19	Kyle
87	Ted

24

Viewing the Output

PROC SQL Output

Which Employees Have Completed Training A or B?	
ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris
87	Ted
91	Rand

25

Issues with Default Behavior

Use the * operator to specify that all columns from **train_a** and **train_b** will be used in the intermediate result tables.

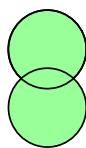
```
title 'Union with Defaults';
proc sql;
select * from work.train_a
union
select * from work.train_b
    where EDate is not missing;
quit;
```

26

s106d02

Viewing the Intermediate Result Sets

Intermediate Result Set 1		
ID	Name	Completion Date
11	Bob	15JUN2012
16	Sam	05JUN2012
14	Pete	21JUN2012
21	Chris	07JUN2012
18	Kim	04JUN2012
17	Pat	22JUN2012
20	Mary	11JUN2012
12	Sue	06JUN2012
87	Ted	05JUN2012
91	Rand	07JUN2012



Intermediate Result Set 2			
Name	ID	Start Date	End Date
Bob	11	09JUL2012	13JUL2012
Pam	15	25JUL2012	27JUL2012
Kyle	19	12JUL2012	20JUL2012
Ted	87	09JUL2012	13JUL2012

Align columns by position in the intermediate result sets. Column 1 from result set 1 is aligned with column 1 from result set 2, and so on. Will this work?

27

Viewing the Log

```
title 'UNION with Defaults';
proc sql;
select * from work.train_a
union
select * from work.train_b
  where EDate is not missing;
WARNING: A table has been extended with null columns to perform
the UNION set operation.
ERROR: Column 1 from the first contributor of UNION is not the
same type as its counterpart from the second.
ERROR: Column 2 from the first contributor of UNION is not the
same type as its counterpart from the second.
367 quit;
```

28

s106d02

Behavior of UNION Operator with Modifiers

Rows with ALL modifier

- Rows from both intermediate result sets are concatenated.
- Duplicate rows are **not** removed from the final result tables.

Columns with CORR modifier

- Columns are matched by name and non-matching columns are removed from the intermediate result sets.

29

Use the ALL keyword when either of the following conditions occurs:

- The presence of duplicates in the final result set will not cause problems.
- Duplicates are not possible. For example, there is a unique or primary key constraint on the column.

Issues with Default Behavior: Corrected

Use the CORR modifier to match columns in the intermediate result sets by name.

```
title 'UNION with CORR Modifier';
proc sql;
select * from work.train_a
union corr
select * from work.train_b
  where EDate is not missing;
quit;
```

Result Set 1

ID
Name
~~Date~~

Result Set 2

Name
ID
~~SDate~~
~~EDate~~

s106d02

30

Viewing the Output

PROC SQL Output

UNION with CORR Modifier	
ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris
87	Ted
91	Rand

31

Setup for the Poll

Open the program **s106a01**, submit the program, and view the report.

Return to the program, add the ALL keyword with the UNION operator, resubmit the program, and view the report. Answer the following question.

Is there any difference in the output with the ALL keyword? Why?

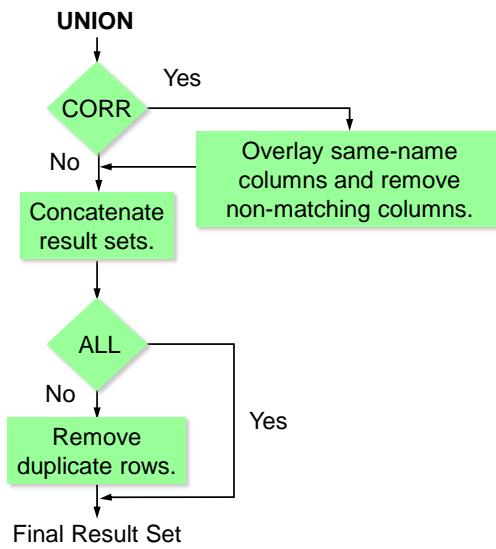
32

6.02 Short Answer Poll

Is there any difference in the output with the ALL keyword? Why?

33

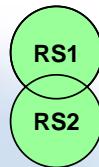
Flow Diagram: UNION Operator



35

Business Scenario

The management team requested a payroll report for Level I, II, and III Orion Star employees.



Payroll Report for Level I, II, and III Employees

Total Paid to ALL Level I Staff	1,234,567
Total Paid to ALL Level II Staff	1,456,789
Total Paid to ALL Level III Staff	2,123,456

36

Business Data

- The **orion.staff** table contains the job title and salary information for all Orion Star employees.
- Submit a separate query for each of the three job title levels.
- Use the UNION set operator twice to combine the intermediate result sets into a single query expression.

Payroll Report for Level I, II, and III Employees

Total Paid to ALL Level I Staff	1,234,567
Total Paid to ALL Level II Staff	1,456,789
Total Paid to ALL Level III Staff	2,123,456

37

UNION Operator

```

title 'Payroll Report for Level I, II,';
title2 'and III Employees';
proc sql;
select 'Total Paid to ALL Level I Staff',
       sum(Salary) format=comma12.
  from orion.staff
 where scan(Job_Title,-1,' ')='I'
union
select 'Total Paid to ALL Level II Staff',
       sum(Salary) format=comma12.
  from orion.staff
 where scan(Job_Title,-1,' ')='II'
union
select 'Total Paid to ALL Level III Staff',
       sum(Salary) format=comma12.
  from orion.staff
 where scan(Job_Title,-1,' ')='III';
quit;

```

38

s106d03

Viewing the Output

PROC SQL Output

Payroll Report for Level I, II, and III Employees	
Total Paid to ALL Level I Staff	3,582,630
Total Paid to ALL Level II Staff	3,569,580
Total Paid to ALL Level III Staff	2,296,425

39

6.03 Multiple Choice

Is it more or less efficient to use the ALL keyword in a set operation?

- more efficient
- less efficient

40

6.3 The OUTER UNION Operator

Objectives

- Describe the SQL process when you use the OUTER UNION set operator and keywords.
- Use the OUTER UNION set operator.
- Compare the SQL OUTER UNION operator to traditional SAS programming.

44

Business Scenario: Question 2

The education manager requested a report that lists which employees have completed training A and/or B, and on what dates?

ID	Name	Completion Date	Start Date	End Date
11	Bob	15JUN2012	:	:
16	Sam	05JUN2012	:	:
...				
12	Sue	06JUN2012	.	.
11	Bob	.	09JUL2012	13JUL2012
15	Pam	.	25JUL2012	27JUL2012
19	Kyle	.	12JUL2012	20JUL2012
87	Ted	.	09JUL2012	13JUL2012

45

OUTER UNION Operator

The report needs to keep all rows and all columns from the two intermediate result sets. The OUTER UNION operator will accomplish this.

```
proc sql;
  select * from train_a
  outer union
  select * from train_b
    where EDate is not missing;
quit;
```

```
SELECT ...
OUTER UNION <CORR>
SELECT ...
```

46

s106d04

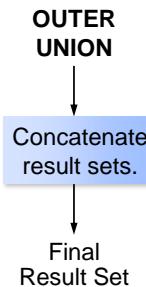
Default Behavior of the OUTER UNION Operator

Rows

- Rows from both intermediate result sets are concatenated.
- Duplicate rows are not removed from the final result set.

Columns

- All columns from both result sets are selected. Common columns are not overlaid.



47

Viewing the Output

PROC SQL Output

ID	Name	Completion Date		ID	Start Date		End Date
		Name	Date		Name	Date	
11	Bob	15JUN2012	
16	Sam	05JUN2012	
14	Pete	21JUN2012	
21	Chris	07JUN2012	
18	Kim	04JUN2012	
17	Pat	22JUN2012	
20	Mary	11JUN2012	
12	Sue	06JUN2012	
87	Ted	05JUN2012	
91	Rand	07JUN2012					
.		.	Bob	11	09JUL2012	13JUL2012	
.		.	Pam	15	25JUL2012	27JUL2012	
.		.	Kyle	19	12JUL2012	20JUL2012	
.		.	Ted	87	09JUL2012	13JUL2012	

48

Behavior of OUTER UNION with Modifiers

Columns with CORR modifier

- All columns from both result sets are selected and common columns are overlaid.

-  The ALL modifier is not applicable with the OUTER UNION operator. This is the default behavior.

49

OUTER UNION Operator

The previous report can be improved by overlaying the common columns. The OUTER UNION operator with the CORR modifier will accomplish this.

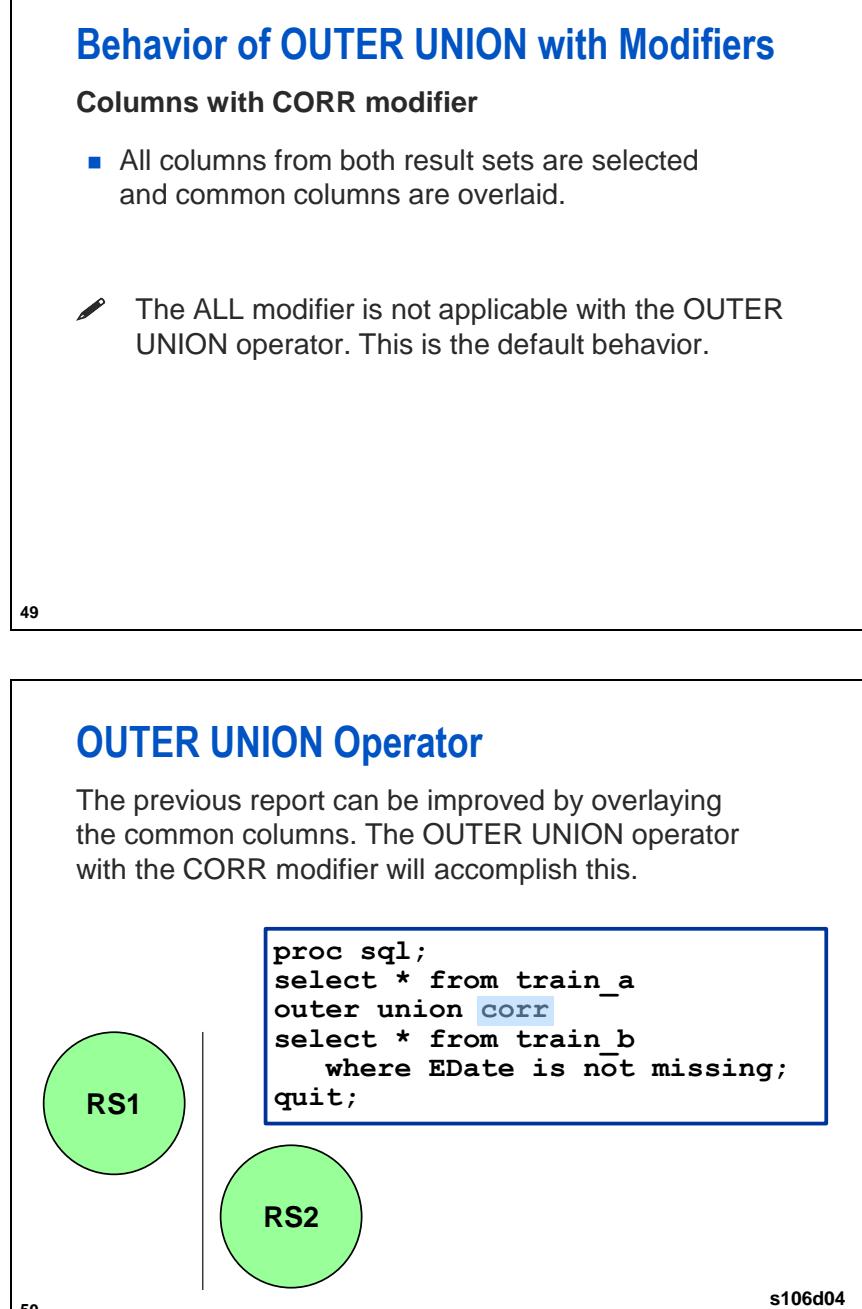
```
proc sql;
  select * from train_a
  outer union corr
  select * from train_b
    where EDate is not missing;
quit;
```

RS1

RS2

50

s106d04

- 
-  OUTER UNION CORR is useful when you have two or more source tables with some columns in common and extra columns that do not exist in all tables. You can use this technique if you want to include at least one of those extra columns in the output.

Viewing the Output

PROC SQL Output

ID	Name	Completion Date	Start Date	End Date
11	Bob	15JUN2012	.	.
16	Sam	05JUN2012	.	.
14	Pete	21JUN2012	.	.
21	Chris	07JUN2012	.	.
18	Kim	04JUN2012	.	.
17	Pat	22JUN2012	.	.
20	Mary	11JUN2012	.	.
12	Sue	06JUN2012	.	.
87	Ted	05JUN2012	.	.
91	Rand	07JUN2012	.	.
11	Bob	.	09JUL2012	13JUL2012
15	Pam	.	25JUL2012	27JUL2012
19	Kyle	.	12JUL2012	20JUL2012
87	Ted	.	09JUL2012	13JUL2012

51

6.04 Multiple Choice Poll

What DATA step statement yields the same results as OUTER UNION CORR?

- a. MERGE
- b. APPEND
- c. SET
- d. STACK

52

SQL versus Traditional SAS Programming

The following programs produce the same report:

```
data trained;
  set train_a train_b;
run;
proc print data=trained label noobs;
run;
```

```
proc sql;
select * from train_a
outer union corr
select * from train_b;
quit;
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Using the UNION Operator

Create a report that displays the total salary for female and male sales representatives and the total number of female and male sales representatives. The **orion.salesstaff** table contains information about all the Orion Star sales representatives, including **Salary** and **Gender**. The query should do the following:

- Create the first row of the report. Use the constant text **Total Paid to All Female Sales Representatives**, **SUM(Salary)**, and the total number of rows using the COUNT(*) function. Summarize data in the **orion.salesstaff** table for those rows that have **Gender** = 'F' and **Job_Title** containing 'Rep'.
- Use the appropriate SET operator.
- Create the second row of the report. Use the constant text **Total Paid to All Male Sales Representatives**, **SUM(Salary)**, and the total number of rows using the COUNT(*) function. Summarize data in the **orion.salesstaff** table for those rows that have **Gender** = 'M' and **Job_Title** containing 'Rep'.
- Provide a title for the report as shown below.

PROC SQL Output

Payroll Report for Sales Representatives		
		Total
Total Paid to All Female Sales Representatives	\$1,898,804	68
Total Paid to All Male Sales Representatives	\$2,655,678	95

2. Using the OUTER UNION Operator with the CORR Keyword

Create a report that displays the sales data for the first and second quarters of 2011. The **orion.qtr1** table contains the sales data for the first quarter, and the **orion.qtr2** table contains the sales data for the second quarter.

Partial PROC SQL Output

First and Second Quarter 2011 Sales						
Order ID	Order Type	Customer ID	Date		Employee ID	
			Order was placed by Customer	Date Delivered		
1241054779	3	24	02JAN2011	05JAN2011	.	
1241063739	1	89	03JAN2011	04JAN2011	.	
.....				
1241731828	1	31	18MAR2011	18MAR2011	.	
1241789227	3	17023	25MAR2011	30MAR2011	.	
1241895594	1	56	05APR2011	09APR2011	121051	
1241909303	0	46966	07APR2011	08APR2011	99999999	
.....				
1242647539	1	45	27JUN2011	27JUN2011	121109	
1242657273	1	90	28JUN2011	28JUN2011	121037	

Level 2**3. Comparing UNION and OUTER UNION Operators**

Stack the **orion.qtr1** and **orion.qtr2** tables to produce a single report.

- a. Use the UNION operator to stack the data in **orion.qtr1** and **orion.qtr2**.
 - b. Use the OUTER UNION operator without the CORR keyword to stack the data in **orion.qtr1** and **orion.qtr2**. Set the LINESIZE system option to 140 to prevent wrapping of the SQL output.
 - c. Were the final results the same? If not, how did they differ? _____
-
- d. Can the UNION operator and OUTER UNION operator yield the same results if you use the CORR keyword in this example? _____

End of Exercises

6.4 The EXCEPT Operator

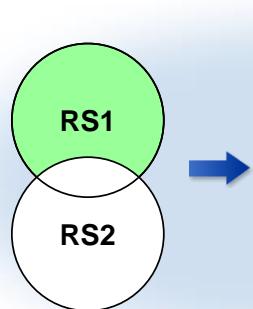
Objectives

- Describe the SQL process when you use the EXCEPT set operator.
- Use the EXCEPT set operator.
- Use the modifiers ALL and CORR.

58

Business Scenario: Question 3

The education manager requested a report that lists which employees have taken training A, but not training B.



Which Employees Have Completed Training A, But Not Training B	
ID	Name
12	Sue
14	Pete
16	Sam
17	Pat
18	Kim
20	Mary
21	Chris
91	Rand

59

EXCEPT Operator

This report lists the employees who have completed training A, but not training B. The EXCEPT operator will accomplish this.

```
title 'Which Employees Have Completed';
title2 'Training A, But Not Training B';
proc sql;
select ID, Name from train_a
except
select ID, Name from train_b
  where Edate is not missing;
quit;
```

SELECT ...
 EXCEPT <ALL><CORR>
 SELECT ...

60

s106d07

Default Behavior of the EXCEPT Operator

Rows

- Duplicate rows are removed from the intermediate result sets.
- Rows from the first intermediate result set that are not in the second intermediate result set are selected.

Columns

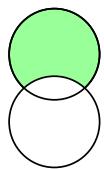
- Columns are matched by position and must be the same data type.
- Column names in the final result set are determined by the first result set.

61

Viewing the Intermediate Result Sets

EXCEPT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



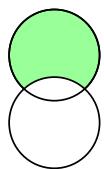
Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

62

Removing Duplicate Rows

EXCEPT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



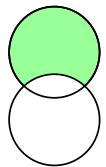
Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

63

Removing Matching Rows

EXCEPT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

64

Viewing the Output

PROC SQL Output

Which Employees Have Completed
Training A, But Not Training B

ID	Name
12	Sue
14	Pete
16	Sam
17	Pat
18	Kim
20	Mary
21	Chris
91	Rand

65

Behavior of EXCEPT Operator with Modifiers

Rows with the ALL modifier

- Duplicate rows are **not** removed from the intermediate result sets.
- Rows from the first intermediate result set that are not in the second intermediate result set are selected.

Columns with the CORR modifier

- Columns are matched by name and non-matching columns are removed from the intermediate result sets.

66

Using a Set Operator in an In-Line View

The previous query can easily become an in-line view used to determine how many employees that have completed training A have not completed training B.

```
proc sql;
select count(ID) as Count
  from (select ID, Name from train_a
        except
        select ID, Name from train_b
          where Edate is not missing);
quit;
```

PROC SQL Output

Count
8

s106d08

67

Setup for the Poll

View the PROC SQL code below and respond to the following True/False statement: Adding the CORR keyword to this code will change the results.

```
proc sql;
  select ID, Name from train_a
except
  select ID, Name from train_b;
quit;
```

Partial train_a

ID	Name	Date
11	Bob	15JUN2012
16	Sam	5JUN2012
14	Pete	21JUN2012

Partial train_b

Name	ID	SDate	EDate
Bob	11	9JUL2012	13JUL2012
Pam	15	25JUL2012	27JUL2012
Kyle	19	12JUL2012	20JUL2012
Chris	21	29JUL2012	.

68

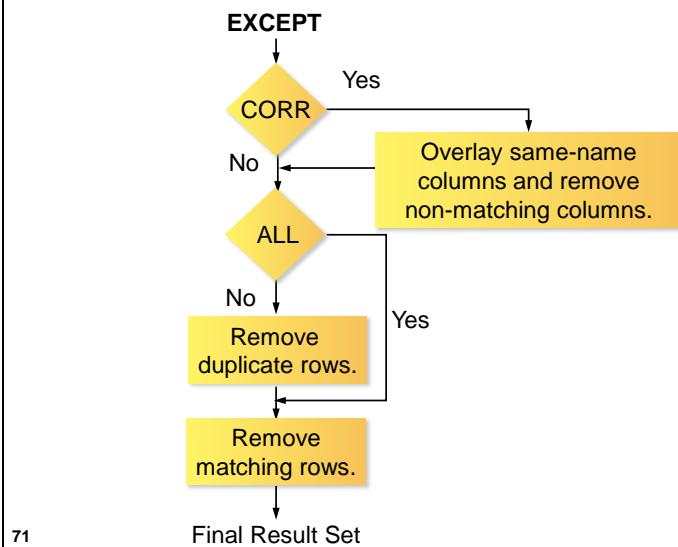
6.05 Poll

Adding the CORR keyword to this code will change the results.

- True
- False

69

Flow Diagram: EXCEPT Operator



71

6.06 Poll

By default, the EXCEPT set operator selects ***all*** of the rows from the first result set that are not in the second result set.

- True
- False

72

6.5 The INTERSECT Operator

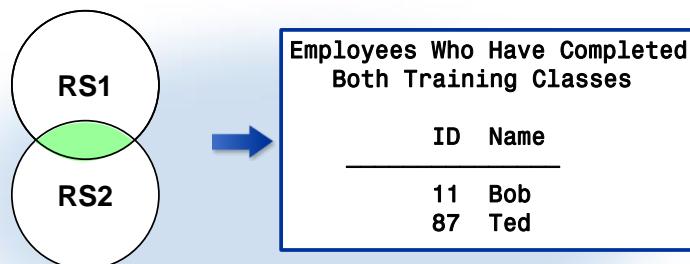
Objectives

- Describe the SQL process when you use the INTERSECT set operator.
- Use the INTERSECT set operator.
- Use the modifiers ALL and CORR.

76

Business Scenario: Question 4

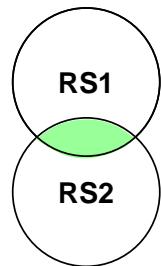
The education manager has requested a report that lists which employees have completed both training classes.



77

INTERSECT Operator

This report requires rows that exist in both **train_a** and **train_b**. The INTERSECT operator will accomplish this.



```
title 'Employees Who Have Completed';
title2 'Both Training Classes';
proc sql;
select ID, Name from train_a
intersect
select ID, Name from train_b
  where EDate is not missing;
quit;
```

SELECT ...
 INTERSECT <ALL><CORR>
 SELECT ...

78

s106d09

Default Behavior of the INTERSECT Operator

Rows

- Duplicate rows are removed from the intermediate result sets.
- Rows from the first intermediate result set that are also in the second intermediate result set are selected.

Columns

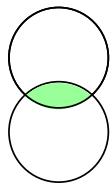
- Columns are matched by position and must be the same data type.
- Column names in the final result set are determined by the first result set.

79

Viewing the Intermediate Result Sets

INTERSECT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



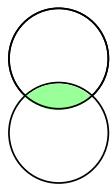
Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

80

Removing Duplicate Rows

INTERSECT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pat
20	Mary
12	Sue
87	Ted
91	Rand



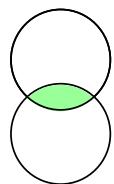
Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

81

Removing Non-Matching Rows

INTERSECT

Intermediate Result Set 1	
ID	Name
11	Bob
16	Sam
14	Pete
21	Chris
18	Kim
17	Pet
20	Mary
12	Sue
87	Ted
91	Rand



Intermediate Result Set 2	
ID	Name
11	Bob
15	Pam
19	Kyle
87	Ted

82

Viewing the Output

PROC SQL Output

Employees Who Have Completed Both Training Classes	
ID	Name
11	Bob
87	Ted

83

Behavior of the INTERSECT Operator with Modifiers

Rows with the ALL modifier

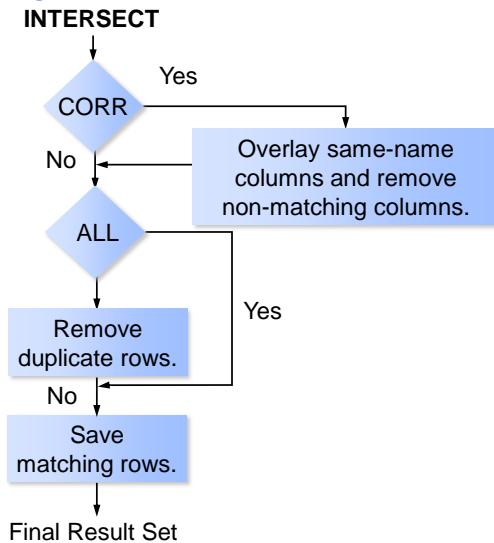
- Duplicate rows are *not* removed from the intermediate result sets.
- Rows from the first intermediate result set that are also in the second intermediate result set are selected.

Columns with the CORR modifier

- Columns are matched by name and non-matching columns are removed from the intermediate result sets.

84

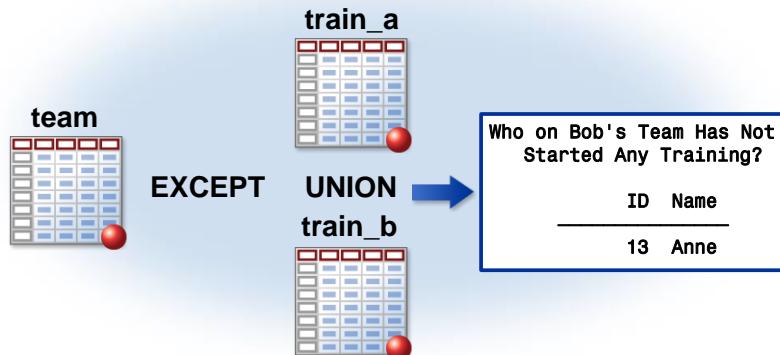
Flow Diagram: INTERSECT Operator



85

Business Scenario

Bob, a team manager, requests your help. He wants to know whether there were any of his team members who had not started either training class. He provided you with a table of his team members.



86

Combining Set Operators

This report needs to return any rows in the **team** table that are not in either **train_a** or **train_b**.

```
title "Who on Bob's Team Has Not";
title2 'Started Any Training';
proc sql;
select ID, Name from team
except
(select ID, Name from train_a
union
select ID, Name from train_b);
quit;
```

87

s106d10

- Set operators follow this order of precedence unless they are overridden by parentheses in the expressions: INTERSECT is evaluated first. OUTER UNION, UNION, and EXCEPT have the same level of precedence.

Viewing the UNION Result Set

The UNION operator will return an intermediate result set of all employees that have started either training class.

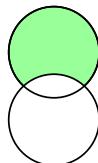
Intermediate Result Set UNION	
ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris
87	Ted
91	Rand

88

Removing the Duplicates

Intermediate Result Bob's Team	
ID	Name
11	Bob
11	Bob
12	Sue
13	Anne
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris

EXCEPT



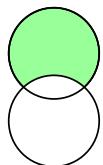
Intermediate Result Set UNION	
ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris
87	Ted
91	Rand

89

Removing the Matching Rows

Intermediate Result Bob's Team	
ID	Name
11	Bob
11	Bob
12	Sue
13	Anne
14	Pete
15	Pam
16	Sam
17	Pet
18	Kim
19	Kyle
20	Mary
21	Chris

EXCEPT



Intermediate Result Set UNION

ID	Name
11	Bob
12	Sue
14	Pete
15	Pam
16	Sam
17	Pat
18	Kim
19	Kyle
20	Mary
21	Chris
87	Ted
91	Rand

90

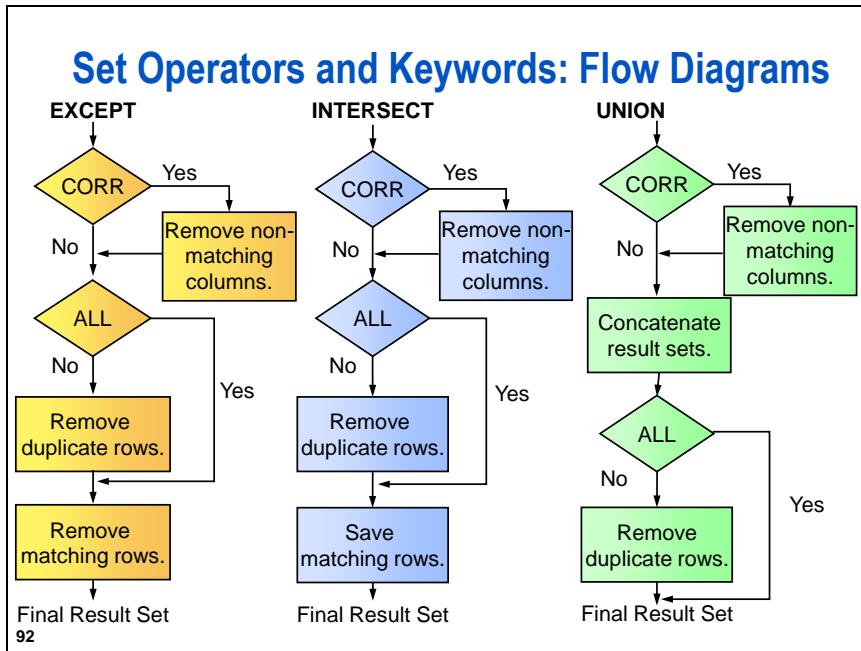
Viewing the Output

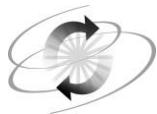
PROC SQL Output

Who on Bob's Team Has Not Started Any Training?

ID	Name
13	Anne

91





Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

4. Using the EXCEPT Operator

Create a report that displays the employee identification numbers of employees who have phone numbers, but do not appear to have address information. The **orion.employee_phones** table contains **Employee_ID** and **Phone_Number**. If an employee's address is on file, the **orion.employee_addresses** table contains the **Employee_ID** value and address information.

The query should

- use the column **Employee_ID** from **orion.employee_phones**
- use the appropriate SET operator
- use the column **Employee_ID** from **orion.employee_addresses**.

PROC SQL Output

Employee IDs with Phone Numbers But Not Address Information
Employee_ID
12099
20683
21149

5. Using the INTERSECT Operator

Create a report that shows **Customer_ID** for all customers who placed orders. The **orion.order_fact** table contains information about the orders that were placed by Orion Star customers, including **Customer_ID**. The **orion.customer** table contains information about all Orion Star customers, including **Customer_ID**. The query should

- use the column **Customer_ID** from **orion.order_fact**
- use the appropriate SET operator
- use the column **Customer_ID** from **orion.customer**.

Partial PROC SQL Output

Customers Who Placed Orders
Customer ID
4
5
9
....

Level 2

6. Using the EXCEPT Operator to Count Rows

Create a report that displays the total count of employees who did not make any charitable donations. The **orion.employee_organization** table contains a record for every employee in the Orion Star organization and includes the employee identification numbers. The **orion.employee_donations** table contains records only for employees who made charitable donations, including the **Employee_ID** value.

PROC SQL Output

No. Employees w/ No Charitable Donations
300

7. Using the INTERSECT Operator to Count Rows

Create a report that shows the total number of customers who placed orders. The **orion.order_fact** table contains information about the orders that were placed by Orion Star customers, including **Customer_ID**. The **orion.customer** table contains information about all Orion Star customers, including **Customer_ID**.

PROC SQL Output

No. Customers w/ Orders
75

Challenge

8. Using the EXCEPT Operator with a Subquery

Create a report that displays the employee identification numbers and names of sales representatives who did not sell any products in 2011. The **orion.sales** table contains the **Employee_ID** values of all sales representatives. The **orion.order_fact** table contains the **Employee_ID** value of the salesperson and other information about all sales that were made. The **orion.employee_addresses** table contains the **Employee_ID** and **Employee_Name** values of all Orion Star employees. Provide a title for the report as indicated in the sample output, and include the row number as part of the report.

Partial PROC SQL Output

Sales Reps Who Made No Sales in 2011		
Row	Employee_ID	Employee_Name
1	121044	Abbott, Ray
2	120145	Aisbitt, Sandy
3	121038	Anstey, David
4	121030	Areu, Jeryl
5	121062	Armant, Debra
6	120144	Barbis, Viney
7	120168	Barcoe, Selina
...		

9. Using the INTERSECT Operator with a Subquery

Create a report that includes **Customer_ID** and **Customer_Name** for all customers who placed orders. The **orion.order_fact** table contains information about the orders that were placed by Orion Star customers, including **Customer_ID**. The **orion.customer** table contains information about all Orion Star customers, including **Customer_ID** and **Customer_Name**.

Partial PROC SQL Output

Name of Customers Who Placed Orders	
Customer ID	Customer Name
4	James Kvarnig
5	Sandrina Stephano
9	Cornelia Krahl
10	Karen Ballinger
11	Elke Wallstab
12	David Black
13	Markus Sepke
16	Ulrich Heyde
17	Jimmie Evans
18	Tonie Asmussen
19	Oliver S. Füßling
20	Michael Dineley
23	Tulio Devereaux
24	Robyn Klem
...	

End of Exercises

6.6 Solutions

Solutions to Exercises

1. Using the UNION Operator

```
*** s106s01 ***;

proc sql;
title 'Payroll Report for Sales Representatives';
select 'Total Paid to All Female Sales Representatives',
      sum(Salary) format=dollar14., count(*) label='Total'
   from orion.salesstaff
  where gender='F' and Job_Title like '%Rep%'
union
select 'Total Paid to All Male Sales Representatives',
      sum(Salary) format=dollar14., count(*) label='Total'
   from orion.salesstaff
  where gender='M' and Job_Title like '%Rep%';
quit;
title;
```

2. Using the OUTER UNION Operator with the CORR Keyword

```
*** s106s02 ***;

proc sql;
title 'First and Second Quarter 2011 Sales ';
select *
   from orion.qtr1
outer union corr
select *
   from orion.qtr2;
quit;
title;
```

3. Comparing UNION and OUTER UNION Operators

- a. The **Customer_ID** column contains some **Employee_ID** values because the columns are overlaid.

```
*** s106s03 ***;

/* a */
proc sql;
title 'Results with UNION operator';
select *
   from orion.qtr1
union
select *
   from orion.qtr2;
quit;
title;
```

Partial PROC SQL Output

Results with UNION operator						
Order ID	Order Type	Customer ID	Date			Date Delivered
			Order was placed by Customer	Date Order was Delivered	Date Order was Delivered	
.	1	120127	20JUN1960	13MAY2011	13MAY2011	.
.	3	99999999	14NOV2151	11MAY2011	14MAY2011	.
1241054779	3	24	02JAN2011	05JAN2011	.	.
1241063739	1	89	03JAN2011	04JAN2011	.	.
1241066216	1	171	04JAN2011	04JAN2011	.	.
...						
1241686210	1	10	13MAR2011	19MAR2011	.	.
1241715610	1	92	16MAR2011	16MAR2011	.	.
1241731828	1	31	18MAR2011	18MAR2011	.	.
1241789227	3	17023	25MAR2011	30MAR2011	.	.
1241895594	1	121051	26FEB1960	05APR2011	09APR2011	.
1241909303	0	99999999	02AUG2088	07APR2011	08APR2011	.
1241930625	3	99999999	28JAN1960	09APR2011	14APR2011	.
1241977403	1	120152	20JUN1960	15APR2011	15APR2011	.

- b. With the OUTER UNION operator, all rows from both tables are selected. This includes both unique and non-unique rows. The columns are not overlaid. Examine the column for **Employee_ID**.

```
/* b */
options ls=140;
proc sql;
title 'Results with OUTER UNION operator';
select *
  from orion.qtr1
outer union
select *
  from orion.qtr2;
quit;
title;
```

Partial PROC SQL Output

Results with OUTER UNION operator											
Order ID	Order Type	Customer ID	Date			Order ID	Order Type	Employee ID	Date		
			Order was placed by	Customer	Date Delivered				Customer ID	Customer	Date Delivered
1241054779	3	24	02JAN2011	05JAN2011
1241063739	1	89	03JAN2011	04JAN2011
1241066216	1	171	04JAN2011	04JAN2011
1241086052	3	53	06JAN2011	09JAN2011
1241147641	1	53	13JAN2011	13JAN2011
1241235281	1	171	23JAN2011	30JAN2011
1241244297	1	111	24JAN2011	24JAN2011
1241263172	3	3959	25JAN2011	26JAN2011
1241286432	3	27	28JAN2011	02FEB2011
1241298131	2	2806	29JAN2011	08FEB2011
1241359997	1	12	05FEB2011	05FEB2011
1241371145	1	171	07FEB2011	07FEB2011
1241390440	1	41	09FEB2011	09FEB2011
1241461856	1	18	16FEB2011	17FEB2011
1241561055	1	171	28FEB2011	28FEB2011
1241623505	3	24	06MAR2011	09MAR2011
1241645664	2	70100	09MAR2011	13MAR2011
1241652707	3	27	09MAR2011	14MAR2011
1241686210	1	10	13MAR2011	19MAR2011
1241715610	1	92	16MAR2011	16MAR2011
1241731828	1	31	18MAR2011	18MAR2011
1241789227	3	17023	25MAR2011	30MAR2011
.	1241895594	1	121051	56	05APR2011	09APR2011	.
.	1241909303	0	99999999	46966	07APR2011	08APR2011	.
.	1241930625	3	99999999	27	09APR2011	14APR2011	.
.	1241977403	1	120152	171	15APR2011	15APR2011	.
.	1242012259	1	121040	10	18APR2011	12APR2011	.
.	1242012269	1	121040	45	18APR2011	18APR2011	.
.	1242035131	1	120132	183	21APR2011	21APR2011	.
.	1242076538	3	99999999	31	25APR2011	29APR2011	.
.	1242130888	1	121086	92	01MAY2011	01MAY2011	.
.	1242140006	4	99999999	5	02MAY2011	07MAY2011	.
.	1242140009	2	99999999	90	02MAY2011	04MAY2011	.
.	1242149082	1	121032	90	03MAY2011	03MAY2011	.
.	1242159212	3	99999999	5	04MAY2011	09MAY2011	.
.	1242161468	3	99999999	2550	04MAY2011	09MAY2011	.
.	1242162201	3	99999999	46966	05MAY2011	06MAY2011	.
.	1242173926	3	99999999	1033	06MAY2011	10MAY2011	.
.	1242185055	1	120136	41	08MAY2011	08MAY2011	.
.	3	99999999	70079	11MAY2011	14MAY2011	.
.	1	120127	171	13MAY2011	13MAY2011	.
.	1242259863	2	99999999	70187	16MAY2011	21MAY2011	.
.	1242265757	1	121105	10	16MAY2011	16MAY2011	.
.	1242449327	3	99999999	27	26JUL2011	26JUL2011	.
.	1242458099	1	121071	10	06JUN2011	06JUN2011	.
.	1242467585	3	99999999	34	07JUN2011	13JUN2011	.
.	1242477751	3	99999999	31	08JUN2011	12JUN2011	.
.	1242493791	1	121056	5	10JUN2011	10JUN2011	.
.	1242502670	1	121067	31	11JUN2011	11JUN2011	.
.	1242515373	3	99999999	17023	13JUN2011	18JUN2011	.
.	1242534503	3	99999999	70165	15JUN2011	22JUN2011	.
.	1242557584	2	99999999	89	17JUN2011	21JUN2011	.
.	1242559569	1	120130	171	18JUN2011	18JUN2011	.

- c. Were the final results the same? If not, how did they differ? **No, the final results are different.**

When the UNION operator is used, the Customer_ID column contains some Employee_ID values because the columns are overlaid. When the OUTER UNION operator is used, all of the rows from both tables are selected. These rows include both unique and non-unique rows. Also, notice that there is a separate column for Employee_ID, because the columns are not overlaid when you use the OUTER UNION operator.

- d. Can the UNION operator and OUTER UNION operator yield the same results if you use the CORR keyword in this example? No, the UNION operator with the CORR keyword cannot yield the same results as the OUTER UNION operator with the CORR keyword for this example. This is because UNION CORR discards columns that do not have matching names whereas OUTER UNION CORR retains all columns.

```
/* d - NO*/
proc sql;
title 'Results with UNION operator and CORR modifier';
select *
   from orion.qtr1
union corr
select *
   from orion.qtr2;
quit;
title;
proc sql;
title 'Results with OUTER UNION operator and CORR modifier';
select *
   from orion.qtr1
outer union corr
select *
   from orion.qtr2;
quit;
title;
```

4. Using the EXCEPT Operator

```
*** s106s04 ***;

proc sql;
title "Employee IDs with Phone Numbers But Not Address Information";
select Employee_ID
   from orion.employee_phones
except
select Employee_ID
   from orion.employee_addresses;
quit;
title;

/* Alternative Solution using the CORR keyword */
proc sql;
title "Employee IDs with Phone Numbers But Not Address Information";
select *
   from orion.employee_phones
except corr
select *
   from orion.employee_addresses;
quit;
title;
```

5. Using the INTERSECT Operator

```
*** s106s05 ***;

proc sql;
title 'Customers Who Placed Orders';
select Customer_ID
   from orion.order_fact
intersect
select Customer_ID
   from orion.customer;
quit;
title;
```

6. Using the EXCEPT Operator to Count Rows

```
*** s106s06 ***;

proc sql;
select count(*) label='No. Employees w/ No Charitable Donations'
   from (select Employee_ID
         from orion.employee_organization
      except
      select Employee_ID
         from orion.employee_donations);
quit;
```

7. Using the INTERSECT Operator to Count Rows

```
*** s106s07 ***;

proc sql;
select count(*) label='No. Customers w/ Orders'
   from (select Customer_ID
         from orion.order_fact
      intersect
      select Customer_ID
         from orion.customer);
quit;
title;
```

8. Using the EXCEPT Operator with a Subquery

```
*** s106s08 ***;

proc sql number;
title 'Sales Reps Who Made No Sales in 2011';
select Employee_ID, Employee_Name
  from orion.employee_addresses
 where Employee_ID in
    (select Employee_ID
      from orion.sales
     where Job_Title like '%Rep%'
 except all
    select Employee_ID
      from orion.order_fact
     where year(Order_Date)=2011);
quit;
title;
```

9. Using the INTERSECT Operator with a Subquery

```
*** s106s09 ***;

proc sql;
title 'Name of Customers Who Placed Orders';
select Customer_ID, Customer_Name
  from orion.customer
 where Customer_ID in
    (select Customer_ID
      from orion.order_fact
 intersect
    select Customer_ID
      from orion.customer);
quit;
title;
```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

6.01 Poll – Correct Answer

By default, the UNION, INTERSECT, and EXCEPT set operators remove duplicate rows from the query output.

- True
- False

16

6.02 Short Answer Poll – Correct Answer

Is there any difference in the output with the ALL keyword? Why?

Yes, Bob and Ted are included twice. This is because they have completed both training classes. Without the ALL keyword, duplicates are removed.

34

6.03 Multiple Choice – Correct Answer

Is it more or less efficient to use the ALL keyword in a set operation?

- more efficient
- less efficient

A second pass of the data is not required to determine whether duplicates exist.

41

6.04 Multiple Choice Poll – Correct Answer

What DATA step statement yields the same results as OUTER UNION CORR?

- a. MERGE
- b. APPEND
- c. SET
- d. STACK

53

6.05 Poll – Correct Answer

Adding the CORR keyword to this code will change the result.

- True
- False

The columns in the two intermediate result sets are determined by the individual queries. They are the same in both intermediate result sets, so the CORR modifier will not change the result.

70

6.06 Poll – Correct Answer

By default, the EXCEPT set operator selects ***all*** of the rows from the first result set that are not in the second result set.

- True
- False

By default, the EXCEPT operator eliminates duplicate rows first. It selects only *unique* rows from the first result set that are not in the second result set.

73

Chapter 7 Creating Tables and Views

7.1	Creating Tables with the SQL Procedure	7-3
	Exercises	7-16
7.2	Creating Views with the SQL Procedure	7-19
	Exercises	7-31
7.3	Solutions	7-34
	Solutions to Exercises	7-34
	Solutions to Student Activities (Polls/Quizzes)	7-39

7.1 Creating Tables with the SQL Procedure

Objectives

- Create a new table and add data using a single query.
- Create a new empty table by copying column structure from an existing table.
- Create a new empty table by defining the column structure.
- Add data into a table.

3

Creating a New Table

There are three ways to create new tables in PROC SQL.

Method 1 Copy columns and rows from existing table(s).



Method 2 Copy columns but no rows from an existing table.



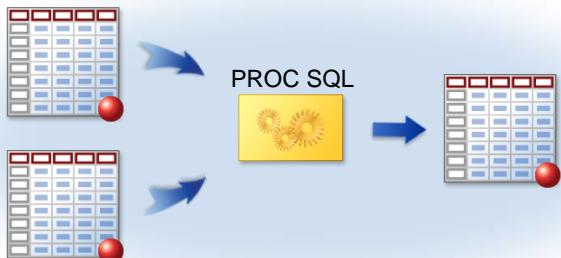
Method 3 Define only the columns in the PROC SQL code.



4

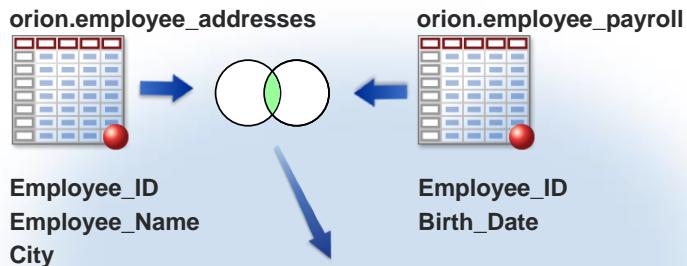
Business Scenario: Method 1

Management wants to recognize employee birthdays.
You need to write code to generate a table with each employee's birth month. Existing tables contain the columns and rows that you need.



5

Business Data



`orion.birthmonths`

Name	City	Birth Month
Aisbitt, Sandy	Melbourne	1
Sheedy, Sherie	Melbourne	1
Tannous, Cos	Melbourne	1

6

Creating a Table: Method 1

For this task, use method 1. Create a table that contains columns and rows returned by a query on existing tables.

```
proc sql;
  create table orion.birthmonths as
    select Employee_Name as Name format=$25.,
           City format=$25.,
           month(Birth_Date) as BirthMonth
             'Birth Month' format=3.
      from orion.employee_payroll as p,
           orion.employee_addresses as a
     where p.Employee_ID=a.Employee_ID
       and Employee_Term_Date is missing
     order by BirthMonth,City,Name;
quit;
```

7

s107d01

This method is most often used to create subsets or supersets of tables. This is the only method that *creates and populates* a table in one statement.

The CREATE TABLE statement does not produce a report. A separate SELECT statement is required to produce a report based on the new table.

Viewing the Log

```
proc sql;
  create table orion.birthmonths as
    select Employee_Name as Name format=$25.,
           City format=$25.,
           month(Birth_Date) as BirthMonth
             'Birth Month' format=3.
      from orion.employee_payroll as p,
           orion.employee_addresses as a
     where p.Employee_ID=
           a.Employee_ID
       and Employee_Term_Date is missing
     order by BirthMonth,City,Name;
NOTE: Table ORION.BIRTHMONTHS created, with 308 rows and 3 columns.
quit;
```

8

Verifying the New Table

```
proc sql;
describe table orion.birthmonths;
select * from orion.birthmonths;
quit;
```

The DESCRIBE statement writes information about the table to the SAS log.

```
proc sql;
describe table orion.birthmonths;
NOTE: SQL table ORION.BIRTHMONTHS was created like:

create table ORION.BIRTHMONTHS( bufsize=8192 )
(
  Name char(40) format=$25.,
  City char(30) format=$25.,
  BirthMonth num format=3. label='Birth Month'
);
```

9

s107d01

Verifying the New Table

```
proc sql;
describe table orion.birthmonths;
select * from orion.birthmonths;
quit;
```

The SELECT statement creates a report that lists the contents of the table.

Partial PROC SQL Output

Name	City	Birth Month
Aisbitt, Sandy	Melbourne	1
Sheedy, Sherie	Melbourne	1
Tannous, Cos	Melbourne	1
Boocks, Michael. R.	Miami-Dade	1
Chinnis, Kumar	Miami-Dade	1

10

s107d01

7.01 Short Answer Poll

What other Base SAS procedures are similar to the PROC SQL DESCRIBE and SELECT statements?

11

Business Scenario: Method 2

Management wants a table for new sales staff that is structured like the **orion.sales** table.



13

Creating a Table: Method 2

Copy the table structure from an existing table with the LIKE clause.

```
CREATE TABLE table-name-2
LIKE table-name-1;

proc sql;
create table work.new_sales_staff
  like orion.sales;
quit;
```

Name the source table.

Partial SAS Log

```
proc sql;
create table work.new_sales_staff
  like orion.sales;
NOTE: Table WORK.NEW_SALES_STAFF created, with 0 rows and 9 columns.
quit;
```

14

s107d02

Use this method when the table that you want to create is structurally identical to an existing table.

Business Scenario: Method 3

You need to create a new table to contain discount information. The structure and data needed are not in an existing table.



15

Creating a Table: Method 3

Define the columns in the CREATE TABLE statement.

```
CREATE TABLE table-name
  (column-name type(length)
   <, ...column-name type(length)>);

proc sql;
  create table discounts
    (Product_ID num format=z12.,
     Start_Date date,
     End_Date date,
     Discount num format=percent.);
quit;
```

The table definition is enclosed in parentheses.

Individual column definitions are separated by commas.

16

s107d03

Use this method when you create a table unlike any other existing table.

Viewing the Log

Partial SAS Log

```
proc sql;
  create table discounts
    (Product_ID num format=z12.,
     Start_Date date,
     End_Date date,
     Discount num format=percent.);
NOTE: Table WORK.DISCOUNTS created, with 0 rows and 4 columns.
quit;
```

17

Defining Columns

For ANSI compliance, PROC SQL accepts the following data types in table definitions:

ANSI Type	Resulting SAS Type	Default Length	Default Format
CHAR(n)	Character	8	\$w.
VARCHAR(n)	Character	8	\$w.
INTEGER	Numeric	8	BEST.
SMALLINT	Numeric	8	BEST.
DECIMAL	Numeric	8	BEST.
NUMERIC	Numeric	8	BEST.
FLOAT	Numeric	8	BEST.
REAL	Numeric	8	BEST.
DOUBLE PRECISION	Numeric	8	BEST.
DATE	Numeric	8	DATE.

18

If you define a column of type INTEGER, SMALLINT, REAL, or DOUBLE PRECISION and you specify a length, do not include the length value in parentheses. For those data types, the syntax is as follows:

```
column-name type n,...
```

Setup for the Poll

Submit the program **s107a01** and review the SAS log.

```
proc sql;
create table discounts
  (Product_ID num format=z12.,
   Start_Date date,
   End_Date date,
   Discount num format=percent.);
describe table discounts;
quit;
```

19

s107a01

7.02 Multiple Choice Poll

Based on the query in program **s107a01**, which different data types does the **work.discounts** table have?

- a. All columns are numeric type.
- b. Some columns are date type and some are numeric type.
- c. Some columns are character type and some are numeric type.
- d. Some columns are character type, some are date type, and some are numeric type.

20

Creating Tables with PROC SQL: Job Aid

Multiple techniques are used to create tables with SQL.

Method	Description	Syntax
1	Create a table and add data all in one step, using a query.	CREATE TABLE <i>table-name</i> AS SELECT... ;
2	Create an empty table by copying column attributes from an existing table using a LIKE clause.	CREATE TABLE <i>table-name</i> LIKE <i>old-table-name</i> ;
3	Create an empty table by manually specifying all column attributes.	CREATE TABLE <i>table-name</i> (<i>column-name type(length)</i> <, ... <i>column-name type(length)</i> >);

22

Loading the Data

Now that you have created the empty **discounts** table, it is time to add data to it.

You need to choose the appropriate technique for adding data into this table.



24

Adding Data to a Table

The **INSERT** statement can be used to add data to an empty table, or to append data to a table that already contains data, using one of three methods.

Method	Description	Syntax
A	One clause per row using column-value pairs	INSERT INTO <i>table-name</i> SET <i>column-name=value, column-name=value,...;</i>
B	One clause per row using positional values	INSERT INTO <i>table-name</i> <i><(column list)></i> VALUES <i>(value,value,...);</i>
C	A query returning multiple rows based on positional values	INSERT INTO <i>table-name</i> <i><(column list)></i> SELECT <i>columns</i> FROM <i>table-name;</i>

25

The order of the values in the query expression must match the order of the column names in the insert column list or, if no list was specified, the order of the columns in the table.

Method A: Adding Data with a SET Clause

The SET clause requires that you add data using column name-value pairs:

```
INSERT INTO table-name
SET column-name=value,
column-name=value,...;
```

```
insert into discounts
  set Product_ID=230100300006,
      Start_Date='01MAR2013'd,
      End_Date='15MAR2013'd,Discount=.33
  set Product_ID=230100600018,
      Start_Date='16MAR2013'd,
      End_Date='31MAR2013'd, Discount=.15;
```

26

s107d04

Method B: Adding Data with a VALUES Clause

The VALUES clause adds data to the columns in a single row of data.

```
INSERT INTO table-name
<(column list)>
VALUES (value,value,...);
```

```
insert into discounts
  values (230100300006,'01MAR2013'd,
          '15MAR2013'd,.33)
  values (230100600018,'16MAR2013'd,
          '31MAR2013'd,.15);
```



The VALUES clause must list values in the same order in which the columns are defined in the table.

27

s107d05

Method B: Adding Data with a VALUES Clause

Optionally, the INSERT statement can list the columns into which data is to be inserted, in the order in which the VALUES clause will provide the data.

```
INSERT INTO table-name
  <(column list)>
VALUES (value,value,...);
```

```
insert into discounts
(Start_Date,End_Date, Product_ID, Discount)
values ('01MAR2013'd,'15MAR2013'd,
        230100300006,.33)
values ('16MAR2013'd,'31MAR2013'd,
        230100600018,.15);
```

The order of the columns in the column list is independent of the order of the columns in the table.

28

s107d05

Method C: Adding Data with a Query

Rows returned by the query are inserted into the table. If the table has existing rows, the new rows are appended.

```
INSERT INTO table-name
  <(column list)>
SELECT columns
  FROM table-name;
```

```
proc sql;
insert into discounts
(Product_ID,Discount,Start_Date,End_Date)
select distinct Product_ID,.35,
        '01MAR2013'd,'31MAR2013'd
  from orion.Product_Dim
 where Supplier_Name contains
        'Pro Sportswear Inc';
quit;
```

⚠ The query must list values in the same order as the INSERT statement column list.

29

s107d06

7.03 Quiz

Locate three syntax errors in the following code:

```
proc sql;
create table discounts
  (Product_ID num format=z12.,
   Start_Date date, End_Date date,
   Discount num format=percent.)
insert into discounts
  (Product_ID,Start_Date,End_Date,Discount)
  values (220200200022,'01Mar2013'd,
          '31Mar2013'd,.35)
  values (220200200024,'01Mar2013'd,
          '31Mar2013'd, '.35');
quit;
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Creating a Table and Adding Data Using a Query

- a. Create a table containing the following columns and name it **orion.employees**:
- **Employee_ID**
 - **Hire_Date**
 - **Salary**
 - **Birth_Date**
 - **Gender**
 - **Country**
 - **City**

Format all date columns with **MMDDYY10.** and **Salary** with **COMMA12.2**. Include only current employees (rows where **Employee_Term_Date** is missing). Order the output by the year value of **Hire_Date** (a variable that is not included in the output table) and then by descending **Salary**. The data that you need can be obtained from these tables:

- **orion.employee_addresses**
- **orion.employee_payroll**

Column Sourcing Information

Employees	Employee_Addresses	Employee_Payroll
Employee_ID	Employee_ID	Employee_ID
Hire_Date	-	Employee_Hire_Date
Salary	-	Salary
Birth_Date	-	Birth_Date
Gender	-	Employee_Gender
Country	Country	-
City	City	-

- b. Query the new table **orion.employees** to produce the following report:

Partial PROC SQL Output

Employee_	ID	Hire_Date	Salary	Birth_Date	Gender	Country	City
-----------	----	-----------	--------	------------	--------	---------	------

121141	01/01/1978	194,885.00	06/19/1948	M	US	Philadelphia
120659	01/01/1978	161,290.00	07/16/1953	M	US	Philadelphia
120103	01/01/1978	87,975.00	01/22/1953	M	AU	Sydney
120712	01/01/1978	63,640.00	06/12/1953	F	US	Miami-Dade
120804	01/01/1978	55,400.00	02/11/1948	M	US	Miami-Dade

Level 2

2. Creating a Table by Defining Its Structure and Adding Data

- Create a table named **orion.rewards** with four columns:
 - Purchased** – a numeric column, formatted to be nine characters wide, including commas and two decimal places
 - Year** – a numeric column, formatted to be four characters wide with no decimal places
 - Level** – a character column, which holds the value *Silver*, *Gold*, or *Platinum*
 - Award** – a character column (maximum of 50 characters)
- Add six rows of data to the **orion.rewards** table, and then write a query to display the contents of the **Rewards** table. The results should resemble the following:

PROC SQL Output

Purchased	Year	Level	Award
200.00	2012	Silver	25% Discount on one item over \$25
300.00	2012	Gold	15% Discount on one order over \$50
500.00	2012	Platinum	10% Discount on all 2012 purchases
225.00	2013	Silver	25% Discount on one item over \$50
350.00	2013	Gold	15% Discount on one order over \$100
600.00	2013	Platinum	10% Discount on all 2013 purchases

Challenge

3. Creating a Table and Inserting Data Using a Complex Query

Create a table named **direct_compensation** in the **work** library. The table should contain the following information for all *non-managerial* Sales staff (those with a level listed in their titles), formatted as indicated:

Employee_ID	Name	Level	Salary	Commission	Direct_Compensation
12345	First Last	(I, II, III, or IV)	12,345.00	1,234.00	13,579.00

- Sales information is available in **orion.order_fact**.
- The table **orion.sales** contains **Employee_ID**, **First_Name**, **Last_Name**, **Job_Title**, and **Salary** information for all Sales staff.
- Job_Title** contains level information for each employee.
- To calculate **Commission**, add the **Total_Retail_Price** values for all sales made by an employee with an order date in 2011, and take 15% of that total value as commission.

- To calculate **Direct_Compensation**, add **Commission** plus **Salary**.

Partial PROC SQL Output

Employee_ID	Name	Level	Salary	Commission	Direct_Compensation
121029	Kuo-Chung Mcelwee	I	27,225.00	3.53	27,228.53
121135	Tammy Ruta	I	27,010.00	24.12	27,034.12
120131	Marinus Surawski	I	26,910.00	59.82	26,969.82
120130	Kevin Lyon	I	26,955.00	11.10	26,966.10
121086	John-Michael Plybon	I	26,820.00	24.08	26,844.08
120136	Atul Leyden	I	26,605.00	13.88	26,618.88
120124	Lucian Daymond	I	26,480.00	132.02	26,612.02
121028	William Smades	I	26,585.00	10.41	26,595.41
120152	Sean Dives	I	26,515.00	2.87	26,517.87
121109	Harold Boulus	I	26,035.00	18.25	26,053.25
121051	Glorina Myers	I	26,025.00	18.76	26,043.77
121106	James Hilburger	I	25,880.00	2.54	25,882.54

End of Exercises

7.2 Creating Views with the SQL Procedure

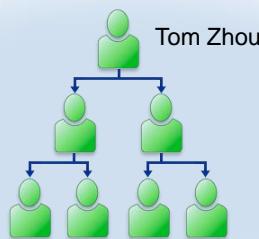
Objectives

- Create a PROC SQL view.
- Use PROC SQL views in SQL queries.
- Use PROC SQL views in other SAS procedures.
- Make a PROC SQL view portable.

37

Business Scenario

Tom Zhou is a sales manager who needs access to personnel information for his staff.



38

Business Data

The data that Tom needs is name, job title, salary, and years of service. This data is contained in three tables.



orion.employee_addresses



orion.employee_payroll



orion.employee_organization

39

Considerations

What is the best way to help Tom, given the following requirements:

- He should not be allowed access to personnel data for any employee that is not his direct report.
- He can write simple PROC SQL queries and use basic SAS procedures, but cannot write complex joins.

A PROC SQL view accessing data for Tom Zhou's direct reports can provide the information that Tom needs in a secure manner.



40

What Is a PROC SQL View?

A *PROC SQL view*

- is a stored query
- contains no actual data
- can be derived from one or more tables, PROC SQL views, DATA step views, or SAS/ACCESS views
- extracts underlying data each time it is used and accesses the most current data
- can be referenced in SAS programs in the same way as a data table
- cannot have the same name as a data table stored in the same SAS library.



41

PROC SQL views are sometimes referred to as *virtual tables* because they are referenced in SAS programs in the same manner as actual data tables, but they are not data tables. They contain no actual data. Instead, they store the instructions required to retrieve and present the data to which they refer.

Creating a PROC SQL View

To create a PROC SQL view, use the CREATE VIEW statement.

```
CREATE VIEW view-name AS
SELECT ...;
```

```
proc sql;
create view orion.tom_zhou as
  select Employee_Name as Name format=$25.0,
         Job_Title as Title format=$15.0,
         Salary 'Annual Salary' format=comma10.2,
         int((today()-Employee_Hire_Date)/365.25)
              as YOS 'Years of Service'
  from employee_addresses as a,
       employee_payroll as p,
       employee_organization as o
 where a.Employee_ID=p.Employee_ID and
       o.Employee_ID=p.Employee_ID and
       Manager_ID=120102;
quit;
```

s107d07

42

View the Log

Partial SAS Log

```

16 proc sql;
17 create view orion.tom_zhou as
18   select Employee_Name as Name format=$25.0,
19         Job_Title as Title format=$15.0,
20         Salary 'Annual Salary' format=comma10.2,
21         int((today()-Employee_Hire_Date)/365.25)
22           as YOS 'Years of Service'
23   from employee_addresses as a,
24       employee_payroll as p,
25       employee_organization as o
26  where a.Employee_ID=p.Employee_ID and
27        o.Employee_ID=p.Employee_ID and
28        Manager_ID=120102;
NOTE: SQL view ORION.TOM_ZHOU has been defined.

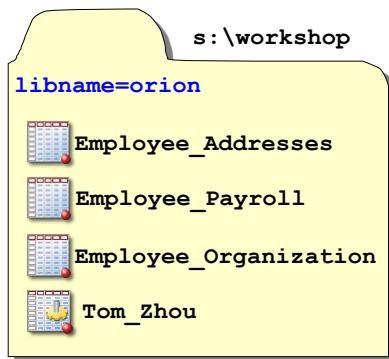
```

43

Within any SAS library, a view cannot have the same name as an existing table, nor can a table be created with the same name as an existing view.

Location of a PROC SQL View

ANSI standards specify that the view must reside in the same SAS library as the contributing table or tables.



44

Location of the Source Tables: ANSI

In PROC SQL, the default libref for the table (or tables) in the FROM clause is the libref of the library that contains the view. When the view and data source are in the same location, you specify a one-level name for the table (or tables) in the FROM clause.

```
create view orion.tom_zhou as
...
from employee_addresses as a,
      employee_payroll as p,
      employee_organization as o
```

45

Using the one-level name in this instance does not indicate the tables are in the SAS **work** library. As this is PROC SQL, this indicates that the tables will be in the same physical location as the view.

A PROC SQL view can be used when the physical location of the view and the underlying data are not collocated. The code in the view will use the traditional two-level names for the tables, and an additional clause, the USING clause, will be added.

7.04 Short Answer Poll

Open and submit **s107a02**. View the results in the log.

Why was the SQL view not created?

46

Using a View

Tom can use the view to produce simple reports.

```
title "Tom Zhou's Direct Reports";
title2 "By Title and Years of Service";
select *
  from orion.tom_zhou
 order by Title desc, YOS desc;
```

Partial PROC SQL Output (executed 05SEP2012)

Tom Zhou's Direct Reports By Title and Years of Service			
Name	Title	Annual Salary	Years of Service
Nowd, Fadi	Sales Rep. IV	30,660.00	34
Hofmeister, Fong	Sales Rep. IV	32,040.00	29
Phoumirath, Lynelle	Sales Rep. IV	30,765.00	22
Platts, Alexei	Sales Rep. IV	32,490.00	10

48

s107d08



Years of Service will change because the TODAY function was used when creating the view.

Using a View

Tom can also use the view to produce simple descriptive statistics to help him better manage his group.

```
title "Tom Zhou's Group - Salary Statistics";
proc means data=orion.tom_zhou min mean max;
  var salary;
  class title;
run;
```

PROC MEANS Output

Tom Zhou's Group - Salary Statistics				
The MEANS Procedure				
Analysis Variable : Salary Annual Salary				
Title	N Obs	Minimum	Mean	Maximum
Sales Rep. I	18	25185.00	26466.67	27260.00
Sales Rep. II	13	26165.00	27123.46	28100.00
Sales Rep. III	12	28135.00	29895.42	36605.00
Sales Rep. IV	5	30660.00	31369.00	32490.00

49

s107d08

Business Scenario

You created a PROC SQL view to provide Tom Zhou access to personnel data for his direct reports.

Tom copied his view to a folder on his hard drive.

Now Tom reports that the view does not work anymore, and he asked for your help to resolve the problem.



50

Exploring the Problem

Tom submitted the following:

```
libname orion 'c:\temp';
proc sql;
title "Tom Zhou's Direct Reports";
title2 "By Title and Years of Service";
select *
from orion.tom_zhou
order by Title_desc, YOS_desc;
quit;
title;
```

51

s107d09

Viewing the Log

Partial SAS Log

```

libname orion 'c:\temp';
NOTE: Libref ORION was successfully assigned as follows:
      Engine:      V9
      Physical Name: c:\temp
proc sql;
title "Tom Zhou's Direct Reports";
title2 "By Title and Years of Service";
select *
      from orion.tom_zhou
      order by Title desc, YOS desc;
ERROR: File ORION.EMPLOYEE_ADDRESSES.DATA does not exist.
ERROR: File ORION.EMPLOYEE_PAYROLL.DATA does not exist.
ERROR: File ORION.EMPLOYEE_ORGANIZATION.DATA does not
exist.
quit;
title;
NOTE: The SAS System stopped processing this step because
of errors.

```

52

Considerations

Tom moved his view to his C:\TEMP folder and redefined the **orion** library there. This violated the one-level naming convention in the FROM clause in the view code.

```

libname orion 'c:\temp';
proc sql;
title "Tom Zhou's Direct Reports";
title2 "By Title and Years of Service";
select *
      from orion.tom_zhou
      order by Title desc, YOS desc;
quit;

```

53

Making a View Portable

```
CREATE VIEW view AS SELECT...
<USING LIBNAME-clause<, ...LIBNAME-clause>>;
```

```
create view orion.Tom_Zhou as
  select Employee_Name as Name format=$25.0,
         Job_Title as Title format=$15.0,
         Salary "Annual Salary" format=comma10.2,
         int((today()-Employee_Hire_Date)/365.25)
              as YOS 'Years of Service'
    from orion.employee_addresses as a,
         orion.employee_payroll as p,
         orion.employee_organization as o
   where a.Employee_ID=p.Employee_ID
     and o.Employee_ID=p.Employee_ID
     and Manager_ID=120102
  using libname orion "s:\workshop";
```

54

s107d10

two-level data set names A USING clause names the location of the tables.

If tables referenced in the FROM clause reside in different libraries, multiple LIBNAME statements can be specified in the USING clause, separated by commas. For example:

```
create view xyz as
  select <...>
  from one.a, two.b
  where <...>
  using libname one 'c:\...',
        libname two 'd:\...';
```

Validating the View

```
describe view orion.tom_zhou;
```

Partial SAS Log

```
describe view orion.tom_zhou;
NOTE: SQL view ORION.TOM_ZHOU is defined as:

      select Employee_Name as Name format=$25.0,
             Job_Title as Title format=$15.0,
             Salary label='Annual Salary' format=COMMA10.2,
             INT((TODAY() - Employee_Hire_Date) / 365.25)
                  as YOS label='Years of Service'
        from orion.employee_addresses a,
             orion.employee_payroll p,
             orion.employee_organization o
       where a.Employee_ID = p.Employee_ID and
             o.Employee_ID = p.Employee_ID and
             Manager_ID= 120102
  using libname orion "s:\workshop";
```

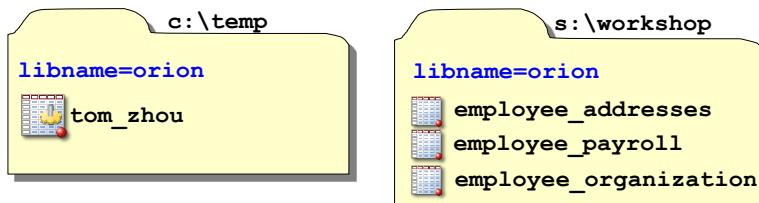
55

s107d10

Two-Level Table Names in Permanent Views

```
CREATE VIEW proc-sql-view AS SELECT ...
<USING LIBNAME-clause<, ...LIBNAME-clause>>;
```

- The USING clause libref is local to the view, and it will not conflict with an identically named libref in the SAS session.
- When the query finishes, the libref is disassociated.



56

7.05 Quiz

What differences are there between the SQL code written to the SAS log by the DESCRIBE VIEW statement and the CREATE VIEW code, which actually created the view **orion.tom_zhou**?

Original

```
select Employee_Name as Name format=$25.0,
       Job_Title as Title format=$15.0,
       Salary 'Annual Salary' format=comma10.2,
       int((today()-Employee_Hire_Date)/365.25)
       as YOS 'Years of Service'
  from Employee_Addresses as a,Employee_Payroll as p,
       Employee_Organization as o
 where a.Employee_ID=p.Employee_ID and
       o.Employee_ID=p.Employee_ID and
       Manager_ID=120102;
```

DESCRIBE VIEW

```
select Employee_Name as Name format=$25.0,
       Job_Title as Title format=$15.0,
       Salary label='Annual Salary' format=COMMA10.2,
       int((today()-Employee_Hire_Date)/365.25)
       as YOS label='Years of Service'
  from EMPLOYEE_ADDRESSES a, EMPLOYEE_PAYROLL p,
       EMPLOYEE_ORGANIZATION o
 where (a.Employee_ID=p.Employee_ID) and
       (o.Employee_ID=p.Employee_ID) and
       (Manager_ID=120102);
```

57

7.06 Poll

Considering the differences discussed in the previous quiz, if you submit the code produced by the DESCRIBE VIEW statement, would the view produced be identical to the original view?

- Yes
- No

61

Views: Advantages

You can use views to do the following:

- avoid storing copies of large tables.
- avoid a frequent refresh of table copies. When the underlying data changes, a view surfaces the most current data.
- pull together data from multiple database tables and multiple libraries or databases.
- simplify complex queries.
- prevent other users from inadvertently altering the query code.

63

Views: Disadvantages

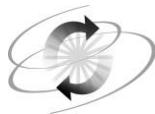
- Because views access the most current data in changing tables, the results might be different each time you access the view.
- Views can require significant resources each time they execute. With a view, you save disk storage space at the cost of extra CPU and memory usage.
- When accessing the same data several times in a program, use a table instead of a view. This ensures consistent results from one step to the next and can significantly reduce the resources that are required.

64

General Guidelines for Creating Views

- Avoid ORDER BY clauses in view definitions, which force data sorting each time that the view is referenced.
- When you create a permanent view with permanent tables in the FROM clause, use a USING clause to specify the location of the libraries to make your view portable.

65



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

4. Creating and Using a View

- Create a view named **orion.phone_list** containing the following columns:
 - **Department** (format=\$25.)
 - **Name** (format=\$25.)
 - **Phone_Number** (label='Home Phone' and format=\$16.)

Use data found in the following tables. (Columns of interest are in parentheses.)

- **orion.employee_addresses** (**Employee_ID**, **Employee_Name**)
- **orion.employee_organization** (**Employee_ID**, **Department**)
- **orion.employee_phones** (**Employee_ID**, **Phone_Number**, **Phone_Type**)

Include only those phone number records where **Phone_Type**="Home".

- Use the new **orion.phone_list** view as the source for a query that produces the phone list for the Engineering Department, sorted by **Name**. Add this title to the report: **Engineering Department Home Phone Numbers**.

PROC SQL Output

Engineering Department Home Phone Numbers		
Name		Home Phone
Arizmendi, Gilbert		+1(619)551-0293
Elleman, Lal		+61(2)5555-4127
Elmoslamy, Wilson		+1(619)551-0291
Hargrave, Seco		+1(215)551-0289
Hartshorn, Darshi		+61(2)5555-2265
Liguori, Donelle		+1(305)551-0290
Mccleary, Bill		+61(3)5555-9767
Peiris, Krishna		+61(2)5555-5585
Sullivan, Lutezenia		+1(305)551-0292

Level 2

5. Creating and Using a View to Provide Consolidated Information

- a. Create a view named **orion.t_shirts** containing the following columns:

- **Product_ID**
- **Supplier_Name** (formatted as \$20.)
- **Product_Name**
- **Price** (label it **Retail Price.**)

Use data from the following tables. (Columns of interest are in parentheses.)

- **orion.product_dim (Product_ID, Product_Name, Supplier_Name)**
- **orion.price_list (Product_ID, Unit_Sales_Price)**

Select only those records where **Product_Name** includes the word **T-Shirt**.

- b. Write a query to display the data from your new **orion.t_shirts** view. Sort the report by **Supplier_Name** and **Product_ID**. Supply a useful title.

Partial PROC SQL Output

Available T-Shirts			
Product ID	Supplier Name	Product Name	Retail Price
210201000050	3Top Sports	Kid Children's T-Shirt	\$19.60
240500200016	3Top Sports	T-Shirt	\$31.70
220101300001	A Team Sports	T-Shirt, Short-sleeved, Big Logo	\$33.60
220101300012	A Team Sports	Men's T-Shirt Small Logo	\$19.40
220101300017	A Team Sports	Toncot Beefy-T Emb T-Shirt	\$16.60

- c. Write a query using your new **orion.t_shirts** view to display **Product_ID**, **Product_Name**, and **Price** for all T-shirts priced less than \$20.00. Sort the report by price. Supply a useful title.

PROC SQL Output

T-Shirts under \$20		
Product ID	Product Name	Retail Price
220101300017	Toncot Beefy-T Emb T-Shirt	\$16.60
220101300012	Men's T-Shirt Small Logo	\$19.40
210201000050	Kid Children's T-Shirt	\$19.60

Challenge

6. Creating and Using a View That Updates Itself over Time

- a. Create a view named **orion.current_catalog**. The view must contain all of the columns in **orion.product_dim** and a new column named **Price** (labeled **Current Retail Price**). Use data found in the following tables:
- **orion.product_dim**
 - **orion.price_list**

The value of the current price (**Price**) is determined by multiplying the original value of **Unit_Sales_Price** by the number of years since the product was first offered and then by the inflation factor (**Factor**). Calculate the number of years that a product was offered by subtracting the year in which the product was first offered from the current year. **Factor** is a percentage of the original price. The formula is shown below:

$$\text{Price} = \text{Unit_Sales_Price} * \\ (\text{Factor}^{**}(\text{year(Today())}-\text{year(Start_Date)}))$$

Round the calculated price to the nearest cent.

- b. Write a query to display **Supplier_Name**, **Product_Name**, and **Price** from your new **orion.current_catalog** view for products with **Roller Skate** in the product name. Sort the output by supplier name and price. Supply a useful title for your report. If the current year is later than 2012, your results will differ because prices have increased.

Partial PROC SQL Output

Current Roller Skate Prices		
Supplier Name	Product Name	Current Retail Price
Magnifico Sports	Children's Roller Skates	\$102.40
Magnifico Sports	Pro-roll Lazer Roller Skates	\$113.20
Magnifico Sports	Pro-roll Sabotage-Rp Roller Skates	\$175.90

- c. Write a query to display **Product_Name**, the original value of **Unit_Sales_Price** (formerly **Price**), the current value of **Price**, and the amount of increase (**Increase** calculated as **Price - Unit_Sales_Price**) for all products having an increase greater than \$5.00. Sort the report by decreasing **Increase**. Supply a useful title. If the current year is later than 2012, your results will differ because prices have increased.

Partial PROC SQL Output

2012 prices > \$5.00 higher than original price			
Product Name	Old Price	New Price	Increase
Twain Ac7/Ft7 Men's Roller Skates	\$181.30	\$186.74	\$5.44
Top R&D Long Jacket	\$536.30	\$541.66	\$5.36

End of Exercises

7.3 Solutions

Solutions to Exercises

1. Creating a Table and Adding Data Using a Query

a.

```
*** s107s01 ***;

/* a. */
proc sql;
create table orion.employees as
  select a.Employee_ID,
         Employee_Hire_Date as Hire_Date
    format=mmddyy10.,
         Salary format=comma12.2,
         Birth_Date format=mmddyy10.,
         Employee_Gender as Gender,
         Country, City
    from orion.employee_addresses as a,
         orion.employee_payroll as p
   where a.Employee_ID=p.Employee_ID
     and Employee_Term_Date is missing
   order by year(Hire_Date), salary desc;
quit;
```

b.

```
/* b. */
title;
proc sql;
select *
  from orion.employees;
quit;
```

2. Creating a Table by Defining Its Structure and Adding Data

a.

```
*** s107s02 ***;
/* a. */
proc sql;
create table orion.rewards
(Purchased num format=comma9.2,
 Year num format=4., Level char(9),
 Award char(50));
quit;
```

b.

```
/* b. */
proc sql;
insert into orion.rewards (Year,Level,Purchased,Award)
```

```

values (2012,'Silver',200,'25% Discount on one item over $25')
values (2012,'Gold',300,'15% Discount on one order over $50')
values (2012,'Platinum',500,'10% Discount on all 2012
      purchases')
values (2013,'Silver',225,'25% Discount on one item over $50')
values (2013,'Gold',350,'15% Discount on one order over $100')
values (2013,'Platinum',600,'10% Discount on all 2013
      purchases');

quit;
proc sql;
select *
  from orion.rewards;
quit;
 ****Alternate method - add the data using set clauses
proc sql;
insert into orion.rewards
  set year=2012,Level='Silver',Purchased=200,
      Award='25% Discount on one item over $25'
  set year=2012,Level='Gold',Purchased=300,
      Award='15% Discount on one order over $50'
  set year=2012,Level='Platinum',Purchased=500,
      Award='10% Discount on all 2012 purchases'
  set year=2013,Level='Silver',Purchased=225,
      Award='25% Discount on one item over $50'
  set year=2013,Level='Gold',Purchased=350,
      Award='15% Discount on one order over $100'
  set year=2013,Level='Platinum',Purchased=600,
      Award='10% Discount on all 2013 purchases';
quit;
 ****

```

3. Creating a Table and Inserting Data Using a Complex Query

```

*** s107s03 ***;

proc sql;
create table work.direct_compensation as
  select hr.Employee_ID,
         catx(' ',First_Name,Last_Name) as Name,
         scan(Job_Title,-1) format=$3. as Level,
         Salary format=comma12.2,
         Sales*0.15 format=comma7.2 as Commission,
         Salary+calculated Commission format=comma12.2
               as Direct_Compensation
  from orion.sales as hr,
       (select Employee_ID,
              sum(Total_Retail_Price) as Sales
       from orion.order_fact
       where year(Order_Date)=2011
             group by Employee_ID) as s
  where hr.Employee_ID=s.Employee_ID

```

```

        and Job_Title like ('%Rep%');

select *
  from work.direct_compensation
  order by Level, Direct_Compensation desc;
quit;

```

4. Creating and Using a View

a.

```

*** s107s04 ***;

/* a. */
proc sql;
create view orion.phone_list as
  select Department format=$25.,
         Employee_Name as Name format=$25.0,
         Phone_Number 'Home Phone' format=$16.
    from orion.employee_addresses as a,
         orion.employee_phones as p,
         orion.employee_organization as o
   where a.Employee_ID=p.Employee_ID and
         o.Employee_ID=p.Employee_ID and
         Phone_Type="Home";
quit;

```

b.

```

/* b. */
proc sql;
title "Engineering Department Home Phone Numbers";
select Name, Phone_Number
  from orion.lhone_list
 where Department="Engineering"
  order by Name;
quit;
title;

```

5. Creating and Using a View to Provide Consolidated Information

a.

```

*** s107s05 ***;

/* a. */
proc sql;
create view orion.t_shirts as
  select d.Product_ID, Supplier_Name format=$20.,
         Product_Name,
         Unit_Sales_Price as Price 'Retail Price'
    from orion.product_dim as d, orion.price_list as l
   where d.Product_ID=l.Product_ID
         and lowcase(Product_Name) like '%t-shirt%';
quit;

```

b.

```
/* b. */
proc sql flow=6 35;
title "Available T-Shirts";
select *
  from orion.t_shirts
  order by Supplier_Name, Product_ID;
quit;
title;
```

c.

```
/* c. */
proc sql;
title "T-Shirts under $20";
select Product_ID, Product_Name, Price format=dollar6.2
  from orion.t_shirts
 where Price < 20
  order by Price;
quit;
title;
```

6. Creating and Using a View That Updates Itself over Time

a.

```
*** s107s06 ***
/* a. */
proc sql;
create view orion.current_catalog as
/*orion.product_dim includes duplicate records*/
select distinct d.*,
      round((l.Unit_Sales_Price *
             factor**year(Today())-year( Start_Date )),.01)
      'Current Retail Price' format=dollar13.2 as Price
  from orion.product_dim as d,
       orion.price_list as l
 where d.Product_ID=l.Product_ID;
quit;
```

b.

```
/* b. */
proc sql;
title "Current Roller Skate Prices";
select Supplier_Name, Product_Name, Price
  from orion.current_catalog
 where lowcase(Product_Name) like '%roller skate%'
  order by Supplier_Name, Price;
quit;
title;
```

c.

```
/* c. */
proc sql;
title "2012 prices > $5.00 higher than original price";
select c.Product_Name, Unit_Sales_Price 'Old Price',
      Price 'New Price',
      Price-Unit_Sales_Price as Increase format=dollar9.2
   from orion.current_catalog as c,
        orion.price_list as p
  where c.Product_ID=p.Product_id
        and calculated Increase gt 5
  order by Increase Desc;
quit;
title;
```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

7.01 Short Answer Poll – Correct Answer

What other Base SAS procedures are similar to the PROC SQL DESCRIBE and SELECT statements?

PROC CONTENTS provides information similar to the **PROC SQL DESCRIBE** statement.

PROC PRINT produces a report similar to the **PROC SQL SELECT** statement.

12

7.02 Multiple Choice Poll – Correct Answer

Based on the query in program **s107a01**, which different data types does the **work.discounts** table have?

- a. All columns are numeric type.
- b. Some columns are date type and some are numeric type.
- c. Some columns are character type and some are numeric type.
- d. Some columns are character type, some are date type, and some are numeric type.

SAS has only two data types: character and numeric.

In PROC SQL, if you specify the ANSI data type DATE when you create a table, the actual data type in the underlying SAS data set is numeric.

- ✍ When the DATE data type defines the data, Enterprise Guide displays DATE in the Properties window.

21

7.03 Quiz – Correct Answer

Locate three syntax errors in the following code:

1. missing semicolon in the CREATE TABLE statement
2. missing comma in the column name list
3. quotation marks around numeric value in the second VALUES clause

```
proc sql;
create table Discounts
  (Product_ID num format=z12.,
   Start_Date date, End_Date date,
   Discount num format=percent.);
insert into discounts
  (Product_ID,Start_Date,End_Date,Discount)
  values (220200200022, '01Mar2013'd,
          '31Mar2013'd, .35)
  values (220200200024, '01Mar2013'd,
          '31Mar2013'd, .35) ;
```

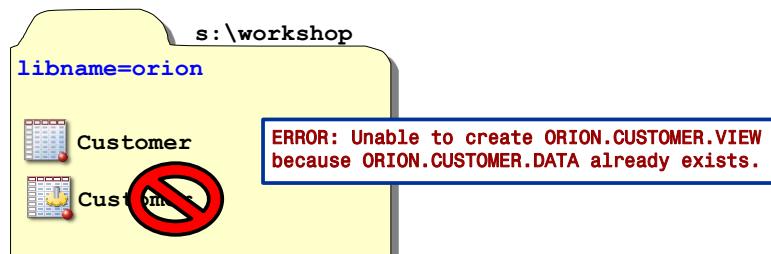
33

7.04 Short Answer Poll – Correct Answer

Open and submit **s107a02**. View the results in the log.

Why was the SQL view not created?

The PROC SQL view cannot have the same name as the data table stored in the same library.



47

7.05 Quiz – Correct Answer

Differences between the SQL code produced by the DESCRIBE VIEW statement and the actual CREATE VIEW code, which created `orion.tom_zhou`:

Original

```
select Employee_Name as Name format=$25.0,
       Job_Title as Title format=$15.0,
       Salary 'Annual Salary' format=comma10.2
       int((today()-Employee_Hire_Date)/365.25)
           as YOS 'Years of Service'
  from Employee_Addresses as a, Employee_Payroll as p,
       Employee_Organization as o
 where a.Employee_ID=p.Employee_ID and
       o.Employee_ID=p.Employee_ID and
       Manager_ID=120102;
```

ANSI labels in the
CREATE VIEW code

DESCRIBE VIEW

```
select Employee_Name as Name format=$25.0,
       Job_Title as Title format=$15.0,
       Salary label='Annual Salary' format=0,
       int((today()-Employee_Hire_Date)/365.25)
           as YOS label='Years of Service'
  from EMPLOYEE_ADDRESSES a, EMPLOYEE_PAYROLL p,
       EMPLOYEE_ORGANIZATION o
 where (a.Employee_ID=p.Employee_ID) and
       (o.Employee_ID=p.Employee_ID) and
       (Manager_ID=120102);
```

DESCRIBE VIEW
produced SAS (LABEL=)
syntax

continued...

7.05 Quiz – Correct Answer

Differences between the SQL code produced by the DESCRIBE VIEW statement and the actual CREATE VIEW code, which created `orion.tom_zhou`:

Original

```
select Employee_Name as Name format=$25.0
       Job_Title as Title format=$15.0
       Salary 'Annual Salary' format=0
       int((today()-Employee_Hire_Date) as AS
           'Years of Service'
  from Employee_Addresses as a, Employee_Payroll as p,
       Employee_Organization as o
 where a.Employee_ID=p.Employee_ID and
       o.Employee_ID=p.Employee_ID and
       Manager_ID=120102;
```

The original code included
int((today()-Employee_Hire_Date) as AS
as YOS 'Years of Service'

DESCRIBE VIEW

```
select Employee_Name as Name format=$25.0
       Job_Title as Title format=$15.0
       Salary label='Annual Salary' format=0
       int((today()-Employee_Hire_Date) as AS
           'Years of Service'
  from EMPLOYEE_ADDRESSES a, EMPLOYEE_PAYROLL p,
       EMPLOYEE_ORGANIZATION o
 where (a.Employee_ID=p.Employee_ID) and
       (o.Employee_ID=p.Employee_ID) and
       (Manager_ID=120102);
```

The DESCRIBE VIEW statement
omitted the keyword AS.

continued...

7.05 Quiz – Correct Answer

Differences between the SQL code produced by the DESCRIBE VIEW statement and the actual CREATE VIEW code, which created `orion.tom_zhou`:

Original

```
select Employee_Name as Name format=$25.0,
       Job_Title as Title format=$15.0,
       Salary 'Annual Salary' format=comma10.2,
       int((today()-Employee_Hire_Date)/365.25)
             as YOS 'Years of Service'
  from Employee_Addresses as a, Employee_Payroll as p,
       Employee_Organization as o
 where a.Employee_ID=p.Employee_ID and
       o.Employee_ID=p.Employee_ID and
       Manager_ID=120102;
```

DESCRIBE VIEW

```
select Employee_Name as Name format=$25.0
       Job_Title as Title format=$15.0
       Salary label='Annual Salary' / inserted parentheses.
              int((today()-Employee_Hire_Date))
                    as YOS label='Years of Service'
  from EMPLOYEE_ADDRESSES a, EMPLOYEE_PAYROLL p,
       EMPLOYEE_ORGANIZATION o
 where (a.Employee_ID=p.Employee_ID) and
       (o.Employee_ID=p.Employee_ID) and
       (Manager_ID=120102);
```

7.06 Poll – Correct Answer

Considering the differences discussed in the previous quiz, if you submit the code produced by the DESCRIBE VIEW statement, would the view produced be identical to the original view?

- Yes
- No

Chapter 8 Advanced PROC SQL Features

8.1 Dictionary Tables and Views	8-3
Exercises	8-13
8.2 Using SQL Procedure Options.....	8-16
8.3 Interfacing PROC SQL with the Macro Language.....	8-23
Demonstration: Using Macro Variables to Make Your Program Dynamic	8-38
Exercises	8-40
8.4 Solutions	8-44
Solutions to Exercises	8-44
Solutions to Student Activities (Polls/Quizzes)	8-51

8.1 Dictionary Tables and Views

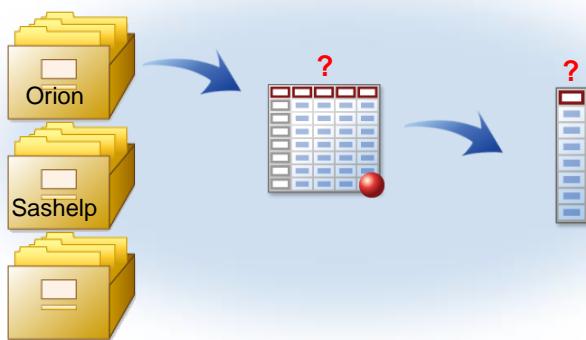
Objectives

- Use dictionary tables and views to obtain information about SAS files.

3

Business Scenario

You have inherited many different data tables and want to become familiar with their content.



4

Dictionary Tables: Overview

Dictionary tables are Read-Only metadata views that contain session metadata, such as information about SAS libraries, data sets, and external files in use or available in the current SAS session.

Dictionary tables are

- created at SAS session initialization
- updated automatically by SAS
- limited to Read-Only access.

You can query dictionary tables with PROC SQL.



5



Metadata is data that provides information about other data.

Querying Metadata about SAS Libraries

There can be more than 30 dictionary tables. We will focus on using data from three of the tables.

DICTIONARY.TABLES

- detailed information about tables

DICTIONARY.COLUMNS

- detailed information about all columns in all tables

DICTIONARY.MEMBERS

- general information about SAS library members

6



SAS librefs are limited to eight characters. The libref **dictionary** is an automatically assigned, reserved libref that is accessible only from within PROC SQL.

Exploring Dictionary Tables

You can use a DESCRIBE statement to explore the structure of dictionary tables:

```
describe table dictionary.tables;
```

Partial Log

NOTE: SQL table DICTIONARY.TABLES was created like:

```
create table DICTIONARY.TABLES
(
  libname char(8) label='Library Name',
  memname char(32) label='Member Name',
  ...
  crdate num format=DATETIME informat=DATETIME label='Date Created',
  modate num format=DATETIME informat=DATETIME label='Date Modified',
  nobs num label='Number of Physical Observations',
  obslen num label='Observation Length',
  nvar num label='Number of Variables', ...);
```

7

s108d01

The DESCRIBE TABLE statement is a good tool for exploring dictionary tables. The complete log notes from the DESCRIBE statement are shown below:

```
create table DICTIONARY.TABLES
(
  libname char(8) label='Library Name',
  memname char(32) label='Member Name',
  memtype char(8) label='Member Type',
  dbms_memtype char(32) label='DBMS Member Type',
  memlabel char(256) label='Dataset Label',
  typemem char(8) label='Dataset Type',
  crdate num format=DATETIME informat=DATETIME label='Date Created',
  modate num format=DATETIME informat=DATETIME label='Date Modified',
  nobs num label='Number of Physical Observations',
  obslen num label='Observation Length',
  nvar num label='Number of Variables',
  protect char(3) label='Type of Password Protection',
  compress char(8) label='Compression Routine',
  encrypt char(8) label='Encryption',
  npage num label='Number of Pages',
  filesize num label='Size of File',
  pcompress num label='Percent Compression',
  reuse char(3) label='Reuse Space',
  bufsize num label='Bufsize',
  delobs num label='Number of Deleted Observations',
  nlobs num label='Number of Logical Observations',
  maxvar num label='Longest variable name',
  maxlabel num label='Longest label',
  maxgen num label='Maximum number of generations',
  gen num label='Generation number',
  attr char(3) label='Dataset Attributes',
  indxtype char(9) label='Type of Indexes',
  datarep char(32) label='Data Representation',
  sortname char(8) label='Name of Collating Sequence',
  sorttype char(4) label='Sorting Type',
```

```

sortchar char(8) label='Charset Sorted By',
reqvector char(24) format=$HEX48 informat=$HEX48 label='Requirements Vector',
datarepname char(170) label='Data Representation Name',
encoding char(256) label='Data Encoding',
audit char(3) label='Audit Trail Active?',
audit_before char(3) label='Audit Before Image?',
audit_admin char(3) label='Audit Admin Image?',
audit_error char(3) label='Audit Error Image?',
audit_data char(3) label='Audit Data Image?'
);

```

Querying Dictionary Information

Display information about the tables in the **orion** library.

```

title 'Tables in the ORION Library';
proc sql;
select memname 'Table Name',
       nobs,nvar,crdate
  from dictionary.tables
 where libname='ORION';
quit;

```

Library names are stored in uppercase in dictionary tables.

8

s108d01



SAS library and table names are stored in uppercase in the dictionary tables. Using SAS functions, such as UPCASE() or LOWCASE(), when querying dictionary tables dramatically degrades query performance. For example, using **upcase(Libname)='ORION'** will cause *all* assigned librefs assigned to the SAS session to be opened to return this information. This prevents the PROC SQL Query Optimizer from seeing many conditions that could be optimized. This is especially apparent when librefs are assigned to a DBMS using a SAS/ACCESS engine.

When you query dictionary tables, you supply values to the WHERE clause in the appropriate case, and match the known case for library and table names (uppercase) and for column names (mixed case). (Know your data!)

The PRESERVE_TAB_NAMES=YES and PRESERVE_COL_NAMES=YES options change how some table and column names are seen by SAS in the dictionary tables. These options might require further investigation to maximize the efficiency of your queries.

Viewing the Output

Partial PROC SQL Output

Tables in the ORION Library				
Table Name	Number of Observations	Physical	Number of Variables	Date Created
CUSTOMER	77	12	14DEC07:08:05:44	
CUSTOMER_TYPE	8	4	14DEC07:08:05:42	
EMPLOYEE_ADDRESSES	424	9	23OCT07:10:12:23	
EMPLOYEE_DONATIONS	124	7	14JAN08:10:34:21	
EMPLOYEE_INFORMATION	424	11	05JUL12:11:25:08	
EMPLOYEE_ORGANIZATION	424	4	12OCT07:14:06:32	
EMPLOYEE_PAYROLL	424	8	14DEC07:08:08:06	

9



This report will vary based on the tables created during class.

Querying Dictionary Information

Display information about the columns in `orion.employee_addresses`.

```
title 'Columns in the orion.employee_addresses '
      'Table';
proc sql;
select Name,Type,Length
  from dictionary.columns
 where libname='ORION'
       and memname='EMPLOYEE_ADDRESSES';
quit;
```

Table names (*memnames*)
are also stored in uppercase
in dictionary tables.

10

s108d01

Viewing the Output

PROC SQL Output

Columns in the orion.employee_addresses Table		
Column Name	Column Type	Column Length
Employee_ID	num	8
Employee_Name	char	40
Street_ID	num	8
Street_Number	num	8
Street_Name	char	40
City	char	30
State	char	2
Postal_Code	char	10
Country	char	2

Column names are stored in mixed case.

11

Using Dictionary Information

Which tables contain the Employee_ID column?

```
title 'Tables Containing an Employee_ID '
      'Column';
proc sql;
select memname 'Table Names', name
  from dictionary.columns
 where libname='ORION' and
       upcase(name)='EMPLOYEE_ID';
quit;
```

Because different tables might use different cases for same-named columns, you can use the UPCASE function for comparisons. However, this significantly degrades the performance of the query.

12

s108d01

Viewing the Output

Partial PROC SQL Output

Tables Containing an Employee_ID Column	
Table Names	Column Name
EMPLOYEE_ADDRESSES	Employee_ID
EMPLOYEE_DONATIONS	Employee_ID
EMPLOYEE_INFORMATION	Employee_ID
EMPLOYEE_ORGANIZATION	Employee_ID
EMPLOYEE_PAYROLL	Employee_ID
EMPLOYEE_PHONES	Employee_ID
ORDER_FACT	Employee_ID
SALESSTAFF	Employee_ID
STAFF	Employee_ID

All Employee_ID column names are stored in uniform mixed case, so the UPCASE function is not needed the next time that a query such as this is executed.

13



The tables identified in this report will vary depending on the files created in your SAS session.

8.01 Short Answer Poll

What is wrong with the following code?

```
proc print data=dictionary.tables label;
  var memname nobs nvar;
run;
```

14

Using Dictionary Tables in Other SAS Code

SAS provides views based on the dictionary tables in the **Sashelp** library.

Most of the **Sashelp** library dictionary view names are similar to dictionary table names, but are shortened to eight characters or less. They begin with the letter v and do not end in s. For example:

```
dictionary.tables = sashelp.vtable
```

The following code executes successfully:

```
title 'Tables in the ORION Library';
proc print data=sashelp.vtable label;
  var memname nobs nvar;
  where libname='ORION';
run;
```

16

s108d02

8.02 Short Answer Poll

In your SAS environment, navigate to the **Sashelp** library. Scroll down to examine the **Sashelp** views.

Which view shows the names and data types of all the columns in every table available in the SAS session?

18

Business Scenario

An analyst asked for a report showing the SAS tables, views, catalogs, and other SAS objects that are in the various libraries.



20

Using Dictionary Information

Use **sashelp.vmember** to extract information from **dictionary.members** in a PROC TABULATE step.

```

title 'SAS Objects by Library';
proc tabulate data=sashelp.vmember format=8.;
  class libname memtype;
  keylabel N=' ';
  table libname, memtype/rts=10
    misstext='None';
  where libname in ('ORION','SASUSER',
                    'SASHHELP');
run;
  
```

21

s108d02

Viewing the Output

PROC TABULATE Output

SAS Objects by Library					
	Member Type				
	CATALOG	DATA	ITEMSTOR	MDDB	VIEW
Library Name					
ORION	None	19	None	None	4
SASHHELP	132	191	2	2	38
SASUSER	3	101	1	None	1

22



The numbers on this report might vary depending on the files that were created in your SAS session.



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Using PROC SQL Options and Displaying the Contents of a Dictionary Table

- a. Write a query that retrieves **memname** (table name) and **memlabel** (description of the table) from **dictionary.dictionaries**. Include only a single row per table name displayed.
Use the FLOW= option to ensure that columns wrap between 6 – 35 characters. Title the report **Dictionary Tables**.

Partial PROC SQL Output

Dictionary Tables	
Member Name	Data Set Label
CATALOGS	Catalogs and catalog-specific information
CHECK_CONSTRAINTS	Check constraints
COLUMNS	Columns from every table
CONSTRAINT_COLUMN_USAGE	Constraint column usage
CONSTRAINT_TABLE_USAGE	Constraint table usage
DATAITEMS	Information Map Data Items
DESTINATIONS	Open ODS Destinations
DICTIONARIES	DICTIONARY tables and their columns
ENGINES	Available engines
EXTFILES	Files defined in FILENAME statements, or implicitly

- b. Use the columns **memname**, **type**, and **length** from **dictionary.columns** where **libname="ORION"** and the table contains a column (**Name**) named **Customer_ID** (**upcase(Name)='CUSTOMER_ID'**). Title the report **Tables containing Customer_ID**.

PROC SQL Output

Tables containing Customer_ID			
Member Name	Column Type	Column Length	
CUSTOMER	num	8	
ORDER_FACT	num	8	
QTR1	num	8	
QTR2	num	8	

Level 2

2. Using PROC SQL Options and Displaying Dictionary Table Information

- a. Produce a report that includes **memname** (table name), **memlabel** (table description), and a count of the number of columns in each table, by querying **dictionary.dictionaries**. Display only one row per table name. Use the FLOW= option to wrap columns between 6 – 35 characters. Title the report **Dictionary Tables**.

Partial PROC SQL Output

Dictionary Tables		
Member Name	Data Set Label	Columns
CATALOGS	Catalogs and catalog-specific information	10
CHECK_CONSTRAINTS	Check constraints	4
COLUMNS	Columns from every table	18
CONSTRAINT_COLUMN_USAGE	Constraint column usage	7

- b. List the table name (**memname**), number of rows (**nobs**), number of columns (**nvar**), file size (**filesize**), length of the widest column (**maxvar**), and length of the widest column label (**maxlabel**) by querying **dictionary.tables**. Limit the list to tables in the **orion** library (**libname="ORION"**) and exclude views (**memtype ne 'VIEW'**) from your report. Order the report by table name. Give the columns labels as indicated by the sample output. Name the report **Orion Library Tables**. If you wrote additional tables to the **orion** library, your output might differ from the sample PROC SQL output.

PROC SQL Output

Orion Library Tables					
Table	Rows	Columns	File Size	Widest Column	Widest Label
CUSTOMER	77	12	49152	18	19
EMPLOYEE_ADDRESSES	424	9	86016	13	0
EMPLOYEE_DONATIONS	124	7	36864	11	11
EMPLOYEE_ORGANIZATION	424	4	49152	11	0
EMPLOYEE_PAYROLL	424	8	40960	18	0
EMPLOYEE_PHONES	923	3	229376	12	11
ORDER_FACT	617	12	73728	18	48
PRICE_LIST	259	6	24576	16	24
PRODUCT_DIM	481	8	131072	16	16
QTR1	22	5	8192	13	33
QTR2	36	6	8192	13	33
SALES	165	9	32768	11	0
SALESSTAFF	163	10	36864	13	25
STAFF	424	10	57344	13	25

Challenge

3. Using PROC SQL Options, SAS System Options, and Dictionary Tables to Document Data Tables in the Orion Star Library

- a. Use PROC SQL options, SAS system options, and **dictionary.tables** to document the data tables in the **orion** library. In the report, exclude views and output only one row per table. For each numeric column, the query should place an asterisk beside each of the maximum values. (See the sample output below.) If you created tables in the **orion** library during this class, your output might have additional entries.

Partial PROC SQL Output

ORION Library Table Information				
Table	Rows	Columns	File Size (Bytes)	Widest Column
CUSTOMER	77	12*	49152	18*
EMPLOYEE_	424	9	86016	13
ADRESSES				
EMPLOYEE_	124	7	36864	11
...				
SALESSTAFF	163	10	36864	13
STAFF	424	10	65536	13

* Largest in the Library

Hint: Consider using a CASE statement to generate the asterisks in the report.

End of Exercises

8.2 Using SQL Procedure Options

Objectives

- Use SQL procedure options to control processing details.

26

PROC SQL Options

PROC SQL options give you finer control over your SQL processes by providing the following features:

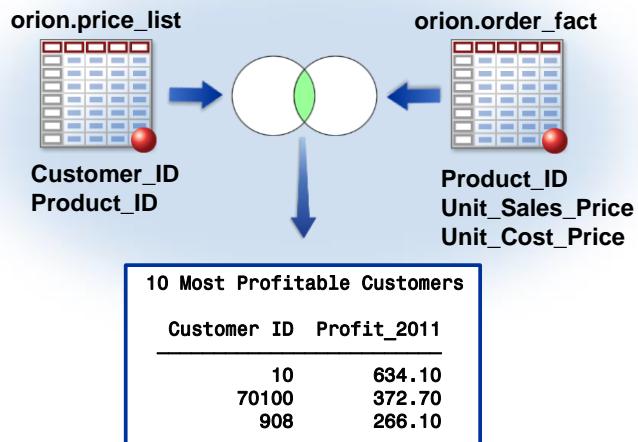
- syntax checking without executing your code
- expanding SQL statements to their fully qualified values
- restricting the number of rows processed
- providing system utilization statistics for query tuning

PROC SQL options;

27

Business Scenario

You have been asked to produce a report of the 10 most profitable customers.



28

OUTOBS= PROC SQL Option

Join all rows, but limit output to the 10 most profitable customers.

```

proc sql outobs=10;
  title "10 Most Profitable Customers";
  select Customer_ID,
         sum(Unit_Sales_Price-Unit_Cost_Price)
           as Profit_2011 format=comma8.2
    from orion.price_list as p,
         orion.order_fact as o
   where p.Product_ID=o.Product_id
     and year(Order_date)=2011
   group by Customer_ID
   order by Profit_2011 desc;
quit;
  
```

OUTOBS=n

29

s108d03

Viewing the Output

PROC SQL Output

10 Most Profitable Customers

Customer ID	Profit_2011
10	634.10
70100	372.70
908	266.10
31	225.35
171	207.50
27	204.95
70201	197.85
2806	177.95
34	142.65
46966	138.80

30

8.03 Multiple Answer Poll

Specifying the NOEXEC option in a PROC SQL statement does which of the following?

- a. prevents statement execution for the current invocation of PROC SQL
- b. applies only to the SELECT statement
- c. checks SQL query syntax without actually executing the statements
- d. displays rewritten PROC SQL statements after references are expanded and certain other transformations are made

31

SQL Options: Controlling Processing

Option	Effect
OUTOBS= <i>n</i>	Restricts the number of rows that a query outputs (displays or writes to a table).
INOBS= <i>n</i>	Sets a limit of <i>n</i> rows from each source table that contributes to a query.
NOEXEC	Checks syntax for all SQL statements without executing them.

33

continued...

SQL Options: Controlling Display

Option	Effect
PRINT NOPRINT	Controls whether the results of a SELECT statement are displayed as a report.
NONUMBER NUMBER	Controls whether the row number is displayed as the first column in query output.

34

continued...

The default value appears first in the slides.

The NONUMBER|NUMBER option affects only PROC SQL reports displayed in an output destination. It cannot be used to add observation numbers to data tables created using PROC SQL.

SQL Options: Controlling Display

Option	Effect
NOSTIMER STIMER	Controls whether PROC SQL writes resource utilization statistics to the SAS log.
NODOUBLE DOUBLE*	Controls whether the report is double-spaced.

* This option applies to the SAS LISTING destination only.

35

continued...

SQL Options: Controlling Display

Option	Effect
NOFLOW FLOW * FLOW=n FLOW=n <m>	controls the appearance of wide character columns. The FLOW option causes text to flow in its column rather than wrapping an entire row. Specifying <i>n</i> determines the width of the flowed column. Specifying <i>n</i> and <i>m</i> floats the width of the column between the limits to achieve a balanced layout.

* This option applies to the SAS LISTING destination only.

36

INOBS= PROC SQL Option

For testing purposes, read 10 rows from `orion.price_list`.

```
proc sql inobs=10;
title "orion.price_list - INOBS=10";
select Product_ID,
       Unit_Cost_Price format=comma8.2,
       Unit_Sales_Price format=comma8.2,
       Unit_Sales_Price-Unit_Cost_Price
       as Margin format=comma8.2
  from orion.price_list;
quit;
```

INOBS=n

- ✍ The INOBS= option applies to each input table.
Use the RESET statement with INOBS=MAX
to read all observations in the table without re-invoking
the procedure.

37

s108d04

Viewing the Output

PROC SQL Output

orion.price_list - INOBS=10			
Product ID	Unit Cost Price	Unit Sales Price	Margin
210200100009	15.50	34.70	19.20
210200100017	17.80	40.00	22.20
210200200023	8.25	19.80	11.55
210200600067	28.90	67.00	38.10
210200600085	17.85	39.40	21.55
210200600112	9.25	21.80	12.55
210200900033	6.45	14.20	7.75
210200900038	9.30	20.30	11.00
210201000050	9.00	19.60	10.60
210201000126	2.30	6.50	4.20

38

Resetting Options

Use the RESET statement to add or change PROC SQL options without re-invoking the procedure.

```
RESET option(s);
```

- ✍ An option specified in the PROC SQL statement remains in effect until you re-invoke PROC SQL or issue a RESET statement.

39

8.04 Quiz

Open the program **s108a01**. This SQL query joins **orion.employee_addresses** and **orion.employee_donations** to calculate each employee's total charitable contribution for 2011. Output rows are numbered and are limited to 10 observations.

Without re-invoking PROC SQL, add a statement before the second query that does the following:

- displays output rows without row numbers
- ensures that only nine rows are output

40

8.3 Interfacing PROC SQL with the Macro Language

Objectives

- State the purpose of the SAS macro language.
- Create and use SAS macro variables in PROC SQL.
- Insert information from dictionary tables into SAS macro variables to create a self-modifying SQL query.

44

Business Scenario

A manager asked you to include some report data in a report title.

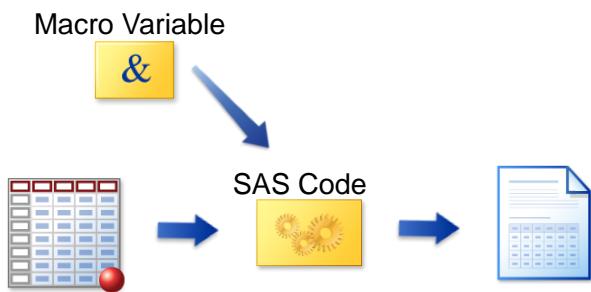
Sales Department Employees Earning More than the Department Average of 27503.06		
Row	Employee_ID	Salary
57	121081	30235
58	121082	28510

Including this type of information in a report title requires an understanding of SAS macro variables and their use in PROC SQL.

45

Creating and Referencing Macro Variables

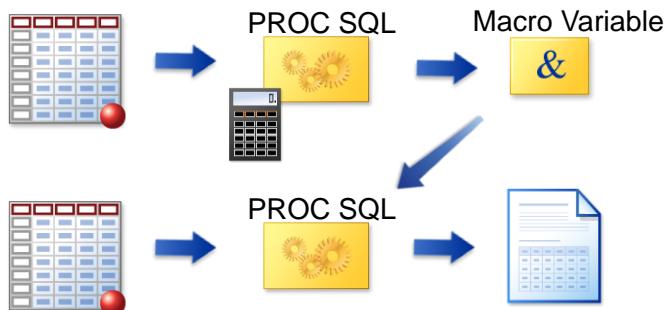
Create macro variables and then submit code to create a report using the defined macro variables.



46

PROC SQL and Macro Variables

In PROC SQL, you can use *macro variables* to store values returned by a query. You can then reference macro variables in other PROC SQL statements and steps.



47

SAS Macro Facility Overview

The SAS macro facility enables you to use the SAS macro language to do the following:

- create macro variables that contain text, and reference them anywhere in a SAS program.



- write special programs (macros) that generate customized SAS code.



48

Macro Variables

SAS macro variables are stored in an area of memory referred to as the *global symbol table*.

There are two types of macro variables: automatic (created and updated by SAS) and user-defined.

Partial Global Symbol Table

Name	Value
...	...
SYSLAST	NULL
SYSSCP	WIN
SYSTIME	09:00
SYSVER	9.3
DSN	employee_payroll

automatic

user-defined

49

The following macro statement will display all macro variables in the global symbol table. This report will be written to the SAS log.

```
%put _all_;
```

To create a report of automatic macro variables, submit the following:

```
%put _automatic_;
```

Creating and Referencing Macro Variables

Taking advantage of macro variables requires two steps.

1. Create and assign a value to a macro variable using one of these methods:
 - %LET statement in SAS code
 - INTO clause in a PROC SQL query
 - CALL SYMPUTX routine in SAS code
2. Reference the macro variable in SAS code so that SAS can resolve the macro variable value.
 - Use `¯o-name`

50

Creating Macro Variables: %LET Statement

Use the %LET macro statement to create and assign a value to a macro variable.

```
%LET variable=value;
```

```
%let DSN=employee_payroll;
%let bigsalary=100000;
%let libname='ORION';
```

variable is any valid SAS variable name, and *value* is any text string.

In SAS macro statements, most keywords are preceded by a percent sign (%).

51

The %LET statement is a global statement.

-  In the %LET statement, *value* can be any string.
- Maximum length is 64K characters.
 - Minimum length is 0 characters (*null value*).
 - Numeric tokens are stored as character strings.
 - Mathematical expressions are not evaluated.
 - The case of *value* is preserved.

- Quotation marks are stored as part of *value*.
- Leading and trailing blanks are removed from *value* before the assignment is made.

Creating Macro Variables: %LET Statement

If the macro variable does not exist, then the %LET creates it and assigns it a value.

If the macro variable exists (it is in the Global Symbol table), then the %LET assigns the new value.

```
%let DSN=employee_payroll;
%let bigsalary=100000;
%let libname='ORION' ;
```

Partial Global Symbol Table

Name	Value
DSN	employee_payroll
BIGSALARY	100000
LIBNAME	'ORION'

52

Referencing Macro Variables

Macro variable references begin with an ampersand (&) followed by a macro variable name.

Partial Global Symbol Table

Name	Value
DSN	employee_payroll
BIGSALARY	100000
LIBNAME	'ORION'

```
proc sql;
select *
  from orion.&dsn
 where Salary > &bigsalary;
```

53

...

Resolving Symbolic References

When the code is submitted, the first macro variable is resolved. The value is obtained from the symbol table and substituted in the program before the syntax is evaluated.

Partial Global Symbol Table

Name	Value
DSN	employee_payroll
BIGSALARY	100000
LIBNAME	'ORION'

```
proc sql;
select *
  from orion.employee_payroll
 where Salary > &bigsalary;
```

55

...

Resolving Symbolic References

When the code is submitted, the second macro variable is resolved. The value is obtained from the symbol table and substituted in the program before the syntax is evaluated.

Partial Global Symbol Table

Name	Value
DSN	employee_payroll
BIGSALARY	100000
LIBNAME	'ORION'

```
proc sql;
select *
  from orion.employee_payroll
 where Salary > 100000;
```

57

Referencing a Macro Variable in Quotation Marks

To reference a macro variable within a quoted text string, enclose the reference in ***double*** quotation marks.

Partial Global Symbol Table

Name	Value
CITY	San Diego

A macro variable resolves within double quotation marks.

```
title "Report for &city";
```

58

Referencing a Macro Variable in Quotation Marks

To reference a macro variable within a quoted text string, enclose the reference in ***double*** quotation marks.

Partial Global Symbol Table

Name	Value
CITY	San Diego

A macro variable resolves within double quotation marks.

```
title "Report for &city";
```

```
title "Report for San Diego";
```

Output

```
Report for San Diego
```

59

8.05 Quiz

What happens if the macro variable value includes quotation marks?

```
%let city='San Diego';
```

Global Symbol Table

Name	Value
CITY	'San Diego'
DATE	05JAN2007
AMOUNT	975

```
where city="&city";
```

```
where city="??";
```

60

s108a02

Referencing a Macro Variable in Quotation Marks

A macro variable within **single** quotation marks will not resolve. The text is treated as literal, and no attempt is made to reference the global symbol table.

```
title 'Report for &city';
```

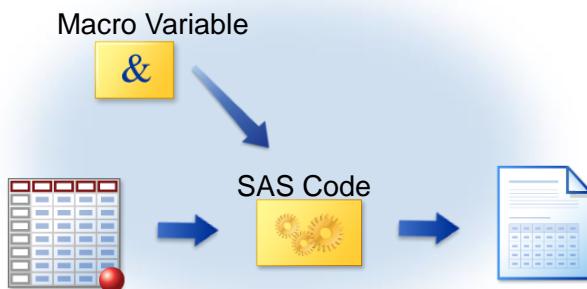
Output

```
Report for &city
```

62

Business Scenario

Create macro variables and then a report using the defined macro variables. You want to validate that the macro variables work in code.



63

Displaying Macro Variable Values

Use the %PUT statement to display the resolved macro variable value along with descriptive text in the SAS log.

```
%let DSN=employee_payroll;
%let bigsalary=100000;
%put DSN is &DSN;
%put bigsalary is &botsalary;
```

%PUT text;

Partial SAS Log

```
%put DSN is &DSN;
DSN is employee_payroll
%put bigsalary is &botsalary;
botsalary is 100000
```

64

s108d05

Displaying Macro Variable Values

When you submit code containing macro variable references, use the SYMBOLGEN system option to see the value that was substituted in the code echoed in the SAS log.

```
options symbolgen; OPTIONS SYMBOLGEN;
proc sql;
title "Salaries > &bigsalary";
select Employee_ID, Salary
  from orion.&DSN
 where Salary > &bigsalary;
quit;
```

Partial SAS Log

```
68 options symbolgen;
69 proc sql;
SYMBOLGEN: Macro variable BIGSALARY resolves to 100000
70 title "Salaries > &bigsalary";
```

65

s108d05



NOSYMBOLGEN is the default system option setting.

Validating Macro Variable Resolution

Display the results of a resolved macro variable reference in the SAS log with the SYMBOLGEN system option.

```
%let DSN=employee_payroll;
%let bigsalary=100000;
%put DSN is &DSN;
%put bigsalary is &bigsalary;
options symbolgen;
proc sql;
title "Salaries > &bigsalary";
select Employee_ID, Salary
  from orion.&DSN
 where Salary > &bigsalary;
quit;
```

66

s108d05

Viewing the Log

Partial SAS Log

```
%let DSN=employee_payroll;
%let bigsalary=100000;

%put DSN is &DSN;
DSN is employee_payroll
%put bigsalary is &bigsalary;
bigsalary is 100000

options symbolgen;
proc sql;
SYMBOLGEN: Macro variable BIGSALARY resolves to 100000
title "Salaries > &bigsalary";
select Employee_ID, Salary
from orion.&DSN
SYMBOLGEN: Macro variable DSN resolves to employee_payroll
where Salary
SYMBOLGEN: Macro variable BIGSALARY resolves to 100000
!
      > &bigsalary;
```

67

Viewing the Output

PROC SQL Output

Salaries > 100000	
Employee_ID	Salary
120101	163040
120102	108255
120259	433800
120260	207885
120261	243190
120262	268455
120659	161290
121141	194885
121142	156065

68

Creating Macro Variables with PROC SQL

In PROC SQL, you use an INTO clause to create macro variables and assign a value to them or to update existing macro variable values. The INTO clause must follow the SELECT clause.

```
SELECT ...
  INTO ...
  FROM table/view ...
  <additional clauses>
```

The INTO clause has three syntaxes, and each produces a different result.

70

The INTO clause occurs between the SELECT list and the FROM clause. It cannot be used in a CREATE TABLE or CREATE VIEW statement. Use the NOPRINT option if you do not want to display the query result.

Business Scenario

Create a report listing all employees in the Sales Department with salaries above the department average. Include the average salary for the department in the report title.

Sales Department Employees Earning More than the Department Average of 27503.06		
Row	Employee_ID	Salary
57	121081	30235
58	121082	28510
59	121085	32235
60	121087	28325
61	121089	28095

71

Creating Macro Variables: Syntax 1

Syntax 1 places values from the *first row* returned by an SQL query into one or more macro variables. Data from additional rows returned by the query is ignored.

```
SELECT column-1 format=format-name.  
<, ...column-n>  
    INTO :macvar_1<, ... :macvar_n>  
    FROM table/view ...
```

The value from the first column in the SELECT clause is placed in the first macro variable listed in the INTO clause, and so on.

72

This method is most often used with queries that return only one row.

- ✍ If a format is not applied to large numeric values, then the values can be stored in scientific notation.

Creating One Macro Variable: Syntax 1

Step 1 Calculate the average Sales Department salary and store the value in a macro variable.

```
%let Dept=Sales;  
proc sql noprint;  
select avg(Salary)  
  into :MeanSalary  
  from orion.employee_payroll as p,  
       orion.employee_organization as o  
  where p.Employee_ID=o.Employee_ID  
    and Department=propcase("&Dept");
```

73

s108d06

Creating One Macro Variable: Syntax 1

Step 2 List employees in the Sales Department with salaries greater than the average. Include the average salary in the title.

```
reset print number;
title "&Dept Department Employees Earning";
title2 "More than the Department Average "
      "of &MeanSalary";
select p.Employee_ID, Salary
  from orion.employee_payroll as p,
       orion.employee_organization as o
 where p.Employee_ID=o.Employee_ID
   and Department=Propcase("&Dept")
   and Salary > &MeanSalary;
```

74

s108d06

Viewing the Output

Partial PROC SQL Output (Rows 57–69)

Sales Department Employees Earning		
More than the Department Average of 27503.06		
Row	Employee_ID	Salary
57	121081	30235
58	121082	28510
59	121085	32235
60	121087	28325
61	121089	28095
62	121095	28010
63	121099	32725
64	121100	28135
65	121104	28315
66	121105	29545
67	121107	31380
68	121129	30945
69	121139	27700

8.06 Quiz

How many changes must be made to the program to generate a report showing how many Engineering Department employees earn above-average salaries?

```
%let Dept=Sales;
proc sql noprint;
select avg(Salary)
  into :MeanSalary
  from orion.employee_payroll as p,
       orion.employee_organization as o
 where p.Employee_ID=o.Employee_ID
   and Department=Propcase("&Dept");
reset print number;
title "&Dept Department Employees Earning";
title2 "More than the Department Average "
      "of &MeanSalary";
select p.Employee_ID, Salary
  from orion.employee_payroll as p,
       orion.employee_organization as o
 where p.Employee_ID=o.Employee_ID
   and Department=Propcase("&Dept")
   and Salary>&MeanSalary;
quit;
title;
```



Using Macro Variables to Make Your Program Dynamic

1. Retrieve the program **s108d06**.
2. Change the value assigned to the macro variable **Dept** in the %LET statement to **Engineering** and submit the program.

```
%let Dept=Engineering;
proc sql noprint;
select avg(Salary)
  into :MeanSalary
  from orion.employee_payroll as p,
       orion.employee_organization as o
  where p.Employee_ID=o.Employee_ID
    and Department=propcase("&Dept");
reset print number;
title "&Dept Department Employees Earning";
title2 "More than the Department Average "
      "of &MeanSalary";
select p.Employee_ID, Salary
  from orion.employee_payroll as p,
       orion.employee_organization as o
  where p.Employee_ID=o.Employee_ID
    and Department=Propcase("&Dept")
    and Salary > &MeanSalary;
quit;
title;
```

3. Examine the results in the Output window.

PROC SQL OUTPUT

Engineering Department Employees Earning More Than The Department Average Of 30698.33
--

Row	Employee_ID	Salary
1	120117	31670
2	121016	48075

The average and the results from the query are populated dynamically by making one simple change from **Sales** to **Engineering** for the value of **Dept**. To generate another report with different values, simply make one change to the value of **Dept** and the program is capable of making the substitutions dynamically at the appropriate places within the code.

End of Demonstration

Creating Multiple Macro Variables: Syntax 1

Using the INTO clause, create multiple macro variables that contain the minimum, average, and maximum salary for the entire company.

```
select avg(Salary),min(Salary),max(Salary)
      into :MeanSalary, :MinSalary, :MaxSalary
      from orion.Employee_payroll;
%put Mean: &MeanSalary Min: &MinSalary
          Max: &MaxSalary;
```

Partial SAS Log

```
Mean: 38041.51 Min: 22710 Max: 433800
```



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

4. Creating and Using Macro Variables

- Write a query for the **employee_payroll** table that returns the highest value of **Salary** (**max(Salary)**). Title the report as indicated in the sample output.

PROC SQL Output

Highest Salary in employee_payroll
—————
433800

- Use %LET statements to create and assign values to two macro variables:

DataSet = employee_payroll
VariableName = Salary

Use a %PUT statement to write their values back to the SAS log in this context:

SAS Log

NOTE: DataSet=Employee_payroll, VariableName=Salary

- Modify the query that you wrote in 4.a. Replace each hardcoded (typed) reference to **employee_payroll** (once in the title and once in the query) with a reference to the macro variable **DataSet (&DataSet)**. Replace each hardcoded (typed) reference to **Salary** (once in the title and once in the query) with a reference to the macro variable **VariableName (&VariableName)**. Resubmit the query, and verify that the results are identical to the results obtained in a. above.
- Use the %LET statements to change the values of your macro variables:

DataSet = Price_List
VariableName = Unit_Sales_Price

Run the report again. The results appear below:

PROC SQL Output

Highest Unit_Sales_Price in Price_List
—————

630.4

Level 2

5. Creating a Macro Variable from an SQL Query

- a. Produce a report of **Country** and a new column named **Purchases** (**SUM (Total_Retail_Price)**). Group the report by **Country**. Include only orders placed in 2011. Order the report so that the highest values of **Purchases** sums are at the top. The data is in the following tables (columns of interest in parenthesis):
- **Order_fact (Customer_ID, Total_Retail_Price)**
 - **Customer (Customer_ID, Country)**

Name the report **2011 Purchases by Country**. Label the columns as indicated in the sample report:

PROC SQL Output

2011 Purchases by Country		
Customer	Country	Purchases
	US	\$10,655.97
	CA	\$5,210.5
	AU	\$3,902.49
	ZA	\$1,954.1
	TR	\$961.4
	DE	\$429.6
	IL	\$194.6

- b. Write a query similar to the first but modified to produce a report of **Purchases** by **Customer_Name** for the year 2011. (Keep a copy of your first query for use in 5.c.below.) Subset the query so that only customers from the top-buying country (the one listed at the top of the last report; in this case, **US**) are included. Order the report so that customers with the highest purchases are at the top. The data is in the following tables:
- **Order_fact (Customer_ID, Total_Retail_Price)**
 - **Customer (Customer_ID, Customer_Name, Country)**

Name the report **2011 US Customer Purchases** with a second title line of **Total US Purchases: \$10,655.97** as indicated by the previous query. When you run the report, it should produce the following results:

Partial PROC SQL Output

2011 US Customer Purchases
Total US Purchases: \$10,655.97

Customer Name	Purchases
Karen Ballinger	\$3,479.09
Cynthia Martinez	\$1,777.60
Cynthia McCluney	\$1,093.60
Wynella Lewis	\$736.60
Alvan Goheen	\$728.80

- c. Automate the report that you wrote in **5.b.**
- 1) Modify the query from **5.a.** so that instead of producing output, it merely writes the values for **Country** and **Purchases** for the first returned row into macro variables named **Country** and **Country_Purchases**, respectively.
 - 2) Modify the query that you wrote in **5.b.** by substituting the macro variable reference **&Country** for each instance where you typed the value **US** in the query, and the macro reference **&Country_Purchases** in the second title line in place of the **\$10,655.97** value that you typed previously. When you are finished, run the modified queries. The results should be exactly as produced in **5.b.**

Hint: Do not forget to use double quotation marks around macro variable references. Review syntax 1 for inserting values into macro variables, if necessary.

- d. Starting with the code from **5.c.**, modify the first query so that the country with the lowest total purchases is read into the macro variable instead of the highest. Then rerun both queries. The queries should produce the following results without further modification:

PROC SQL Output

2007 IL Customer Purchases	
Total IL Purchases: \$194.60	
Customer Name	Purchases
Avinoam Zweig	\$194.60

Challenge

6. Using Dictionary Tables and Macro Variables to Build Low-Maintenance Code

Create a program that, when provided with a libref and the name of a table column, locates in the specified library all of the tables that contain the named column, and writes a custom SQL query that joins them all, producing an output table.

Use the program **s108e06** as a starting point. This is a partially completed utility program that will, when completed, identify all of the tables in any user-specified library. Then join them based on a user-specified common key column to produce a single output table. Your job is to complete the first three steps of the program, which create and populate the macro variables required by the fourth (already completed) section of the program. The fourth section of the program demonstrates SAS macro programming techniques that are outside the scope of an SQL class. Do not modify this section until you successfully completed the exercise.

In **6.a.** below, you write a query to create a macro variable and store the number of tables in the specified library that contains the common key column. In **6.b.**, you write a second query that creates a series of macro variables storing the names of the individual tables. In **6.c.**, you write a third query that creates a single macro variable containing a comma-delimited list of all the tables that contain the common key column.

The two %LET statements near the top of the program assign values to two macro variables: **Library** and **JoinVar**. When you write your portion of the program, reference those values (**&Library**, **&JoinVar**) wherever you type the SAS library libref or the column name in your program code.

Following the %LET statements, there are three sections reserved to write each of the three queries necessary to complete the program. They are clearly labeled with comments: **Query 1:**, **Query 2:**, and **Query 3:**. Remember not to edit the already completed macro program code that follows these sections.

- In the section reserved for Query 1, write a PROC SQL query that counts the number of tables in **Library** that contains the column **JoinVar**. Write that value to a macro variable named **Rows**.
- In the section reserved for Query 2, write a PROC SQL query that creates a series of macro variables named **Table1**, **Table2**,..., **TableN** (that is, one for each of the tables in **Library** that contains the column **JoinVar**). You should use the macro variable **Rows**, created in the previous query, to determine the number of macro variables that you need (**N**).
- In the section reserved for Query 3, write a PROC SQL query that creates a macro variable named **SourceTables**. **SourceTables** must contain a list of the fully qualified names of all of the tables in **Library** that contain the column **JoinVar**. The names should be separated by commas. The contents should resemble the following: **LIBRARY.Table1**, **LIBRARY.Tables**,..., **LIBRARY.TableN**.
- After you create your queries, run the entire program. Review the SAS log. There should be only one error message located near the bottom of the SAS log.

Partial SAS Log

```
ERROR: ***** JoinTheTables Macro *****
ERROR: Column Some_Column_Name not found in any of the ORION tables
ERROR: ***** JoinTheTables Macro *****
```

- Near the top of the code, change the value assigned to **JoinVar** in the %LET statement to **Customer_ID** and rerun the program. Review the SAS log and explore the table produced by the program to answer the following questions:
 - What is the name of the table produced? _____
 - How many columns does the table contain? _____
 - How many tables contain the column **Customer_ID**? _____
- Change the value assigned to **JoinVar** in the %LET statement to **Product_ID** and rerun the program. Review the SAS log and explore the table produced by the program to answer the following questions:
 - What is the name of the table produced? _____
 - How many columns does the table contain? _____
 - How many tables contain the column **Product_ID**? _____

End of Exercises

8.4 Solutions

Solutions to Exercises

1. Using PROC SQL Options and Displaying the Contents of a Dictionary Table

a.

```
*** s108s01 ***;
proc sql flow=6 35;
title "Dictionary Tables";
select distinct memname,memlabel
  from dictionary.dictionaries;
quit;
title;
```

b.

```
proc sql;
title "Tables containing Customer_ID";
select Memname, Type, Length
  from dictionary.columns
 where libname="ORION"
       and upcase(Name)="CUSTOMER_ID";
quit;
title;
```

2. Using PROC SQL Options and Displaying Dictionary Table Information

a.

```
*** s108s02 ***;
proc sql flow=6 35;
title "Dictionary Tables";
select Memname as Table, Memlabel as Contents,
      count(*) as Columns
  from dictionary.dictionaries
 group by Table, Contents;
quit;
title;
```

b.

```
proc sql flow=6 35;
title "Orion Library Tables";
select memname "Table",
      nobs "Rows",
      nvar "Columns",
      filesize "File Size",
      maxvar 'Widest Column',
      maxlabel 'Widest Label'
  from dictionary.tables
 where libname='ORION'
       and memtype ne 'VIEW'
 order by memname;
```

```
quit;
title;
```

3. Using PROC SQL Options, SAS System Options, and Dictionary Tables to Document Data Tables in the Orion Star Library

```
*** s108s03 ***;
/*************************************
 This query remerges summary data and will not work
 if OPTIONS SQLREMERGE=NO is in effect.
 *****/
options ls=85 ps=20;

proc sql flow=2 25;
title "ORION Library Table Information";
footnote "* Largest in the Library";
select memname "Table",
       cats(nobs,
             case
               when nobs=max(nobs) then "*"
               else ""
             end) "Rows",
       cats(nvar,
             case
               when nvar=max(nvar) then "*"
               else ""
             end) "Columns",
       cats(put(filesize,commal2.),
             case
               when filesize=max(filesize) then "*"
               else ""
             end) "File Size (Bytes)",
       cats(maxvar,
             case
               when maxvar=max(maxvar) then "*"
               else ""
             end) 'Widest Column'
from dictionary.tables
where libname='ORION'
      and memtype ne 'VIEW';
quit;
title;
footnote;
```

4. Creating and Using Macro Variables

a.

```
*** s108s04 ***;

proc sql;
title "Highest Salary in Employee_payroll";
select max(Salary)
```

```
from orion.employee_payroll;
quit;
title;
```

b.

```
%let DataSet=employee_payroll;
%let VariableName=Salary;
%put NOTE: DataSet=&DataSet, VariableName=&VariableName;
```

c.

```
proc sql;
title "Highest &VariableName in &DataSet";
select max(&VariableName)
      from orion.&DataSet;
quit;
title;
```

d.

```
%let DataSet=Price_List;
%let VariableName=Unit_Sales_Price;
proc sql;
title "Highest &VariableName in &DataSet";
select max(&VariableName)
      from orion.&DataSet;
quit;
title;
```

5. Creating a Macro Variable from an SQL Query

a.

```
*** s108s05 ***;
proc sql;
title "2011 Purchases by Country";
select Country,
       sum(Total_Retail_Price) format=dollar10.2 as Purchases
  from orion.customer as c,
       orion.order_fact as o
 where c.Customer_ID=o.Customer_ID
   and year(Order_Date)=2011
 group by Country
 order by Purchases desc;
quit;
title;
```

b.

```
proc sql;
title "2011 US Customer Purchases";
title2 "Total US Purchases: $10,655.97";
select Customer_Name,
       sum(Total_Retail_Price) format=dollar10.2 as Purchases
  from orion.customer as c,
       orion.order_fact as o
```

```

    where c.Customer_ID=o.Customer_ID
        and year(Order_Date)=2011
        and Country="US"
    group by Customer_Name
        order by Purchases desc;
quit;
title;

```

c.

```

proc sql noprint;
select Country,
       sum(Total_Retail_Price) format=dollar10.2
       into :Country, :Country_Purchases
from orion.customer as c,
     orion.order_fact as o
where c.Customer_ID=o.Customer_ID
      and year(Order_Date)=2011
group by Country
order by Purchases desc;
reset print;

title "2011 &Country Customer Purchases";
title2 "Total &Country Purchases: &Country_Purchases";
select Customer_Name,
       sum(Total_Retail_Price) format=dollar10.2 as Purchases
from orion.customer as c,
     orion.order_fact as o
where c.Customer_ID=o.Customer_ID
      and year(Order_Date)=2011
      and Country="&Country"
group by Customer_Name
order by Purchases desc;
quit;
title;

```

d.

```

proc sql noprint;
select Country,
       sum(Total_Retail_Price) format=dollar10.2
       into :Country, :Country_Purchases
from orion.customer as c,
     orion.order_fact as o
where c.Customer_ID=o.Customer_ID
      and year(Order_Date)=2011
group by Country
order by Purchases /* ascending requires no keyword */;
reset print;
title "2011 &Country Customer Purchases";
title2 "Total &Country Purchases: &Country_Purchases" ;
select Customer_Name,
       sum(Total_Retail_Price) As Purchases
format=dollar10.2

```

```

from orion.Customer as c,
      orion.Order_fact as o
where c.Customer_ID=o.Customer_ID
      and year(Order_Date)=2011
      and Country=&Country"
group by Customer_Name
order by Purchases desc;
quit;
title;

```

6. Using Dictionary Tables and Macro Variables to Build Low-Maintenance Code

a.

```

*** s108s06 ***
/* Warning: Case sensitive! */
%let JoinVar=Some_Column_Name;
%let Library=%upcase(ORION);
proc sql noprint;
select strip(put(count(*),5.))
      into :Rows
      from dictionary.columns
      where libname=&Library
            and name=&JoinVar";
quit;

```

b.

```

proc sql noprint;
select memname
      into :Table1:-:Table&Rows
      from dictionary.columns
      where libname=&Library
            and name=&JoinVar";
quit;

```

c.

```

proc sql noprint;
select catx('.',libname,memname)
      into :SourceTables separated by ','
      from dictionary.columns
      where libname=&Library
            and NAME=&JoinVar";
quit;

*****
For the exercise, do not edit below this line!
*****



*****
This macro program joins all the tables in Library
which contain the column JoinVar by JoinVar, producing
a table called Joined_By_JoinVar. By default, the table
is created in the work library. For example, if

```

```

JoinVar = Employee_ID
then the table produced is
  work.Joined_by_employee_id
******/
```

```

options mprint;
%Macro JoinTheTables(OutLib);
%if &OutLib= %then %let OutLib=work;
%if &Rows gt 1 %then %do;
  %do i=1 %to &rows;
    proc sql noprint;
      select catx('.','&&Table&i",Name)
        into :&&Table&i.._Columns separated by ","
        from Dictionary.Columns
        where libname="&Library"
          and MEMNAME="&&Table&i"
          and Name ne "&JoinVar"
      ;
    quit;
    %end;
  %put _user_;
  proc sql;
    create table &OutLib..Joined_by_&JoinVar AS
      select &&Table1..&JoinVar
    %do i=1 %to &Rows;
      %let ThisColumn=&&Table&i.._Columns;
      , &&ThisColumn
    %end;
    from &SourceTables
    where &Table1..&JoinVar=&Table2..&JoinVar
  %do i=2 %to %eval(&Rows-1);
    %let j=%eval(&i+1);
    and &&Table&i..&JoinVar=&&Table&j..&JoinVar
  %end;
  ;
  quit;
  %put NOTE: ***** JoinTheTables Macro ****;;
  %put NOTE: Column &JoinVar found was found in &Rows tables;
  %put NOTE: ***** JoinTheTables Macro ****;;
%end;
%else %if &Rows=1 %then %do;
  %put NOTE: ***** JoinTheTables Macro ****;;
  %put NOTE: Column &JoinVar found only in ;
  %put NOTE: &Library..&Table1 table ;
  %put NOTE: No join could be performed ;
  %put NOTE: ***** JoinTheTables Macro ****;;
%end;
%else %do;
  %put ERROR: ***** JoinTheTables Macro ****;;
  %put ERROR: Column &JoinVar not found in any ;
  %put ERROR: of the &Library tables ;
;
```

```
%put ERROR: ***** JoinTheTables Macro *****;  
%end;  
%mend;  
  
%JoinTheTables;
```

- d. No solution is required.
- e. For the **Customer_ID** column:
 - 1) What is the name of the table produced? **work.Joined_by_Customer_ID**
 - 2) How many columns does the table contain? **22**
 - 3) How many tables contained the column **Customer_ID**? **4**
- f. For the **Product_ID** column:
 - 1) What is the name of the table produced? **work.Joined_by_Product_ID**
 - 2) How many columns does the table contain? **24**
 - 3) How many tables contained the column **Product_ID**? **5**

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

8.01 Short Answer Poll – Correct Answer

What is wrong with the following code?

```
proc print data=dictionary.tables label;
  var memname nobs nvar;
run;
```

In SAS procedures or the DATA step, the libref
cannot exceed 8 characters.

15

8.02 Short Answer Poll – Correct Answer

In your SAS environment, navigate to the **Sashelp** library.
Scroll down to examine the **Sashelp** views.

Which view shows the names and data types
of all the columns in every table available
in the SAS session?

SASHELP.vcolumn

19

8.03 Multiple Answer Poll – Correct Answer

Specifying the NOEXEC option in a PROC SQL statement does which of the following?

- a. prevents statement execution for the current invocation of PROC SQL
- b. applies only to the SELECT statement
- c. checks SQL query syntax without actually executing the statements
- d. displays rewritten PROC SQL statements after references are expanded and certain other transformations are made

32

8.04 Quiz – Correct Answer

Without re-invoking PROC SQL, add a statement before the second query that does the following:

- displays output rows without row numbers
- ensures that only nine rows are output

```
reset nonumber outobs=9;
```

This statement displays output rows without row numbers and ensures that only nine rows are output without re-invoking PROC SQL.

41

8.05 Quiz – Correct Answer

What happens if the macro variable value includes quotation marks?

```
%let city='San Diego';
```

Global Symbol Table

Name	Value
CITY	'San Diego'
DATE	05JAN2007
AMOUNT	975

A macro variable resolves within double quotation marks.

```
where city="&city";
```

```
where city="??";
```

Partial SAS Log

```
NOTE: No rows were selected.
```

s108a02

61

8.06 Quiz – Correct Answer

How many changes must be made to the program to generate a report showing how many Engineering Department employees earn above-average salaries?

One. Modify the value assigned to the macro variable Dept in the %LET statement.

```
%let Dept=Engineering;
```

77

Chapter 9 Learning More

9.1 SAS Resources	9-3
9.2 Beyond This Course.....	9-6

9.1 SAS Resources

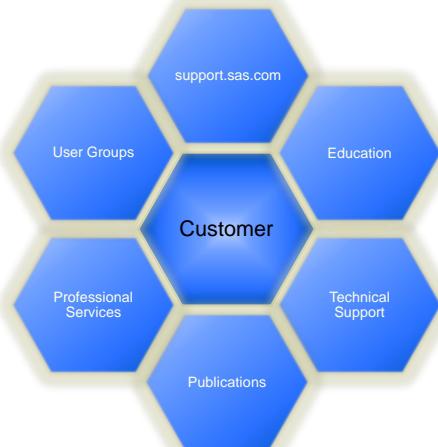
Objectives

- Identify the areas of support that SAS offers.

3

Customer Support

SAS provides a variety of resources to help customers.



<http://support.sas.com/resourcekit/>

4

Education

SAS Education provides comprehensive training, including

- more than 200 course offerings
- world-class instructors
- multiple delivery methods
- worldwide training centers.



<http://support.sas.com/training/>

5

SAS Global Certification Program

SAS Education also provides

- globally recognized certifications
- preparation materials
- practice exams.



<http://support.sas.com/certify/>

6

Networking

Social media channels and user group organizations enable you to

- interact with other SAS users and SAS staff
- learn new programming tips and tricks
- get exclusive discounts.



For training-specific information:

<http://support.sas.com/training/socialmedia>

7

SAS Books

SAS Books offers a complete selection of publications, including

- eBooks
- CD-ROM
- hard-copy books
- books written by outside authors.



<http://support.sas.com/bookstore/>

1-800-727-3228

8

Extended Learning Pages

To grow your SAS skills, remember to activate the ***extended learning page*** for this course.



Individual learning software is available.

9

9.2 Beyond This Course

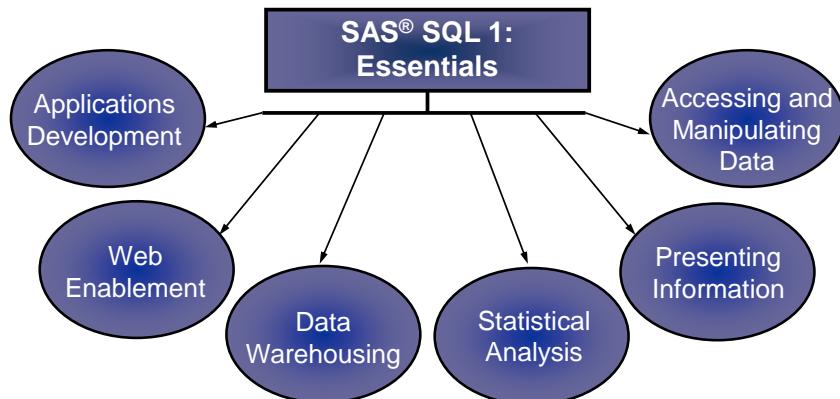
Objectives

- Identify training opportunities to build on the skills you learned in this course.

11

Next Steps

SAS® SQL 1: Essentials provides portable skills that are useful in most of the focus areas shown below.



12

Next Steps

To learn more about this:	Enroll in this course:
Writing self-modifying, environmentally aware, easily maintained SAS programs	SAS® Macro Language 1: Essentials
Efficiently preparing data for data mining and analysis	SAS® Programming 3: Advanced Techniques and Efficiencies
Creating tabular detail and summary reports	SAS® Report Writing 1: Essentials
Writing DATA step programs that manipulate high-precision numeric data and can easily work on multiple observations in parallel to speed up computationally intense processing	DS2 Programming Essentials

13

Appendix A Self-Study Topics

A.1 Correlated Subqueries	A-3
Exercises	A-25
A.2 Introduction to Indexes	A-26
A.3 Creating Indexes.....	A-33
Exercises	A-39
A.4 Maintaining Tables with the SQL Procedure.....	A-42
Exercises	A-58
A.5 Integrity Constraints.....	A-61
A.6 Interfacing PROC SQL with the Macro Language.....	A-70
Demonstration: Creating and Using Macro Variables in PROC SQL	A-79
A.7 Solutions	A-80
Solutions to Exercises	A-80
Solutions to Student Activities (Polls/Quizzes)	A-87

A.1 Correlated Subqueries

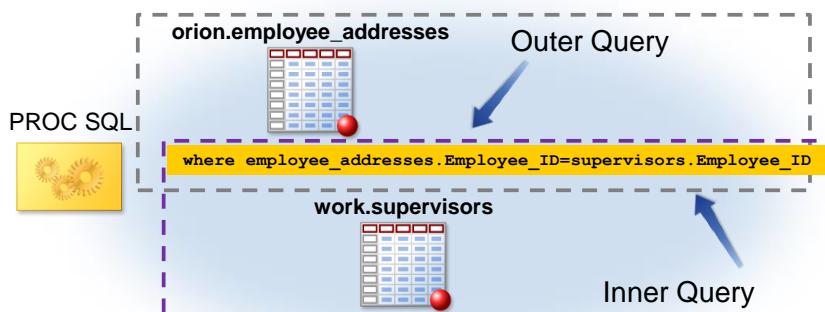
Objectives

- Define correlated subqueries.
- Describe how data is subset using correlated subqueries.

3

Business Scenario

Use a correlated subquery to create a report listing the employee identifier and name for all managers in Australia.



4

Correlated Subqueries

Correlated subqueries

- cannot be evaluated independently
- require values to be passed to the inner query from the outer query
- are evaluated for each row in the outer query.

5

Correlated Subqueries

In a correlated subquery, the outer query provides information so that the subquery resolves successfully.

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
   (select Country
      from work.supervisors
     where employee_addresses.Employee_ID=
          supervisors.Employee_ID);
quit;
```

You must qualify each column with a table name.

6

s1aad01

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID) ;
quit;
```

Step 1: The outer query takes the first row in **orion.employee_addresses** and finds **Employee_ID** and **Employee_Name**.

7

Partial Listing of orion.employee_addresses	
Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors	
Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID) ;
quit;
```

Step 2: In the subquery, try to match **employee_addresses.Employee_ID** of **120145** with the value of **supervisors.Employee_ID** to find a qualifying row in **work.supervisors**.

8

NO MATCH

Partial Listing of orion.employee_addresses	
Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors	
Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

Partial Listing of
orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Steps 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is a match.

MATCH

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

Partial Listing of
orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Step 3: The subquery passes the value of **Country** from the selected row in **work.supervisors** back to the outer query, where the = operator compares this value to 'AU' for selection in the main query. In this case, the main query WHERE expression (where 'AU'='US') resolves to FALSE.

FALSE

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

A.01 Quiz

Given the following query, subquery, and data in **work.supervisors**, what is the maximum number of rows that will be selected by the outer query?

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

work.supervisors	
Employee_ID	Country
120798	US
120800	US
120104	AU
120735	US
121141	US
...	...
120262	US
120679	US
120103	AU
120668	US
121143	US
120260	US
120672	AU

11

The Outer Query Controls the Result Set

The outer query determines which rows cause the inner query to resolve successfully.

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

work.supervisors	
Employee_ID	Country
120798	US
120800	US
120104	AU
120735	US
121141	US
...	...
120262	US
120679	US
120103	AU
120668	US
121143	US
120260	US
120672	AU

13

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Steps 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is no match.

NO MATCH

14

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Step 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is a match.

MATCH

15

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Correlated Subqueries

```

proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID);
quit;

```

Subquery returns 'AU'

Step 3: The subquery passes the value of Country from the selected row in **work.supervisors** back to the outer query, where the = operator compares this value to 'AU' for selection in the main query. In this case, the main query WHERE expression (**WHERE 'AU'='AU'**) resolves to TRUE.

16 **TRUE**

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...
16

Correlated Subqueries

```

proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID);
quit;

```

Step 4: Write **Employee_ID** and **Manager_Name** from **orion.employee_addresses** as the first row in a newly created report.

Build the first row of the report:

Employee_ID	Manager_Name
120104	Billington, Karen

17 ...
17

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
     from work.supervisors
    where employee_addresses.Employee_ID=
          supervisors.Employee_ID) ;
quit;
```

Partial Listing of
orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Steps 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is no match.

NO MATCH

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
     from work.supervisors
    where employee_addresses.Employee_ID=
          supervisors.Employee_ID) ;
quit;
```

Partial Listing of
orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Steps 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is a match.

MATCH

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

19

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
      (select Country
        from work.supervisors
       where employee_addresses.Employee_ID=
             supervisors.Employee_ID);
quit;
```

Subquery returns 'US'

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Step 3: The subquery passes the value of Country from the selected row in **work.supervisors** back to the outer query, where the = operator compares this value to 'AU' for selection in the main query. In this case, the main query WHERE expression (**WHERE 'AU'='US'**) resolves to FALSE.

20 FALSE

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
      (select Country
        from work.supervisors
       where employee_addresses.Employee_ID=
             supervisors.Employee_ID);
quit;
```

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...

Steps 1 and 2 (Repeat): Read the next row from **orion.employee_addresses** and pass the corresponding employee ID to the subquery to look for a matching employee ID in **work.supervisors**. There is a match.

21 MATCH

Correlated Subqueries

```

proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID);
quit;

```

Subquery returns 'AU'

Step 3: The subquery passes the value of Country from the selected row in **work.supervisors** back to the outer query, where the = operator compares this value to 'AU' for selection in the main query. In this case, the main query WHERE expression (**WHERE 'AU'='AU'**) resolves to TRUE.

22 **TRUE**

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

...
23

Correlated Subqueries

```

proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID);
quit;

```

Step 4: Write **Employee_ID** and **Manager_Name** from **orion.employee_addresses** as the first row in a newly created report.

Build the second row of the report:

Employee_ID	Manager_Name
120104	Billington, Karen
120103	Dawes, Wilson

23 ...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID) ;
quit;
```

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

Continue repeating the steps until all rows are read from orion.employee_addresses.

24 MATCH ...

Correlated Subqueries

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
  (select Country
    from work.supervisors
   where employee_addresses.Employee_ID=
        supervisors.Employee_ID) ;
quit;
```

Subquery returns 'AU'

Partial Listing of orion.employee_addresses

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

work.supervisors

Employee_ID	Country
120798	US
120104	AU
120735	US
121141	US
120103	AU
120668	US
120672	AU

Step 3: The subquery passes the value of Country from the selected row in **work.supervisors** back to the outer query, where the = operator compares this value to 'AU' for selection in the main query. In this case, the main query WHERE expression (**WHERE 'AU'='AU'**) resolves to TRUE.

25 TRUE ...

Viewing the Output

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
from orion.employee_addresses
where 'AU'=
      (select Country
       from work.supervisors
       where employee_addresses.Employee_ID=
             supervisors.Employee_ID) ;
quit;
```

**Partial Listing of
orion.employee_addresses**

Employee_ID	Employee_Name
120145	Aisbitt, Sandy
120798	Ardskin, Elizabeth
120656	Amos, Salley
120104	Billington, Karen
121035	Blackley, James
121141	Bleu, Henri Le
120103	Dawes, Wilson
120672	Guscott, Verne

Step 4: Write **Employee_ID** and **Manager_Name** from **orion.employee_addresses** as the first row in a newly created report.

Build the third (and final) row of the report:

Employee_ID	Manager_Name
120104	Billington, Karen
120103	Dawes, Wilson
120672	Guscott, Verne

26

Business Scenario

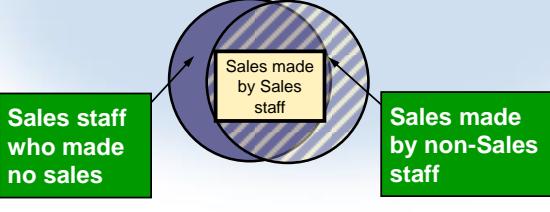
Create a report showing **Employee_ID** and **Job_Title** columns of all sales personnel who did not make any sales.



orion.sales
(all Sales staff)



orion.order_fact
(all sales)



The diagram consists of two overlapping circles. The left circle is labeled "Sales staff who made no sales". The right circle is labeled "Sales made by non-Sales staff". The intersection area is shaded purple and contains the text "Sales made by Sales staff".

28

The EXISTS and NOT EXISTS Condition

The EXISTS condition tests for the existence of a set of values returned by the subquery.

- The EXISTS condition is true if the subquery returns at least one row.
- The NOT EXISTS condition is true if the subquery returns no data.

29

Correlated Subqueries

The table **orion.sales** contains the employee IDs, job titles, and other demographic information about the Orion Star Sales staff.

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
       where sales.Employee_ID=
             order_fact.Employee_ID);
```

orion.sales

s1aad02
...

30

Correlated Subqueries

The **orion.order_fact** table contains a row for each product sold to a customer.

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
       where sales.Employee_ID=
             order_fact.Employee_ID);
quit;
```

31

s1aad02
...

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
       where sales.Employee_ID=
             order_fact.Employee_ID);
quit;
```

32

s1aad02

Testing Concepts: Referencing Columns

Are the highlighted column references equivalent?
Will they result in the same output?

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
      from orion.order_fact
     where sales.Employee_ID =
          order_fact.Employee_ID);
quit;
```

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
      from orion.order_fact
     where Employee_ID=Employee_ID);
quit;
```

33

Setup for the Poll

1. Submit the program **s1aaa01** and review the results.
2. Change the original code to the code shown below.
3. Submit the changed program and review the results.

Your instructor will review the log results with you.

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
      from orion.order_fact
     where Employee_ID=Employee_ID);
quit;
```

34

A.02 Poll

Is it necessary to qualify the column names in the inner WHERE clause as follows?

```
where sales.Employee_ID=order_fact.Employee_ID
```

- Yes
- No

35

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
      from orion.order_fact
     where sales.Employee_ID=
          order_fact.Employee_ID);
quit;
      where 120121=120121
```

MATCH

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

37

Correlated Subqueries

```

proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;      where 120121=120121

```

MATCH

The NOT EXISTS clause is **FALSE**.
No output rows are written.

orion.sales	
Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact		
Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

38

...

Correlated Subqueries

```

proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;      where 120122=120122

```

MATCH

orion.sales	
Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact		
Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

39

...

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
    from orion.order_fact
    where sales.Employee_ID=
          order_fact.Employee_ID);
quit;
      where 120122=120122
```

MATCH

The NOT EXISTS clause is FALSE.
No output rows are written.

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

40

...

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
    from orion.order_fact
    where sales.Employee_ID=
          order_fact.Employee_ID);
quit;
```

NO MATCH

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

41

...

Correlated Subqueries

```

proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;

```

TRUE

NO MATCH

The NOT EXISTS clause evaluates as **TRUE**.
The first output row is written.

orion.sales	
Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact		
Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

Partial PROC SQL Output

Employee_ID	Job_Title
120102	Sales Manager
...	...

42

Correlated Subqueries

```

proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;

```

where 120123=120123

MATCH

orion.sales	
Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact		
Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

43

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;      where 120123=120123
```

MATCH

The NOT EXISTS clause is FALSE.
No output rows are written.

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

44

...

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;
```

NO MATCH

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

45

...

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;
```

NO MATCH

The NOT EXISTS clause evaluates as **TRUE**.
The next output row is written.

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

Partial PROC SQL Output

Employee_ID	Job_Title
120102	Sales Manager
120103	Sales Manager
...	...

46

Correlated Subqueries

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
       (select *
        from orion.order_fact
        where sales.Employee_ID=
              order_fact.Employee_ID);
quit;
      where 120124=120124
```

MATCH

The NOT EXISTS clause is **FALSE**.
No output rows are written.

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

47

Viewing the Output

```
proc sql;
select Employee_ID, Job_Title
  from orion.sales
 where not exists
   (select *
      from orion.order_fact
     where sales.Employee_ID=
          order_fact.Employee_ID);
quit;
```

When EOF is reached for the table in the outer query, PROC SQL stops processing the query.

orion.sales

Employee_ID	Job_Title
...	...
120121	Sales Rep. II
120122	Sales Rep. II
120102	Sales Manager
120123	Sales Rep. I
120103	Sales Manager
120124	Sales Rep. I
...	...

orion.order_fact

Employee_ID	Order_Date	Quantity
...
120122	28MAY2004	1
120121	24JUN2004	1
120124	08OCT2005	1
120123	18AUG2004	1
...

Partial PROC SQL Output

Employee_ID	Job_Title
120102	Sales Manager
120103	Sales Manager



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

1. Creating a Simple Correlated Subquery

Create a report showing **Employee_ID** and the birth month (calculated as **month(Birth_date)**) for all Australian employees, using a correlated subquery.

The table **orion.Employee_Payroll** contains **Employee_ID** and **Birth_Date**.

In the subquery, select only **Country** from **orion.Employee_Addresses**. Use a WHERE clause to return only rows where the **EmployeeID** in **orion.Employee_Addresses** matches the **EmployeeID** in **orion.Employee_Payroll**.

Order the report by birth month.

Partial PROC SQL Output

Australian Employees' Birth Months	
Employee_ID	Birth Month
120145	1
120198	1
120147	1
120127	1
120191	1

Level 2

2. Using a Correlated Subquery

Generate a report that displays **Employee_ID**, **Employee_Gender**, and **Marital_Status** for all employees who donate more than 0.02 of their salary. The table **orion.Employee_donations** contains **Employee_ID**, quarterly donations (**Qtr1-Qtr4**), and charities (**Recipients**).

The **orion.Employee_Payroll** table contains **Employee_ID** and **Salary** information.

Partial PROC SQL Output

Employees With Donations > 0.02 Of Their Salary		
Employee_ID	Employee_Gender	Marital_Status
120267	F	M
120680	F	S
120686	F	M
120689	F	S
120753	M	S

End of Exercises

A.2 Introduction to Indexes

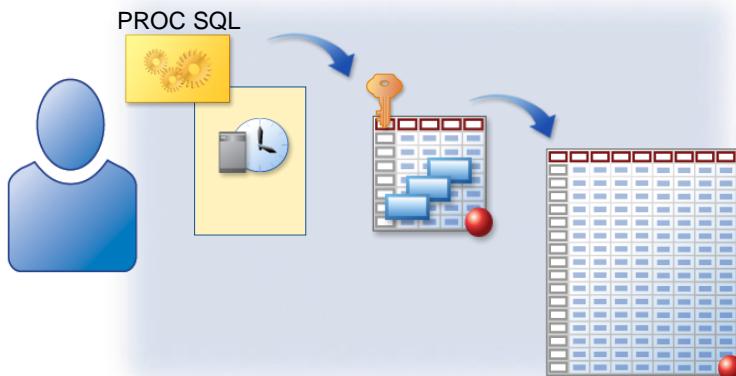
Objectives

- Describe the functions of an index.
- Determine whether the SQL optimizer elected to use an index to process an SQL query.
- Take explicit control of index usage in an SQL query.

52

Business Scenario

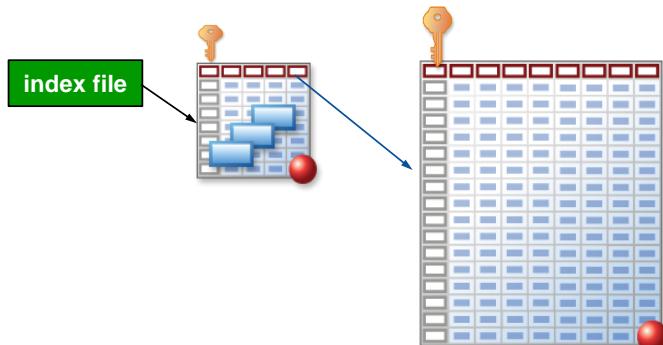
Investigate the use of indexes to optimize data subsetting or joining tasks.



53

What Is an Index?

An *index* is an auxiliary data structure that stores information about the location of indexed table rows, based on the values in one or more **key** columns. You can index both character and numeric columns.



54

Simplified Index File

The index file consists of entries that are stored in order by the key variables

Indexed SAS Data Table

Obs	ID	Loc	Level
1	1001	US	I
2	1002	AU	II
3	1003	AU	III
4	1004	US	III
5	1005	US	II

Index File - Key=Level

Key Value	Location of Obs Page(obs,obs,...)
I	1(1,...) 2(...) ...
II	1(2,5,...) 2(...) ...
III	1(3,4...) 2(...) ...

Query Code

```
proc sql;
select *
  from table
 where Level='II';
quit;
```

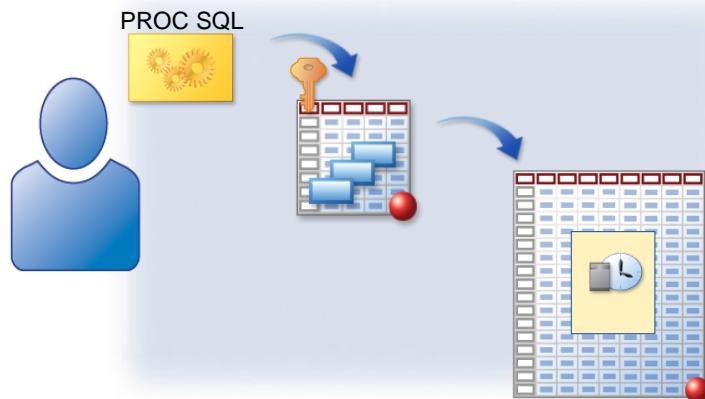
Data Processed

Obs	ID	Loc	Level
2	1002	AU	II
5	1005	US	II

55

Why Use Indexes?

The SQL procedure can use available indexes to optimize data subsetting or joining tasks.



56

Why Use Indexes?

Indexes can provide fast access to small subsets of data.

```
proc sql;
  select *
    from orion.staff
   where Job_Title='Sales Rep. I';
quit;
```

'Sales Rep. I' is one
of many distinct values of
the variable Job_Title.

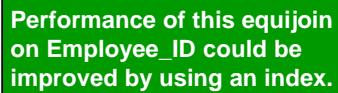
57

Why Use Indexes?

Indexes can enhance join performance, especially equijoins.

```
proc sql;
select *
  from orion.employee_payroll,
       orion.employee_addresses
 where employee_payroll.Employee_ID=
       employee_addresses.Employee_ID;
```

Performance of this equijoin
on Employee_ID could be
improved by using an index.



58

Index Terminology

Two types of indexes are as follows:

- *Simple*
 - is based on values of only one column
 - *Composite*
 - is based on values of more than one column,
concatenated to form a single value (for example,
Product_ID and **Order_ID**)
 - can include mixed variable types; that is,
a composite index can contain both character
and numeric variables
-  A table can have multiple simple and composite indexes.

59

Naming SAS Indexes

A ***simple*** index must use the name of the underlying column.

A ***composite*** index cannot use the name of any existing column.

Index Type	Key Variables	Index Name
Simple	ProductID	ProductID
Simple	ProductGroup	ProductGroup
Composite	OrderID ProductID	SaleID

60

Using an Index

If an index is available, the SQL Optimizer first estimates both the effects of using and not using the index, and then chooses the most efficient method to process a query.

You can use the MSGLEVEL system option to increase the detail level of SAS log messages, alerting you to when an index was used.

```
OPTIONS MSGLEVEL = N | I;
```

- N prints notes, warnings, and error messages only.
N is the default.
- I prints additional notes pertaining to index usage, merge processing, and sort utilities.

61

Using an Index

The compound index **Locale** was created on the **City**, **State**, and **Country** columns of **work.employee_addresses**.

Use the MSGLEVEL=I system option to determine which queries used the index.

Partial SAS Log

```
options msglevel = i;
proc sql;
select *
  from work.employee_addresses
  where State in ('CA','PA');
INFO:Index Locale selected for WHERE clause optimization.

select *
  from work.employee_addresses
  where Postal_Code ='33135';
quit;
```

The INFO message applies only to the preceding WHERE clause.
INFO:Index Locale selected for WHERE clause optimization.
no index for Postal_Code

62

s1aad03

Setup for the Poll

Submit the program **s1aaa02** to do the following:

- Create the table **work.employee_addresses**.
- Add a composite index named **Locale** based on the columns **City**, **State**, and **Country**.
- Run three queries against the table.

Review your SAS log.

63

A.03 Multiple Choice Poll

Which query (or queries) used the **Locale** index?

- a. Query 1 (State)
- b. Query 2 (State and City)
- c. Query 3 (City)
- d. Queries 1 and 2
- e. Queries 1 and 3
- f. all queries used the index
- g. none of the above

64

Controlling Index Usage for WHERE Processing

Data set options enable you to direct SAS when to use an index.



IDXWHERE=YES

IDXWHERE=NO

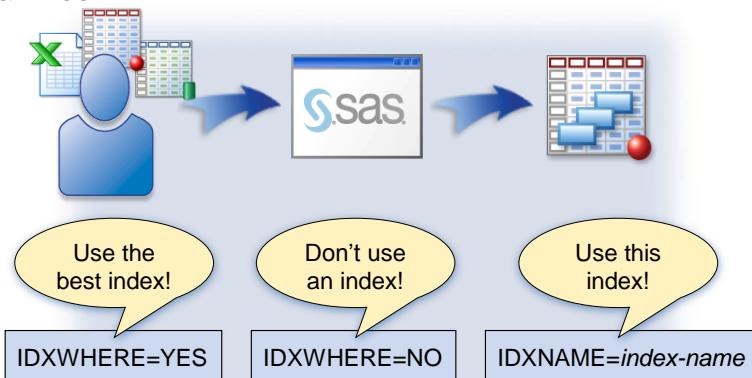
IDXNAME=*index-name*

...

66

Controlling Index Usage for WHERE Processing

Data set options enable you to direct SAS when to use an index.



67

A.3 Creating Indexes

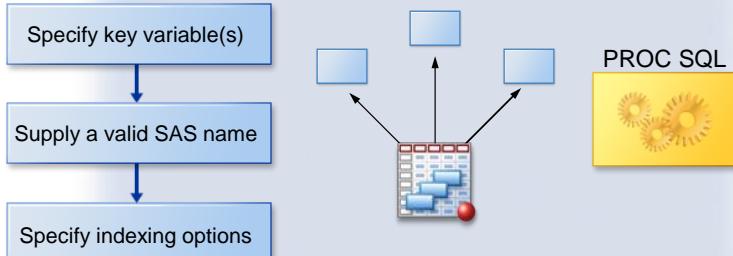
Objectives

- Create an index.
- Explain the costs and benefits of indexes.

70

Business Scenario

Indexes need to be created on an existing data set using PROC SQL to optimize subsetting and join queries.



71

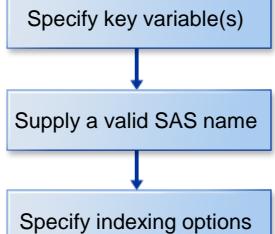
Creating an Index Using PROC SQL

```

proc sql;
create unique index Employee_ID
  on work.employee_addresses
  (Employee_ID)
create index Locale
  on work.employee_addresses
  (State, City);
quit;
    
```

```

PROC SQL;
CREATE <option> INDEX index-name
  ON table-name(column-name-1, ...
                column-name-n);
QUIT;
    
```



72

s1aad04
...

Creating an Index Using PROC SQL

```
proc sql;
create unique index Employee_ID
  on work.employee_addresses
  (Employee_ID)
create index Locale
  on work.employee_addresses
  (State, City);
quit;
```

Specify key variable(s).

```
PROC SQL;
CREATE <option> INDEX index-name
  ON table-name(column-name-1, ...
    column-name-n);
QUIT;
```

73

s1aad04 ...

Creating an Index Using PROC SQL

```
proc sql;
create unique index Employee_ID
  on work.employee_addresses
  (Employee_ID)
create index Locale
  on work.employee_addresses
  (State, City);
quit;
```

Specify key variable(s).

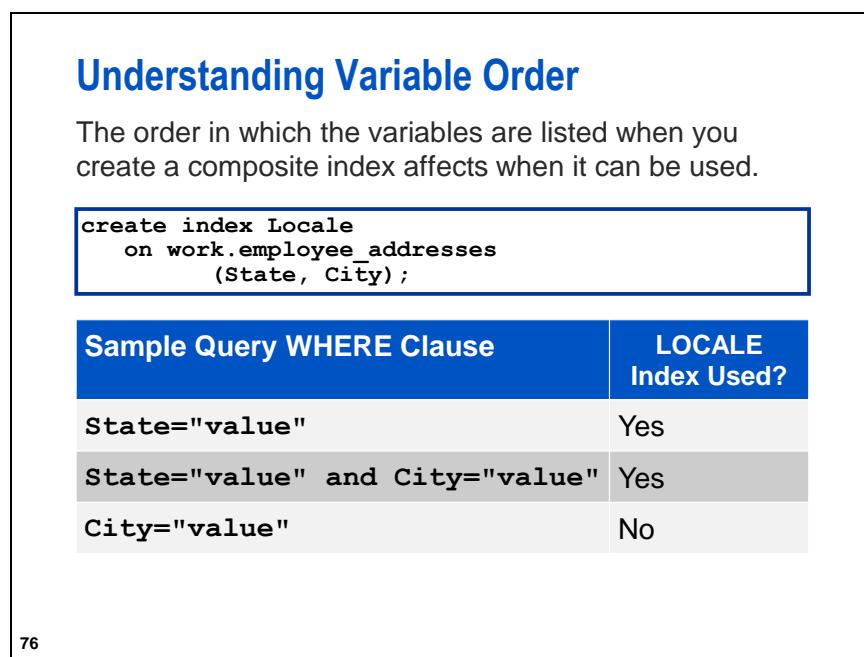
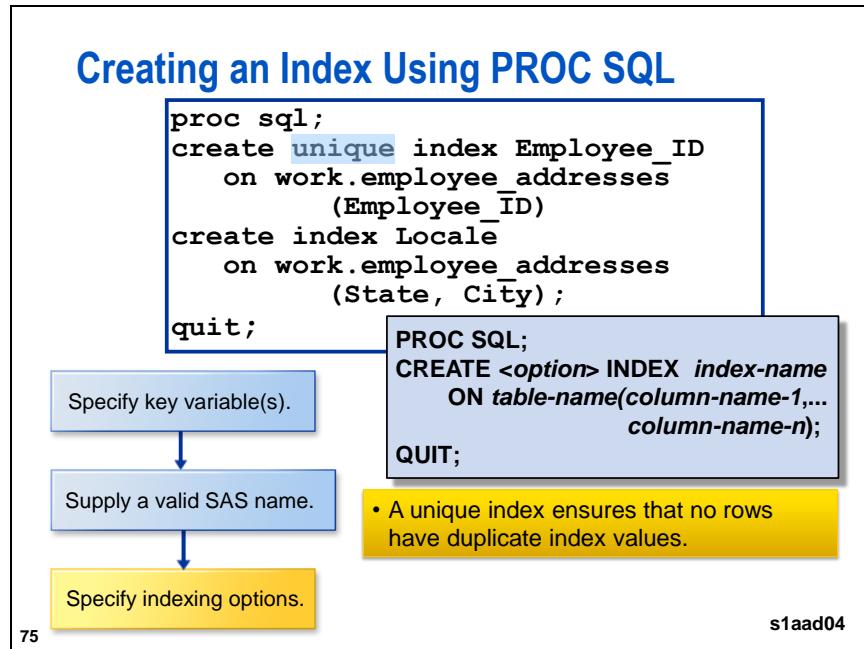
Supply a valid SAS name.

```
PROC SQL;
CREATE <option> INDEX index-name
  ON table-name(column-name-1, ...
    column-name-n);
QUIT;
```

- Simple index names must match the name of the column being indexed.
- Composite index names cannot be the same as any column name in the table.

74

s1aad04 ...



Guidelines for Indexing

DO This	DON'T Do This
Index to retrieve a small subset of observations from a large data set.	Index a data set with fewer than three data file pages.
Index on variables that are discriminating; make the first key variable the most discriminating.	Index on variables with few unique values.
Create and maintain indexes that will be used frequently.	Index frequently changing data unless it is worth the resources.
Sort data before indexing.	Create indexes on rarely used variables.
Minimize the number of indexes.	Try to create a single index for all queries.

77

Performance Tradeoffs

Benefits	Costs
Indexing enables fast access to small ($\leq 15\%$) subsets of data.	CPU cycles and I/O to create indexes
Equijoins can be performed without sorting.	CPU cycles and I/O to update the index whenever data changes
BY-group processing can be performed without sorting.	Memory to load index pages and code for use
Uniqueness can be enforced.	Disk space to store the index file

78

A.04 Quiz

You frequently need to subset the **orion.order_fact** data set, based on **Order_Type**.

Submit the program **s1aaa03** and review the output.

Should you index the **Order_Type** column to improve query performance?



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

3. Creating an Index and Using It in a Query

- a. Open and submit the starter program **s1aae03** to create a temporary table **work.products**, which will contain only unique records from **orion.product_dim**.
- b. Modify the starter program code to create an index on the **work.products** table based on the column **Product_Name**. Remember to give the simple index the same name as the column on which it is based.
- c. Write a query that returns the rows from **work.products** where the **Product_Name** column contains the word “T-Shirt.” Set **options msglevel=I** before running your report to put notes in the SAS log about index usage. Name your report **T-Shirt List**. Review the SAS log. Were any indexes used to process this query? _____

PROC SQL Output

T-Shirt List	
Product Name	Supplier Name
Big Guy Men's T-Shirt	Eclipse Inc
Big Guy Men's T-Shirt Dri Fit	Eclipse Inc
Butch T-Shirt with V-Neck	Luna sastreria S.A.
Eliza T-Shirt	Luna sastreria S.A.
Fred T-Shirt	Luna sastreria S.A.
Instyle T-Shirt	AllSeasons Outdoor Clothing
Kid Children's T-Shirt	3Top Sports
Men's T-Shirt Small Logo	A Team Sports
N.d.gear Basic T-Shirt	Greenline Sports Ltd
O'my Children's T-Shirt with Logo	Luna sastreria S.A.
Osprey Men's King T-Shirt w/Small Logo	Triple Sportswear Inc
Roth T-Shirt	Luna sastreria S.A.
T-Shirt	3Top Sports
T-Shirt, Short-sleeved, Big Logo	A Team Sports
Toncot Beefy-T Emb T-Shirt	A Team Sports
Tony's Cut & Sew T-Shirt	Eclipse Inc
Triffy Logo T-Shirt with V-Neck	Triffy B.V.
Tyfoon Ketch T-Shirt	Typhoon Clothing
Woman's T-Shirt w/Hood	Eclipse Inc
Wyoming Men's T-Shirt with V-Neck	Luna sastreria S.A.

Level 2

4. Creating a Unique Index and Writing a Query Using the Indexed Variable

- Using the program **s1aae04** as starter code, create a temporary table **work.products** that contains only unique records from **orion.product_dim**. Create a unique index on your **work.products** table based on the column **Product_ID**.
- Write a query that joins **work.products** and **orion.order_fact** on **Product_ID** and produces a report of the distinct combinations of **Product_Name** and **Supplier_Name** for records where the year of **Order_Date** is 2011. Name your report as shown in the sample output.

Partial PROC SQL Output

2011 Products Purchased	
Product Name	Supplier Name
A-team Smoothsport Bra	A Team Sports
A-team Sweat Round Neck, Small Logo	A Team Sports
Abdomen Shaper	TrimSport B.V.
Aim4it 80% Tungsten 22 Gram	Royal Darts Ltd

- Name the indexes, if any, selected by the SQL processor to optimize the query:
-
-

Challenge

5. Creating Indexes and Writing a Query Using the Indexed Variables

- Use the program **s1aae05** as starter code.
 - Create temporary tables: **work.customers**, **work.orders**, and **work.products**.
 - Create a unique index on the **work.customers** table based on the column **Customer_ID**.
 - Create an index on the **work.orders** table based on the column **Customer_ID**.
 - Create a unique index on the **work.products** table based on the column **Product_ID**.
- Write a query that joins **work.customers**, **work.orders**, and **work.products** to produce a report that contains **Product_Name** and **Supplier_Name** for orders placed in 2011 by customers from South Africa (country code **ZA**). The report should include only one line for each combination of **Product_Name** and **Supplier_Name**. Name your report appropriately.

PROC SQL Output:

Products Ordered by South African Customers in 2011	
Product Name	Supplier Name
A-team Smoothsport Bra	A Team Sports
Buzz Saw	CrystalClear Optics Inc
Children's Roller Skates	Magnifico Sports
Dartsharpener Key ring	Royal Darts Ltd
Fred T-Shirt	Luna sastreria S.A.
Goodtime Toilet Bag	Luna sastreria S.A.
Hot Mini Backboard Bulls	Van Dammeren International
Pro-roll Hot Rod Roller Skates	Magnifico Sports
Pro-roll Lazer Roller Skates	Magnifico Sports
Proskater Kitalpha Gamma Roller Skates	Roll-Over Inc
Proskater Viablade S Roller Skates	Roll-Over Inc
Rollerskate Roller Skates	Magnifico Sports
Ex9 76mm/78a Biofl	

- c. Name the indexes, if any, selected by the SQL processor to optimize the query.
-

End of Exercises

A.4 Maintaining Tables with the SQL Procedure

Objectives

- Update or delete data values in an existing table.
- Add or drop columns in a table.
- Add, drop, or modify the attributes of columns in a table.
- Delete tables and views.

84

Business Scenario

Your manager asked you to maintain tables.

You need to learn the different PROC SQL techniques to add, modify, or delete data.



85

Review: Adding Data to a Table

The INSERT statement can be used to add data to an empty table, or to append data to a table that already contains data, using one of three methods.

Method	Description	Syntax
A	One clause per row using column-value pairs	INSERT INTO <i>table-name</i> SET <i>column-name=value,</i> <i>column-name=value,...;</i>
B	One clause per row using positional values	INSERT INTO <i>table-name</i> <i><(column list)></i> VALUES (<i>value,value,...</i>);
C	A query returning multiple rows, and based on positional values	INSERT INTO <i>table-name</i> <i><(column list)></i> SELECT <i>columns</i> FROM <i>table-name;</i>

86

Business Scenario

All of the Level I sales representatives have been given a 5% raise. You need to update the **sales** table to reflect the new salary for the Level I sales representatives.



87

Modifying Data Values in Existing Rows

Give Level 1 sales representatives a 5% raise.

```
UPDATE table-name
   SET column-name=expression
      <, ...column-name=expression>
 WHERE expression;
```

```
proc sql;
update work.sales
  set Salary=Salary*1.05
 where Job_Title='Sales Rep. I';
quit;
```



Omitting the WHERE expression causes **all** rows to be updated.

88

s1aad05

Viewing the Log

Partial SAS Log

```
proc sql;
update work.sales
  set Salary=Salary * 1.05
 where Job_Title='Sales Rep. I';
  63 rows were updated in WORK.SALES.
quit;
```

89

Business Scenario

The company is reorganizing, and will use new descriptors for each employee's job mastery level.

A new column, **Level**, was added to the **work.staff** table to contain this information.

Your job is to populate the data, based on the following rules:

If Current Job Title Ends in	The New Level Is
I	Apprentice
II	J Journeyman
III	M Master
IV	M Mentor

90

Business Data

Current Job Title	New Level Value
Manager	
Secretary I	Apprentice
Office Assistant II	J Journeyman
Office Assistant III	M Master
Security Guard IV	M Mentor

91

Review: Conditional Processing

```
UPDATE table-name
  SET column1
  CASE <case-operand>
    WHEN when-condition THEN result-expression
    <WHEN when-condition THEN result-expression>
    <ELSE result-expression>
  END <AS column>;
```

92

Conditional Data Modification

```
proc sql;
update work.staff
  set Level=
    case (scan(Job_Title,-1))
      when 'I' then 'Apprentice'
      when 'II' then 'Journeyman'
      when 'III' then 'Master'
      when 'IV' then 'Mentor'
      else ''
    end;
select Employee_ID, Job_Title, Level
  from work.staff;
quit;
```

93

s1aad06

Viewing the Output

Partial Output

Employee ID	Employee Job Title	Level
120101	Director	
120102	Sales Manager	
120103	Sales Manager	
120104	Administration Manager	
120105	Secretary I	Apprentice
120106	Office Assistant II	Journeyman
120107	Office Assistant III	Master
120108	Warehouse Assistant II	Journeyman
120109	Warehouse Assistant I	Apprentice
120110	Warehouse Assistant III	Master
120111	Security Guard II	Journeyman
120112	Security Guard I	Apprentice
120113	Security Guard II	Journeyman

94

Business Scenario

A Human Resources manager asked you to delete all inactive employees in the **orion.staff** table.



95

Deleting Rows from a Table or View

Delete all terminated employees from the **work.staff** table.

**DELETE FROM table/view
WHERE expression;**

```
proc sql;  
delete from work.staff  
  where Emp_Term_Date is not missing;  
quit;
```

Partial SAS Log

NOTE: 116 rows were deleted from WORK.STAFF.



A DELETE statement without a WHERE expression deletes **all** rows and leaves only the table structure.

96

s1aad07

A.05 Quiz

Submit the program **s1aaa04** and review the SAS log.

How many rows were deleted from the temporary table **work.staff**?

97

Maintaining Table Data: Job Aid

Method	Result
INSERT – SET – VALUES – SELECT	Appends new rows of data to the end of existing tables.
UPDATE – SET	Modifies values in existing data rows as specified in the WHERE expression.
DELETE	Deletes existing rows as specified in the WHERE expression.

99

Business Scenario

Your manager asked you to maintain tables.

Table maintenance includes techniques to add, modify, or delete columns in a table.



101

Altering Table Structure: Adding Columns

Use the ALTER statement to add columns to a table.

```
ALTER TABLE table-name
    ADD column-definition< , ... column-definition>;
```

```
proc sql;
  alter table work.employee_payroll
    add Bonus num format=commal0.2,
    Level char(3);
quit;
```

Partial SAS Log

```
NOTE: Table WORK.EMPLOYEE_PAYROLL has been modified, with 10 columns.
```

102

s1aad08

Altering Table Structure: Dropping Columns

Use the ALTER statement to drop columns from a table.

```
ALTER TABLE table-name
    DROP column-1< , ...column-2>;
```

```
proc sql;
  alter table work.employee_payroll
    drop Birth_Date;
quit;
```

Partial SAS Log

```
NOTE: Table WORK.EMPLOYEE_PAYROLL has been modified, with 9 columns.
```

103

s1aad08

Altering Table Structure: Modifying Columns

Use the ALTER statement to modify column attributes in a table. You can alter a column's length, informat, format, and label.

```
ALTER TABLE table-name
    MODIFY column-definition<, ...column-definition>;
```

```
alter table work.employee_payroll
  modify Bonus num format=comma8.2,
  Level char(10)
  label='Employee Level';
```

Partial SAS Log

```
NOTE: Table WORK.EMPLOYEE_PAYROLL has been modified, with 9 columns.
```

104

s1aad08

A.06 Multiple Answer Poll

Which column attributes cannot be changed with the ALTER statement?

- a. data type
- b. format
- c. label
- d. length
- e. name

105

Business Scenario

A Human Resources manager asked you to add two new columns for an employee's first and last names and to modify the format for the **Birth_Date** column in the **work.staff** table.

When the new columns are available, populate the columns with the appropriate values found in the **Name** column.



107

Business Data

Partial **work.staff**

Name	BirthDate
Abbott, Ray	-386
Aisbitt, Sandy	2943
Anger, Rose	11071
Anstey, David	11731



Partial updated **work.staff**

Name	First	Last	BirthDate
Abbott, Ray	Ray	Abbott	12/11/1958
Aisbitt, Sandy	Sandy	Aisbitt	01/22/1968
Anger, Rose	Rose	Anger	04/24/1990
Anstey, David	David	Anstey	02/13/1992

108

Adding and Modifying Columns

Alter the **work.staff** table as follows:

1. Add new columns named **First** and **Last**.
2. Format the **Birth_Date** column with the **MMDDYY10.** format.

```
proc sql;
  alter table work.staff
    add First char(10), Last Char(20)
    modify Birth_Date date format=mmddyy10.;
```

109

s1aad09

Populating Newly Added Columns

3. Populate columns **First** and **Last**.

```
proc sql;
  alter table work.staff
    add First char(10), Last Char(20)
    modify Birth_Date date format=mmddyy10.;

  update work.staff
    set First=strip(scan(Name,2,',','')),
        Last=strip(scan(Name,1,',',''));
```

STRIP(argument)

The STRIP function returns a character string with all leading and trailing blanks removed.

110

s1aad09

Testing Your Work

4. Test the results.

```
update work.staff
  set First=strip(scan(Name,2,',',')),
      Last=strip(scan(Name,1,',','));
select Name, First, Last
  from work.staff
 where First ne strip(scan(Name,2,',',''))
   or Last ne strip(scan(Name,1,',','));
```

111

s1aad09

Viewing the Output

PROC SQL Output

Name	First	Last
Boocks, Michael. R.	Michael. R	Boocks
Plybon, John-Michael	John-Micha	Plybon



The first name is truncated. The column needs to be wider.

112

Modifying Columns

5. Increase the length of column **First** to 20 characters and reload the **First** column.

```
select Name, First, Last
  from work.staff
 where First ne strip(scan(Name,2,','))
   or Last ne strip(scan(Name,1,','));
alter table work.staff
  modify First char(20);
update work.Sales_staff
  set First=strip(scan(Name,2,','));
```

113

s1aad09

Testing Your Work

6. Re-test the results.

```
update work.staff
  set First=strip(scan(Name,2,','));
select Name, First, Last
  from work.staff
 where First ne strip(scan(Name,2,','))
   or Last ne strip(scan(Name,1,','));
```

Partial SAS Log

NOTE: No rows were selected.

- ✍ No output means that the data in columns **First** and **Last** is as expected. Problem resolved!

114

s1aad09

A.07 Quiz

Both of these code segments change the label for the **Phone_Type** column in **orion.employee_phones**. Which is more efficient?

a.

```
create table orion.employee_phones as
select Employee_ID,
       Phone_Type 'Work, Home, Cell',
       Phone_Number
  from orion.employee_phones;
```

b.

```
alter table orion.employee_phones
    modify Phone_type 'Work, Home, Cell';
```

115

Business Scenario

Your manager asked you to maintain tables.

You also need to know how to delete a table or SQL view.



117

Deleting Tables and Views

Use the DROP statement to delete a table or a view.

```
DROP TABLE table-name<, ...table-name>;
DROP VIEW view-name<, ...view-name>;
```

```
proc sql;
drop table work.employee_addresses;
quit;
```

Partial SAS Log

```
NOTE: Table WORK.EMPLOYEE_ADDRESSES has been dropped.
```

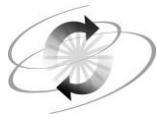
118

s1aad10

Maintaining Tables: Job Aid

Method	Result
INSERT – SET – VALUES	Appends new data rows to the end of existing tables.
UPDATE – SET	Modifies values in existing data rows as specified in the WHERE expression.
DELETE	Deletes existing rows as specified in the WHERE expression.
ALTER – ADD – DROP – MODIFY	Modifies table structure by adding columns, dropping columns, or altering column attributes. (You cannot modify name or type.)
DROP	Deletes an entire table or view.

119



Exercises

If you restarted your SAS session since the last exercise, open and submit the **libname.sas** program found in the data folder.

Level 1

6. Modifying the Table Structure and the Table Data

- a. Marketing determined that the reason that product 210200200022 is not selling well is because it is named “Sunfit Slow Swimming Trunks.” Change the product name and packaging to **Sunfit Speedy Swimming Trunks**.

The program **s1aae06** creates the table **work.products** for you to use in this exercise.

Write a query to update **work.products** by setting **Product_Name** to **Sunfit Speedy Swimming Trunks** where **Product_ID=210200200022**.

- b. Orion Star wants to add a new line of snorkeling products. Add the following rows to the table **work.products**. (Before you start typing, see the comments in the program **s1aae06**.)

Product Information					Supplier Information		
ID	Line	Category	Group	Name	Country	Name	ID
240600100202	Sports	Swim Sports	Snorkel Gear	Coral Dive Mask - Med	AU	Dingo Divers	21001
240600100203	Sports	Swim Sports	Snorkel Gear	Coral Dive Mask - Large	AU	Dingo Divers	21001
240600100212	Sports	Swim Sports	Snorkel Gear	Coral Dive Fins - Med	AU	Dingo Divers	21001
240600100213	Sports	Swim Sports	Snorkel Gear	Coral Dive Fins - Large	AU	Dingo Divers	21001
240600100222	Sports	Swim Sports	Snorkel Gear	Coral Advanced Snorkel	AU	Dingo Divers	21001
240600100223	Sports	Swim Sports	Snorkel Gear	Coral Pro Snorkel	AU	Dingo Divers	21001

- c. Delete the table **work.products**.

Level 2

7. Modifying the Table Structure and the Table Data

- a. Orion Star will drop the entire group of Kid’s Eclipse products from the catalog. The starter program **s1aae07** creates the tables **work.products** and **work.new_products** to use in this exercise. Delete the rows containing information for these products from the table **work.products**. A list of the discontinued products is shown below:

Product Information				
ID	Line	Category	Group	Name
210200300006	Children	Children Sports	Eclipse, Kid's Clothes	Fleece Cuff Pant Kid's
210200300007	Children	Children Sports	Eclipse, Kid's Clothes	Hsc Dutch Player Shirt Junior
210200300052	Children	Children Sports	Eclipse, Kid's Clothes	Tony's Cut & Sew T-Shirt
210200400020	Children	Children Sports	Eclipse, Kid's Shoes	Kids Baby Edge Max Shoes
210200400070	Children	Children Sports	Eclipse, Kid's Shoes	Tony's Children's Deschutz (Bg) Shoes

- b. Add the products in the **work.new_products** table to the **work.products** table.
- c. Top Sports (**Supplier_ID=755**) experienced delays of one week in shipping all sleeping bags (**Product_Group="Sleepingbags"**) from the factory. Modify **work.products** to add a numeric **Shipping_Delay** column. Update the column **Shipping_Delay** in the **work.products** table to show a delay of seven days in shipping Top Sports sleeping bags.
- d. Write a query to display **Supplier_Name**, **Product_Group**, **Product_Name**, and **Shipping_Delay** for all rows in **work.products** where **Shipping_Delay** is not missing. Title the report as shown in the sample output.

PROC SQL Output

Product Shipping Delays		
Product Name	Supplier Name	Shipping_Delay
Basic 10, Left , Yellow/Black	Top Sports	7
Expedition Zero,Medium,Left,Charcoal	Top Sports	7
Expedition Zero,Medium,Right,Charcoal	Top Sports	7
Expedition Zero,Small,Left,Charcoal	Top Sports	7
Expedition Zero,Small,Right,Charcoal	Top Sports	7
Expedition10,Medium,Right,Blue Ribbon	Top Sports	7
Expedition 10,Small,Left,Blue Ribbon	Top Sports	7
Expedition 10,Small,Right,Blue Ribbon	Top Sports	7
Expedition 20,Large,Right,Forestgreen	Top Sports	7
Expedition 20,Medium,Right,Forestgreen	Top Sports	7
Outback Sleeping Bag, Large,Left,Blue/Black	Top Sports	7
Outback Sleeping Bag, Large,Right, Blue/Black	Top Sports	7

- e. Remove the column **Shipping_Delay** from the **work.products** table. Write a query that describes the **work.products** table structure in the log, and review the SAS log to make sure that the column was actually dropped.
- f. Delete both the **work.products** and **work.new_products** tables.

Challenge

8. Maintaining Tables

- a. The program **s1aae08** creates the **work.products**, **work.new_products**, and **work.supplier** tables for you to use in this exercise. Information about new products is in the **work.new_products** table. Information about the supplier of these new products is in the **work.supplier** table. Add all of the products in the **work.new_products** table to the **work.products** table, including the appropriate supplier information for each row.

Hint: There might be data type mismatches between some of the columns in the **work.products** and **work.new_products** tables.

NOTE: The execution of this query involves performing one or more Cartesian product joins that cannot be optimized.

- b. When the data is successfully added to **work.products**, delete the **work.new_products** and **work.supplier** tables.
- c. For products in the category **Swim Products**, query **work.products** to display **Supplier_Country** and the percentage of swim products supplied by each country. The percentage should be calculated as follows:

number of swim products supplied by this country / total number of swim products

Give the report a title as shown in the sample output.

PROC SQL Output

Swim Product Sourcing	
	Supplier Country
AU	67%
US	19%
ES	14%

End of Exercises

A.5 Integrity Constraints

Objectives

- Define integrity constraints.
- Apply integrity constraints to a table.
- Identify and correct integrity constraint violations.

123

Business Scenario

The CFO mandated that no item would be discounted more than 50%. The vice president of Sales requested your help setting this constraint on the **discounts** table.

discounts



Values in the **Discount** column must be less than or equal to .5.

124

Creating Integrity Constraints with PROC SQL

Re-create the **discounts** table with an integrity constraint limiting discounts to 50%.

```
CREATE TABLE table
  (column-specification, ...
   <constraint-specification, ...>);
```

```
create table discounts
  (Product_ID num format=z12.,
   Start_Date date,
   End_Date date,
   Discount num format=percent.,
   constraint ok_discount
   check (Discount le .5));
```

Integrity constraints are assigned as part of the table definition.

125

s1aad11



Integrity constraints must have unique names. Use variable naming conventions when you name an integrity constraint.

- Names consist of 1 to 32 characters.
- Names must begin with a letter or an underscore character. The remaining characters must be a combination of letters, underscores, and numbers.

Integrity Constraints

- Integrity constraints* are rules enforced when data is added to a table to guarantee data validity.
- To preserve the consistency and correctness of your data, specify integrity constraints for the SAS data file.
- SAS uses integrity constraints to validate data values when you insert or update columns for which you defined integrity constraints.

126

Integrity Constraints

Integrity constraints

- follow ANSI standards
- cannot be defined for views
- can be specified when a table is created
- can be added to a table that already contains data
- are commonly found in large database management systems (DBMS) with frequently updated tables.

127

Five Integrity Constraints

General:

- NOT NULL
- CHECK
- UNIQUE

Referential:

- PRIMARY KEY
- FOREIGN KEY

128

NOT NULL	Requires data in this column; missing values are not allowed.
CHECK	Specifies rules to check when values are entered in this column. Attempting to insert or update data in this column that violates these rules causes the update or insertion to be rejected.
UNIQUE	Ensures that every value in a column is unique. Data might be missing, but if one row has a missing value in this column, no other row can have a missing value. All rows must have unique values for this column.

PRIMARY KEY	Identifies a column or columns in a table as the table's primary key. A unique, nonmissing value is required for the primary key. The effect emulates specifying both NOT NULL and UNIQUE constraints.
FOREIGN KEY	Limits the values entered into one or more columns in this table to a list of values provided by the PRIMARY KEY column(s) found in another table. A column with a FOREIGN KEY constraint might have a missing value, but any nonmissing values entered must be one of the values found in the PRIMARY KEY columns of the related table.

Integrity Constraint Violations

Example: Insert three rows of data.

```
insert into Discounts
  values (240500200009, '01Mar2007'd,
          '31Mar2007'd, .45)
  values (220200200036, '01Mar2007'd,
          '31Mar2007'd, .54)
  values (220200200038, '01Mar2007'd,
          '31Mar2007'd, .25);
```

This could be
a costly typo!

129

s1aad11

Integrity Constraint Violations

Partial Log

```
ERROR: Add/Update failed for data set WORK.DISCOUNTS because data
      value(s) do not comply with integrity constraint ok_discount.
NOTE: This insert failed while attempting to add data from VALUES
      clause 2 to the data set.
NOTE: Deleting the successful inserts before error noted above to
      restore table to a consistent state.
```

A constraint violation invokes an UNDO process.

- PROC SQL deletes **all** rows inserted or modified by the errant INSERT INTO or UPDATE statement.
- The table returns to a consistent state, that is, to the condition existing before the statement executed.

130

Controlling the UNDO_POLICY Option

Changing the UNDO_POLICY option in PROC SQL gives you control over how UNDO is performed when integrity constraints are violated.

You can choose to enable rows that do not violate integrity constraints to remain in the table, while you reject only the rows that contain bad data.

131

Controlling the UNDO_POLICY Option

- UNDO_POLICY=REQUIRED (default)
undoes all inserts or updates up to the point of the error. Sometimes the UNDO operation cannot be accomplished reliably.
- UNDO_POLICY=NONE
rejects only rows that violate constraints. Rows that do not violate constraints are inserted.
- UNDO_POLICY=OPTIONAL
operates in a manner similar to REQUIRED when the UNDO operation can be accomplished reliably; otherwise, it operates similar to NONE.

132



The COMMIT and ROLLBACK statements are ANSI standards that are not currently supported by PROC SQL. In PROC SQL, insertions and updates are immediately committed.

UNDO_POLICY=REQUIRED (This is the default.)

PROC SQL reverses all previous row inserts or updates up to the point of the error. Sometimes the UNDO operation cannot be performed reliably. Examples of conditions where changes cannot reliably be reversed include the following:

- When you update or insert data using a SAS/ACCESS view, PROC SQL might not be able to reverse the effects of the statement without also reversing the effects of other changes that were made at the same time.
- When you update or insert data in a SAS data set opened with row-level locking (CNTLLEV=RECORD) through a SAS/SERVEr, other users can edit any row except the one you are currently updating, including rows previously updated by your current INSERT or UPDATE statement.

UNDO_POLICY=NONE

PROC SQL keeps any compliant updates or inserts, and skips updates or inserts that were not compliant.

UNDO_POLICY=OPTIONAL

PROC SQL reverses any updates or inserts that it can reverse reliably. Otherwise, PROC SQL keeps any compliant updates or inserts, and skips updates or inserts that were not compliant.

This option enables PROC SQL to decide which of the first two options to apply, based on current conditions. PROC SQL proceeds as if UNDO_POLICY=REQUIRED were in effect if the UNDO can be reliably performed. Otherwise, PROC SQL proceeds as if UNDO_POLICY=NONE were specified.

Setup for the Poll

Program **s1aaa05** re-creates the table **discounts** with integrity constraints and attempts to insert three rows of data, using UNDO_POLICY=NONE.

- Submit the program **s1aaa05**.
- Review the SAS log.

```
proc sql undo_policy=none;
  insert into Discounts
    values (240500200009,'01Mar2007'd,
            '31Mar2007'd,.45)
    values (220200200036,'01Mar2007'd,
            '31Mar2007'd,.54)
    values (220200200038,'01Mar2007'd,
            '31Mar2007'd,.25);
quit;
```

A.08 Poll

Were the rows containing good data rejected along with the rows containing bad data when you used the PROC SQL option UNDO_POLICY=NONE?

- Yes
- No

134

Troubleshooting Integrity Constraint Violations

When an integrity constraint is violated, the SAS log identifies which VALUES clause contained the error and names the violated integrity constraint.

To correct the problem, you need more information about the violated integrity constraint.

136

Troubleshooting Integrity Constraint Violations

Show the column attributes and integrity constraints for the **discounts** table.

```
proc sql;
  describe table constraints discounts;
quit;
```

```
PROC SQL;
  DESCRIBE TABLE CONSTRAINTS table-name
    <, ...table-name>;
QUIT;
```

DESCRIBE statements produce output in the SAS log.

137

s1aad12

The DESCRIBE TABLE statement (without the CONSTRAINTS keyword) writes a CREATE TABLE statement to the SAS log for the specified table regardless of how the table was originally created (for example, with a DATA step).

If the table contains an index, CREATE INDEX statements for those indexes are also written to the SAS log.

Viewing the Log

Partial SAS Log

-----Alphabetic List of Integrity Constraints-----			
#	Integrity Constraint	Type	Where Clause
1	ok_discount	Check	Discount<=0.5

138

Troubleshooting Integrity Constraint Violations

Statement	Results Produced in the SAS Log
DESCRIBE VIEW	SQL code that would create a view identical to the view being described
DESCRIBE TABLE	SQL code that would create a table identical to the table being described (including indexes) and a description of the table's integrity constraints
DESCRIBE TABLE CONSTRAINTS	A description of the table's integrity constraints

139

Setup for the Poll

The program **s1aaa06** re-creates the table **discounts** with integrity constraints and attempts to insert three rows of data, using UNDO_POLICY=NONE.

- Submit the program **s1aaa06**.
- Review the SAS log.
- Correct the second VALUES clause to provide a discount of 0.45.
- Resubmit **only** the INSERT INTO PROC SQL step.
- Review the SAS log.

140

s1aaa06

A.09 Multiple Choice Poll

After you correct and resubmit only the INSERT INTO query using UNDO_POLICY=NONE, how many rows are in the **discounts** table?

- a. 5
- b. 3
- c. 2

141

A.6 Interfacing PROC SQL with the Macro Language

Objectives

- Create and use SAS macro variables in PROC SQL.
- Insert information from dictionary tables into the SAS macro variables to create a self-modifying SQL query.

145

Business Scenario

Orion Star Club Members are categorized into tiers based on their purchasing volume. The number of tiers increases over time. For each tier, create a report that includes the customer's name and country. Include the tier number and the total number of customers in the tier in the report title.

Tier 1 Customers (48 total)		
Customer Name		Country
Cosi Rim		
Dericka		
...		
Ramesh Tre		
Angel Borw		
...		
Candy Kinsey		AU
Lauren Marx		CA
...		

Tier 2 Customers (21 total)		
Customer Name		Country

Tier 3 Customers (8 total)		
Customer Name		Country

146

Business Data

The **Customer_Type_ID** column is a numeric column in the **customer** table. The first digit of the **Customer_Type_ID** column indicates a customer's tier.

customer	
Customer_Name	Customer_Type_ID
James Kvarniq	1020
Sandrina Stephano	2020
Ulrich Heyde	3010

The diagram shows three green boxes labeled "Tier 1", "Tier 2", and "Tier 3". Arrows point from the value "1020" to "Tier 1", from "2020" to "Tier 2", and from "3010" to "Tier 3".

To extract the first digit, you need to convert **Customer_Type_ID** to a character value.

147

Processing Steps

Creating these reports requires writing multiple queries.

Step 1 Determine the number of tiers.

Step 2 Create macro variables and store tier values and customer counts into macro variables.

Step 3 Use the macro variables to generate reports for each tier.

148

The PUT Function

The PUT function is commonly used to convert numeric values to character. It returns a value using a specified format.

`PUT(source,format.)`

source the SAS variable or constant whose value you want to reformat

format. the SAS format to be applied to *source*

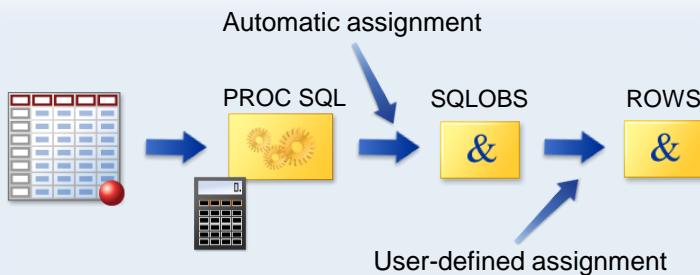
149

SQLOBS Automatic Macro Variable

Executing a PROC SQL statement automatically creates and populates several macro variables.

The **SQLOBS** automatic macro variable records the number of rows (observations) that are processed by an SQL procedure statement. It is assigned a value after the SQL SELECT statement executes.

150



 **SQLOBS** is one of several automatic macro variables that are created when a PROC SQL statement is executed. Refer to *SAS 9.3 SQL Procedure: User's Guide* for specific information about other automatic macro variables.

Creating Dynamic Reports

Step 1 Determine the number of tiers.

```
proc sql;
  select substr(put(Customer_Type_ID,4.),1,1) as Tier,
         count(*)
    from orion.customer
   group by Tier;
%let Rows=&SQLOBS;
%put NOTE: There are &Rows Tiers;
```

 After the first query, **SQLOBS** contains the number of rows (tiers).

151

s1aad13

Viewing the Output

PROC SQL Output

Tier	
1	48
2	21
3	8

1	48
2	21
3	8

Partial SAS Log

%put NOTE: There are &Rows Tiers;
NOTE: There are 3 Tiers

152

Creating a Macro Variable Series: Syntax 2

Syntax 2 extracts values from the first n rows of the query result and inserts these values into a series of n macro variables. Values for rows 1- n in the first column in the SELECT list are placed in a numbered series of macro variables $a1-an$, and so on.

SELECT <i>a, b, ...</i>
INTO <i>:a1-:an, :b1-:bn</i>
FROM <i>table/view ...</i>

153

Creating Dynamic Reports

Step 2 Create macro variables and store tier values and customer counts into macro variables.

```
reset noprint;
select substr(put(Customer_Type_ID, 4.), 1, 1)
      as Tier, count(*)
   into :Tier1-:Tier&Rows, :Count1-:Count&Rows
    from orion.customer
   group by Tier;
%put NOTE: Tier1 is &tier1 Count1 is: &count1;
%put NOTE: Tier2 is &tier2 Count2 is: &count2;
%put NOTE: Tier3 is &tier3 Count3 is: &count3;
```

The second query adds an INTO clause using the value of &Rows to determine the number of macro variables required for the series.

Partial Log

```
NOTE: Tier1 is 1 Count1 is: 48
NOTE: Tier2 is 2 Count2 is: 21
NOTE: Tier3 is 3 Count3 is: 8
```

154

s1aad13

Setup for the Poll

The first query in the program produced output, but the second query produced no output.

```
proc sql;
select substr(put(Customer_Type_ID, 4.), 1, 1)
      as Tier, count(*)
   from orion.customer
  group by Tier;
%let Rows=&SQLOBS;
%put NOTE: There are &Rows Tiers;
reset noprint;
select substr(put(Customer_Type_ID, 4.), 1, 1)
      as Tier, count(*)
   into :Tier1-:Tier&Rows, :Count1-:Count&Rows
    from orion.customer
   group by Tier;
```

155

s1aad13

A.10 Multiple Choice Poll

Why did the first query in the program produce output, but the second query did not?

- a. The INTO clause suppressed output during the execution of the second query.
- b. The RESET NOPRINT statement suppressed output for subsequent queries in this PROC SQL invocation.
- c. The %PUT statement redirected the query results to the SAS log.

156

Creating Dynamic Reports

Step 3 Use the macro variables to generate reports for each tier.

```
reset print;
title "Tier &Tier1 Customers (&Count1 total)";
select Customer_Name, Country 'Country'
  from orion.customer
 where substr(put(Customer_Type_ID,4.),1,1)
       = "&Tier1"
 order by Country, Customer_Name;
```

The query above is copied three times and modified as follows:
 Change &Tier1 to &Tier2 and &Count1 to &Count2 in the second copy.
 Change &Tier1 to &Tier3 and &Count1 to &Count3 in the third copy.

158

s1aad13

Viewing the Output

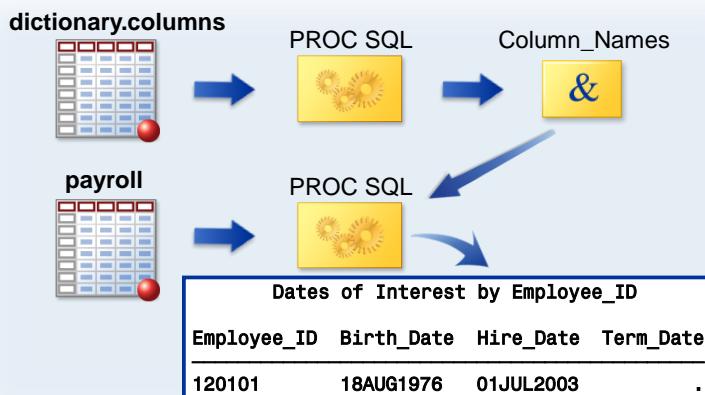
Partial PROC SQL Output

Tier 1 Customers (48 total)	
Customer_Name	Country
Cosi Rimmington	AU
Dericka Pockran	AU
...	
Tier 2 Customers (21 total)	
Customer_Name	Country
Ramesh Trenholme	AU
Angel Borwick	CA
...	
Tier 3 Customers (8 total)	
Customer_Name	Country
Candy Kinsey	AU
Lauren Marx	CA
...	

159

Business Scenario

You maintain a program that produces a report containing **Employee_ID** and all associated date columns from the **orion.employee_payroll** table. Database administrators change date fields in this table. You want to rewrite your program using macro variables to automate the process.



161

Business Data

The payroll data set originally contains three date columns:
Birth_Date, **Hire_Date**, and **Term_Date**.

Dates of Interest by Employee_ID			
Employee_ID	Birth_Date	Hire_Date	Term_Date
120101	18AUG1976	01JUL2003	.

Two additional date columns are added:
Date_Last_Raise and **Promotion_Date**.

Dates of Interest by Employee_ID					
Employee_ID	Birth_Date	Hire_Date	Term_Date	Date_Last_Raise	Promotion_Date
120101	18AUG1976	01JUL2003	.	25JUN04	28DEC03

✍ You want your program to automatically display all dates.

162

Creating Macro Variables: Syntax 3

Syntax 3 extracts values from all rows of the query result and puts them into a single macro variable, separated by the specified delimiter.

```
SELECT column-1<, ... column-2>
  INTO :macvar_1 SEPARATED BY 'delimiter'
        <, ... :macvar_2 SEPARATED BY 'delimiter'>
  FROM table/view ...
```

163



Creating and Using Macro Variables in PROC SQL

1. Open and submit **s1aad14** to create a temporary table for testing.
2. Open and submit **s1aad15**. This program queries the **dictionary.columns** table to store column names in the payroll table that include **date** as a part of the name. A second query produces a report from the payroll tables with the **Employee_ID** column and all other columns that include **date** as a part of the column name.
3. Open and submit **s1aad16**. This program adds and populates two new date columns in the temporary payroll table.
4. Resubmit **s1aad15** to confirm that the program dynamically captures all variables with **date** as a part of the name.

End of Demonstration

A.7 Solutions

Solutions to Exercises

1. Creating a Simple Correlated Subquery

```
*** s1aas01 ***;

proc sql;
title "Australian Employees' Birth Months";
select Employee_ID,
       month(Birth_Date) as BirthMonth format=3.
  from orion.employee_payroll as EP
 where 'AU' =
       (select Country
          from orion.employee_addresses as EA
         where EA.Employee_ID=EP.Employee_ID)
       order by 2;
quit;
title;
```

Alternate Solution with a Noncorrelated Query

```
proc sql;
title "Australian Employees' Birth Months";
select Employee_ID, month(Birth_Date) as BirthMonth
  from orion.employee_payroll
 where Employee_ID in
       (select Employee_ID
          from orion.employee_addresses
         where country='AU')
       order by BirthMonth;
quit;
title;
```

2. Using a Correlated Subquery

```
*** s1aas02 ***:

proc sql;
title "Employees With Donations > 0.02 Of Their Salary";
select Employee_ID, Employee_Gender,
       Marital_Status
  from orion.employee_payroll
 where Employee_ID in
       (select Employee_ID
          from orion.employee_donations
         where employee_payroll.Employee_ID=
               employee_donations.Employee_ID
               and sum(qtr1,qtr2,qtr3,qtr4)/salary > .002)
       order by Employee_ID;
```

```
quit;
title;
```

Alternate Solution Using the EXISTS Condition

```
proc sql;
title "Employees With Donations > 0.02 Of Their Salary";
select Employee_ID, Employee_Gender,
       Marital_Status
  from orion.employee_payroll
 where exists
   (select *
      from orion.employee_donations
     where employee_payroll.Employee_ID=
           employee_donations.Employee_ID
          and sum(qtr1,qtr2,qtr3,qtr4)/salary > .002)
  order by Employee_ID;
quit;
title;
```

3. Creating an Index and Using It in a Query

a.

```
*** s1aas03 ***;

proc sql;
create table work.products as
  select distinct *
    from orion.product_dim;
create index Product_Name
  on products (Product_Name);
quit;
```

b.

```
options msglevel=i;
proc sql;
title "T-Shirt List";
select distinct Product_Name format=$45.,
               Supplier_Name format=$30.
  from products
 where Product_Name like '%T-Shirt%';
quit;
```

c. Yes, the Product_Name index was used to optimize the query.

Partial SAS Log

```
INFO: Index Product_Name selected for WHERE clause optimization.
```

4. Creating a Unique Index and Writing a Query Using the Indexed Variable

a.

```
*** s1aas04 ***;
```

```

proc sql;
create table work.products as
  select distinct *
    from orion.product_dim;
proc sql;
create unique index Product_ID
  on products (Product_ID);
quit;

```

b.

```

options msglevel=i;
proc sql;
title "2011 Products Purchased";
select distinct Product_Name, Supplier_Name
  from products as p, orion.order_fact as o
  where o.product_id=p.product_id
    and year(Order_Date)=2011;
quit;

```

- c. Name the indexes, if any, selected by the SQL processor to optimize the query:

Product_ID on work.Products

5. Creating Indexes and Writing a Query Using the Indexed Variables

a.

```

*** s1aas05 ***;

proc sql;
create table work.customers as
  select distinct *
    from orion.customer;
create table work.orders as
  select *
    from orion.order_fact;
create table work.products as
  select distinct *
    from orion.product_dim;
quit;

proc sql;
create unique index Customer_ID
  on customers (Customer_ID);
create index Customer_ID
  on orders (Customer_ID);
create unique index Product_ID
  on products (Product_ID);
quit;

```

b.

```

options msglevel=i;
proc sql flow=15 40;
title "Products Ordered by South African Customers in 2011";

```

```

select distinct p.Product_Name, p.Supplier_Name
  from customers as c,
       orders as o,
       products as p
 where o.Product_ID=p.Product_ID
   and c.Customer_ID=o.Customer_ID
   and c.Country="ZA"
   and year(o.Order_Date)=2011;
quit;

```

- c. Name the indexes, if any, selected by the SQL processor to optimize the query:

Customer_ID on work.Orders and **Product_ID** on work.Products

6. Modifying the Table Structure and the Table Data

a.

```

*** s1aas06 ***;

proc sql;
update work.products
  set Product_Name ="Sunfit Speedy Swimming Trunks"
  where Product_ID=210200200022;
quit;

```

b.

```

/* Method 1 */
proc sql;
  insert into work.products(Product_ID,Product_Line,
                             Product_Category, Product_Group,
                             Product_Name,Supplier_Country,
                             Supplier_Name, Supplier_ID)
    values(240600100202,"Sports","Swim Sports","Snorkel Gear",
           "Coral Dive Mask - Med","AU","Dingo Divers",21001)
    values(240600100203,"Sports","Swim Sports","Snorkel Gear",
           "Coral Dive Mask - Large","AU","Dingo Divers",21001)
    values(240600100212,"Sports","Swim Sports","Snorkel Gear",
           "Coral Dive Fins - Med","AU","Dingo Divers",21001)
    values(240600100213,"Sports","Swim Sports","Snorkel Gear",
           "Coral Dive Fins - Large","AU","Dingo Divers",21001)
    values(240600100222,"Sports","Swim Sports","Snorkel Gear",
           "Coral Advanced Snorkel","AU","Dingo Divers",21001)
    values(240600100223,"Sports","Swim Sports","Snorkel Gear",
           "Coral Pro Snorkel","AU","Dingo Divers",21001);
quit;

/* Method 2 */
proc sql;
  insert into work.Products
    set Product_ID=240600100202,Product_Line="Sports",
        Product_Category="Swim Sports",Product_Group="Snorkel
        Gear",Product_Name="Coral Dive Mask - Med",

```

```

Supplier_Country="AU", Supplier_Name="Dingo
Divers",Supplier_ID=21001
set Product_ID=240600100203,Product_Line="Sports",
Product_Category="Swim Sports",Product_Group="Snorkel
Gear",Product_Name="Coral Dive Mask - Large",
Supplier_Country="AU",Supplier_Name="Dingo
Divers",Supplier_ID=21001
set Product_ID=240600100212,Product_Line="Sports",
Product_Category="Swim Sports",Product_Group="Snorkel
Gear", Product_Name="Coral Dive Fins - Med",
Supplier_Country="AU",Supplier_Name="Dingo
Divers",Supplier_ID=21001
set Product_ID=240600100213,Product_Line="Sports",
Product_Category="Swim Sports",Product_Group="Snorkel
Gear",Product_Name="Coral Dive Fins - Large",
Supplier_Country="AU",
Supplier_Name="Dingo Divers",Supplier_ID=21001
set Product_ID=240600100222,Product_Line="Sports",
Product_Category="Swim Sports",Product_Group="Snorkel
Gear", Product_Name="Coral Advanced
Snorkel",Supplier_Country="AU",
Supplier_Name="Dingo Divers",Supplier_ID=21001
set Product_ID=240600100223,Product_Line="Sports",
Product_Category="Swim Sports",Product_Group="Snorkel
Gear", Product_Name="Coral Pro Snorkel",
Supplier_Country="AU",
Supplier_Name="Dingo Divers",Supplier_ID=21001;
quit;

```

c.

```

proc sql;
drop table work.products;
quit;

```

7. Modifying the Table Structure and the Table Data

a.

```

*** s1aas07 ***;
proc sql;
delete from work.products
where Product_Group like '%Eclipse, Kid%';
quit;

```

b.

```

proc sql;
insert into work.products
select *
from work.new_products;

```

```

select Product_Name, Supplier_Name, Supplier_Country
  from work.products
 where Product_Category="Swim Sports";
quit;

```

c.

```

proc sql;
alter table work.products
  add Shipping_Delay num;
update work.products
  set Shipping_Delay=7
  where Supplier_ID=755
        and Product_Group="Sleepingbags";
quit;

```

d.

```

proc sql;
title "Product Shipping Delays";
select Product_Name, Supplier_Name, Shipping_Delay
  from work.products
 where not Shipping_Delay is missing;
quit;
title;

```

e.

```

proc sql;
alter table work.products
  drop Shipping_Delay;
describe table work.products;
quit;

```

f.

```

proc sql;
drop table work.products;
drop table work.new_products;
quit;

```

8. Maintaining Tables

a.

```

*** s1aas08 ***;
proc sql;
insert into work.Products
(Product_ID,Product_Line, Product_Category, Product_Group,
 Product_Name ,Supplier_Country, Supplier_Name, Supplier_ID)
select input(Product_ID,12.),Product_Line, Product_Category,
       Product_Group, Product_Name ,Country,
       Name, ID
  from work.new_products as p,
       work.supplier as s;
quit;

```

b.

```
proc sql;
drop table work.new_products;
drop table work.supplier;
quit;
```

c.

```
proc sql;
title "Swim Product Sourcing";
select Supplier_Country,
       Product_Count/SUM(Product_Count) format=percent6.1
  from (select Supplier_Country, Count(*) as Product_Count
         from work.products
        where Product_Category="Swim Sports"
      group by Supplier_Country)
 order by 2 desc;
quit;
```

End of Solutions

Solutions to Student Activities (Polls/Quizzes)

A.01 Quiz – Correct Answer

Given the following query, subquery, and data in **work.supervisors**, what is the maximum number of rows that will be selected by the outer query?

only the three managers where Country='AU' would be selected

```
proc sql;
select Employee_ID,
       Employee_Name as Manager_Name
  from orion.employee_addresses
 where 'AU'=
       (select Country
          from work.supervisors
         where employee_addresses.Employee_ID=
               supervisors.Employee_ID) ;
quit;
```

Employee_ID	Country
120798	US
120800	US
120104	AU
120735	US
121141	US
...	...
120262	US
120679	US
120103	AU
120668	US
121143	US
120260	US
120672	AU

12

A.02 Poll – Correct Answer

Is it necessary to qualify the column names in the inner WHERE clause as follows?

```
where sales.Employee_ID=order_fact.Employee_ID
```

- Yes
- No

36

A.03 Multiple Choice Poll – Correct Answer

Which query (or queries) used the **Locale** index?

- a. Query 1 (State)
- b. Query 2 (State and City)
- c. Query 3 (City)
- d. Queries 1 and 2
- e. **Queries 1 and 3**
- f. all queries used the index
- g. none of the above

65

A.04 Quiz – Correct Answer

Should you index the **Order_Type** column to improve query performance?

No, Order_Type is a low-cardinality column, and none of the values returns a subset less than 15%. A full-table scan yields better performance without the overhead required to maintain an index.

Order_Type	Frequency	Percent
1	324	52.51
2	170	27.55
3	123	19.94

80

A.05 Quiz – Correct Answer

How many rows were deleted from the temporary table work.staff? 424 (all of them)

```
proc sql;
  delete from work.staff;
  where Emp_Term_Date is not missing;
quit;
```

This semicolon ends this statement.

SQL cannot interpret the stand-alone WHERE clause.

```
delete from work.staff;
NOTE: 424 rows were deleted from WORK.STAFF.
      where Emp_Term_Date is not missing;
WARNING: This SAS global statement is not supported in PROC
          SQL. It has been ignored.
```

- ✍ Remember that PROC SQL queries execute immediately.

98

A.06 Multiple Answer Poll – Correct Answer

Which column attributes cannot be changed with the ALTER statement?

- a. data type
- b. format
- c. label
- d. length
- e. name

106

A.07 Quiz – Correct Answer

Both of these code segments change the label for the **Phone_Type** column in **orion.employee_phones**. Which is more efficient?

a.

```
create table orion.employee_phones as
select Employee_ID,
       Phone_Type 'Work, Home, Cell',
       Phone_Number
  from orion.employee_phones;
```

This query reads and writes all the data rows in the table to effect the changes.

b.

```
alter table orion.employee_phones
      modify Phone_type 'Work, Home, Cell';
```

This query modifies only the table metadata.

116

A.08 Poll – Correct Answer

Were the rows containing good data rejected along with the rows containing bad data when you used the PROC SQL option UNDO_POLICY=NONE?

- Yes
 No

Partial Log

```
ERROR: Add/Update failed for data set WORK.DISCOUNTS because data
      value(s) do not comply with integrity constraint ok_discount.
NOTE: This insert failed while attempting to add data from VALUES
      clause 2 to the data set.
NOTE: 3 rows were inserted into WORK.DISCOUNTS --
      of these 1 row was rejected as an ERROR,
      leaving 2 rows that were inserted successfully.
```

135

A.09 Multiple Choice Poll – Correct Answer

After you correct and resubmit only the INSERT INTO query using UNDO_POLICY=NONE, how many rows are in the **discounts** table?

- a. 5
- b. 3
- c. 2

 If you specify UNDO_POLICY=NONE when correcting for constraint violations, ensure that you re-submit only the corrected data rows, or you might inadvertently add unwanted duplicates of the original non-rejected rows to the table.

142

A.10 Multiple Choice Poll – Correct Answer

Why did the first query in the program produce output, but the second query did not?

- a. The INTO clause suppressed output during the execution of the second query.
- b. The RESET NOPRINT statement suppressed output for subsequent queries in this PROC SQL invocation.
- c. The %PUT statement redirected the query results to the SAS log.

157

